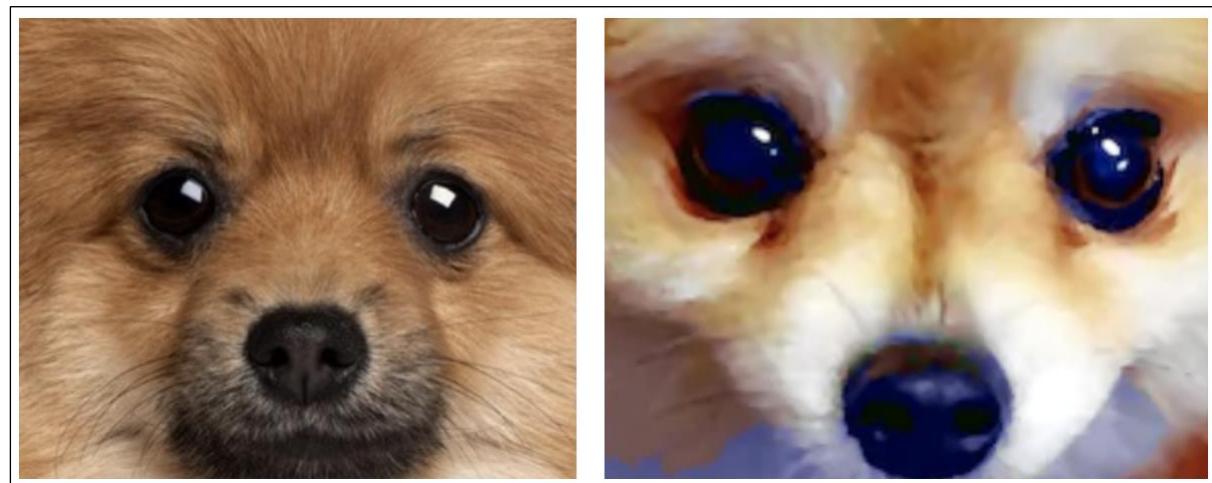
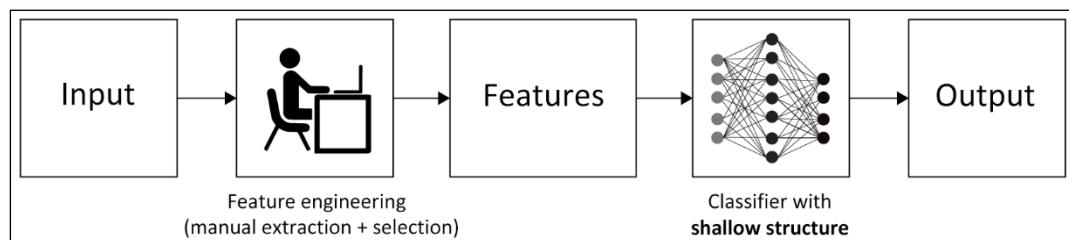
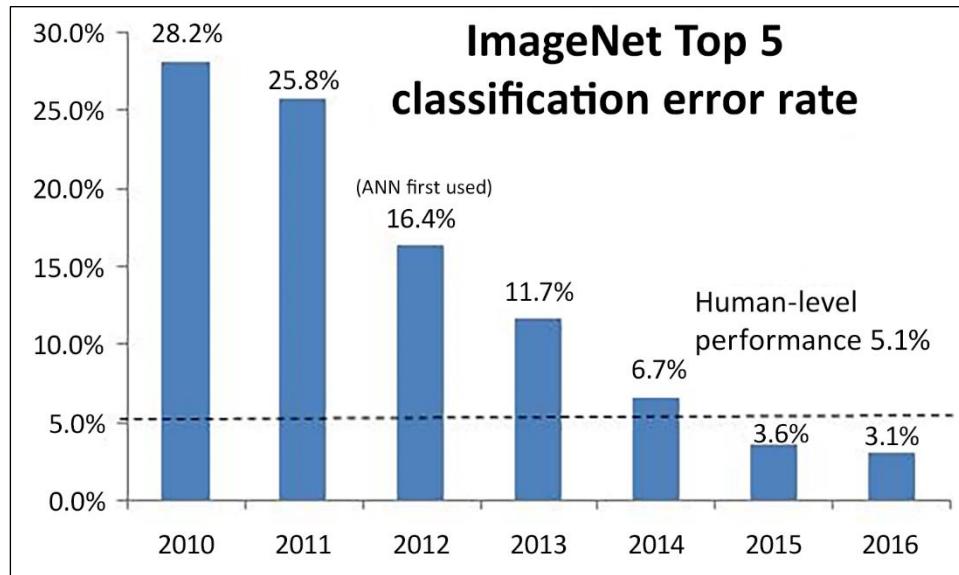
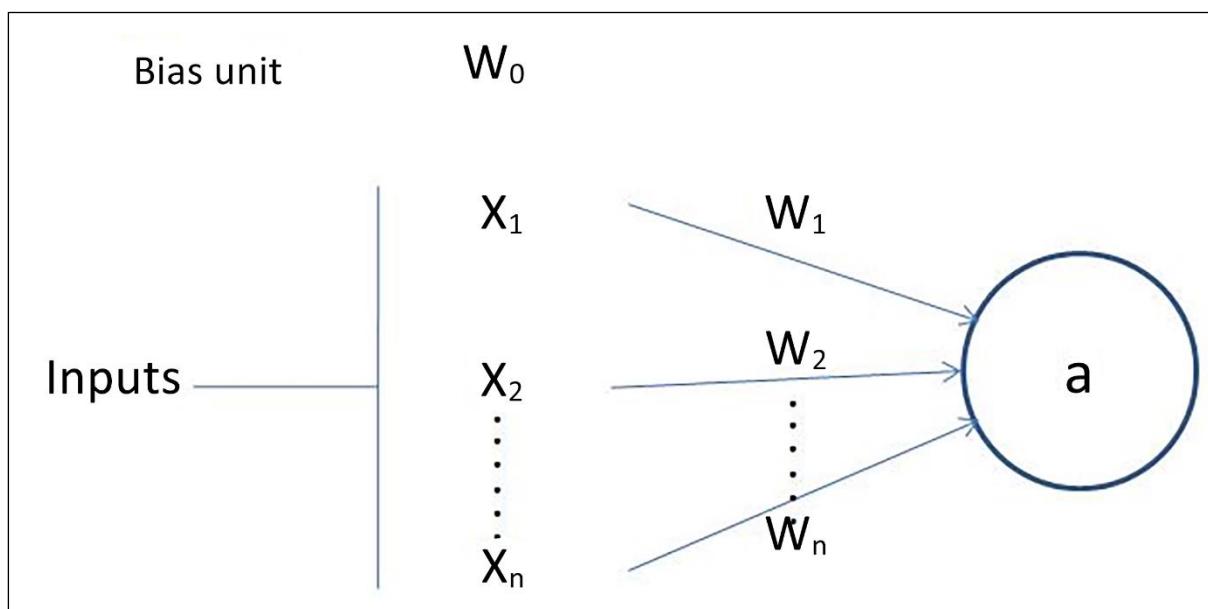
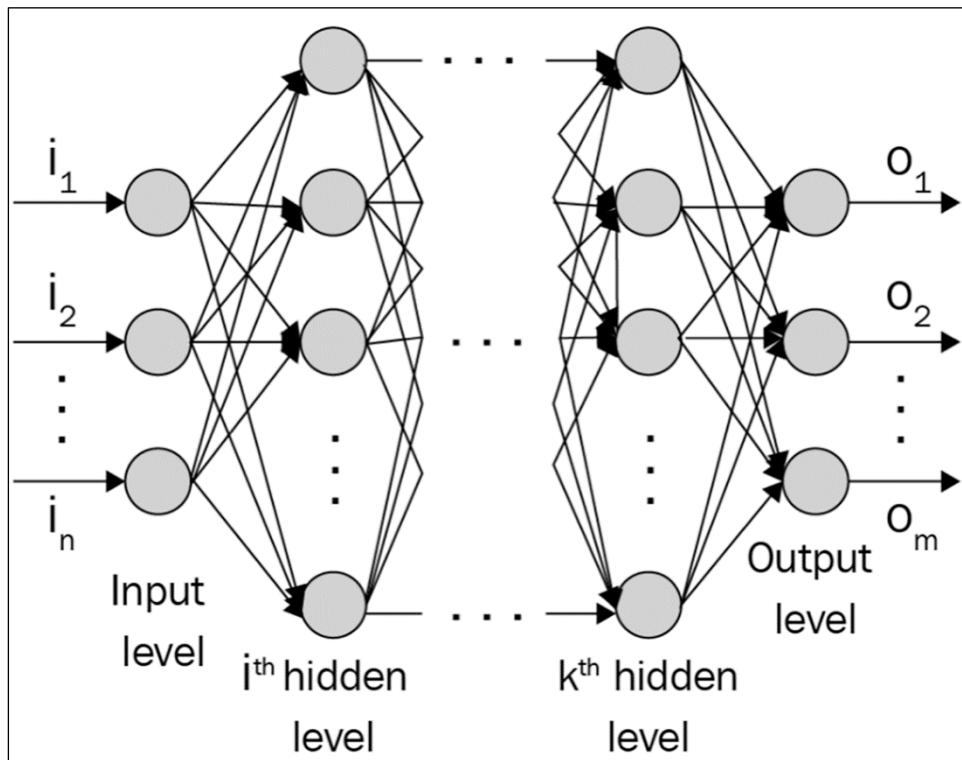
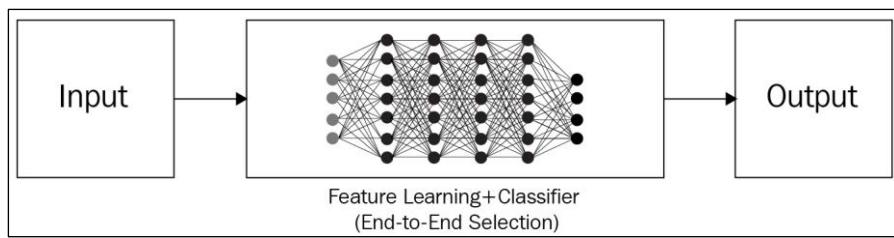
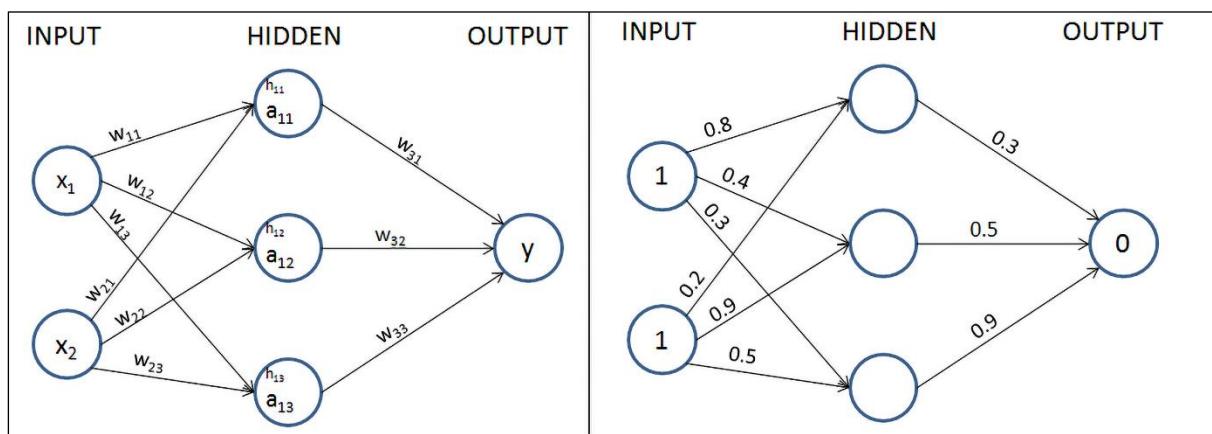
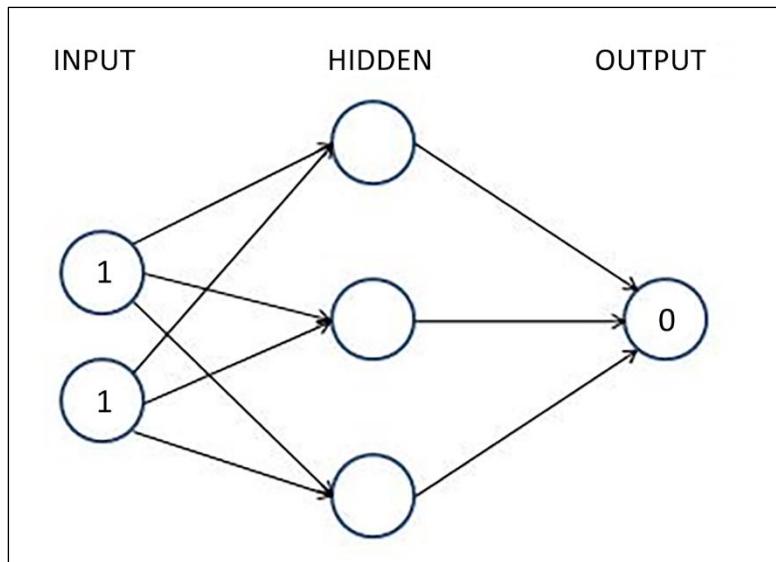
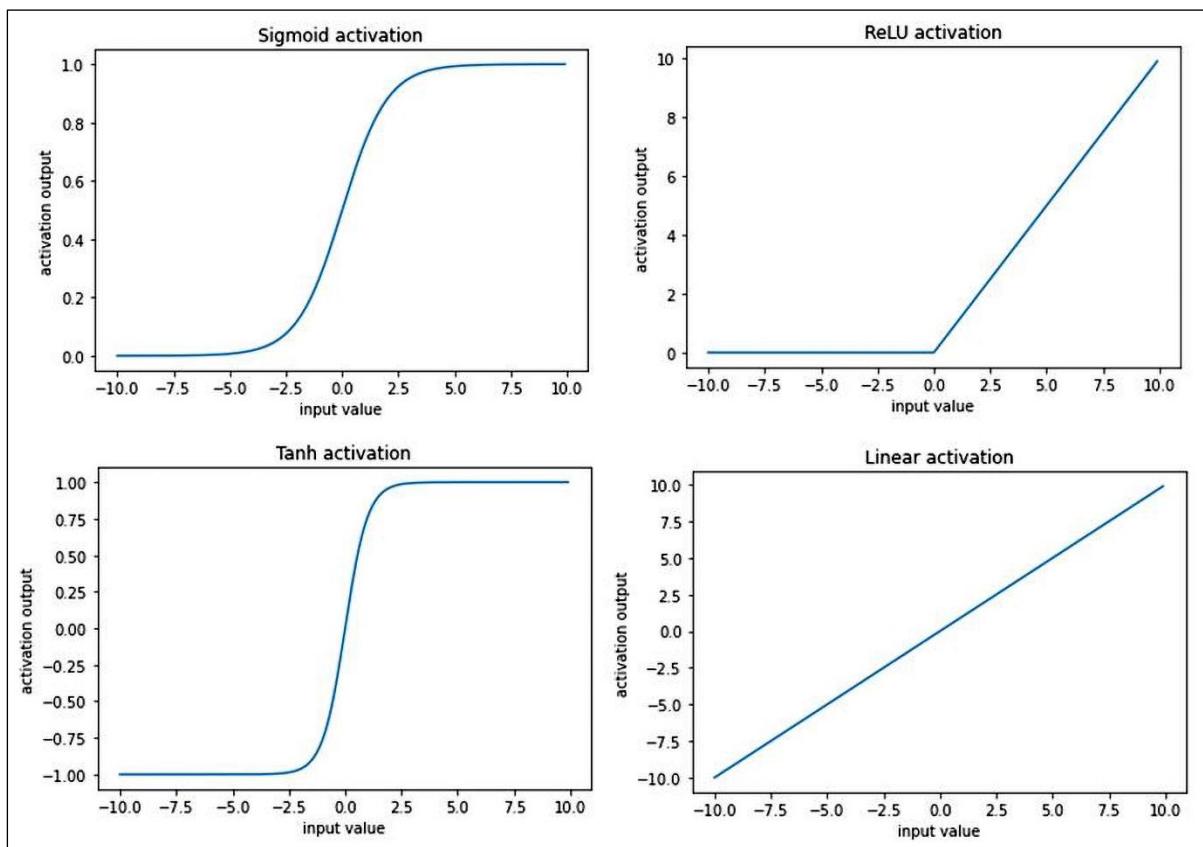
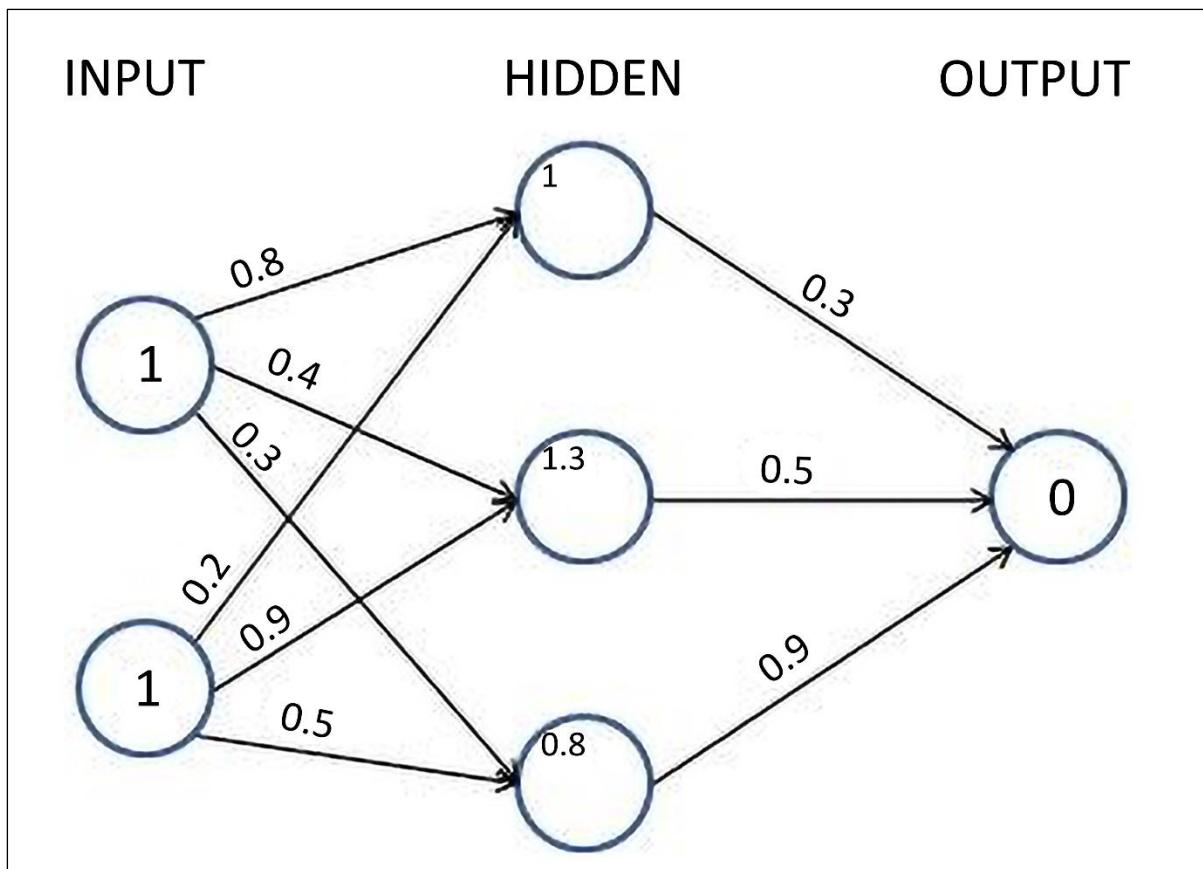


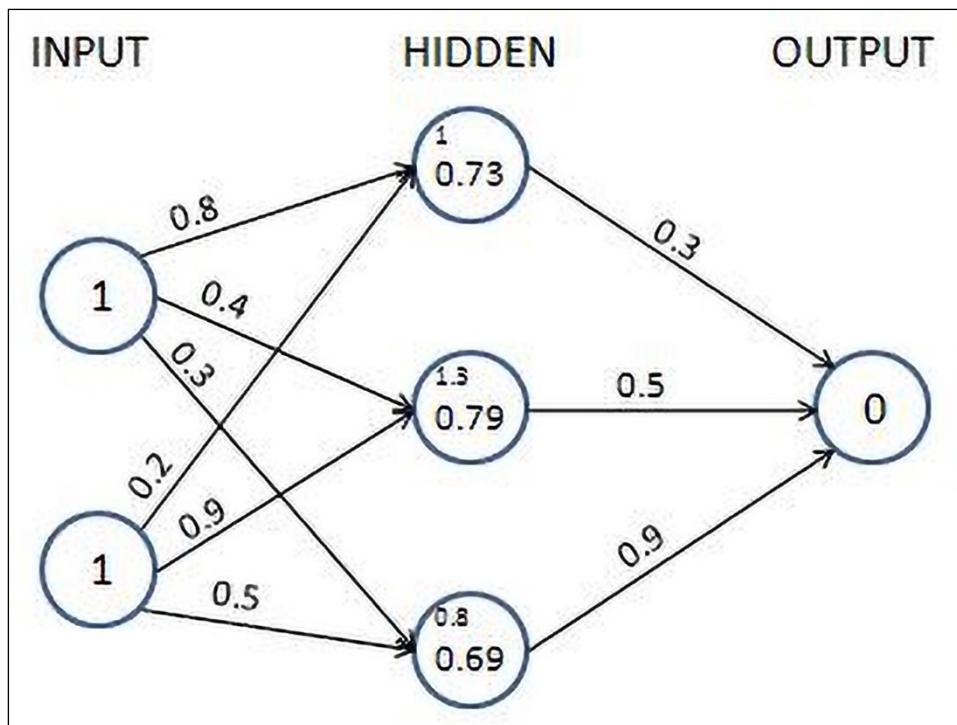
Chapter 1:











Target (correct class)	Prediction probabilities					Cross-entropy loss	
	Dog	Cat	Cow	Hen	Rat		
Dog	0.88	0.02	0.04	0.04	0.02	$-\log(0.88)$	= 0.128
Cat	0.26	0.21	0.17	0.18	0.18	$-\log(0.21)$	= 1.56
Cow	0.01	0.01	0.96	0.01	0.01	$-\log(0.96)$	= 0.04
Hen	0.14	0.09	0.01	0.57	0.19	$-\log(0.57)$	= 0.56
Rat	0.21	0.02	0.05	0.17	0.55	$-\log(0.55)$	= 0.597

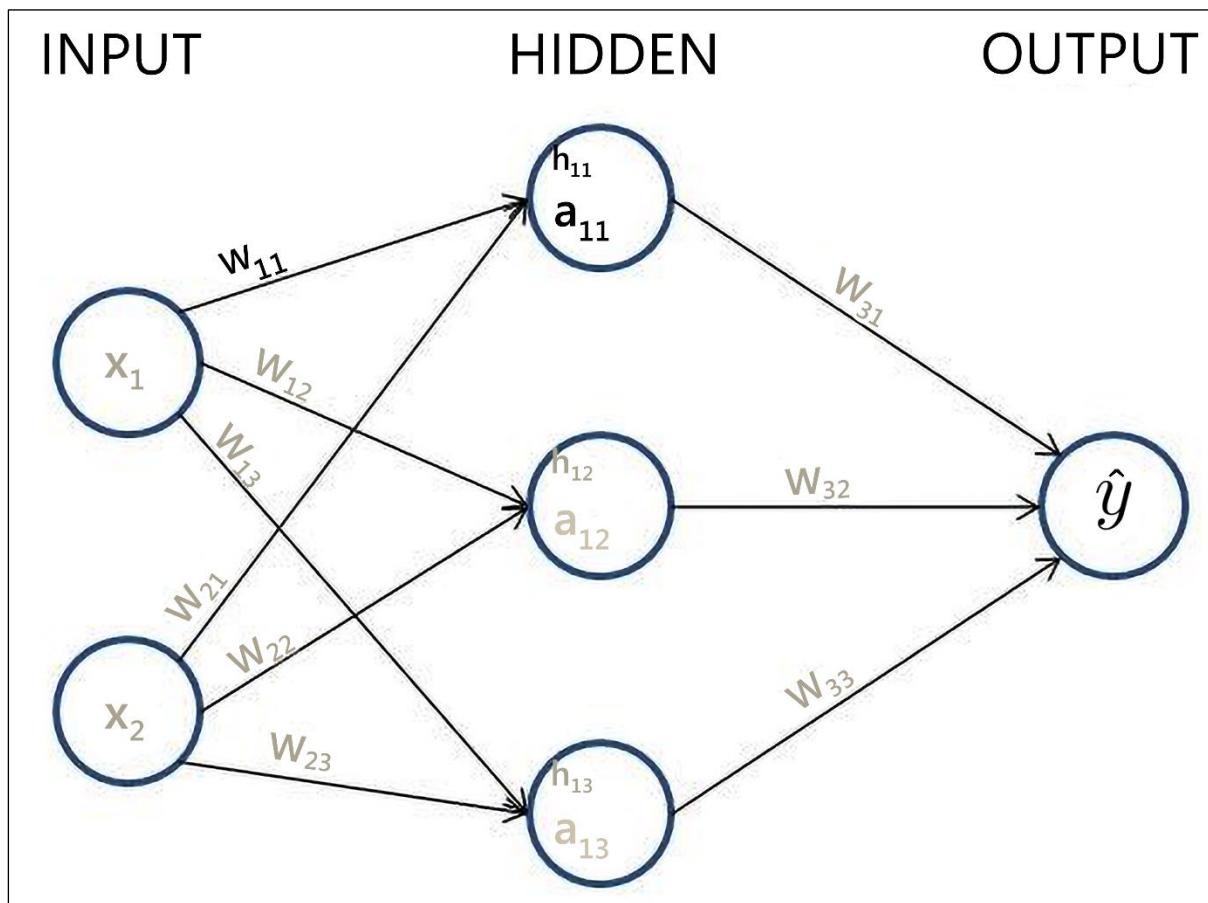
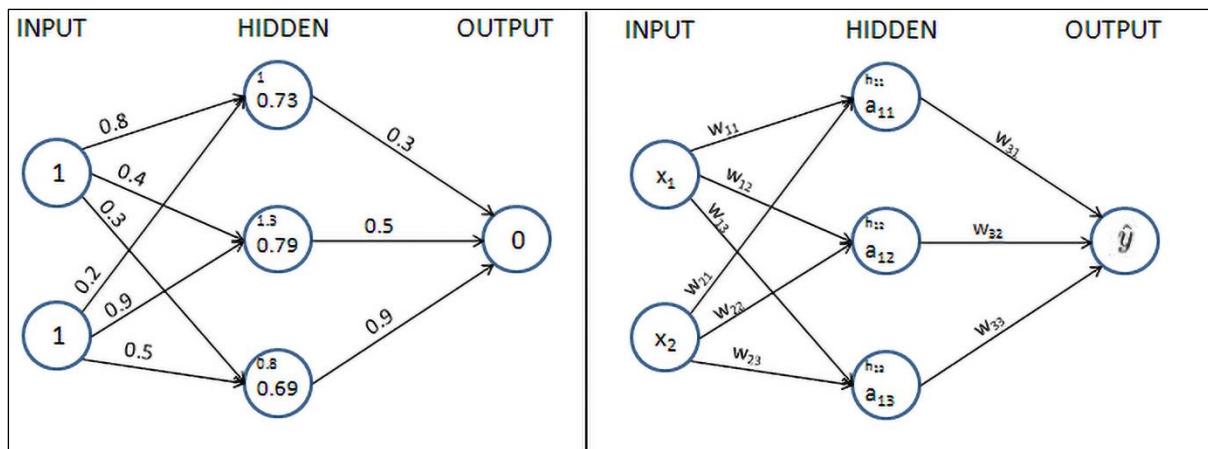
Modern-Computer-Vision-with-PyTorch-2E / Chapter01 / Feed_forward_propagation.ipynb

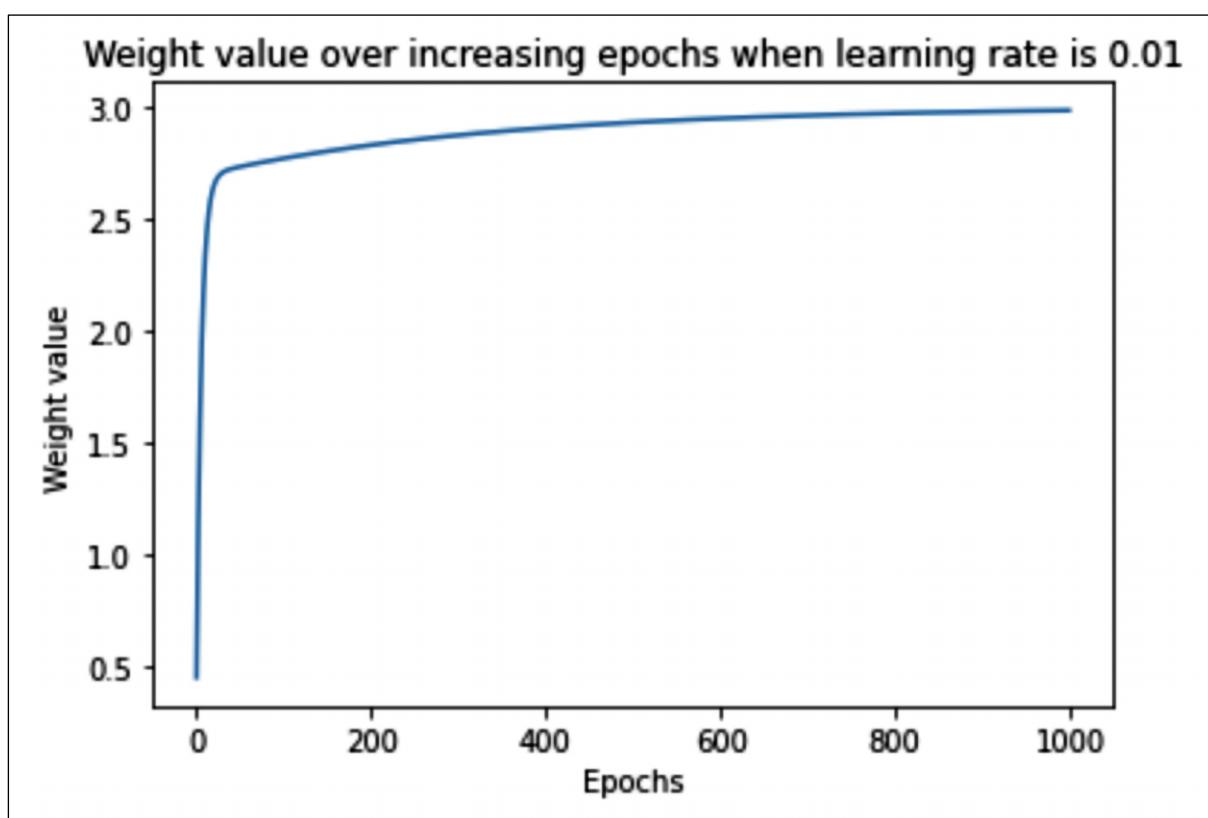
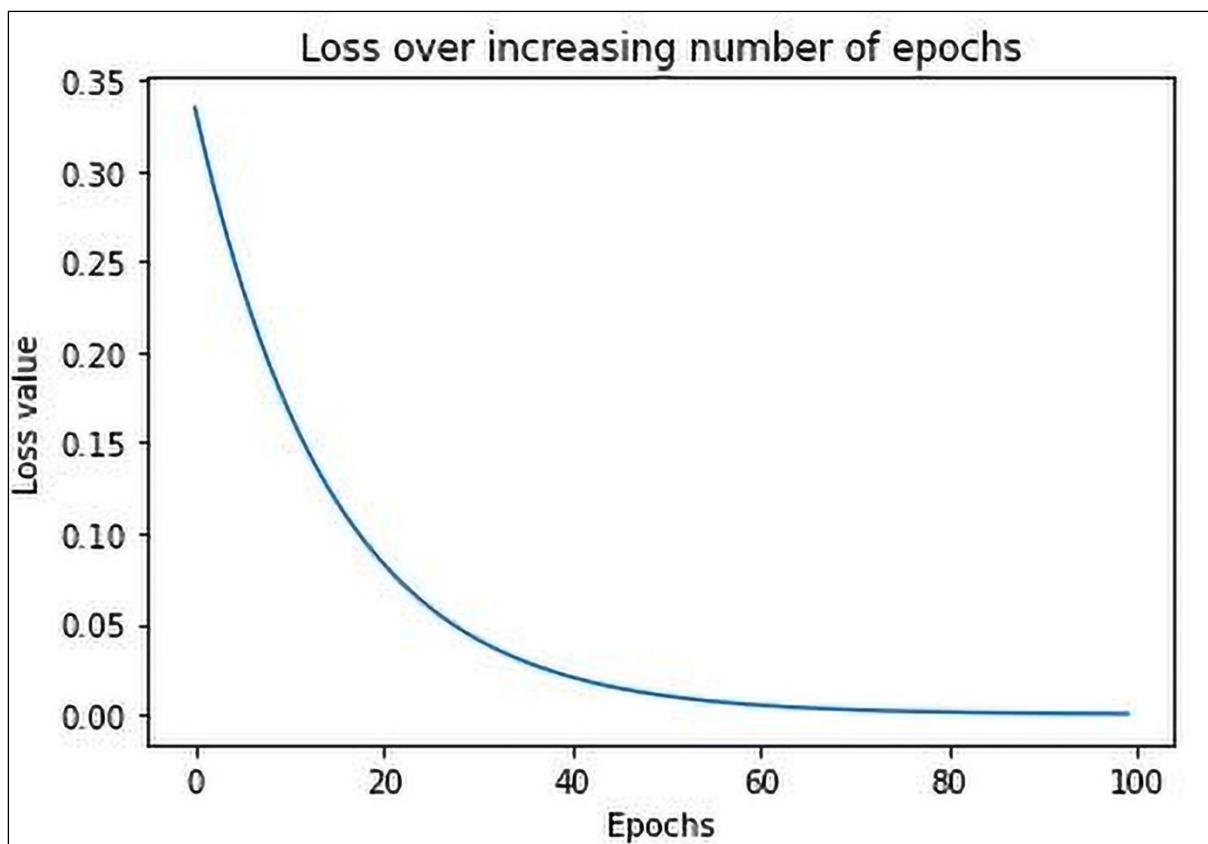
Preview Code Blame 74 lines (74 loc) · 2.02 KB

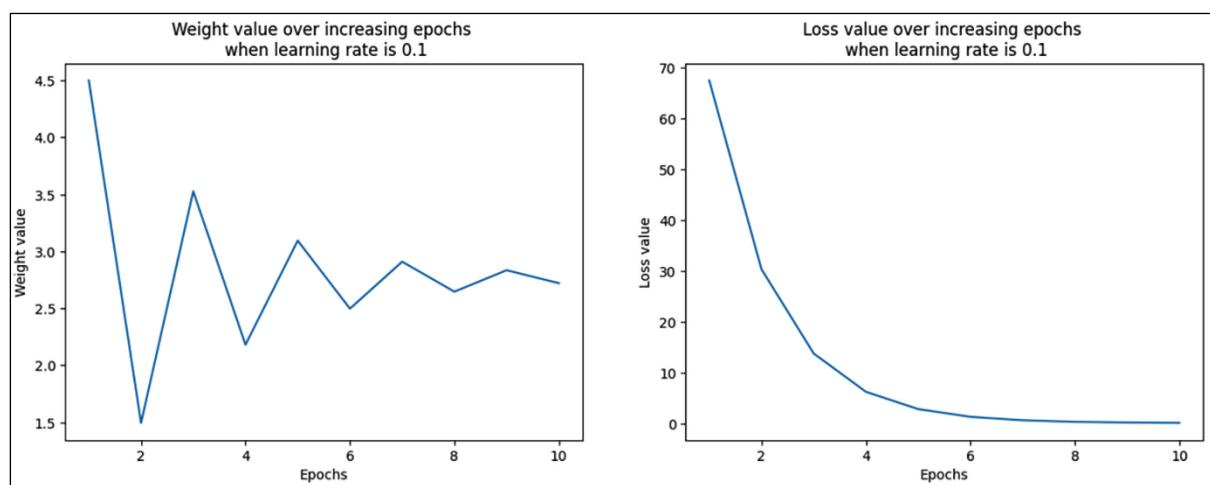
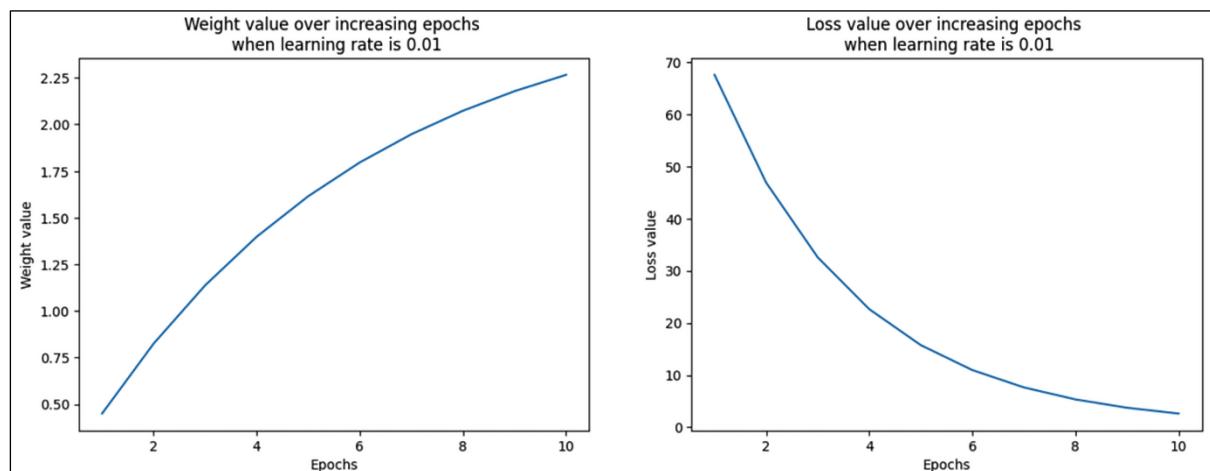
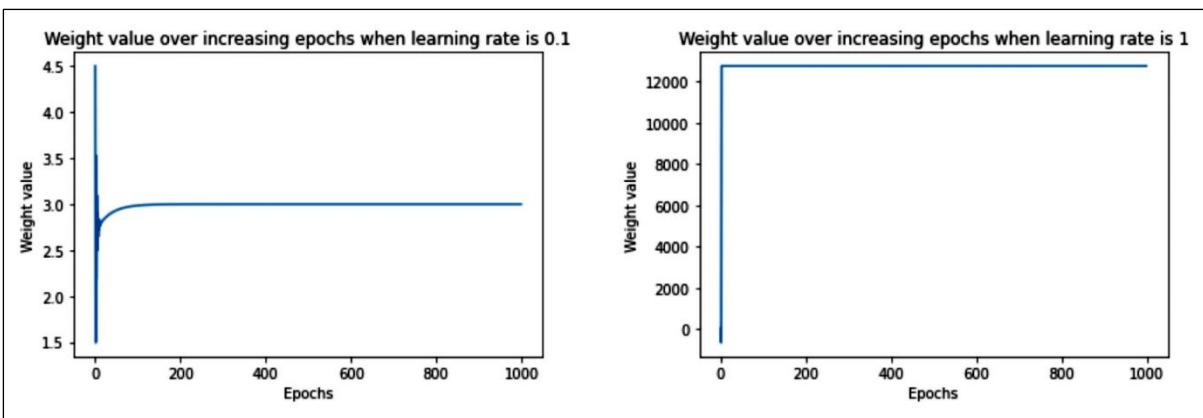
Open in Colab

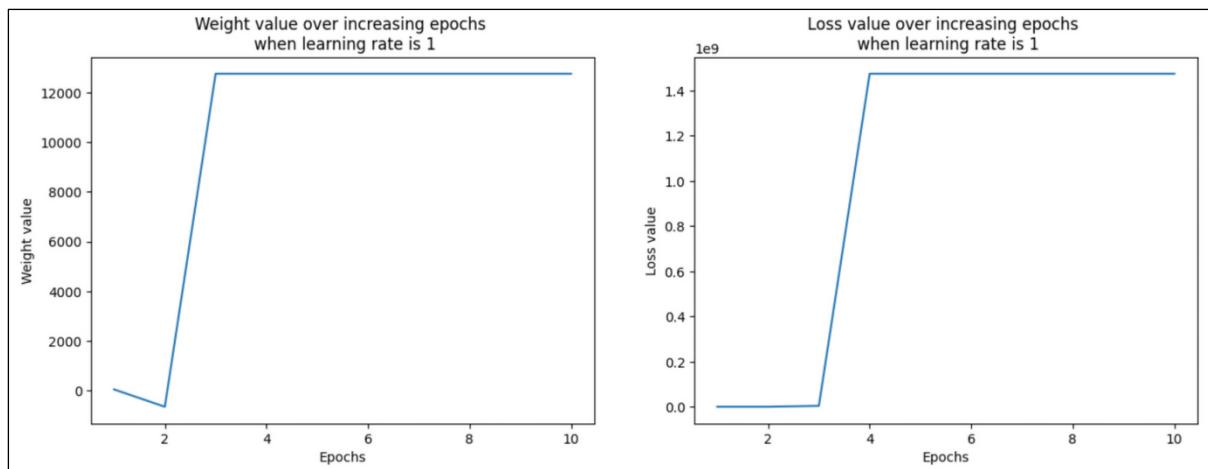
Forward Propagation

```
In [ ]: import numpy as np
def feed_forward(inputs, outputs, weights):
    pre_hidden = np.dot(inputs,weights[0])+ weights[1]
    hidden = 1/(1+np.exp(-pre_hidden))
    pred_out = np.dot(hidden, weights[2]) + weights[3]
    mean_squared_error = np.mean(np.square(pred_out - outputs))
    return mean_squared_error
```



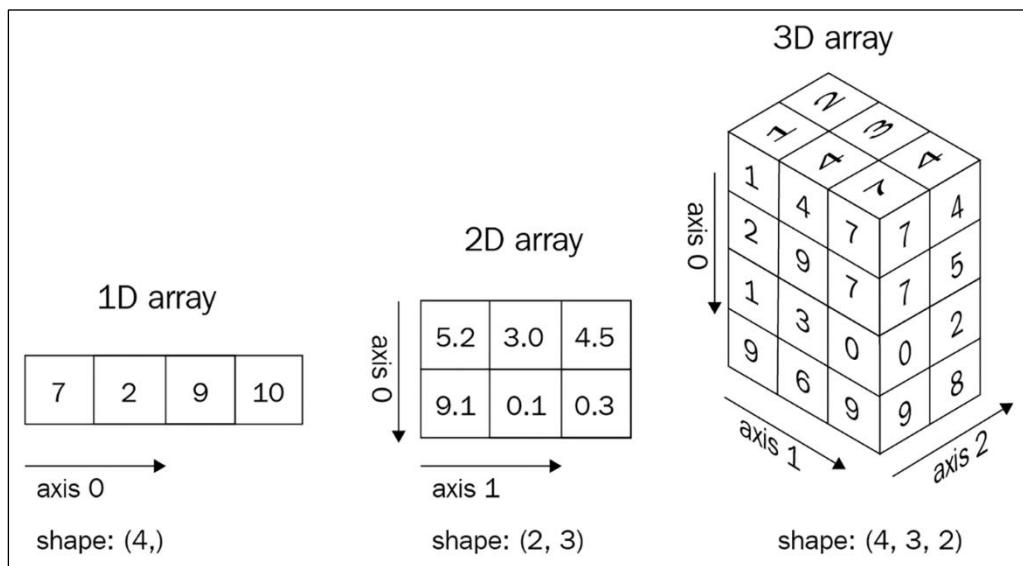






Chapter 2:

PyTorch Build	Stable (2.3.0)	Preview (Nightly)
Your OS	Linux	Mac Windows
Package	Conda	Pip LibTorch Source
Language	Python	C++ / Java
Compute Platform	CUDA 11.8	CUDA 12.1 ROCm 6.0 CPU
Run this Command:	<code>conda install pytorch torchvision torchaudio pytorch-cuda=12.1 -c pytorch -c nvidia</code>	



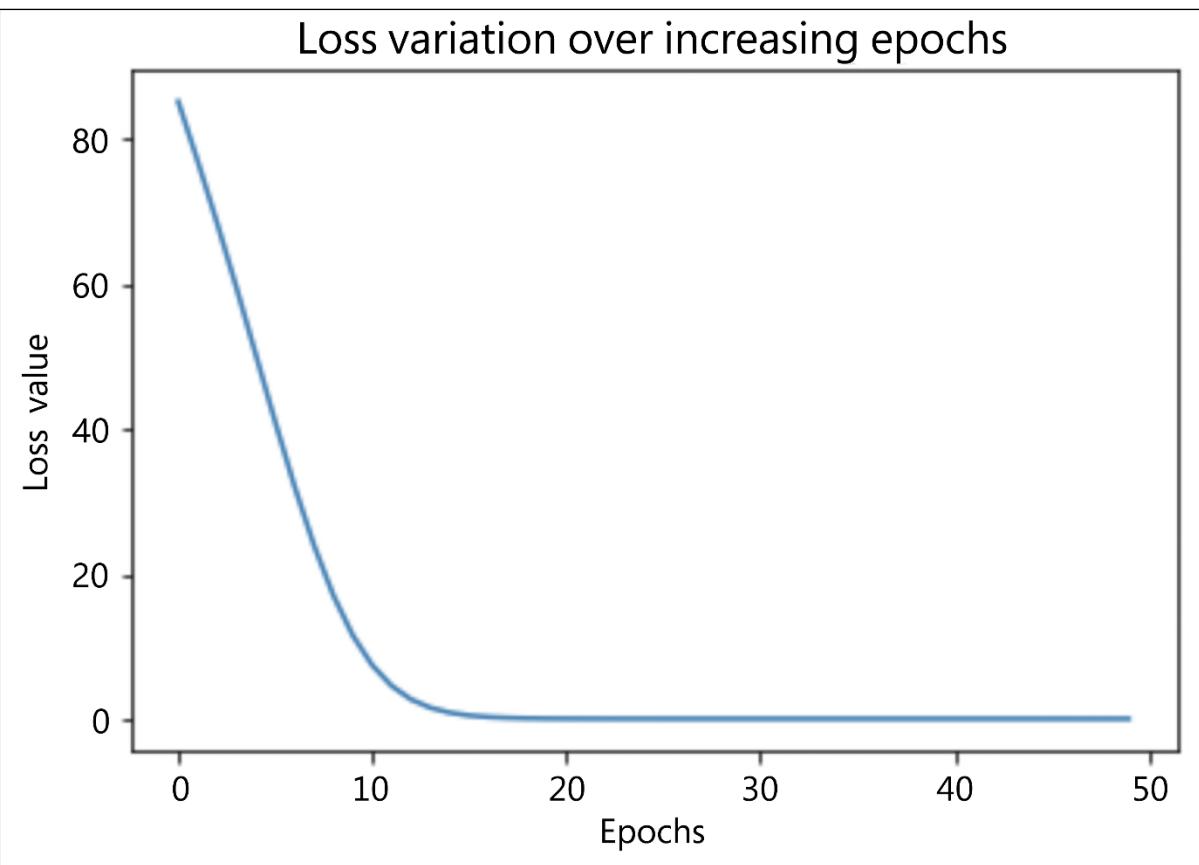
Parameter containing:

```
tensor([[ 0.5670,  0.2775],
       [-0.5525, -0.0506],
       [-0.1226, -0.0549],
       [-0.3667,  0.5775],
       [-0.2847, -0.7009],
       [-0.0449,  0.3303],
       [ 0.2479, -0.1501],
       [-0.4169, -0.0649]], requires_grad=True)
```

```

Parameter containing:
tensor([[ 0.5670,  0.2775],
       [-0.5525, -0.0506],
       [-0.1226, -0.0549],
       [-0.3667,  0.5775],
       [-0.2847, -0.7009],
       [-0.0449,  0.3303],
       [ 0.2479, -0.1501],
       [-0.4169, -0.0649]], requires_grad=True)
Parameter containing:
tensor([-0.7037,  0.4445, -0.4399,  0.6718,  0.2934, -0.6325,  0.2646, -0.5508],
       requires_grad=True)
Parameter containing:
tensor([[ 0.1219, -0.2936,  0.0820,  0.1212, -0.0885, -0.0113,  0.2657,  0.2921]],
       requires_grad=True)
Parameter containing:
tensor([ 0.0119], requires_grad=True)

```



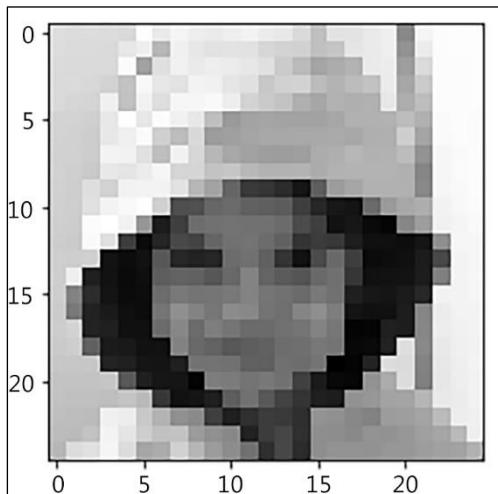
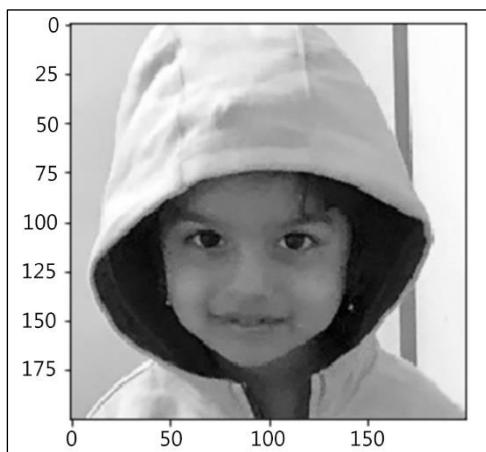
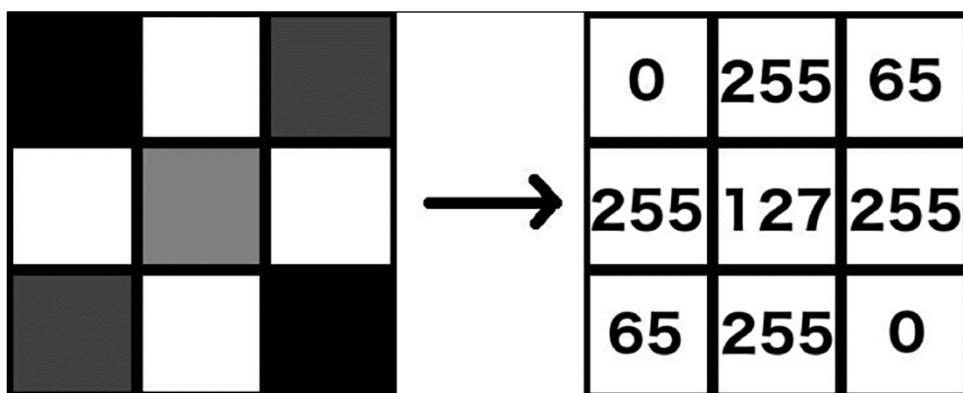
```

=====
Layer (type:depth-idx)          Output Shape      Param #
=====
|-Linear: 1-1                  [-1, 8]           24
|-ReLU: 1-2                     [-1, 8]           --
|-Linear: 1-3                  [-1, 1]            9
=====
Total params: 33
Trainable params: 33
Non-trainable params: 0
Total mult-adds (M): 0.00
=====
Input size (MB): 0.00
Forward/backward pass size (MB): 0.00
Params size (MB): 0.00
Estimated Total Size (MB): 0.00
=====
```

```
1  model.state_dict()

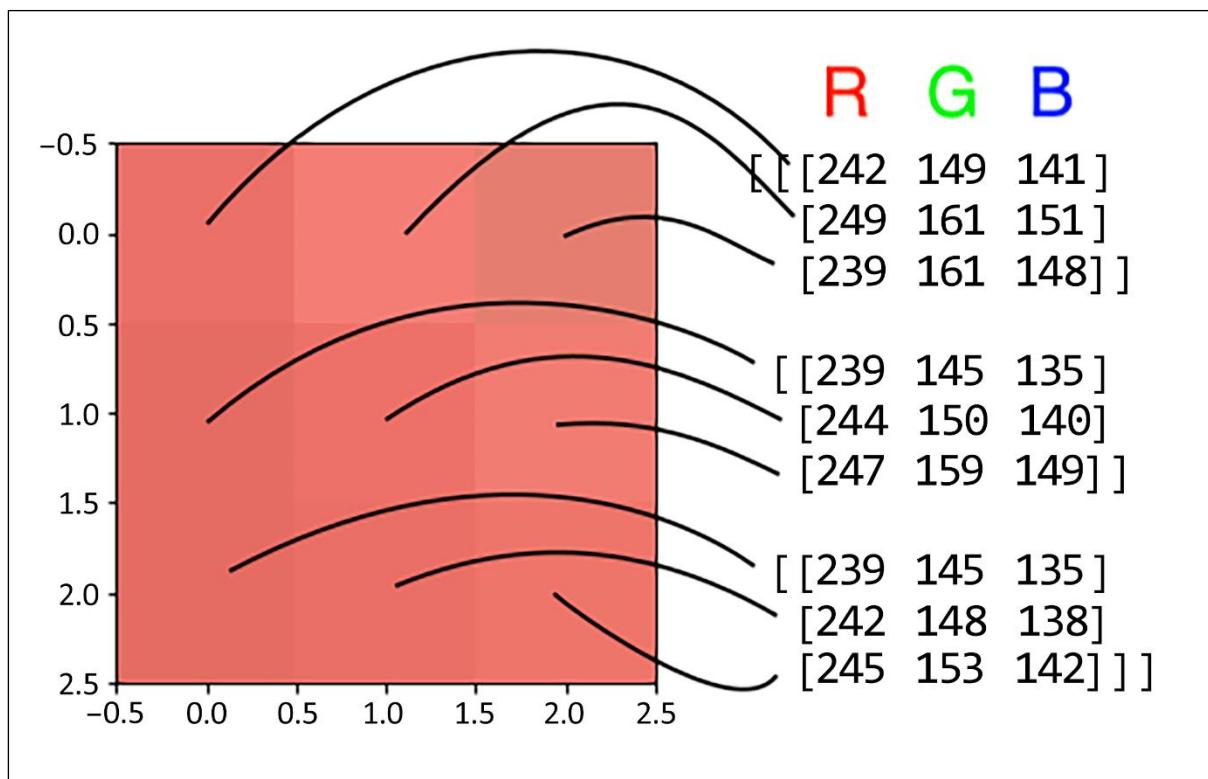
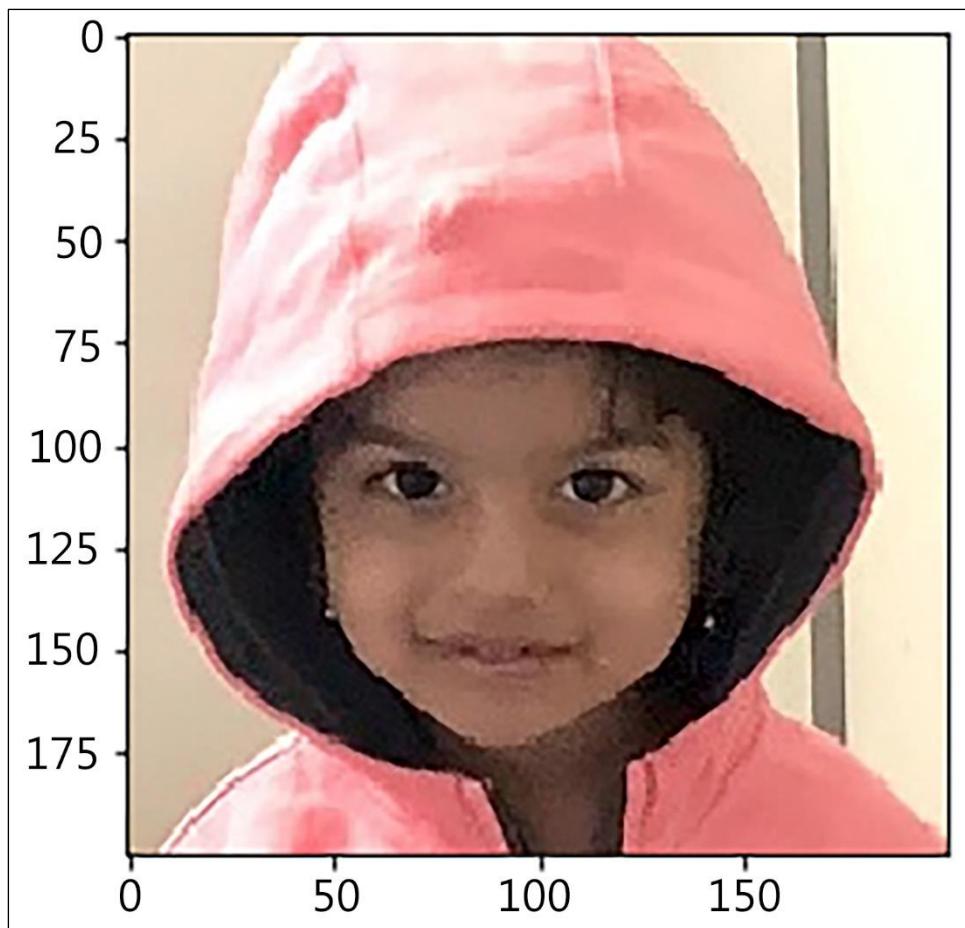
OrderedDict([('0.weight', tensor([[ 0.5090,  0.6708],
                                [-0.5887, -0.2970],
                                [ 0.3078, -0.4445],
                                [-0.3859,  0.0028],
                                [-0.1816,  0.9181],
                                [ 0.1532,  0.6011],
                                [ 0.2814, -0.4834],
                                [-0.6280, -0.5868]])),
('0.bias',
 tensor([ 0.7432, -0.5181, -0.1400,  0.3236, -0.1791, -0.4466, -0.1104, -0.1615])),
('2.weight',
 tensor([[ 0.9044, -0.2407, -0.1512, -0.2253,  0.5417,  0.4821,  0.1548,  0.0964]])),
('2.bias', tensor([0.0956]))])
```

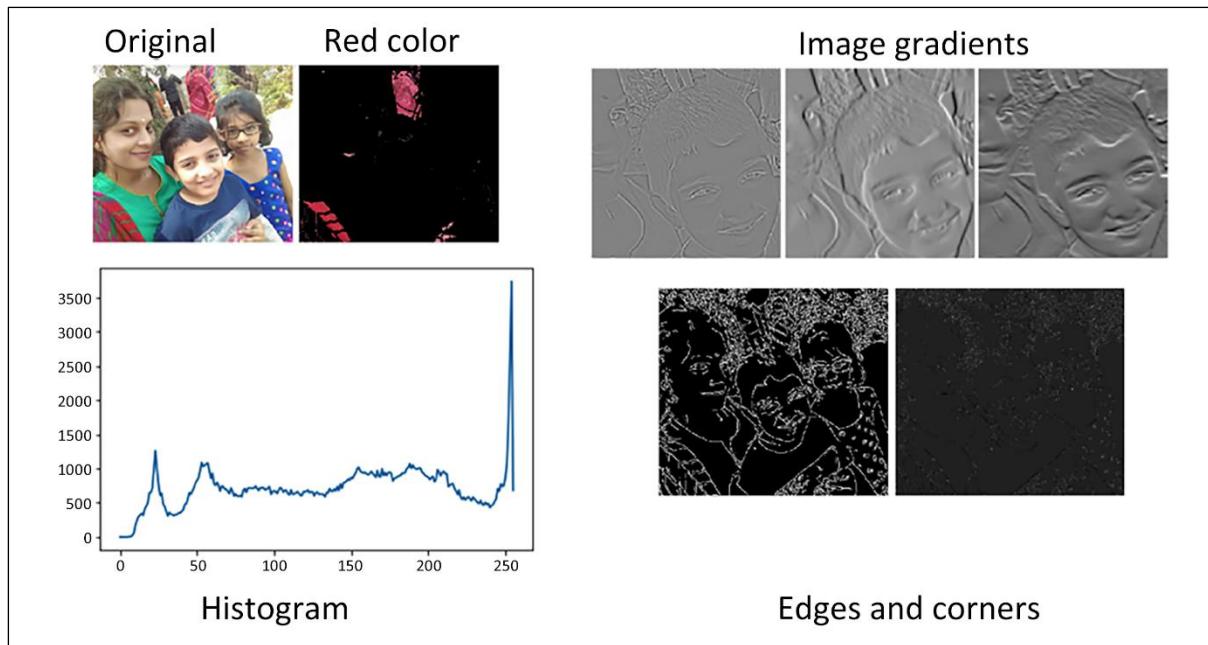
Chapter 3:



```
array([[222, 220, 221, 220, 218, 253, 234, 245, 238, 235, 239, 243, 236,
       232, 218, 193, 228, 228, 234, 239, 139, 245, 252, 253, 253],
      [221, 219, 219, 218, 232, 239, 186, 240, 231, 226, 227, 226, 215,
       212, 209, 193, 199, 229, 234, 239, 150, 236, 252, 253, 253],
      [219, 218, 218, 218, 251, 163, 224, 241, 234, 238, 236, 231, 224,
       204, 188, 166, 173, 180, 234, 236, 159, 219, 252, 252, 253],
      [218, 219, 216, 211, 196, 248, 231, 228, 243, 241, 229, 224, 201,
       209, 210, 189, 181, 189, 196, 235, 168, 204, 252, 252, 253],
```

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y
1	222	220	221	220	218	253	234	245	238	235	239	243	236	232	218	193	228	228	234	239	139	245	252	253	253
2	221	219	219	218	232	239	186	240	231	226	227	226	215	212	209	193	199	229	234	239	150	236	252	253	253
3	219	218	218	218	251	163	224	241	234	238	236	231	224	204	188	166	173	180	234	236	159	219	252	252	253
4	218	219	216	211	196	248	231	228	243	241	229	224	201	209	210	189	181	189	196	235	168	204	252	252	253
5	218	214	213	240	195	242	223	246	246	249	238	211	203	196	177	168	179	176	179	231	175	191	252	252	253
6	212	212	208	232	254	232	252	241	232	192	155	164	166	165	164	163	168	178	178	181	190	178	250	252	251
7	211	209	205	232	240	251	208	191	217	158	161	166	169	169	170	170	171	169	176	177	206	166	250	252	251
8	209	209	205	243	242	225	193	241	215	184	169	163	159	158	160	173	176	184	178	179	189	150	246	250	252
9	210	207	203	232	229	236	246	214	213	196	199	185	179	179	181	172	179	180	180	181	177	136	246	251	252
10	209	206	222	212	242	243	244	226	184	165	104	61	57	48	27	97	158	167	178	178	178	139	246	249	252
11	208	206	225	243	249	254	209	82	85	105	109	100	98	95	65	43	28	24	109	156	169	175	242	248	251
12	208	205	252	255	242	153	33	66	111	116	117	116	115	109	78	66	22	27	14	9	137	159	241	245	249
13	205	204	250	225	63	15	42	77	71	104	115	118	110	101	56	64	60	34	20	20	17	25	145	246	246
14	208	206	209	23	16	22	90	45	39	43	110	115	99	56	23	78	107	65	15	17	20	32	76	244	246
15	208	239	37	22	14	19	97	102	100	90	108	133	104	94	88	108	114	57	21	22	23	33	130	243	246
16	205	133	48	24	15	17	110	124	118	119	124	134	119	116	109	123	116	36	27	25	31	44	242	243	246
17	204	124	38	30	19	16	120	146	133	121	142	138	118	114	135	145	128	24	20	21	33	136	236	243	244
18	205	212	39	37	20	20	110	137	110	128	119	109	109	115	119	133	121	8	9	36	34	137	237	241	243
19	204	206	101	31	29	21	22	132	108	102	91	97	101	111	113	122	118	11	14	38	222	139	233	240	242
20	200	200	200	41	36	25	24	37	117	117	116	101	99	114	111	119	4	8	41	220	219	134	232	239	244
21	197	196	196	199	92	37	26	25	8	127	125	118	122	116	67	31	13	11	150	173	220	131	231	238	242
22	195	193	193	192	198	187	58	22	25	37	97	115	93	70	55	36	33	148	153	165	166	183	233	236	242
23	192	190	189	240	237	202	180	147	140	66	36	52	64	61	51	150	146	138	134	157	159	166	189	237	240
24	189	188	197	244	229	206	194	196	157	146	138	39	63	73	53	144	143	139	150	148	153	161	167	187	239
25	184	220	225	245	221	167	173	209	183	157	143	116	53	74	49	144	150	150	148	153	153	158	162	165	178

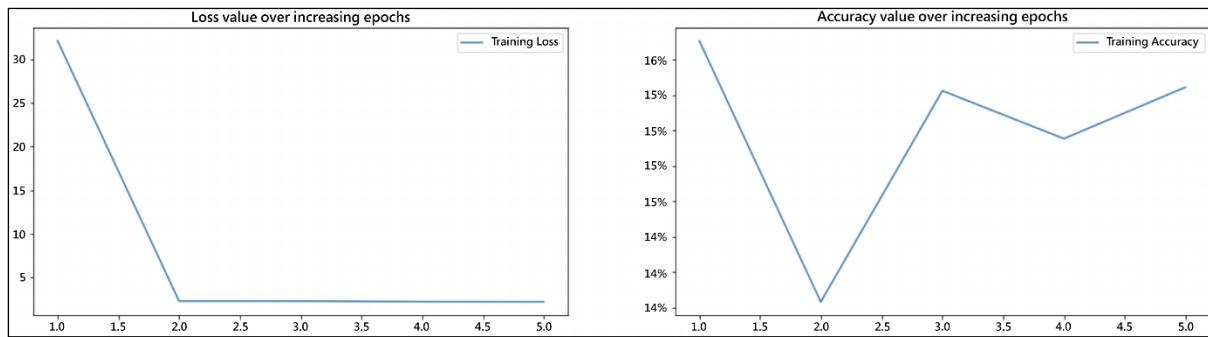


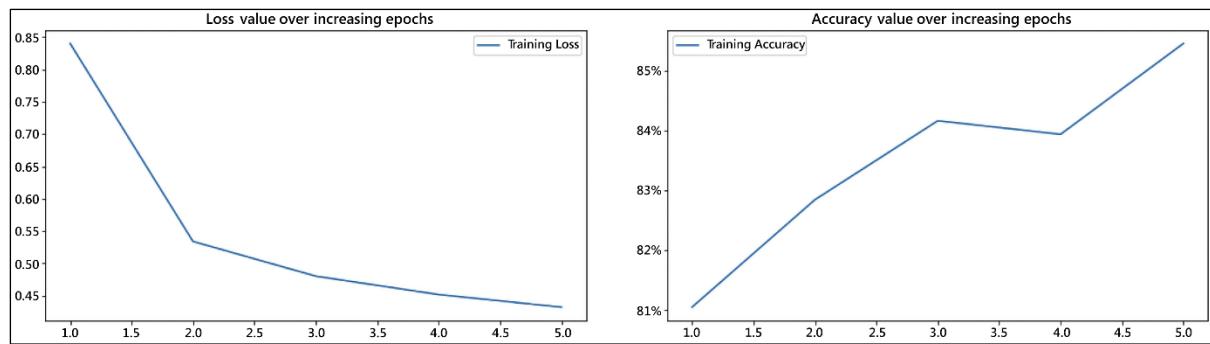


```

tr_images & tr_targets:
  X - torch.Size([60000, 28, 28])
  Y - torch.Size([60000])
  Y - Unique Values : tensor([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
TASK:
  10 class Classification
UNIQUE CLASSES:
  ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

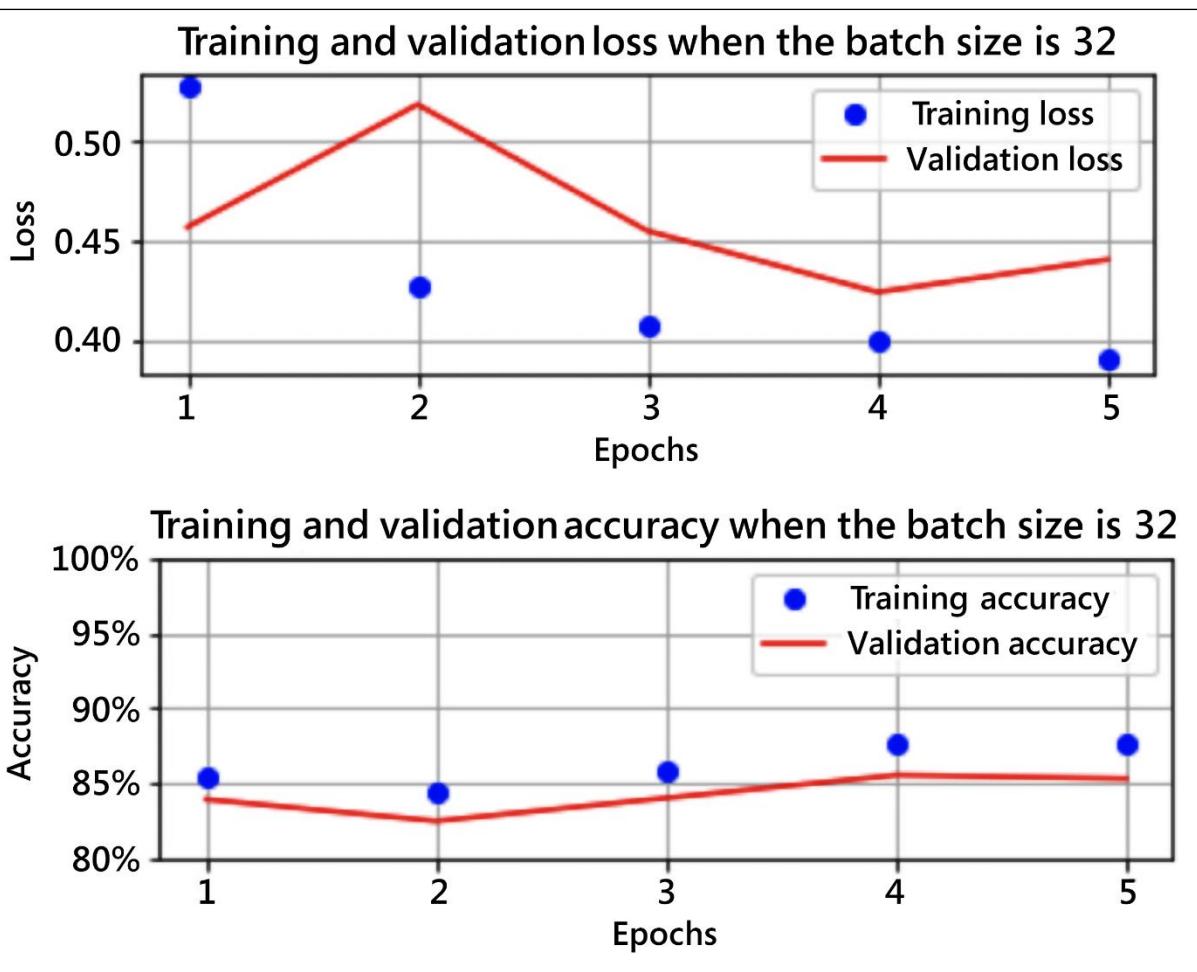
```



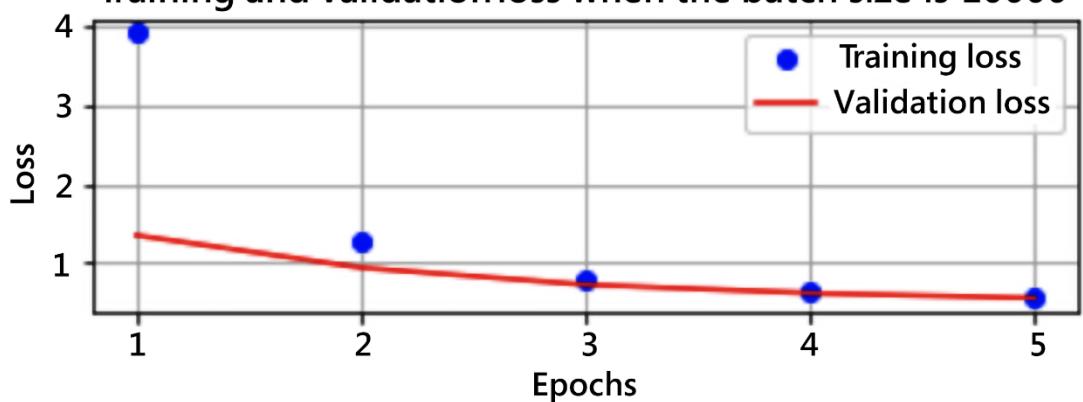


Input	Weight	Sigmoid
255	0.00001	0.501
255	0.0001	0.506
255	0.001	0.563
255	0.01	0.928
255	0.1	1.000
255	0.2	1.000
255	0.3	1.000
255	0.4	1.000
255	0.5	1.000
255	0.6	1.000
255	0.7	1.000
255	0.8	1.000
255	0.9	1.000
255	1	1.000

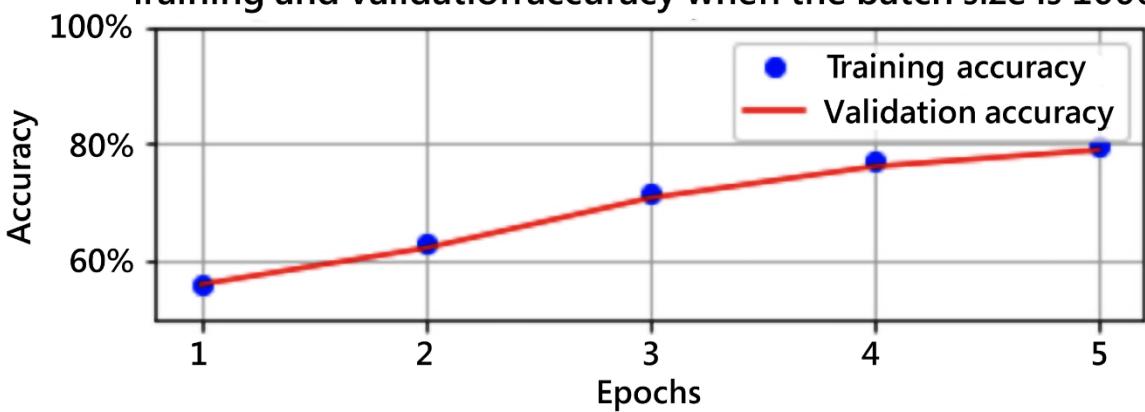
Input	Weight	Sigmoid
1	0.00001	0.500
1	0.0001	0.500
1	0.001	0.500
1	0.01	0.502
1	0.1	0.525
1	0.2	0.550
1	0.3	0.574
1	0.4	0.599
1	0.5	0.622
1	0.6	0.646
1	0.7	0.668
1	0.8	0.690
1	0.9	0.711
1	1	0.731



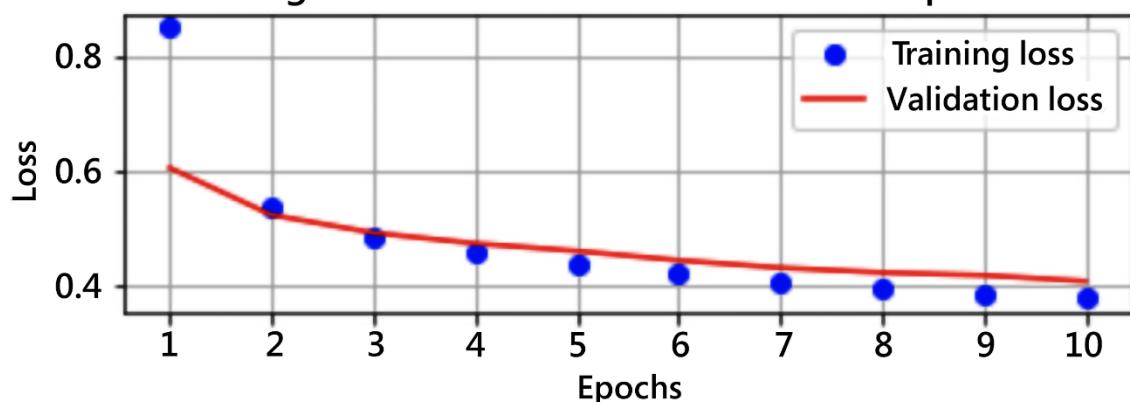
Training and validation loss when the batch size is 10000



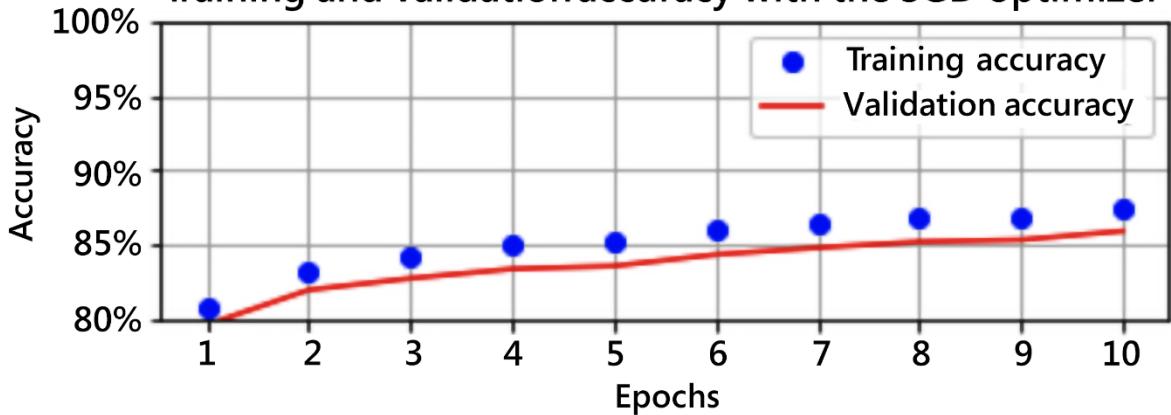
Training and validation accuracy when the batch size is 10000

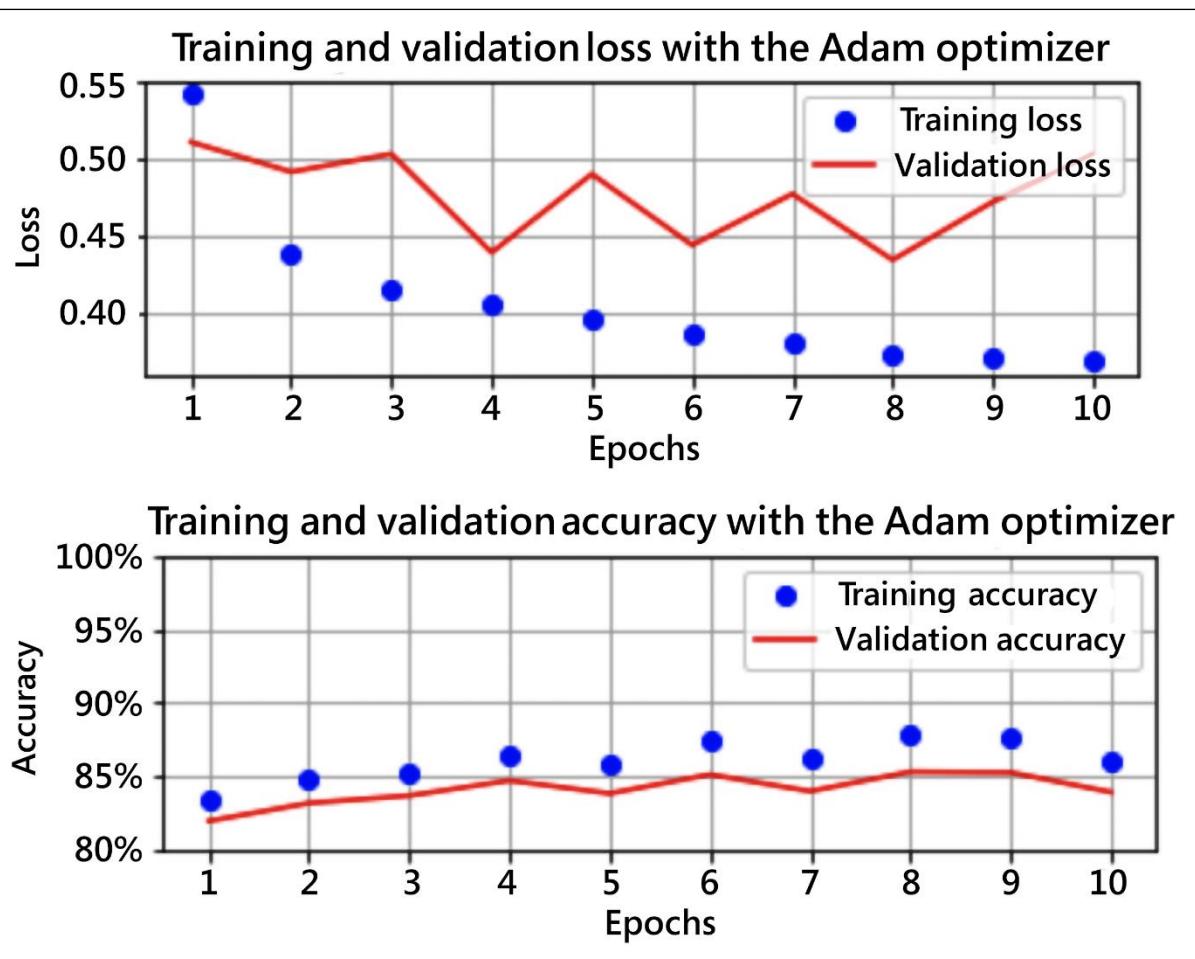


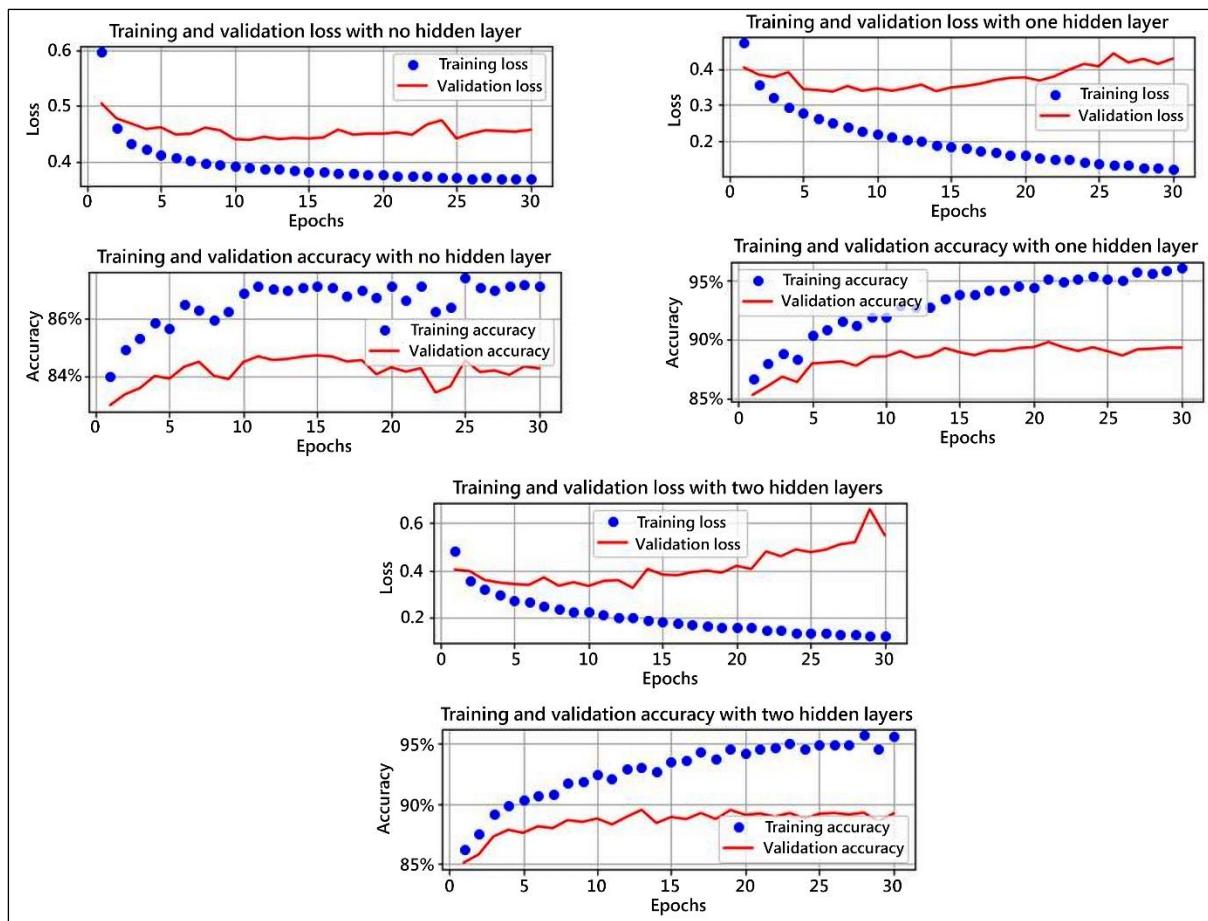
Training and validation loss with the SGD optimizer



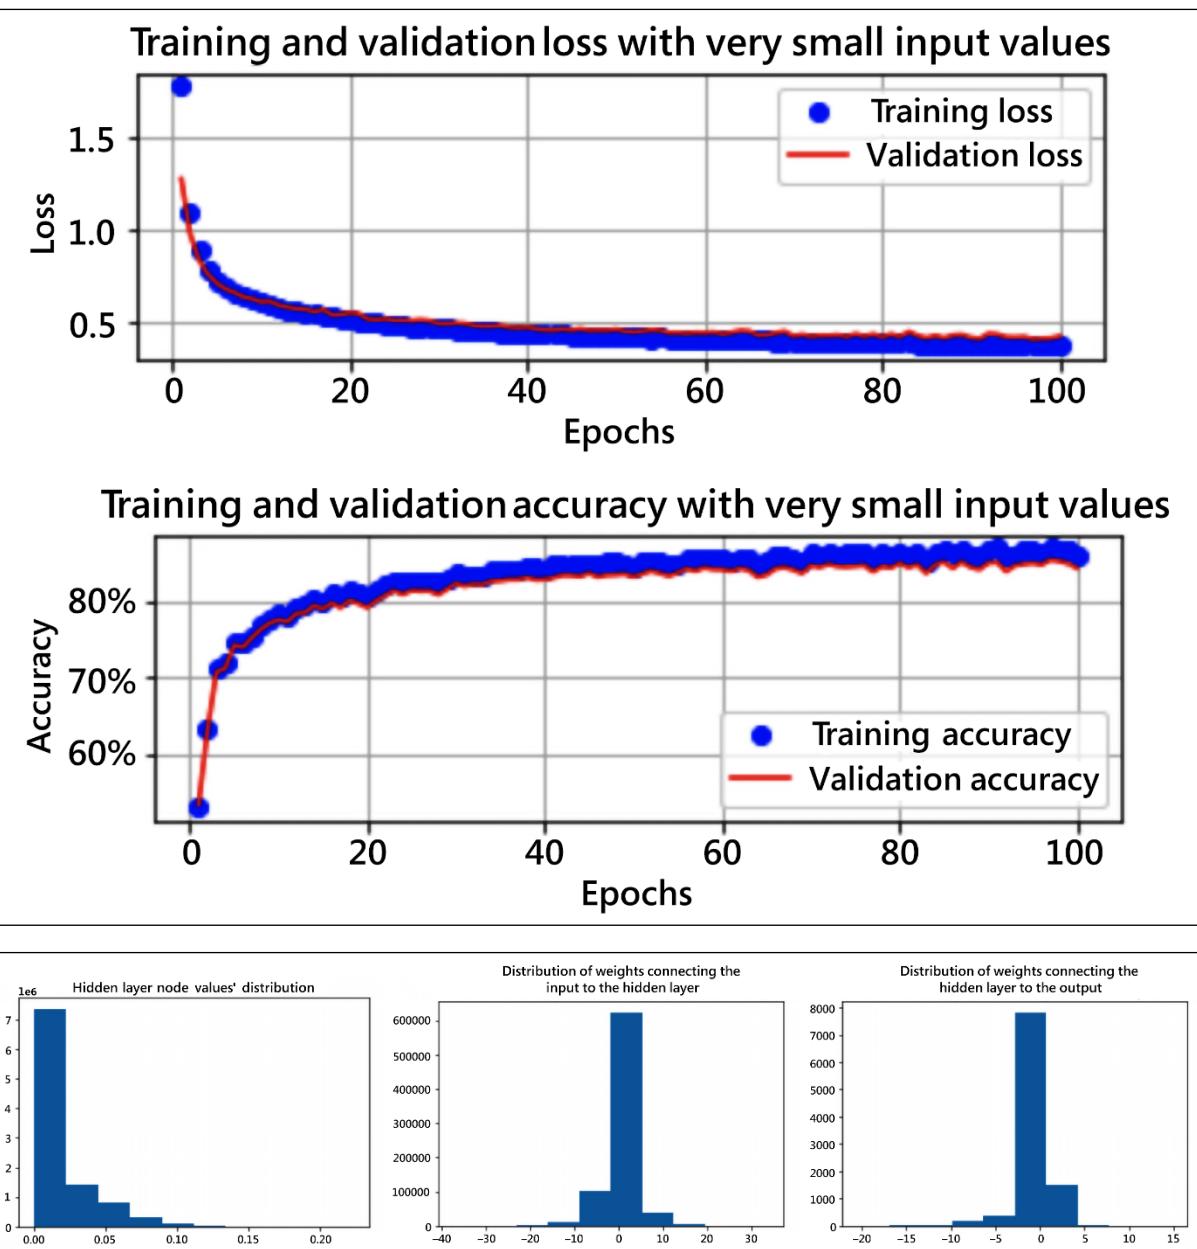
Training and validation accuracy with the SGD optimizer



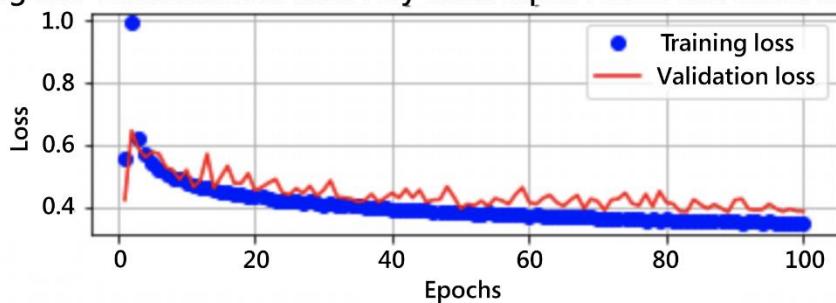




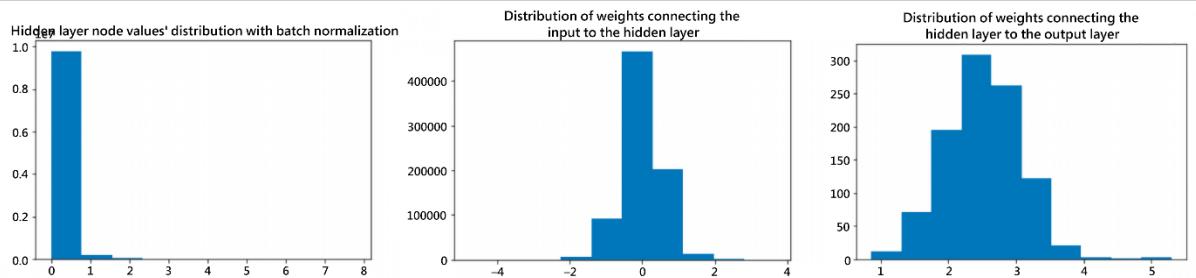
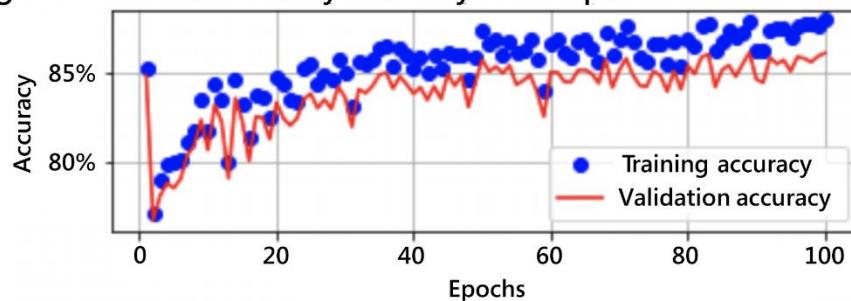
Input	Weight	Sigmoid
0.01	0.00001	0.500
0.01	0.0001	0.500
0.01	0.001	0.500
0.01	0.01	0.500
0.01	0.1	0.500
0.01	0.2	0.500
0.01	0.3	0.501
0.01	0.4	0.501
0.01	0.5	0.501
0.01	0.6	0.501
0.01	0.7	0.502
0.01	0.8	0.502
0.01	0.9	0.502
0.01	1	0.502



Training and validation loss with very small input values and batch normalization



Training and validation accuracy with very small input values and batch normalization



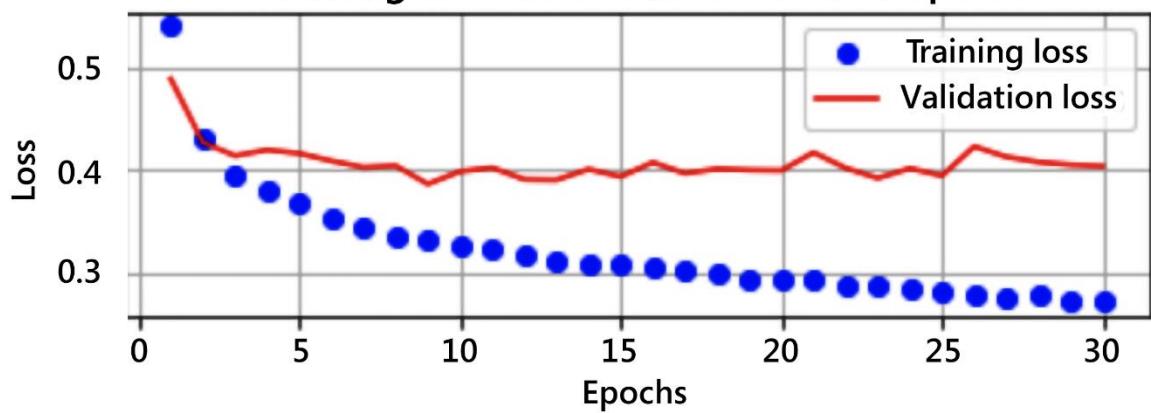
```
[20] 1 model.train()
      2 batch = next(iter(trn_dl))
      3 for i in range(5):
      4     output = model(batch[0])
      5     print(output.mean().item())

      □ -13.275323867797852
      -12.834677696228027
      -11.895054817199707
      -12.713885307312012
      -13.302783012390137

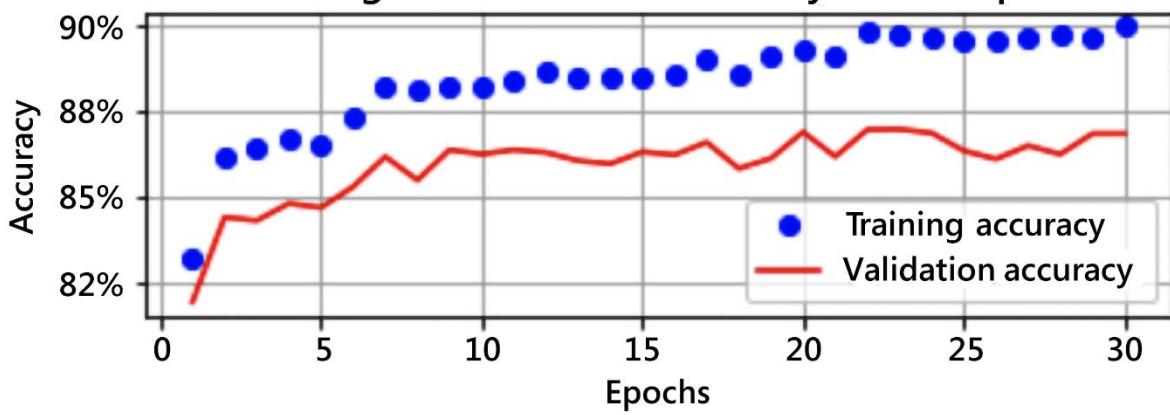
      ⏎ 1 model.eval()
      2 batch = next(iter(trn_dl))
      3 for i in range(5):
      4     output = model(batch[0])
      5     print(output.mean().item())

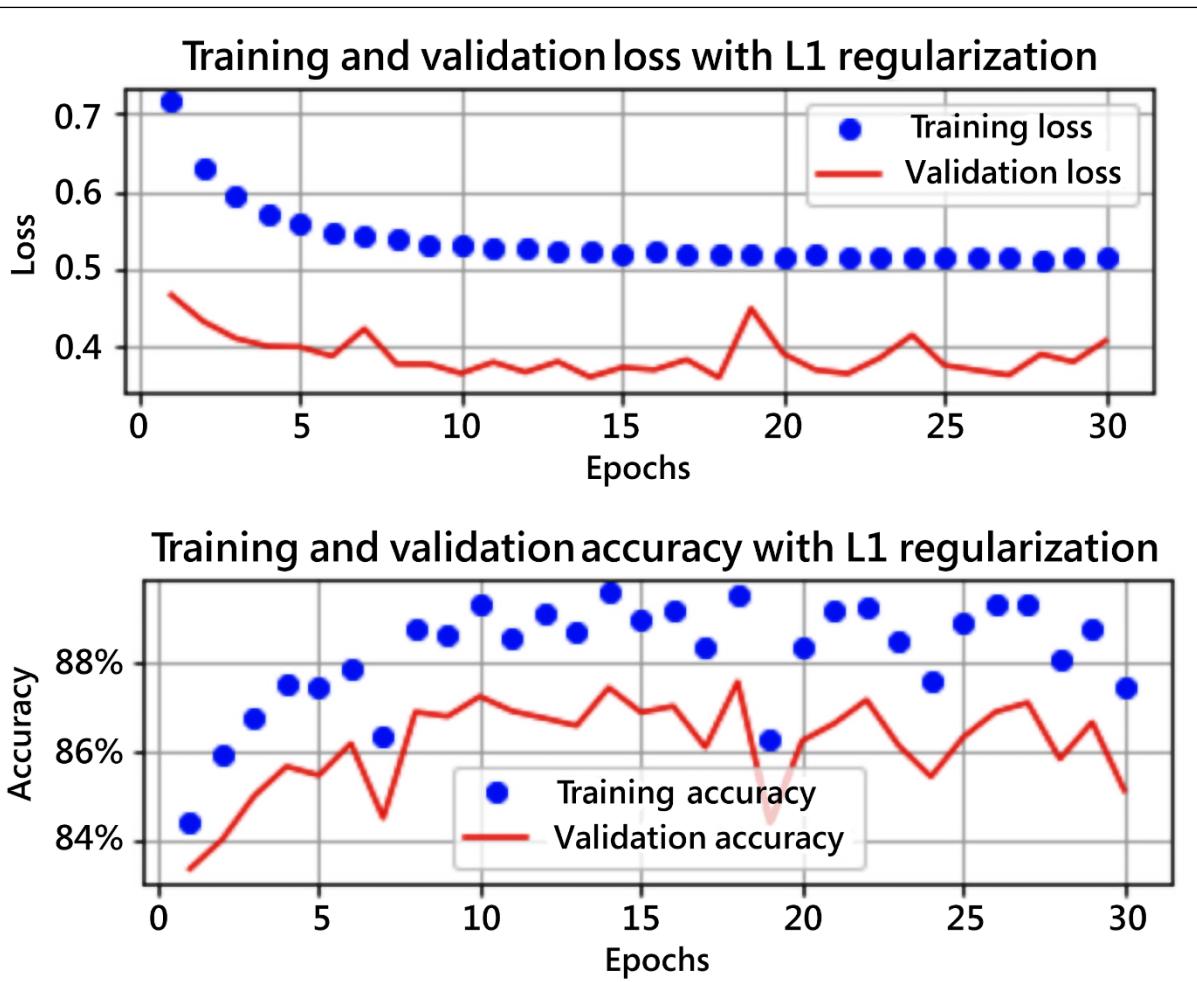
      □ -12.40117359161377
      -12.40117359161377
      -12.40117359161377
      -12.40117359161377
      -12.40117359161377
```

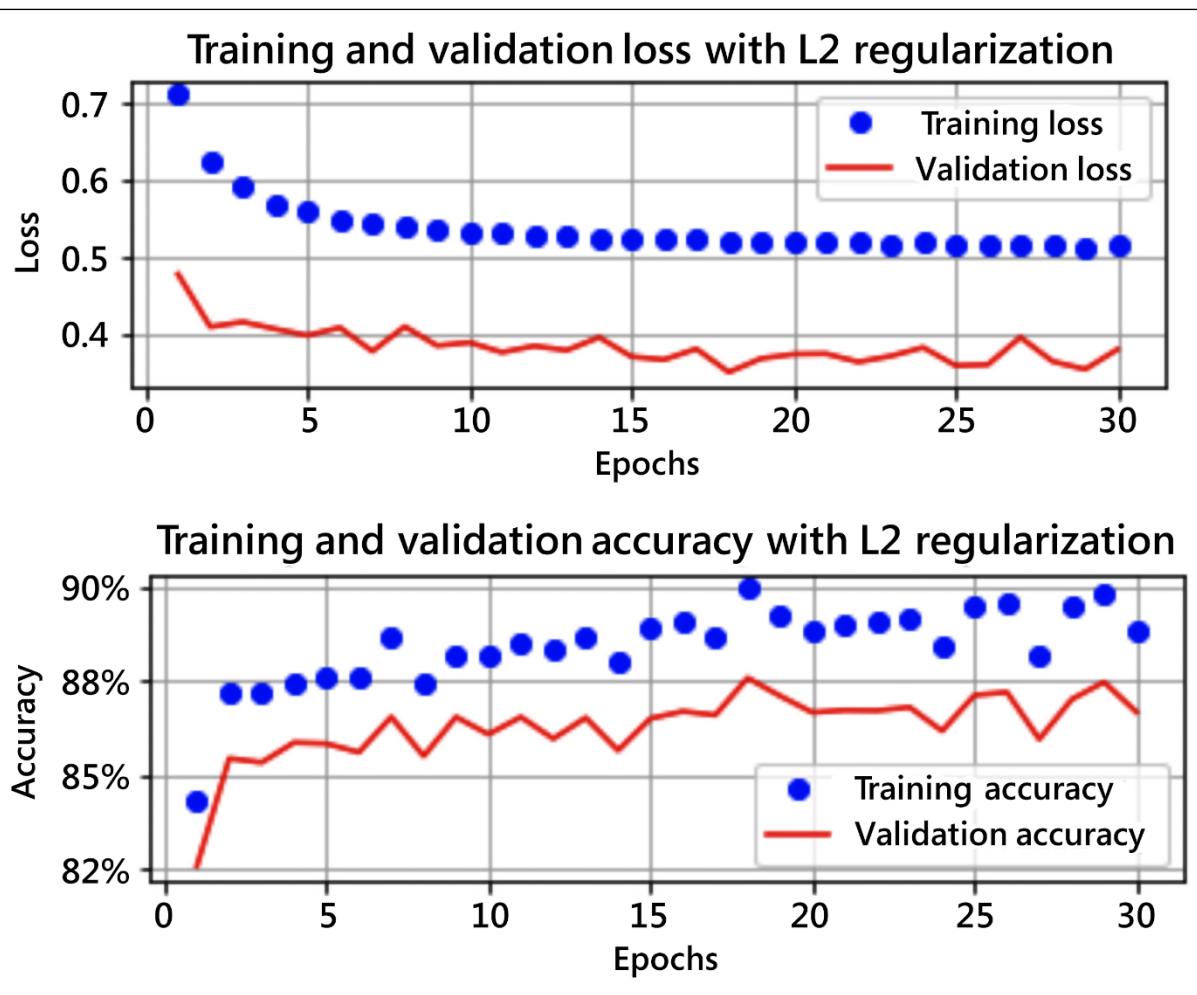
Training and validation loss with dropout



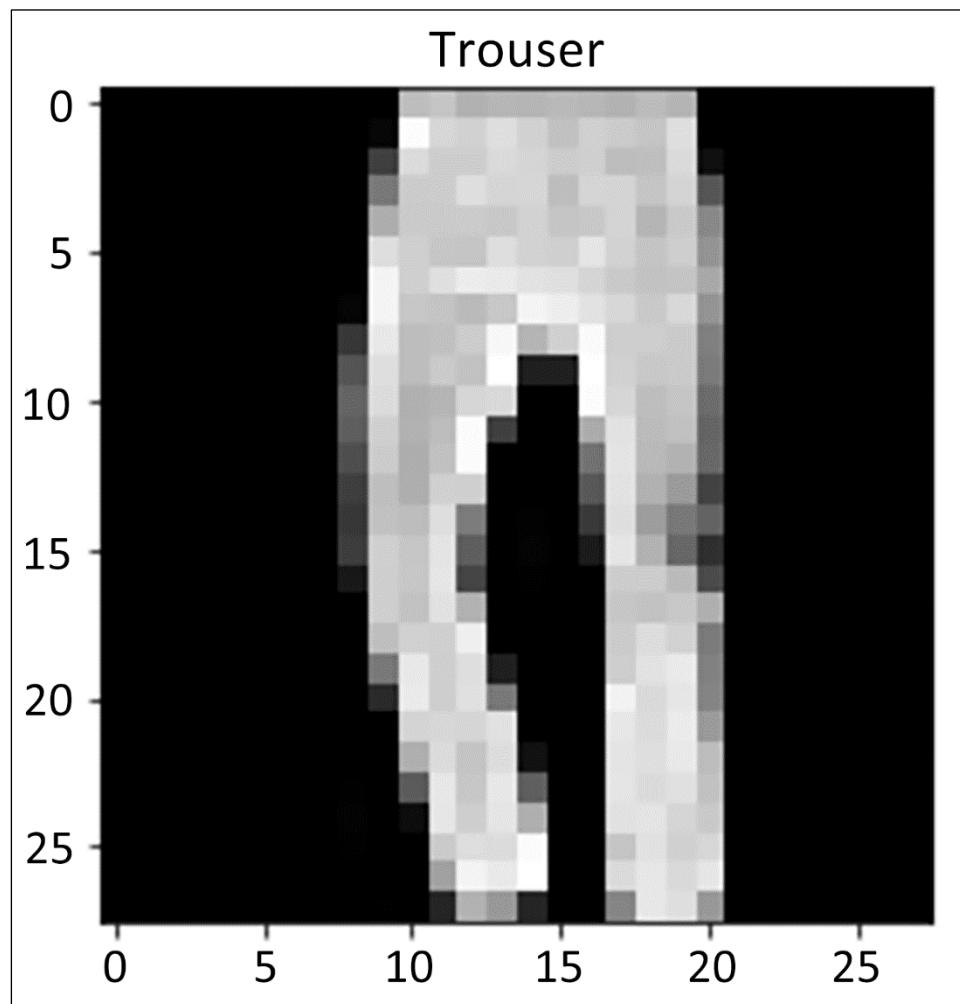
Training and validation accuracy with dropout







Chapter 4:



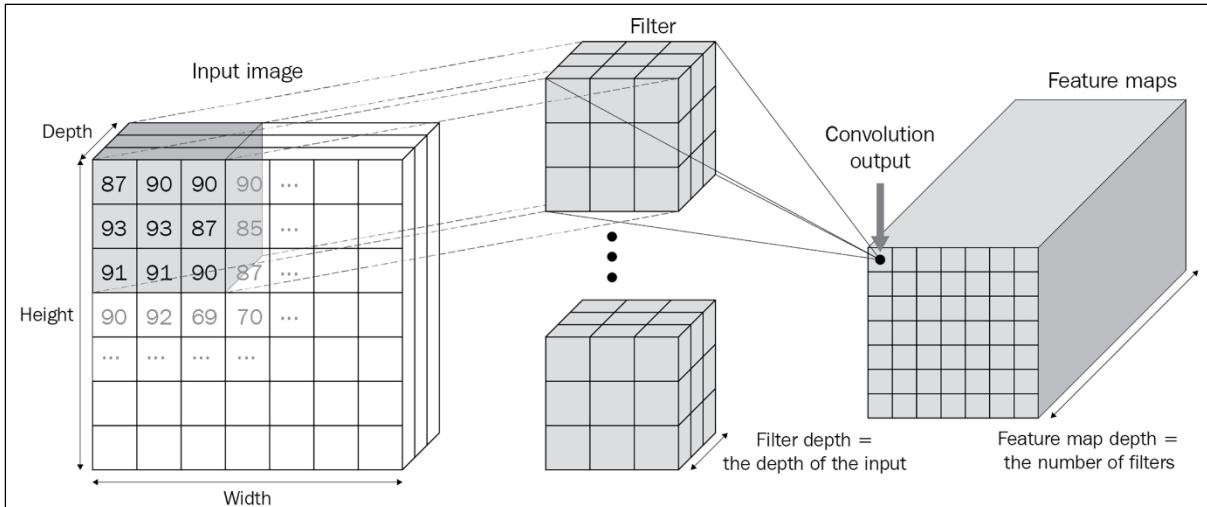
```
array([1.7608897e-08, 1.0000000e+00, 2.6042574e-13, 1.1353759e-10,
       3.1050048e-12, 7.2957764e-16, 8.0109371e-11, 3.8039204e-22,
       1.2800090e-15, 2.8759430e-18], dtype=float32)
```

Probability of each class for various translations										
5 pixels	-5 pixels	-4 pixels	-3 pixels	-2 pixels	-1 pixels	0 pixels	1 pixels	2 pixels	3 pixels	4 pixels
T-shirt/top	0.01	0.00	0.09	0.02	0.00	0.00	0.87	0.00	0.00	0.00
Trouser	0.02	0.00	0.01	0.20	0.02	0.00	0.75	0.00	0.00	0.00
Pullover	0.03	0.06	0.01	0.27	0.13	0.00	0.51	0.00	0.00	0.00
Dress	0.01	0.63	0.00	0.12	0.18	0.00	0.06	0.00	0.00	0.00
Coat	0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Sandal	0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Shirt	0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Sneaker	0.01	0.85	0.00	0.08	0.00	0.00	0.05	0.00	0.00	0.01
Bag	0.00	0.13	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.86
Ankle boot	0.00	0.10	0.00	0.00	0.00	0.01	0.00	0.00	0.00	0.89
	0.00	0.01	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.99

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

1	2
3	4

44	54	64
84	94	104
124	134	144



44	64
124	144

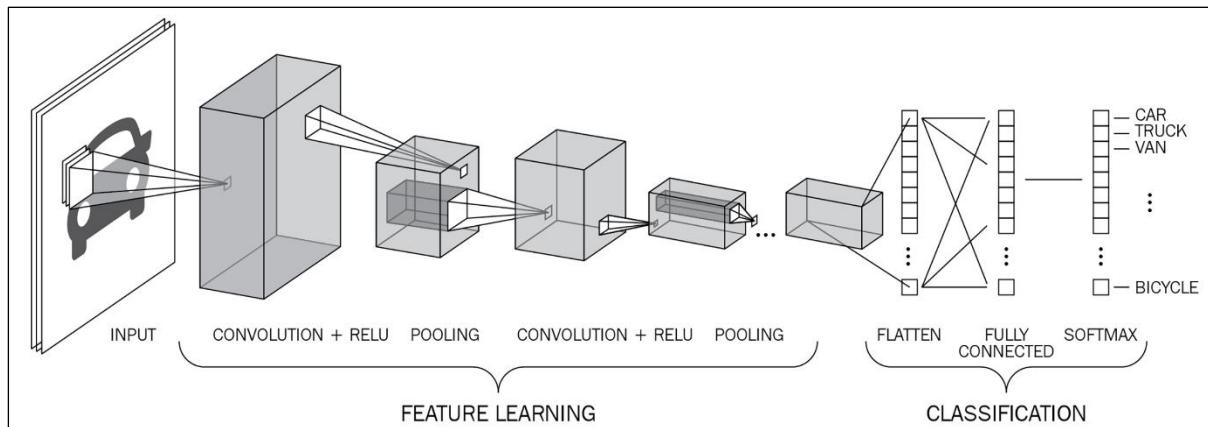
0	0	0	0	0	0
0	1	2	3	4	0
0	5	6	7	8	0
0	9	10	11	12	0
0	13	14	15	16	0
0	0	0	0	0	0

1	2
3	4

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

6	8
14	16



Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 1, 2, 2]	10
MaxPool2d-2	[-1, 1, 1, 1]	0
ReLU-3	[-1, 1, 1, 1]	0
Flatten-4	[-1, 1]	0
Linear-5	[-1, 1]	2
Sigmoid-6	[-1, 1]	0

Total params: 12
 Trainable params: 12
 Non-trainable params: 0

Input size (MB): 0.00
 Forward/backward pass size (MB): 0.00
 Params size (MB): 0.00
 Estimated Total Size (MB): 0.00

```

help(nn.Conv2d)

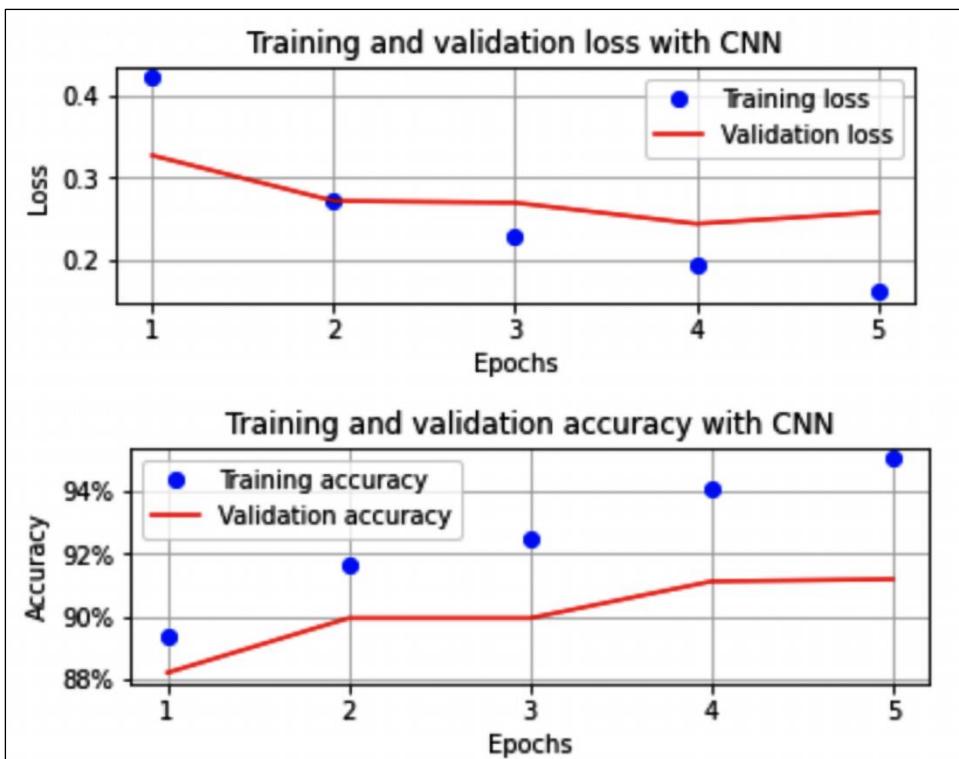
    Args:
        in_channels (int): Number of channels in the input image
        out_channels (int): Number of channels produced by the convolution
        kernel_size (int or tuple): Size of the convolving kernel
        stride (int or tuple, optional): Stride of the convolution. Default: 1
        padding (int or tuple, optional): Zero-padding added to both sides of the input. Default: 0
        padding_mode (string, optional). Accepted values `zeros` and `circular` Default: `zeros`
        dilation (int or tuple, optional): Spacing between kernel elements. Default: 1
        groups (int, optional): Number of blocked connections from input channels to output channels. Default: 1
        bias (bool, optional): If ``True``, adds a learnable bias to the output. Default: ``True``

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 26, 26]	640
MaxPool2d-2	[-1, 64, 13, 13]	0
ReLU-3	[-1, 64, 13, 13]	0
Conv2d-4	[-1, 128, 11, 11]	73,856
MaxPool2d-5	[-1, 128, 5, 5]	0
ReLU-6	[-1, 128, 5, 5]	0
Flatten-7	[-1, 3200]	0
Linear-8	[-1, 256]	819,456
ReLU-9	[-1, 256]	0
Linear-10	[-1, 10]	2,570

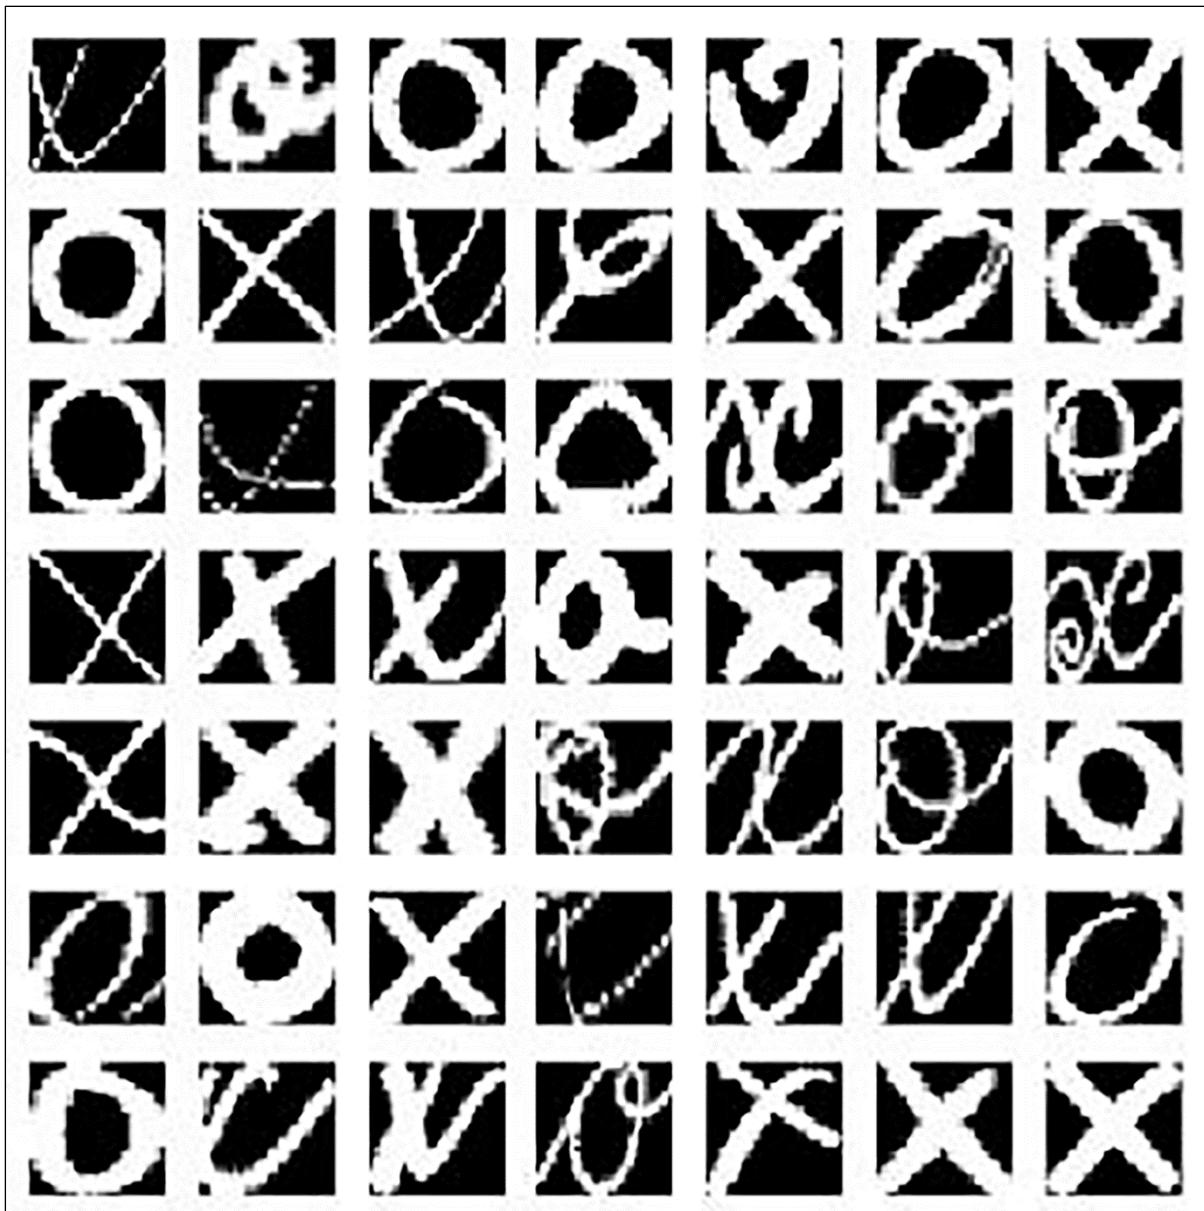
Total params:	896,522
Trainable params:	896,522
Non-trainable params:	0

Input size (MB):	0.00
Forward/backward pass size (MB):	0.69
Params size (MB):	3.42
Estimated Total Size (MB):	4.11



Probability of each class for various translations										
5 pixels	4 pixels	3 pixels	2 pixels	1 pixels	0 pixels	-1 pixels	-2 pixels	-3 pixels	-4 pixels	-5 pixels
0.00	0.00	0.01	0.03	0.16	0.00	0.01	0.00	0.80	0.00	0.00
0.00	0.99	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.01	0.86	0.04	0.00	0.02	0.04	0.02	0.00	0.00	0.00	0.00
0.24	0.11	0.12	0.00	0.00	0.47	0.02	0.03	0.00	0.00	0.00
T-shirt/top	Trouser	Pullover	Dress	Coat	Sandal	Shirt	Sneaker	Bag	Ankle boot	

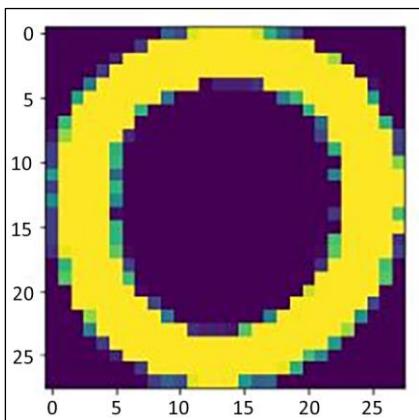
all/o@InterconnectedDemo-Bold@IttL47.png
all/o@Refresh-Regular@LX2MG4.png
all/x@CallistaOllander@7EWgpq.png
all/x@ChristmasSeason@xZ7mjB.png

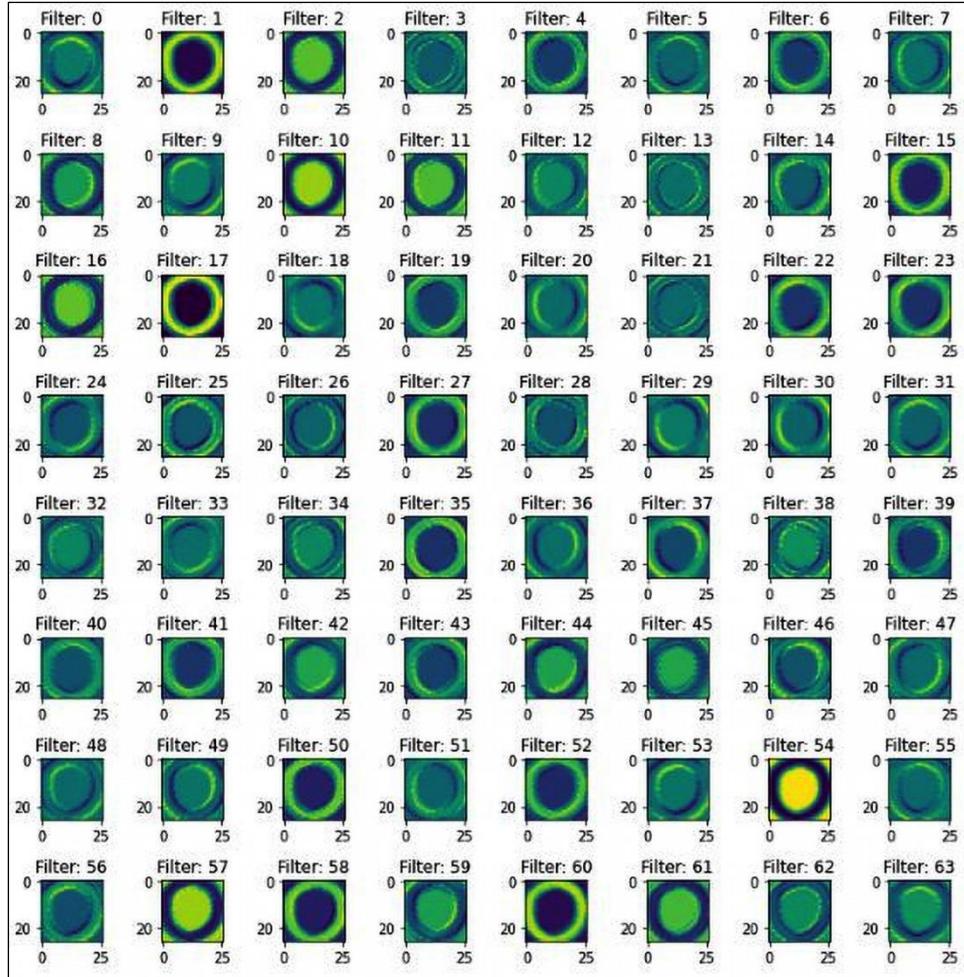


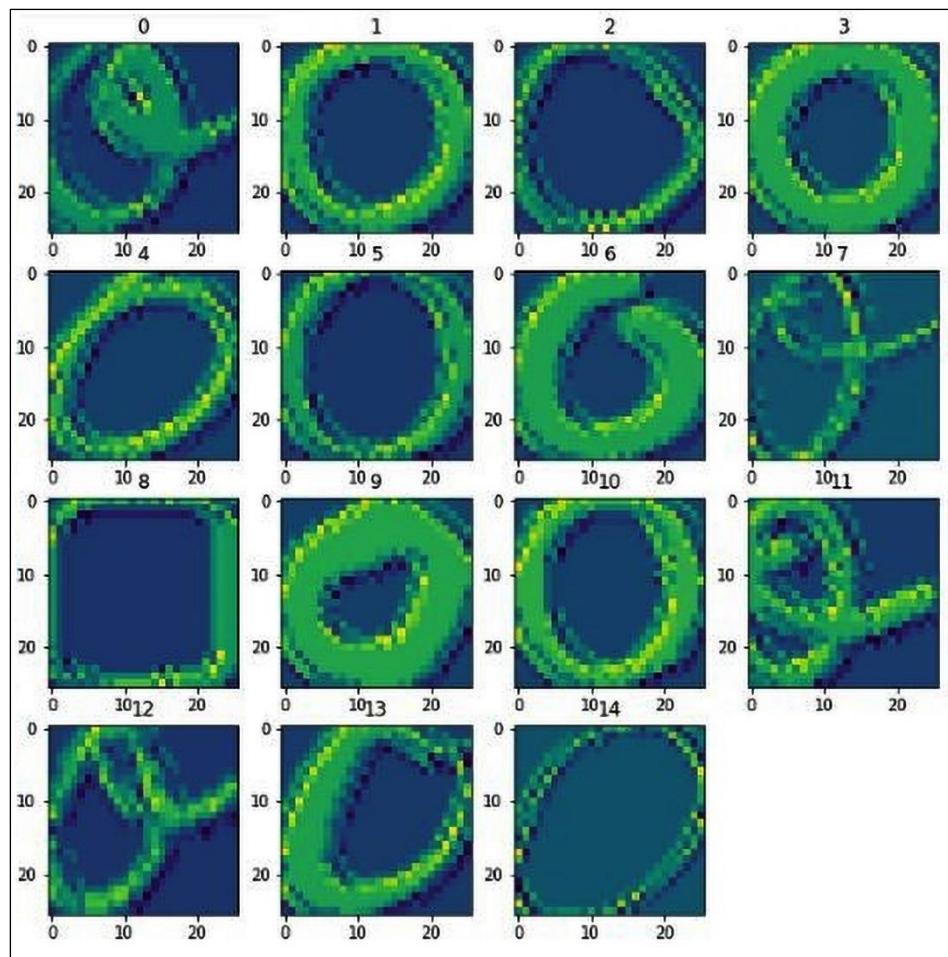
Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 26, 26]	640
MaxPool2d-2	[-1, 64, 13, 13]	0
ReLU-3	[-1, 64, 13, 13]	0
Conv2d-4	[-1, 128, 11, 11]	73,856
MaxPool2d-5	[-1, 128, 5, 5]	0
ReLU-6	[-1, 128, 5, 5]	0
Flatten-7	[-1, 3200]	0
Linear-8	[-1, 256]	819,456
ReLU-9	[-1, 256]	0
Linear-10	[-1, 1]	257
Sigmoid-11	[-1, 1]	0

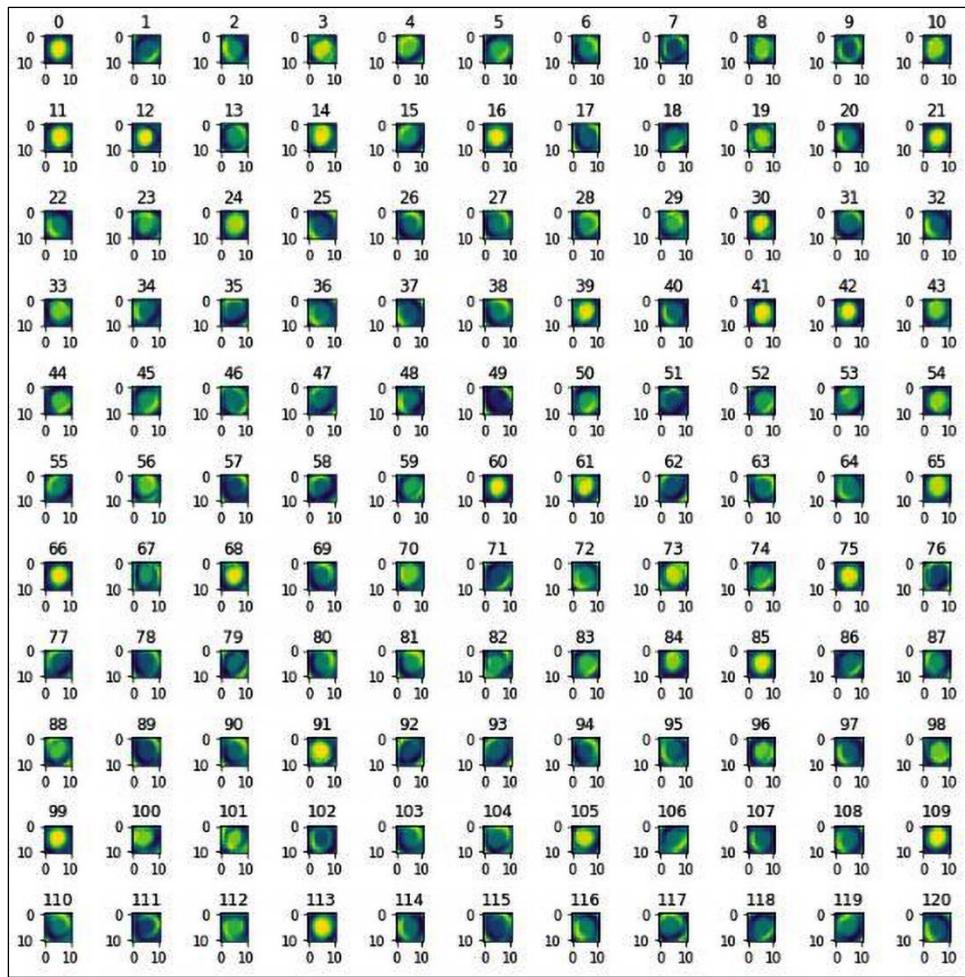
Total params: 894,209
Trainable params: 894,209
Non-trainable params: 0

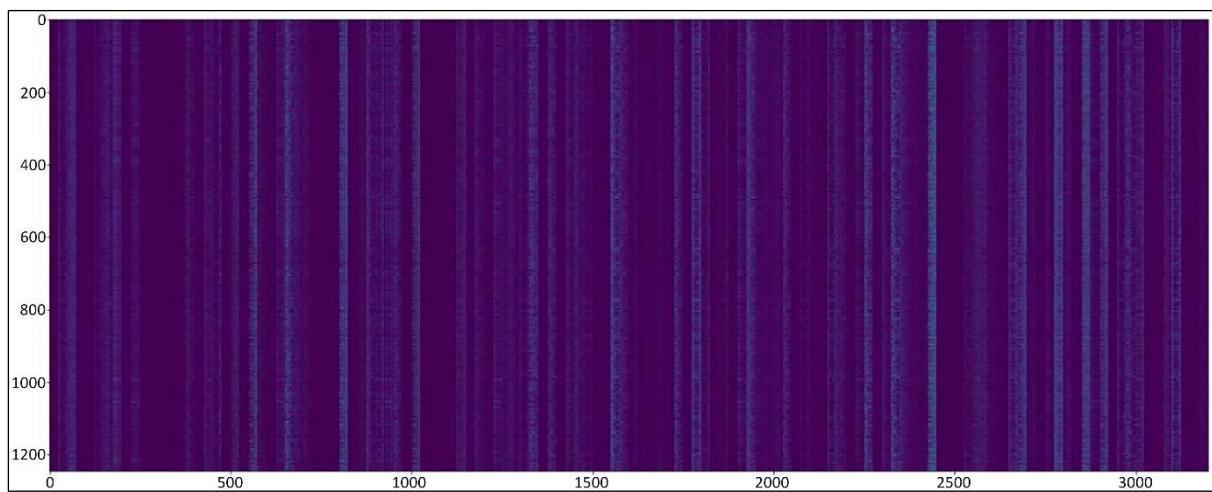
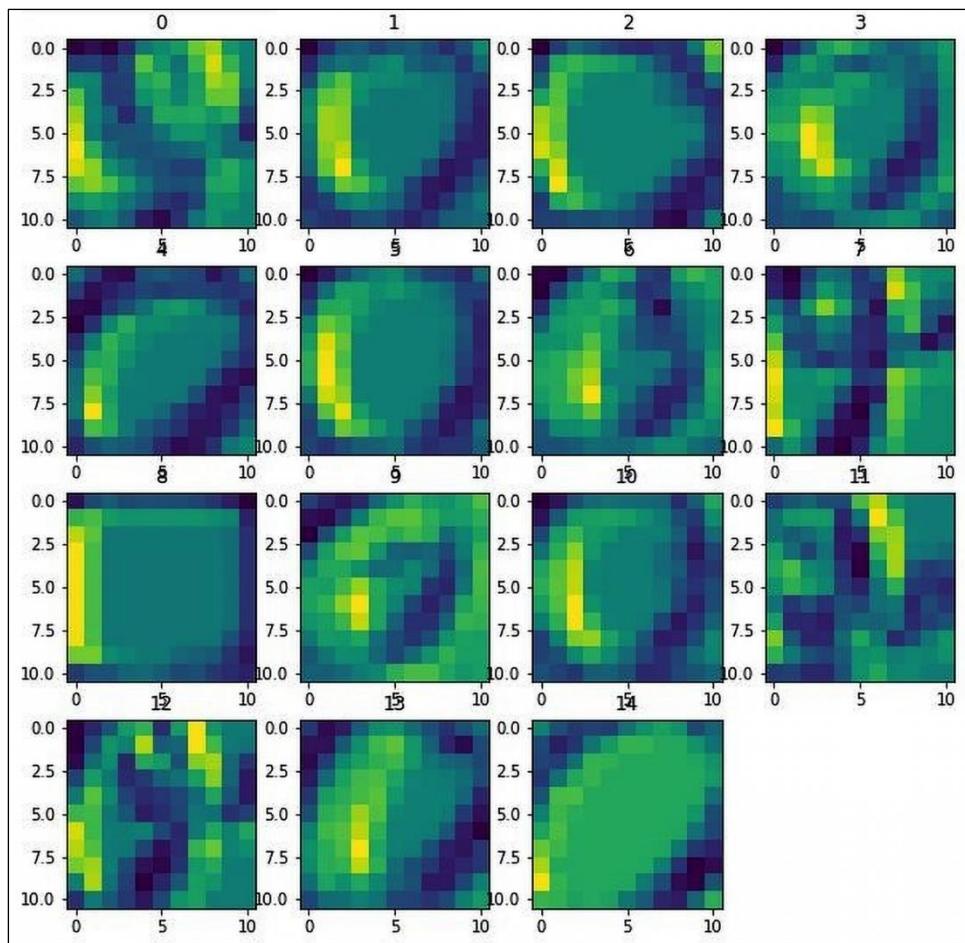
Input size (MB): 0.00
Forward/backward pass size (MB): 0.69
Params size (MB): 3.41
Estimated Total Size (MB): 4.10



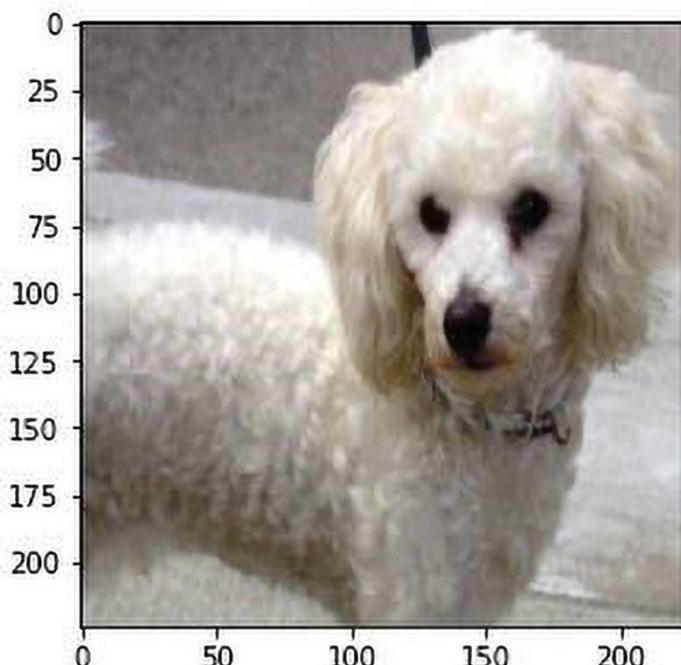








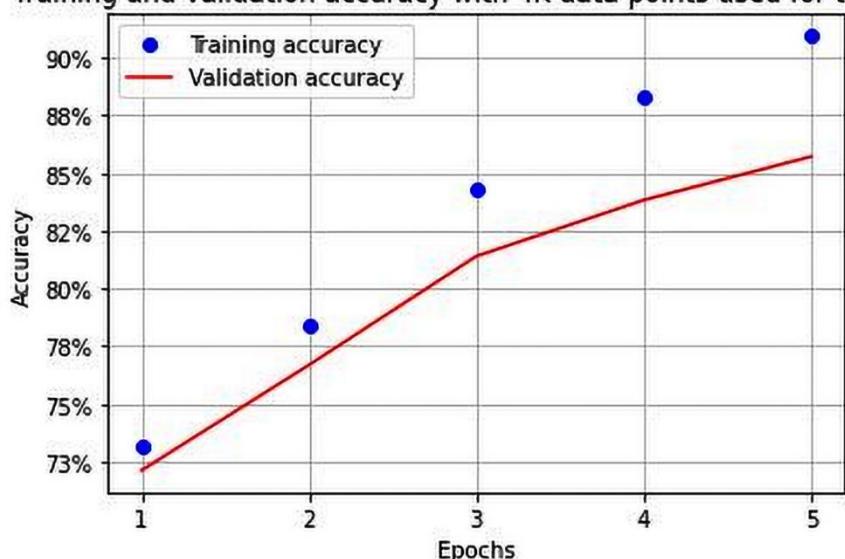
```
tensor([1.], device='cuda:0')
```



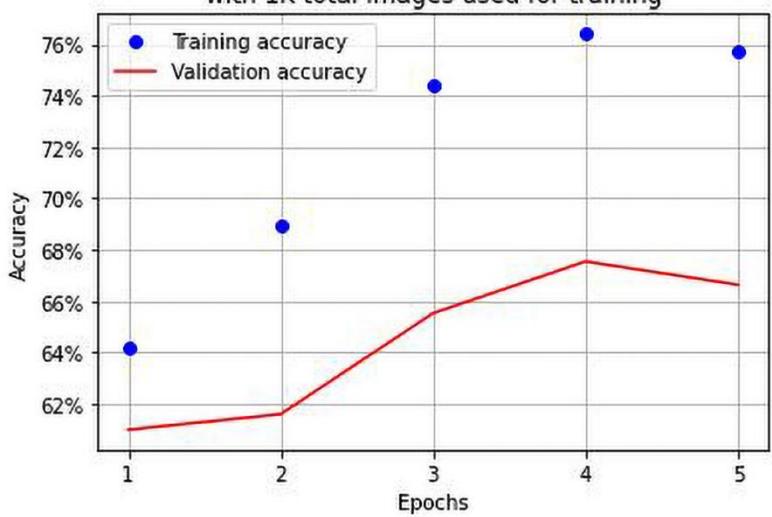
Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 222, 222]	1,792
ReLU-2	[-1, 64, 222, 222]	0
BatchNorm2d-3	[-1, 64, 222, 222]	128
MaxPool2d-4	[-1, 64, 111, 111]	0
Conv2d-5	[-1, 512, 109, 109]	295,424
ReLU-6	[-1, 512, 109, 109]	0
BatchNorm2d-7	[-1, 512, 109, 109]	1,024
MaxPool2d-8	[-1, 512, 54, 54]	0
Conv2d-9	[-1, 512, 52, 52]	2,359,808
ReLU-10	[-1, 512, 52, 52]	0
BatchNorm2d-11	[-1, 512, 52, 52]	1,024
MaxPool2d-12	[-1, 512, 26, 26]	0
Conv2d-13	[-1, 512, 24, 24]	2,359,808
ReLU-14	[-1, 512, 24, 24]	0
BatchNorm2d-15	[-1, 512, 24, 24]	1,024
MaxPool2d-16	[-1, 512, 12, 12]	0
Conv2d-17	[-1, 512, 10, 10]	2,359,808
ReLU-18	[-1, 512, 10, 10]	0
BatchNorm2d-19	[-1, 512, 10, 10]	1,024
MaxPool2d-20	[-1, 512, 5, 5]	0
Conv2d-21	[-1, 512, 3, 3]	2,359,808
ReLU-22	[-1, 512, 3, 3]	0
BatchNorm2d-23	[-1, 512, 3, 3]	1,024
MaxPool2d-24	[-1, 512, 1, 1]	0
Flatten-25	[-1, 512]	0
Linear-26	[-1, 1]	513
Sigmoid-27	[-1, 1]	0

Total params: 9,742,209
Trainable params: 9,742,209
Non-trainable params: 0

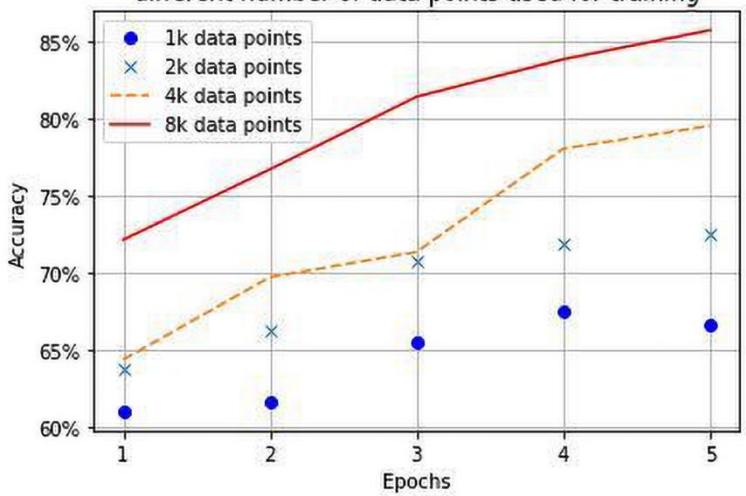
Training and validation accuracy with 4K data points used for training



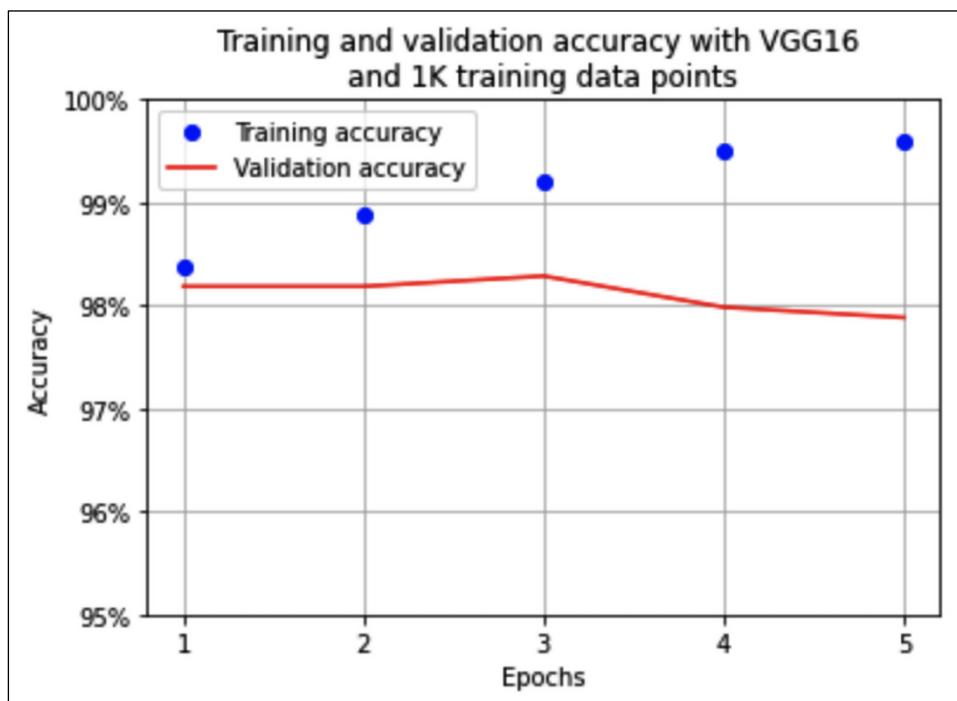
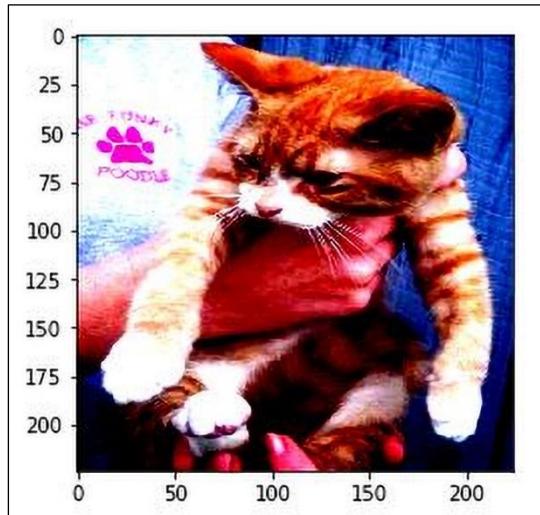
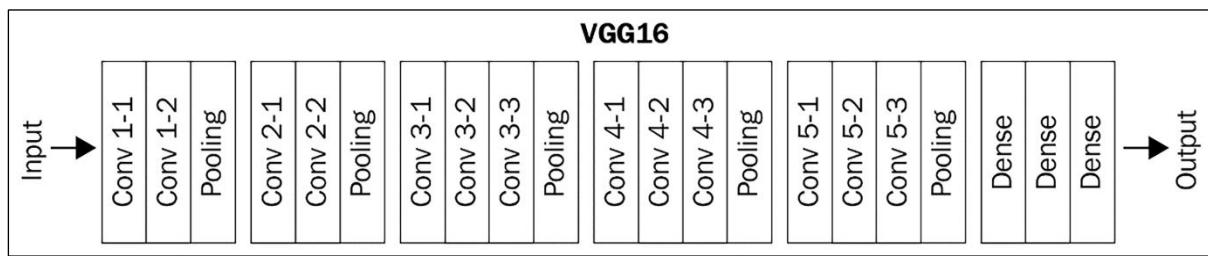
Training and validation accuracy
with 1K total images used for training

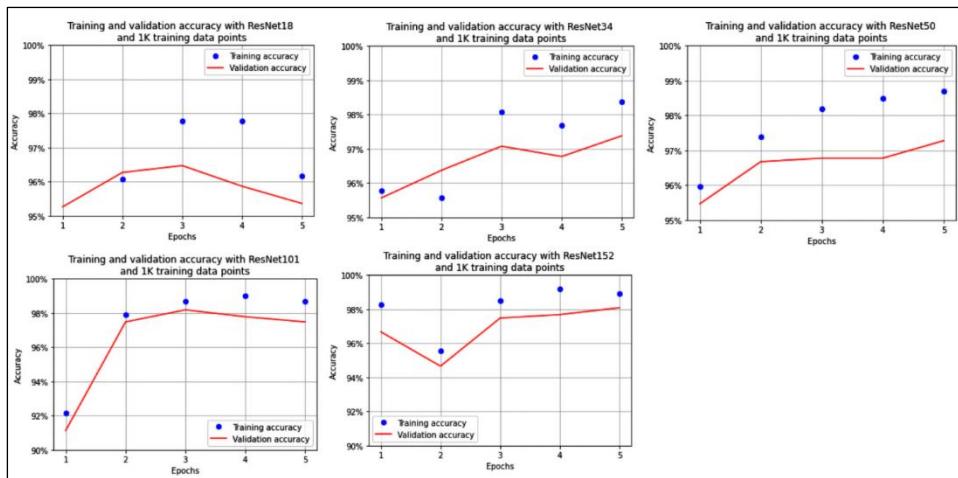
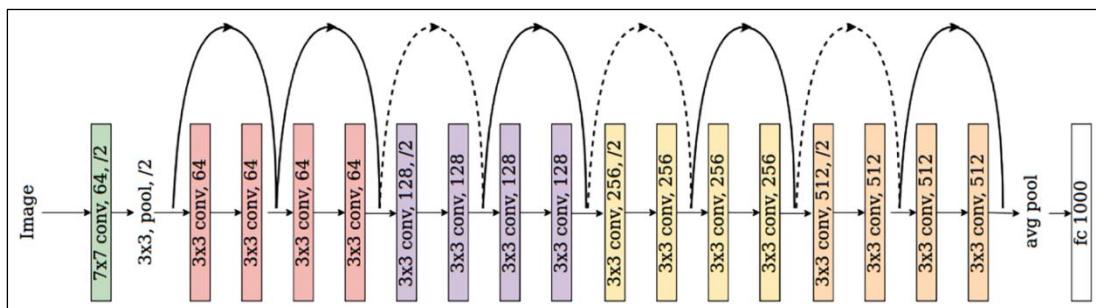
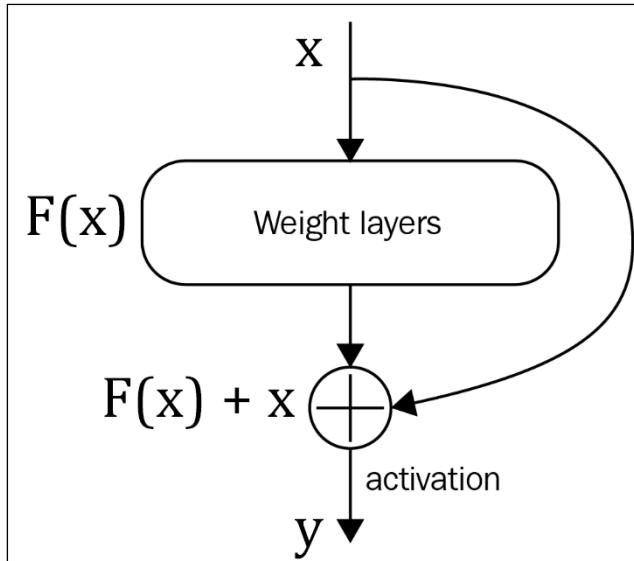
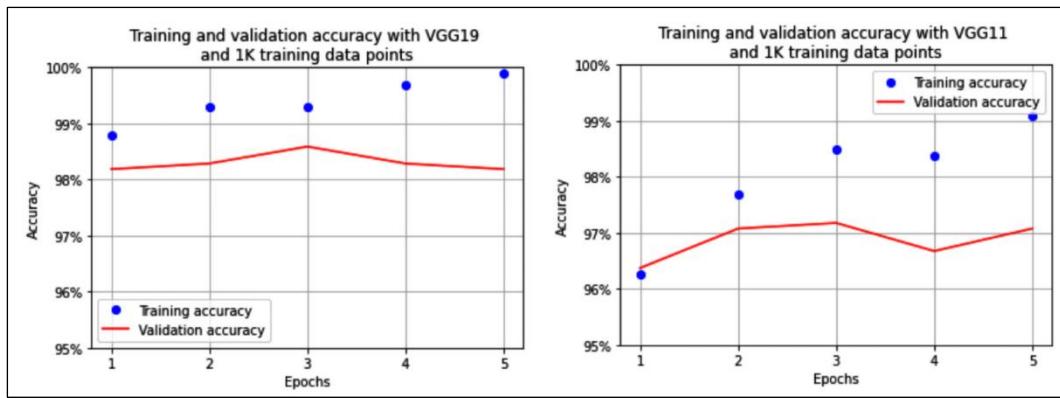


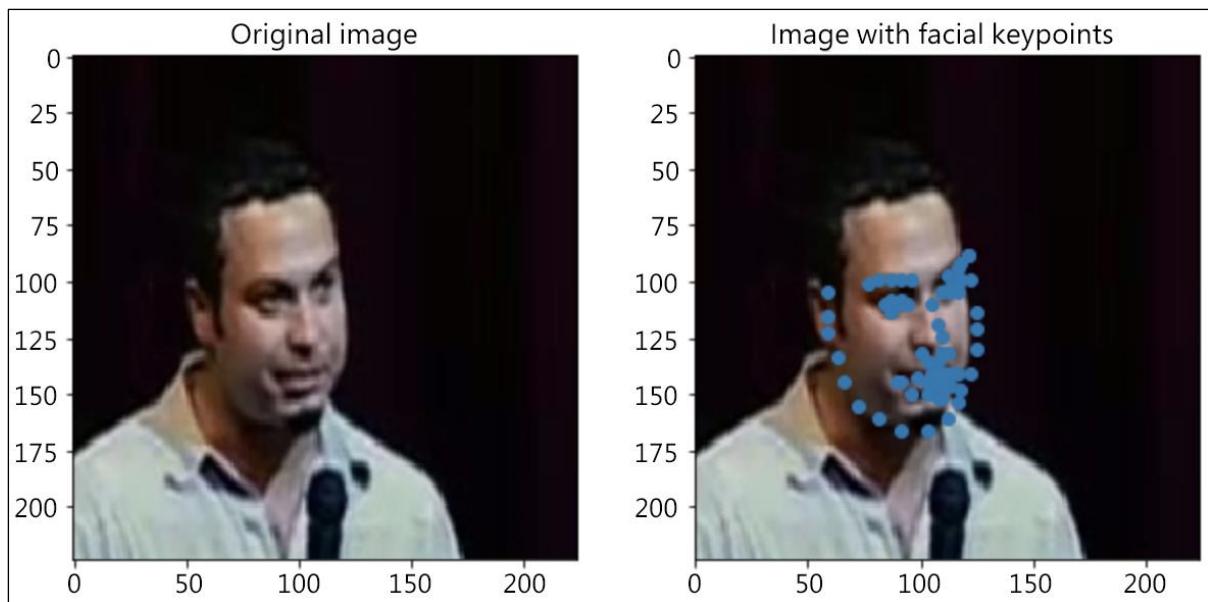
validation accuracy with
different number of data points used for training



Chapter 5:

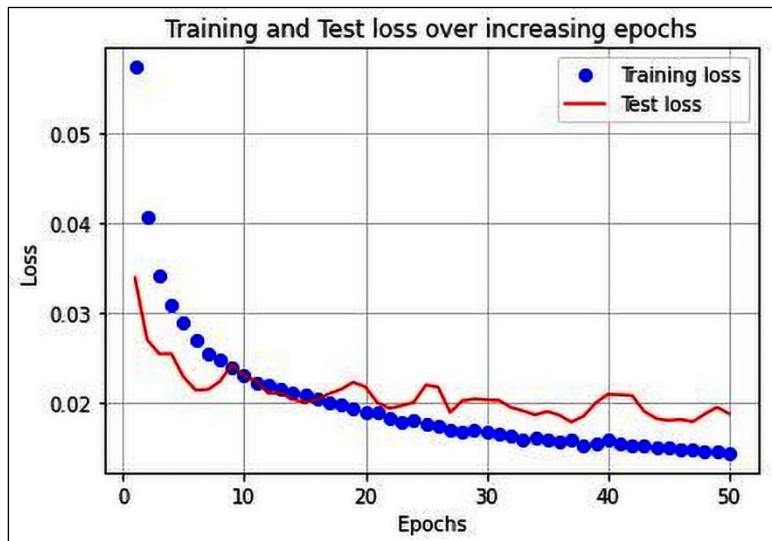


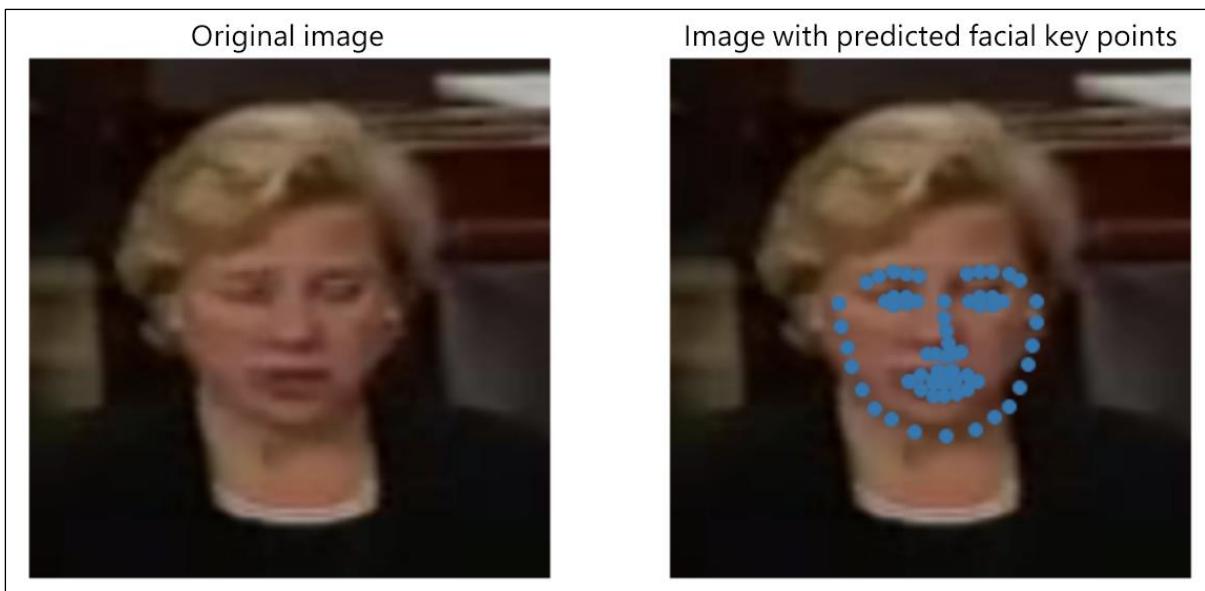


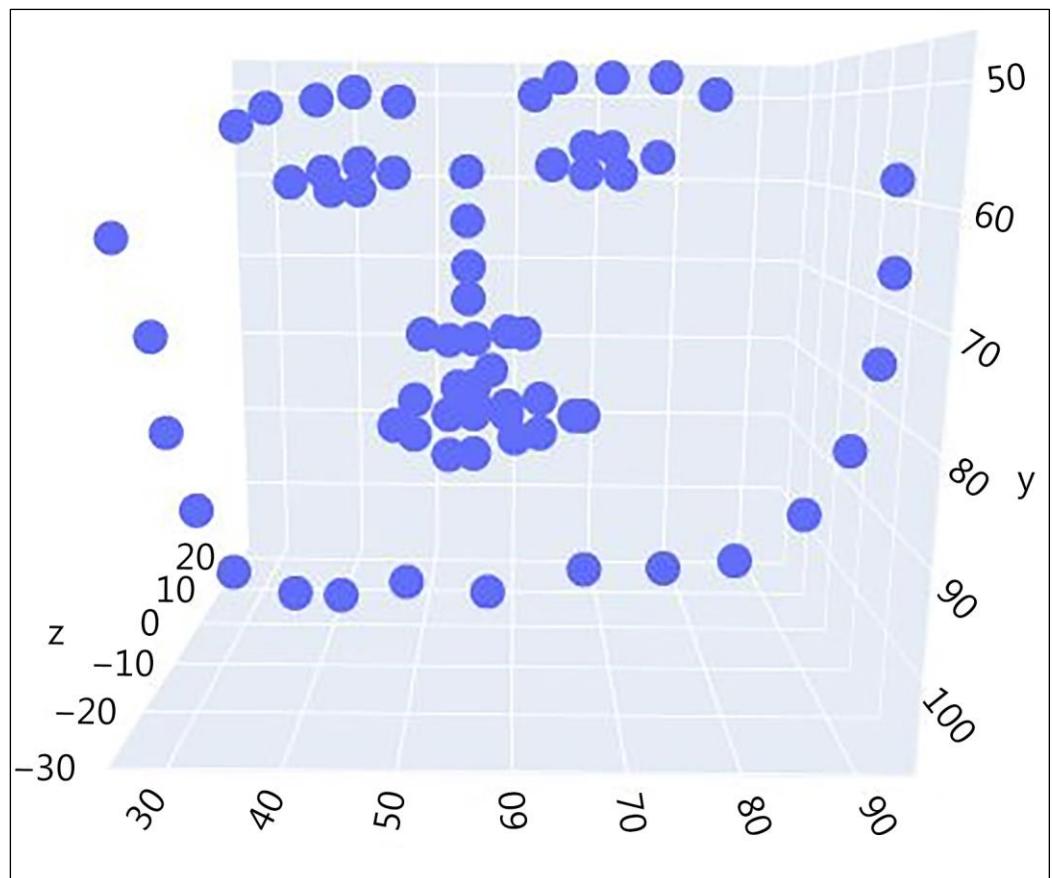


	Unnamed: 0	0	1	2	3	4	5	6
0	Luis_Fonsi_21.jpg	45.0	98.0	47.0	106.0	49.0	110.0	53.0
1	Lincoln_Chafee_52.jpg	41.0	83.0	43.0	91.0	45.0	100.0	47.0
2	Valerie_Harper_30.jpg	56.0	69.0	56.0	77.0	56.0	86.0	56.0
3	Angelo_Reyes_22.jpg	61.0	80.0	58.0	95.0	58.0	108.0	58.0
4	Kristen_Breitweiser_11.jpg	58.0	94.0	58.0	104.0	60.0	113.0	62.0

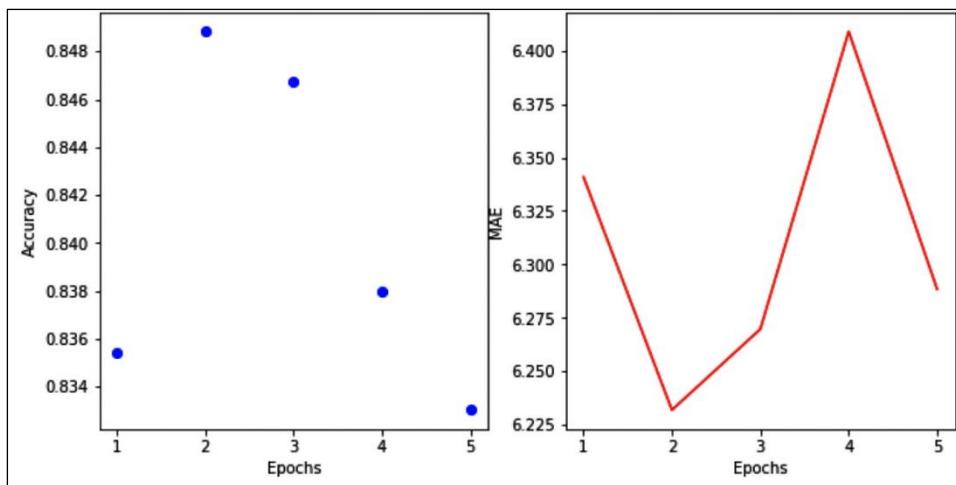
5 rows × 137 columns



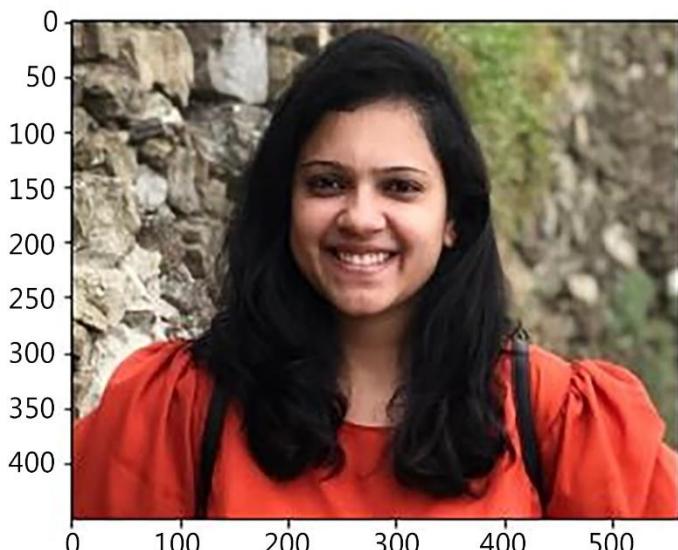




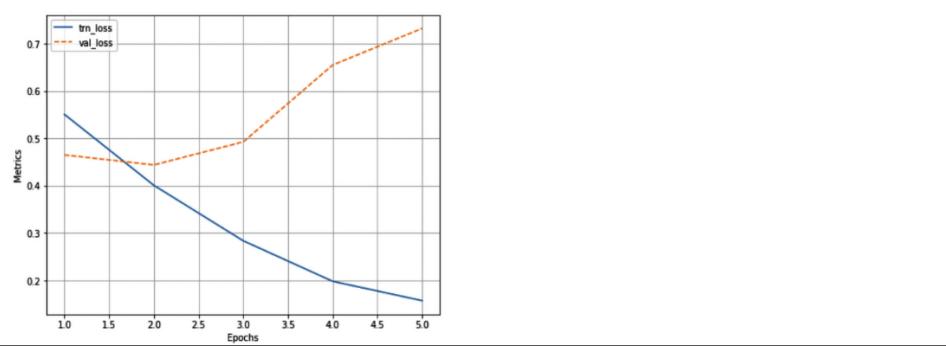
	file	age	gender	race	service_test
0	train/1.jpg	59	Male	East Asian	True
1	train/2.jpg	39	Female	Indian	False
2	train/3.jpg	11	Female	Black	False
3	train/4.jpg	26	Female	Indian	True
4	train/5.jpg	26	Female	Indian	True



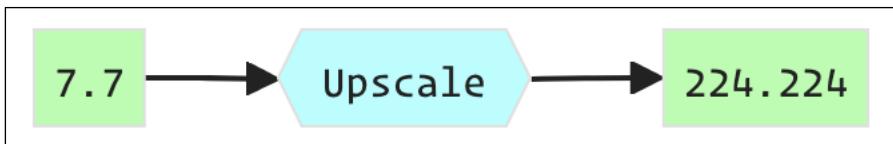
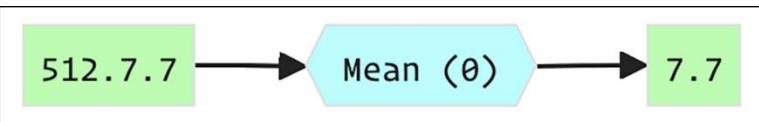
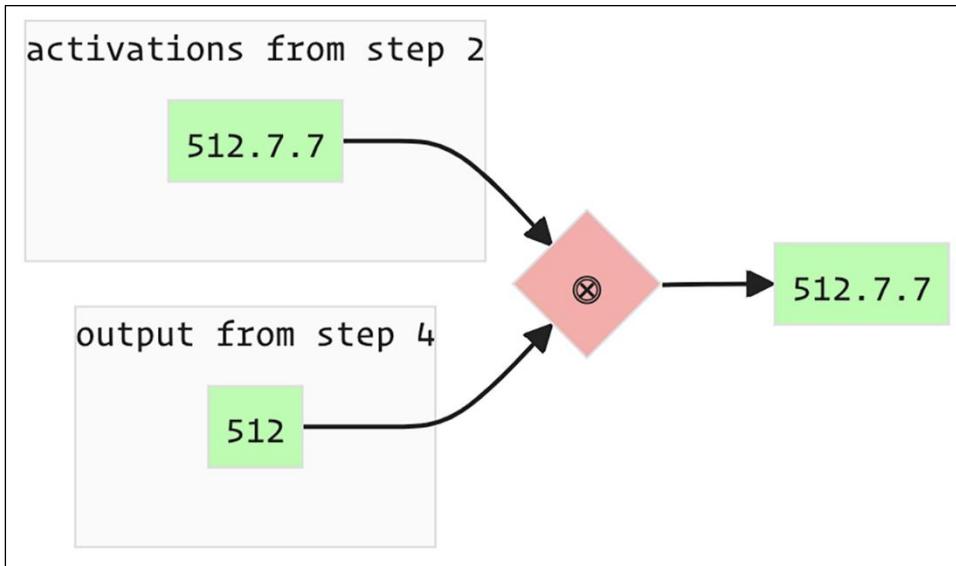
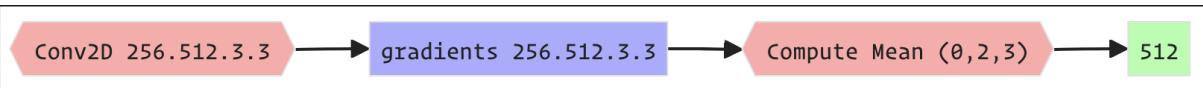
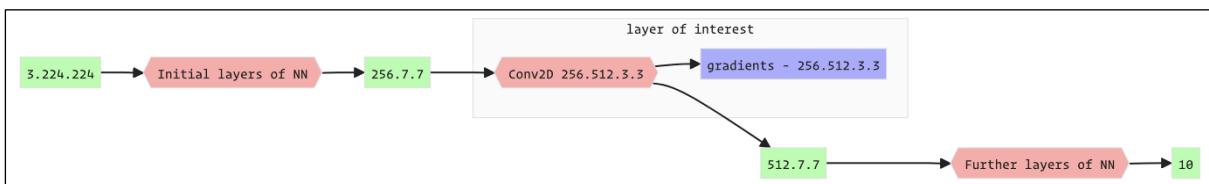
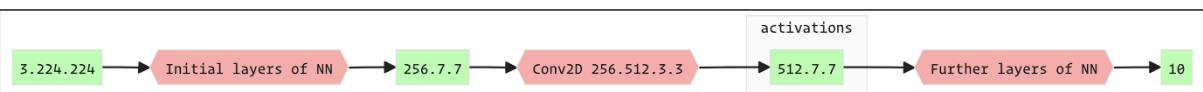
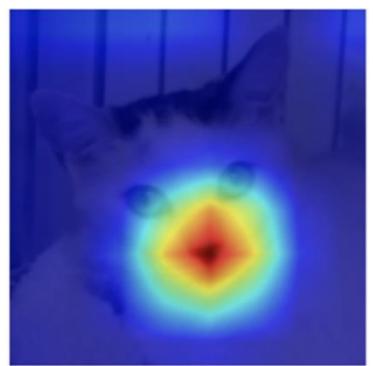
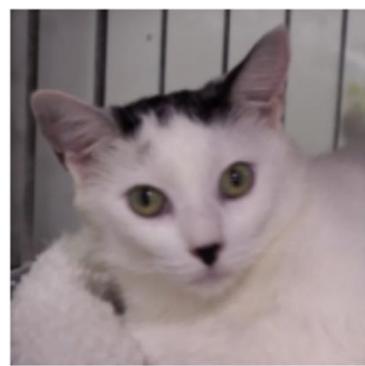
predicted gender: Female ; Predicted age 25

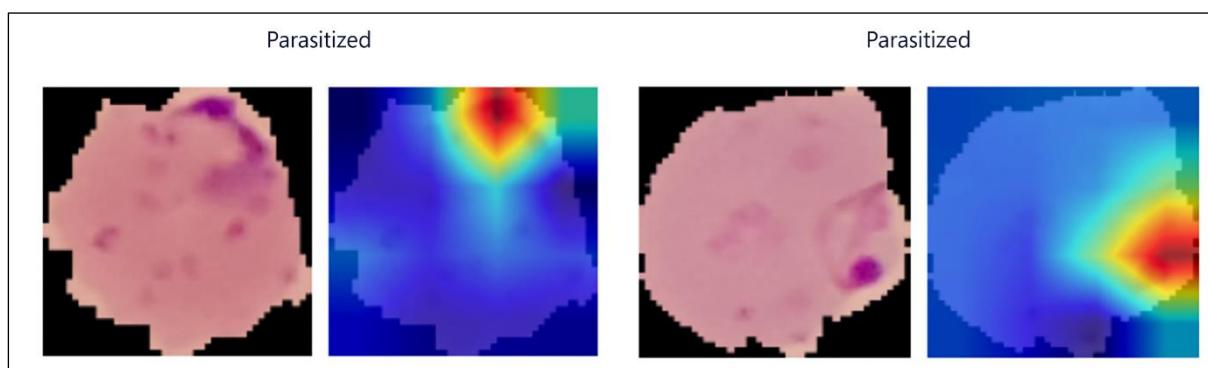
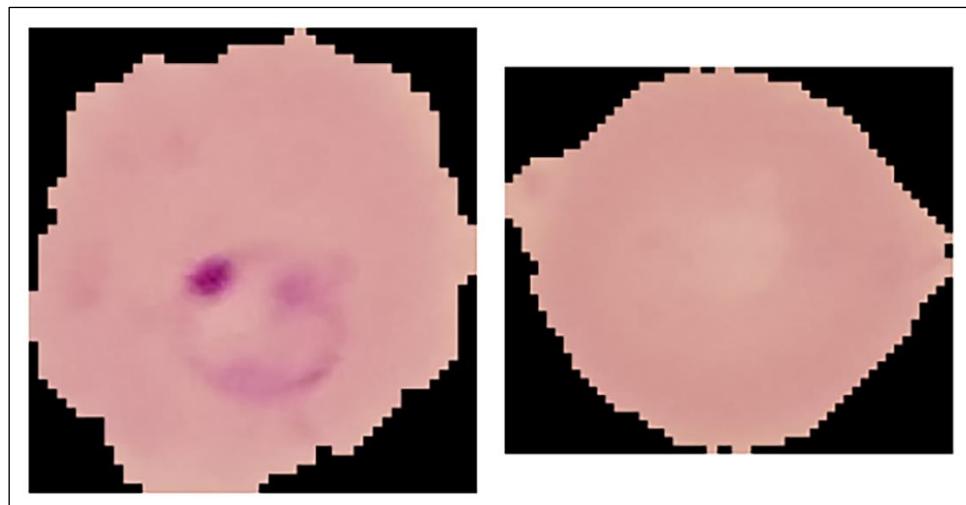
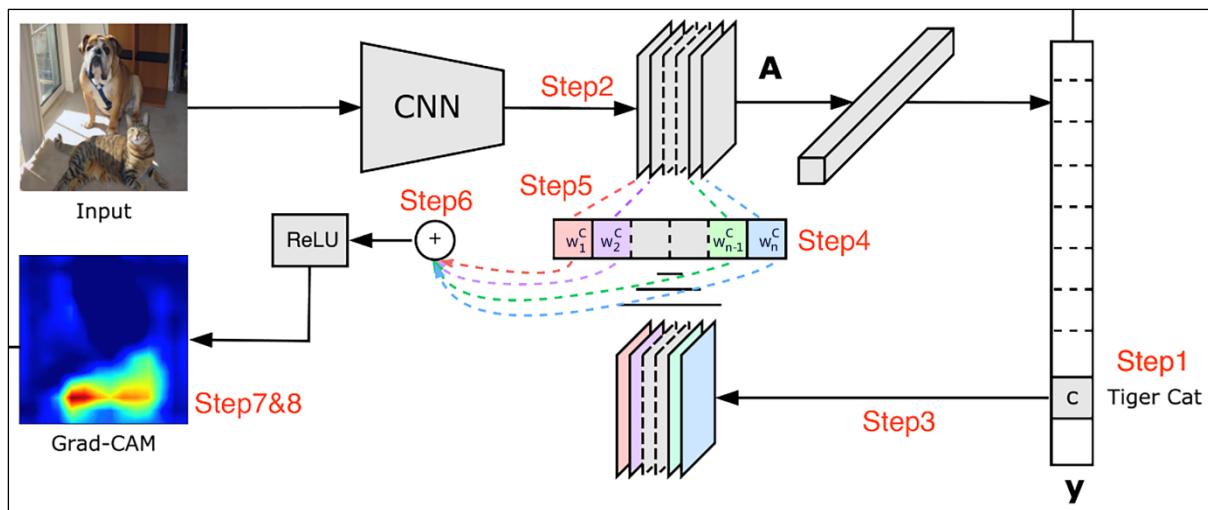


```
EPOCH: 1.000    trn_loss: 0.551 val_loss: 0.465 val_gender_acc: 0.834    val_age_mae: 6.238      (779.37s - 3117.47s remaining)
EPOCH: 2.000    trn_loss: 0.401 val_loss: 0.444 val_gender_acc: 0.847    val_age_mae: 6.229      (1555.70s - 2333.56s remaining)
EPOCH: 3.000    trn_loss: 0.284 val_loss: 0.493 val_gender_acc: 0.846    val_age_mae: 6.340      (2335.09s - 1556.73s remaining)
EPOCH: 4.000    trn_loss: 0.198 val_loss: 0.655 val_gender_acc: 0.842    val_age_mae: 6.339      (3110.83s - 777.71s remaining)
EPOCH: 4.994    trn_loss: 0.375 val_loss: 0.969 val_gender_acc: 0.969    val_age_mae: 6.541      (3887.91s - 4.54s remaining)
 0%| 0 / 6 [00:00<7, ?it/s] /usr/local/lib/python3.6/dist-packages/numpy/core/fromnumeric.py:3335: RuntimeWarning: Mean
  out=out, **kwargs)
/usr/local/lib/python3.6/dist-packages/numpy/core/_methods.py:161: RuntimeWarning: invalid value encountered in double_scalars
ret = ret.dtype.type(ret / rcount)
100%|██████████| 6/6 [00:00<00:00, 291.40it/s] EPOCH: 5.000    trn_loss: 0.157 val_loss: 0.733 val_gender_acc: 0.847    val_age_
```



Chapter 6:

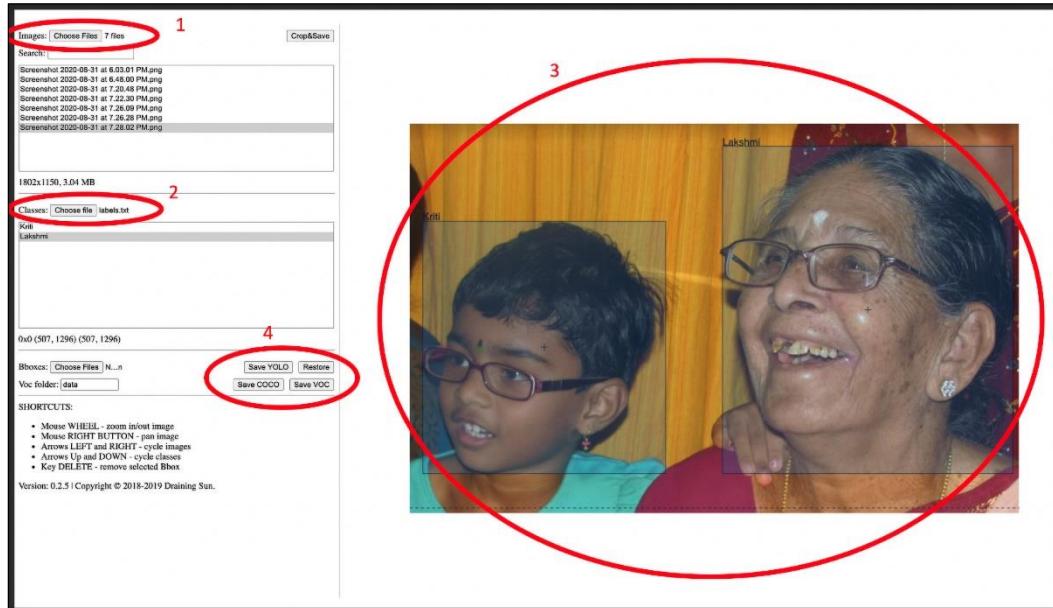
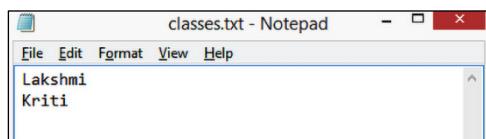
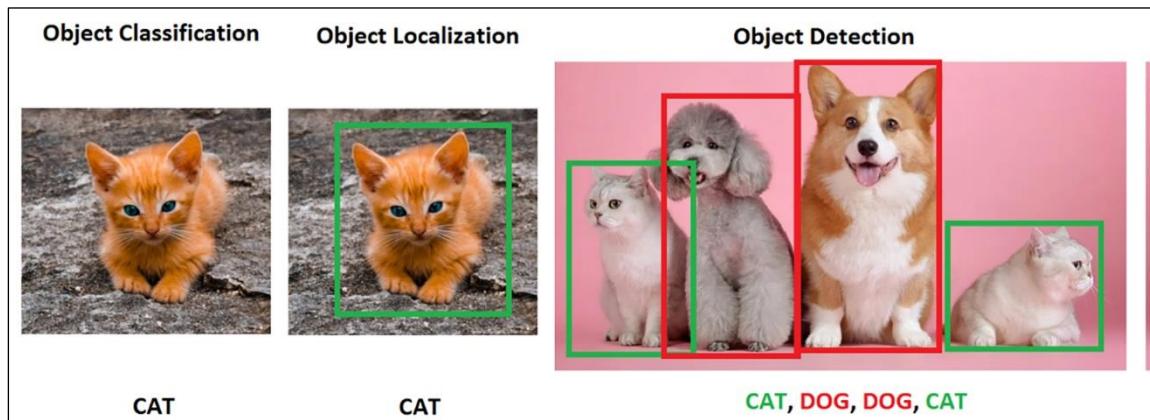


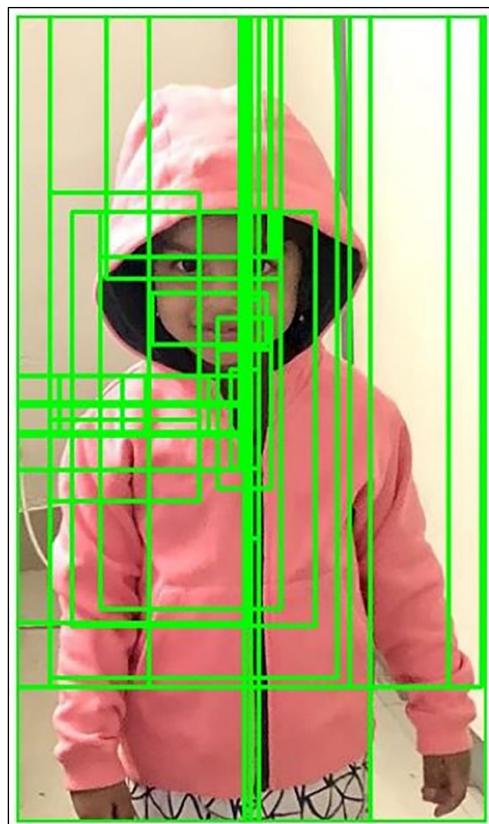
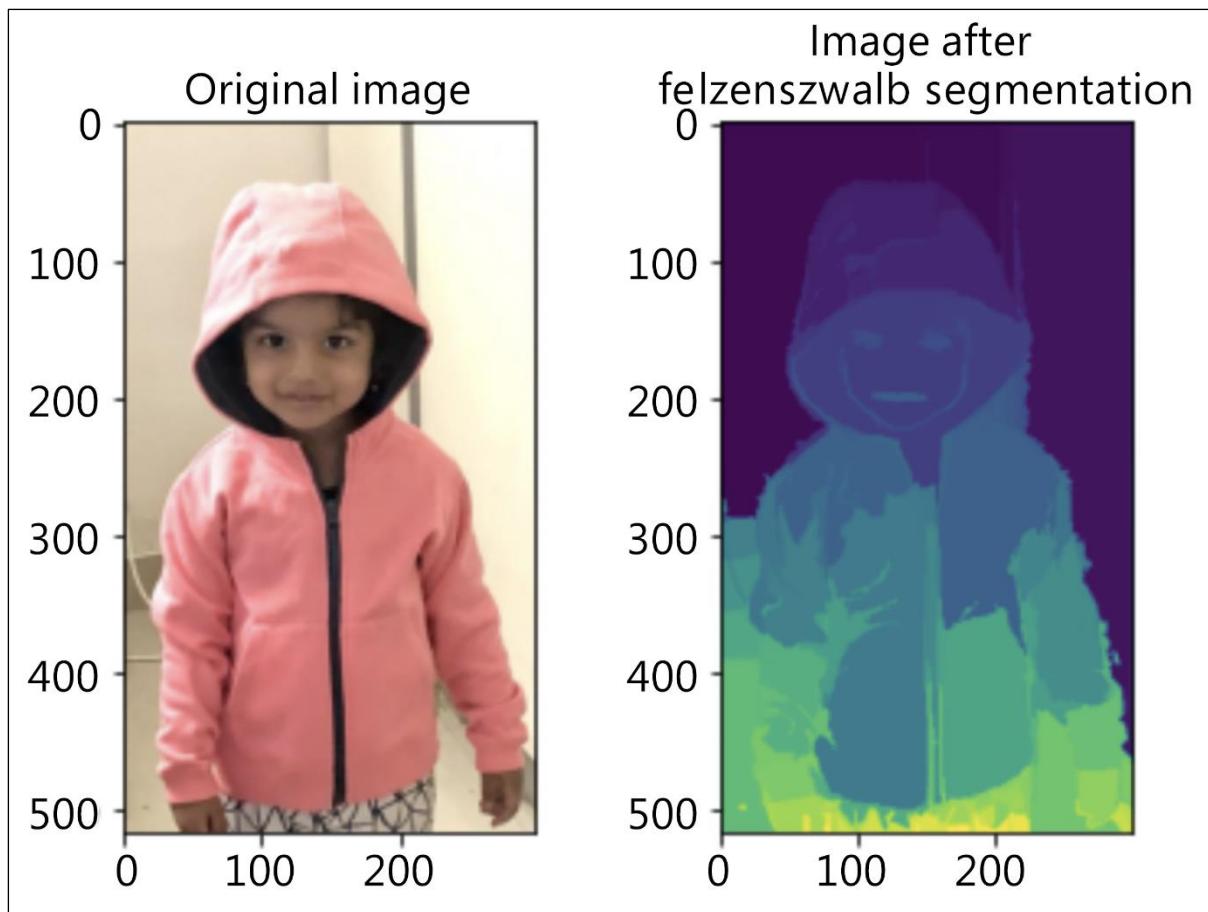


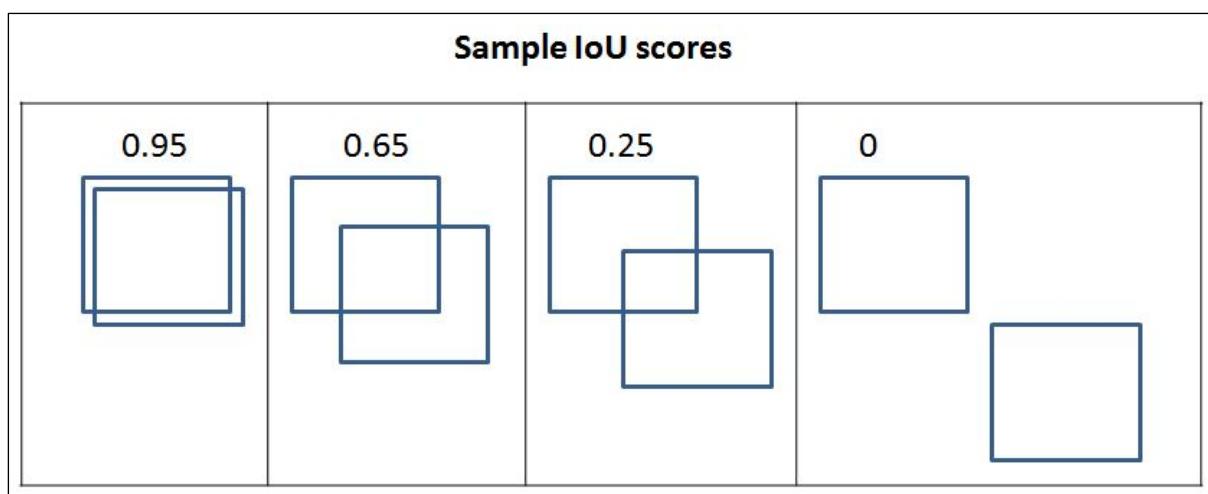
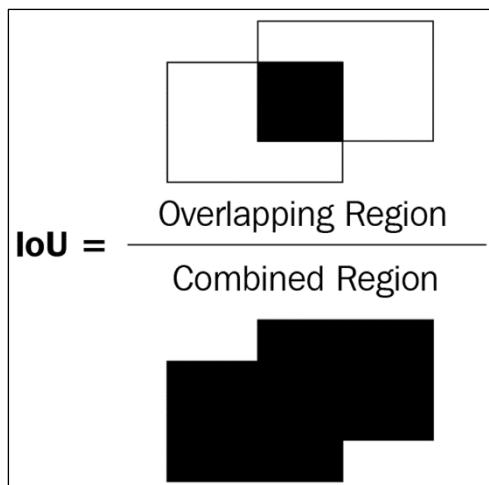
<i>Augment</i>	<i>Batch-Norm</i>	<i>Train Accuracy</i>	<i>Validation Accuracy</i>
No	No	95.9	94.5
No	Yes	99.3	97.7
Yes	Yes	97.7	97.6

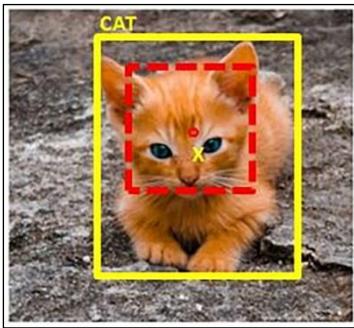
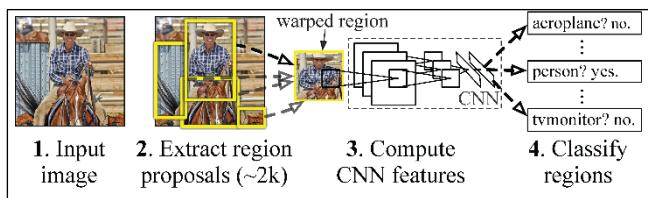
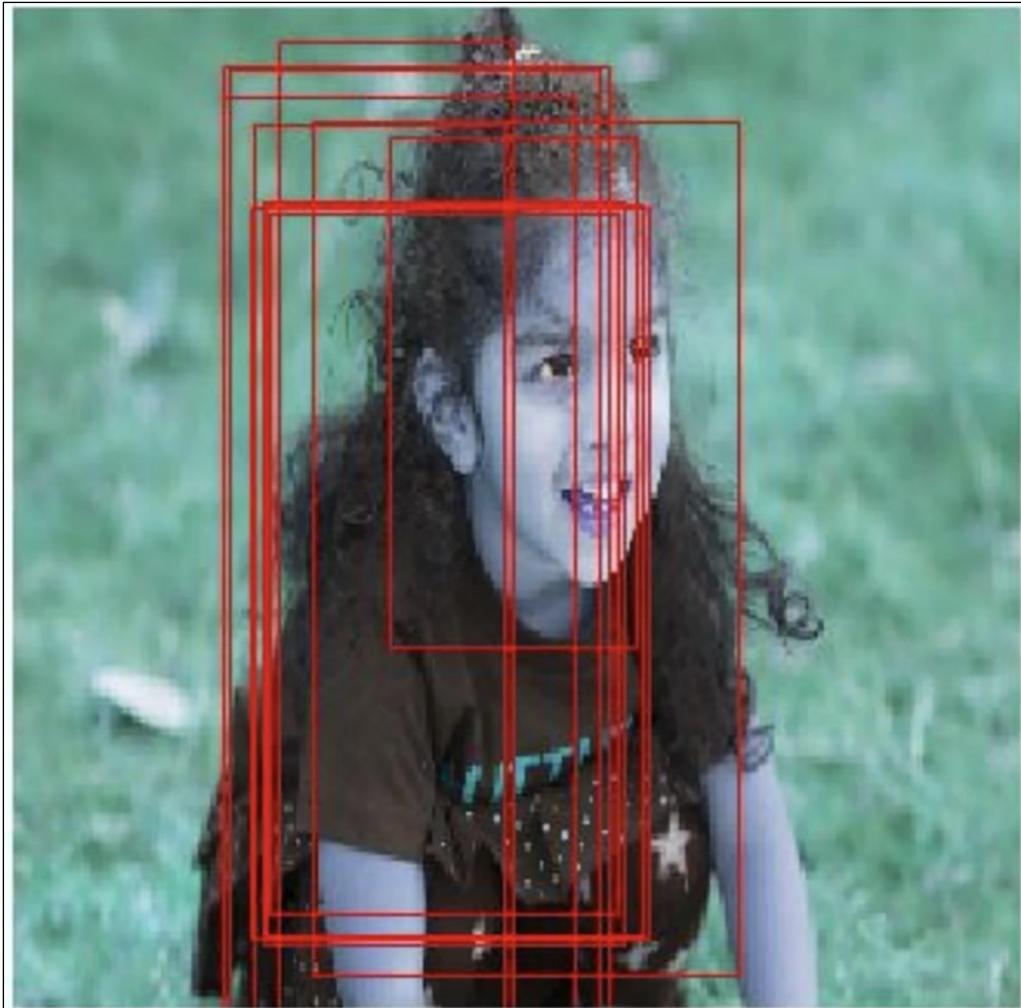
		Predicted	
		0	1
Actual	0	TN	FP
	1	FN	TP

Chapter 7:

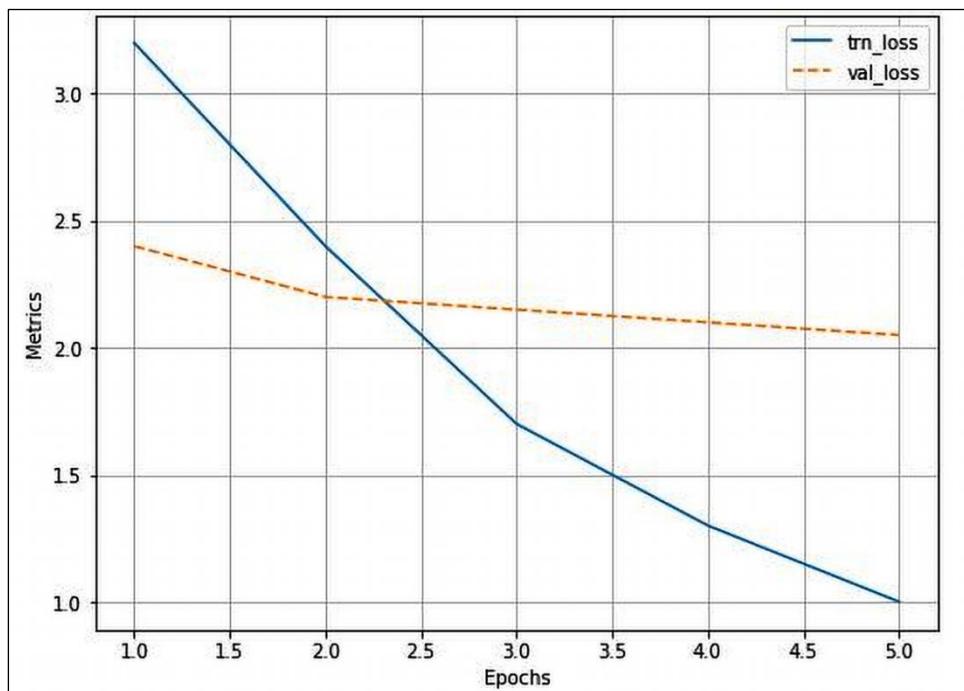


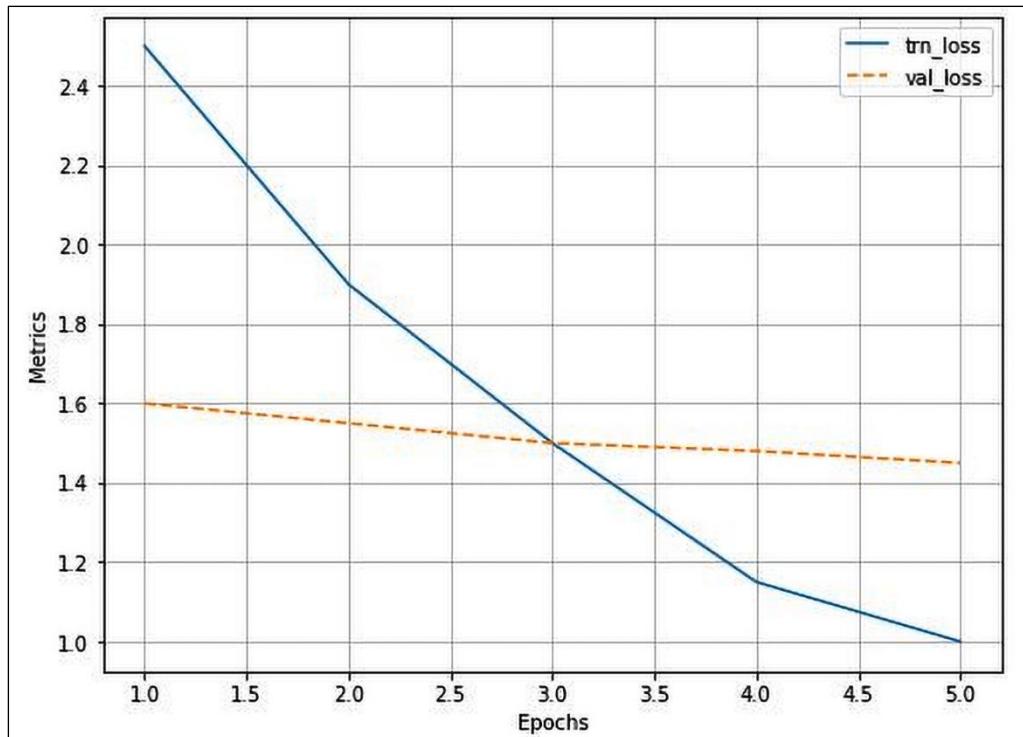
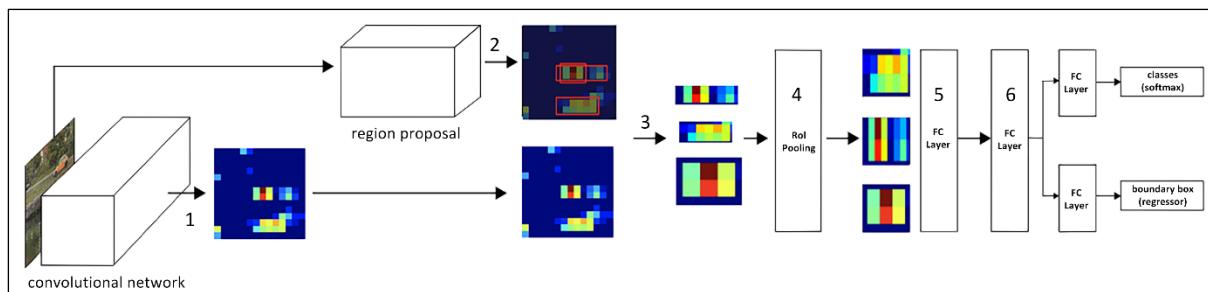




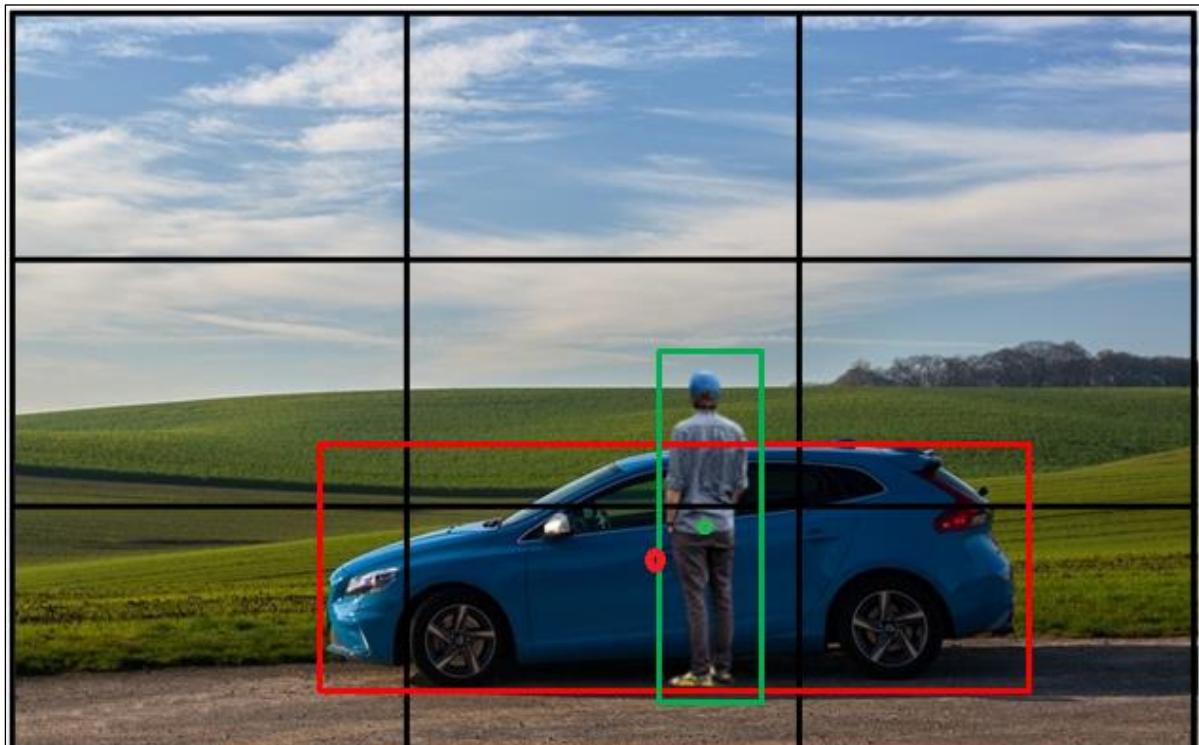


	ImageID	Source	LabelName	Confidence	XMin	XMax	YMin	YMax	IsOccluded	IsTruncated	IsGroupOf	IsDepiction	IsInside
20	002f8241bd829022	xclick	Bus	1	0.257812	0.515625	0.485417	0.891667	1	0	0	0	0
21	002f8241bd829022	xclick	Bus	1	0.535937	0.907813	0.347917	0.997917	1	1	0	0	0
191	013b99371484d3d5	xclick	Bus	1	0.154688	0.920312	0.102083	0.872917	0	0	0	0	0
322	01f8886b50a031a1	xclick	Truck	1	0.012821	0.987179	0.000000	0.969512	0	0	0	0	0
405	02717d30304f4849	xclick	Bus	1	0.106250	0.926562	0.266667	0.635417	0	0	0	0	0





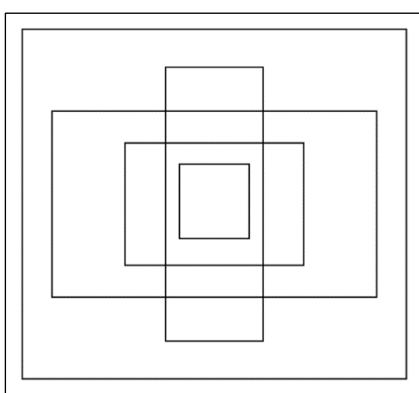
Chapter 8:

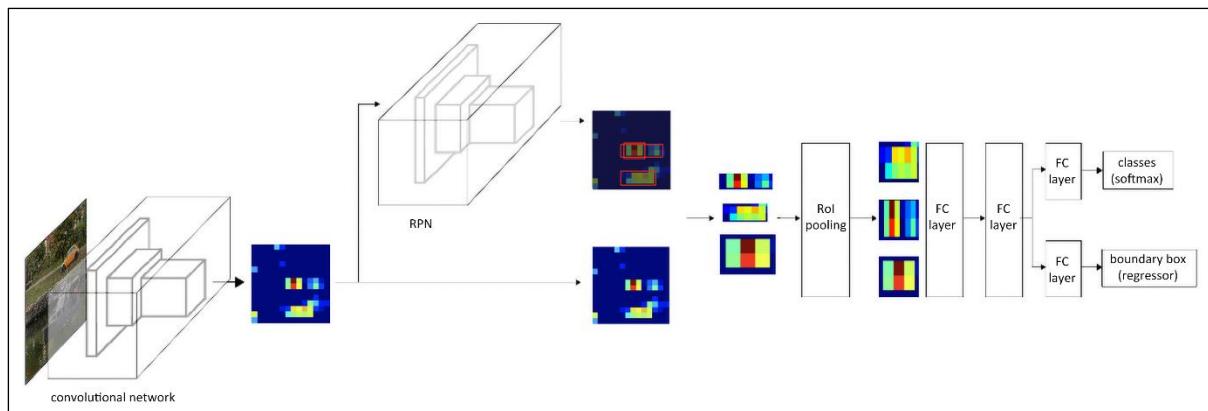
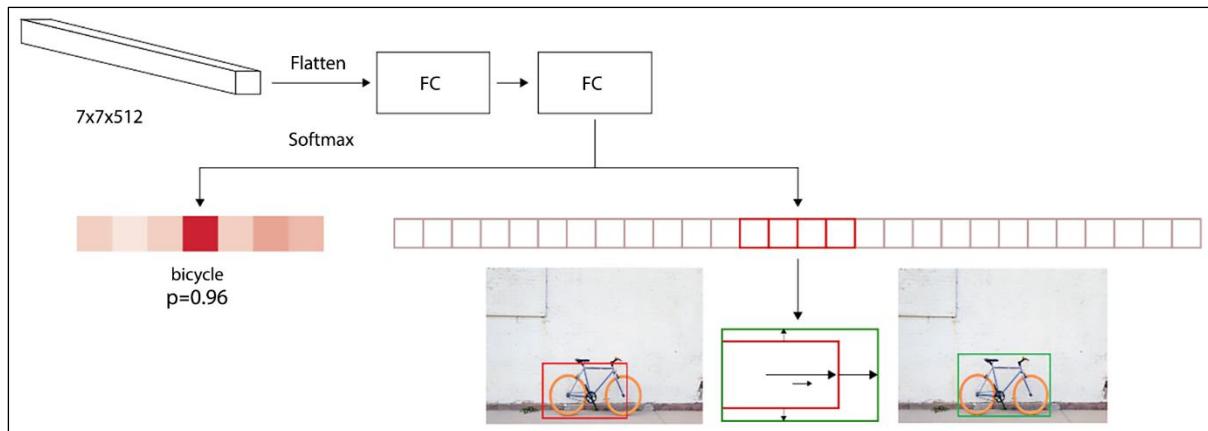


Anchor box 1



Anchor box 2





```
=====
Layer (type:depth-idx)                               Param #
=====
├── GeneralizedRCNNTransform: 1-1                  --
├── BackboneWithFPN: 1-2                           (26,799,296)
├── RegionProposalNetwork: 1-3                     593,935
└── ROIHeads: 1-4                                 13,905,930
=====
Total params: 41,299,161
Trainable params: 14,499,865
Non-trainable params: 26,799,296
=====
```

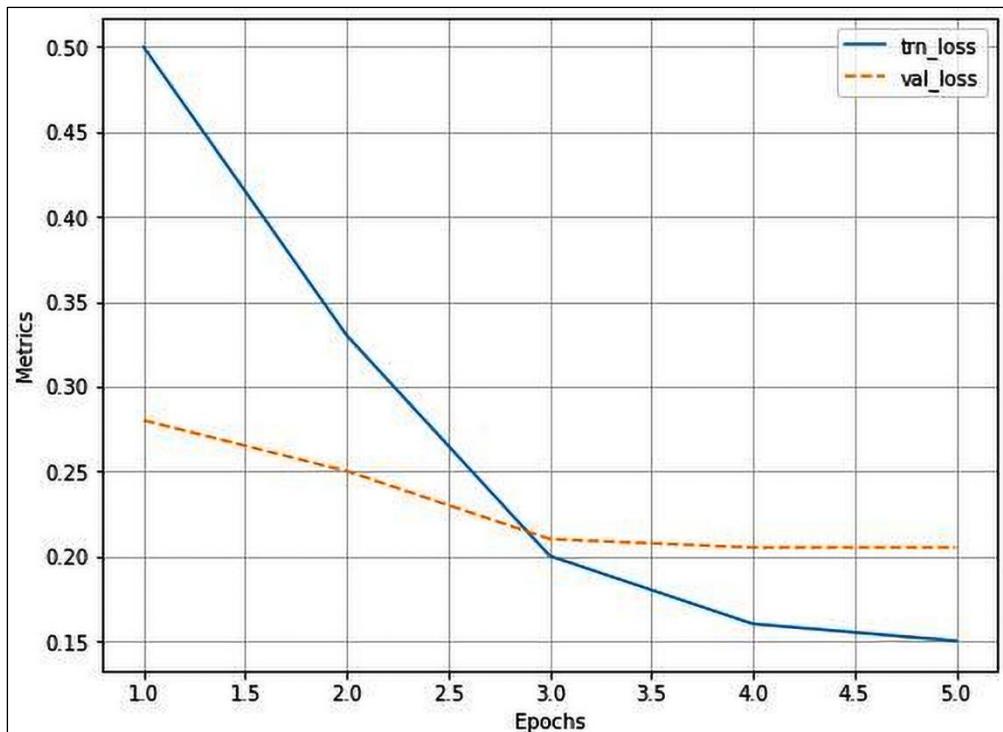
```
(transform): GeneralizedRCNNTransform(
    Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    Resize(min_size=(800,), max_size=1333, mode='bilinear')
)
```

```
(rpn): RegionProposalNetwork(
    (anchor_generator): AnchorGenerator()
    (head): RPNHead(
        (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (cls_logits): Conv2d(256, 3, kernel_size=(1, 1), stride=(1, 1))
        (bbox_pred): Conv2d(256, 12, kernel_size=(1, 1), stride=(1, 1))
    )
)
```

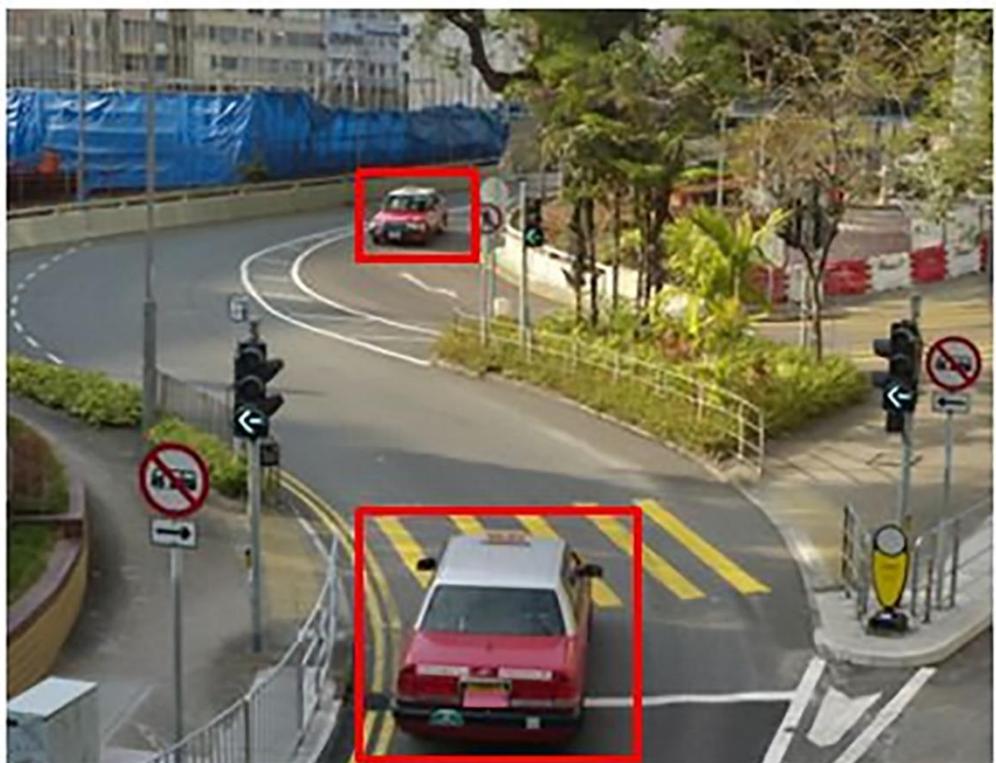
```

(roi_heads): RoIHeads(
    (box_roi_pool): MultiScaleRoIAlign()
    (box_head): TwoMLPHead(
        (fc6): Linear(in_features=12544, out_features=1024, bias=True)
        (fc7): Linear(in_features=1024, out_features=1024, bias=True)
    )
    (box_predictor): FastRCNNPredictor(
        (cls_score): Linear(in_features=1024, out_features=2, bias=True)
        (bbox_pred): Linear(in_features=1024, out_features=8, bias=True)
    )
)
)

```

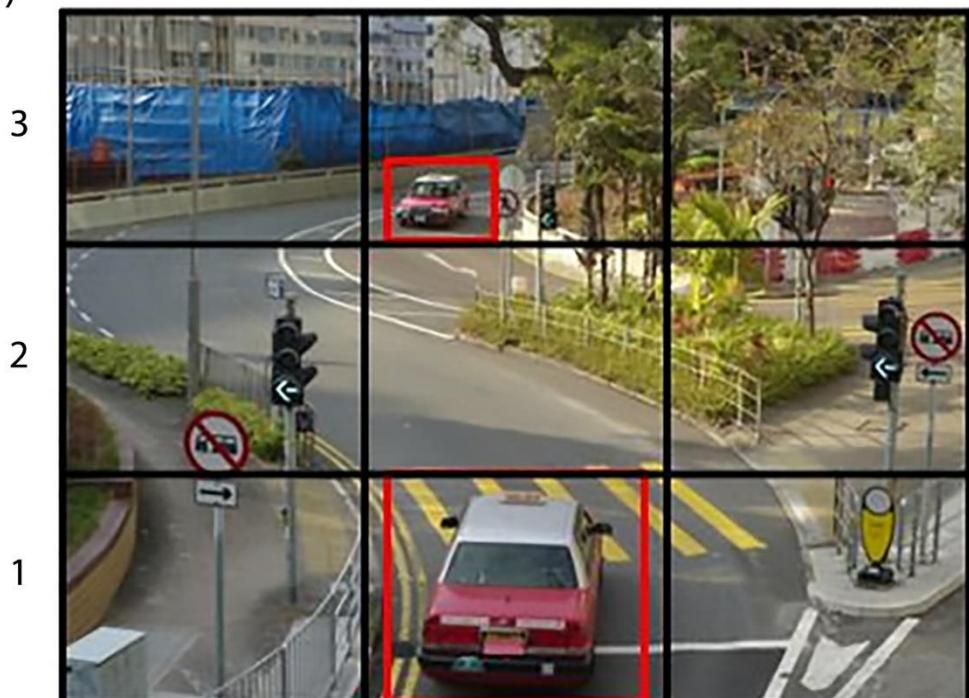


(0,0)



(100,100)

(0,0)



a

b

c

(100,100)

y=

pc

bx

by

bw

bh

c1

c2

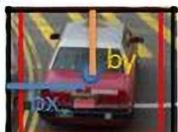
c3

(0,0)



(1,1)

(0,0)



(1,1)

0

?

?

?

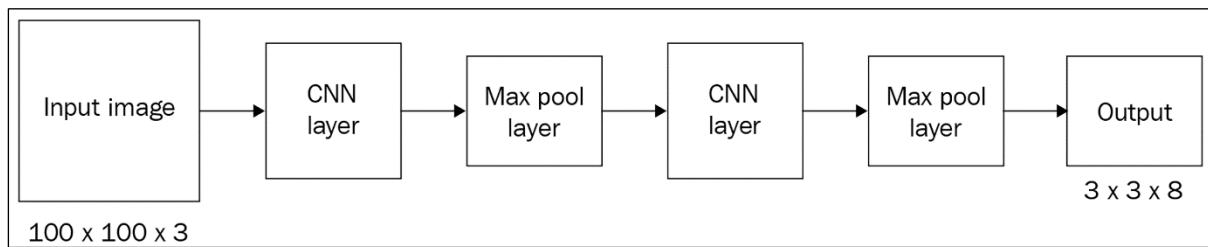
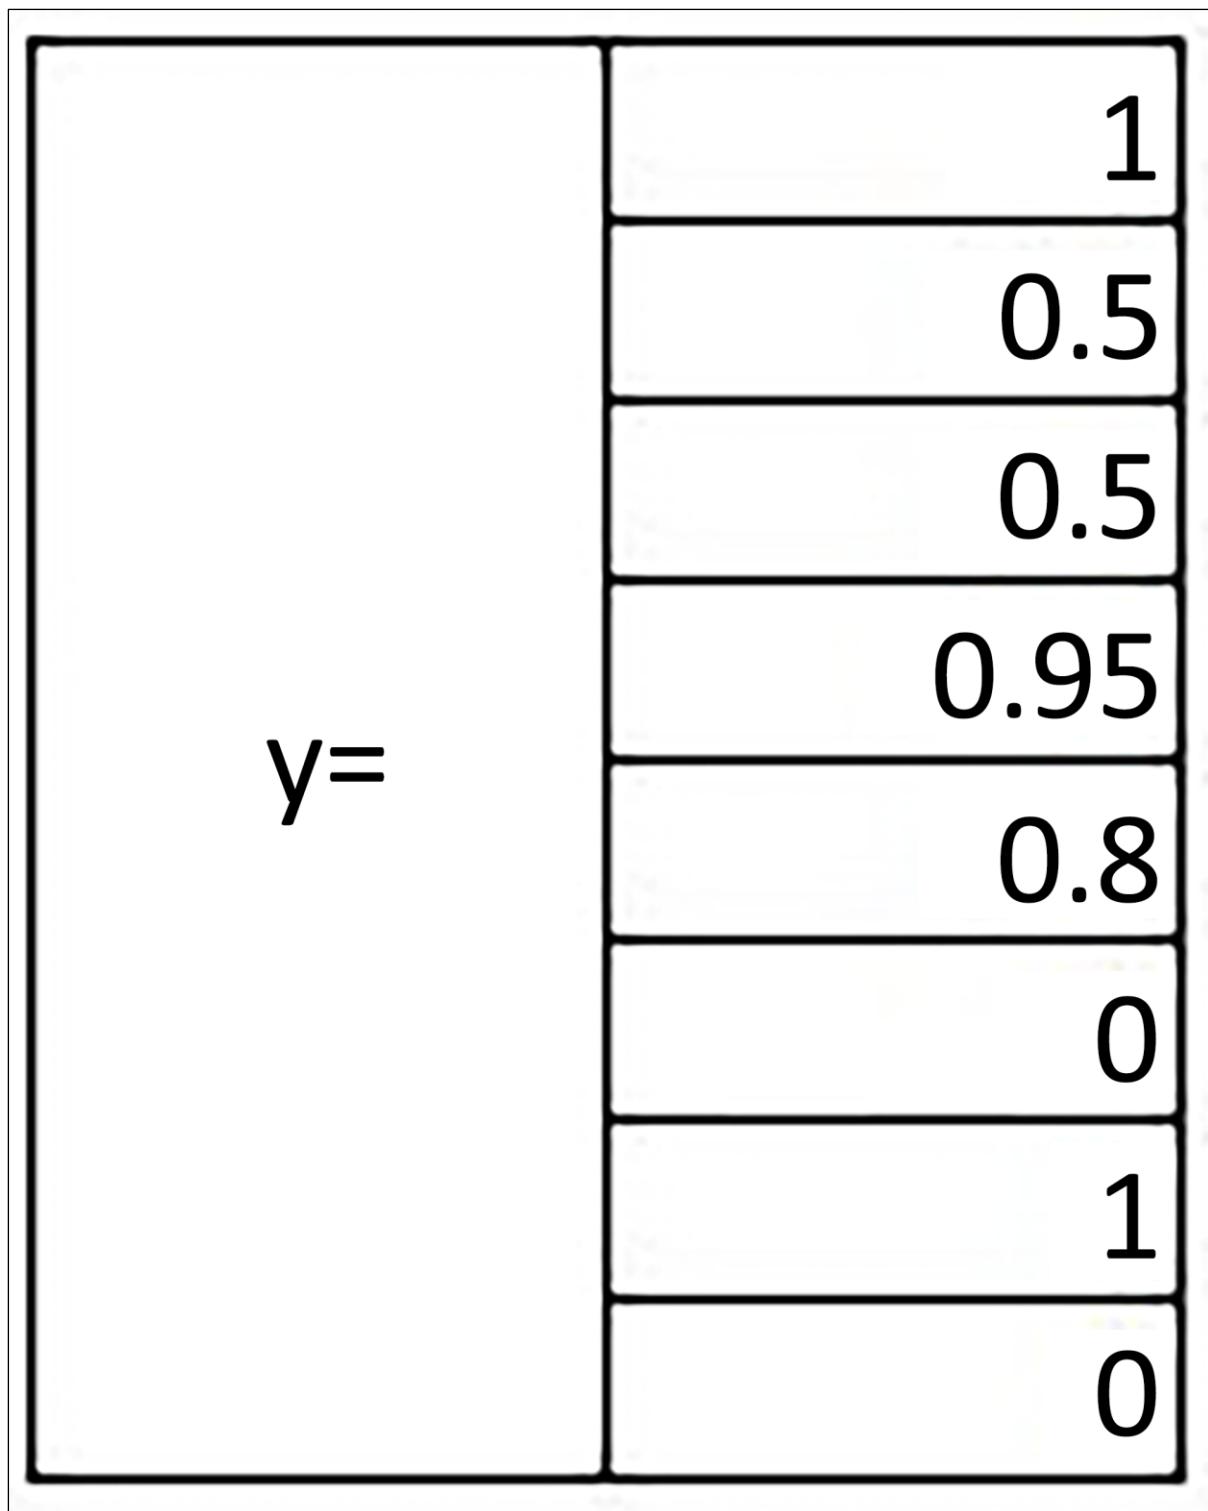
?

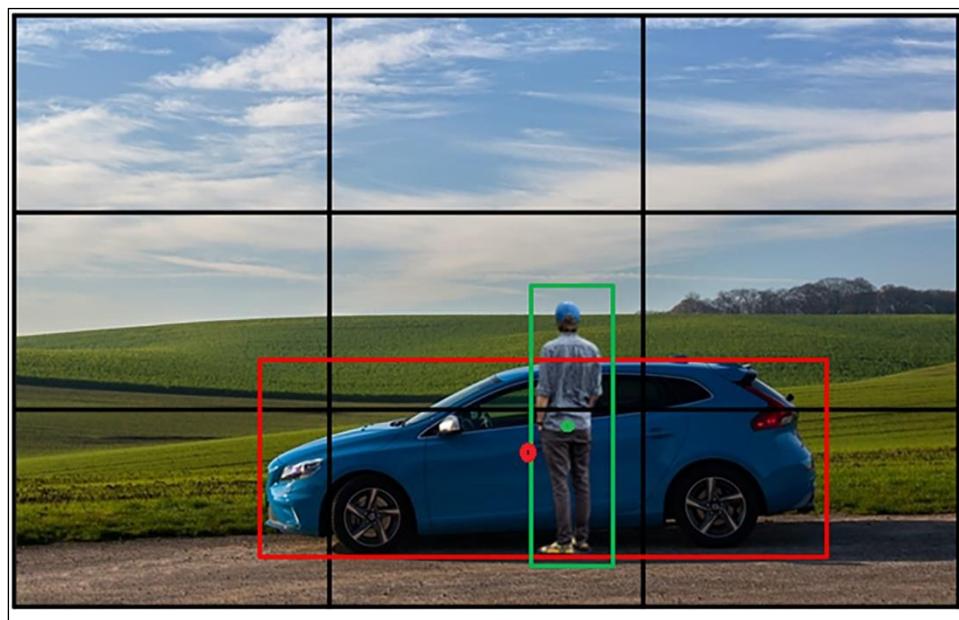
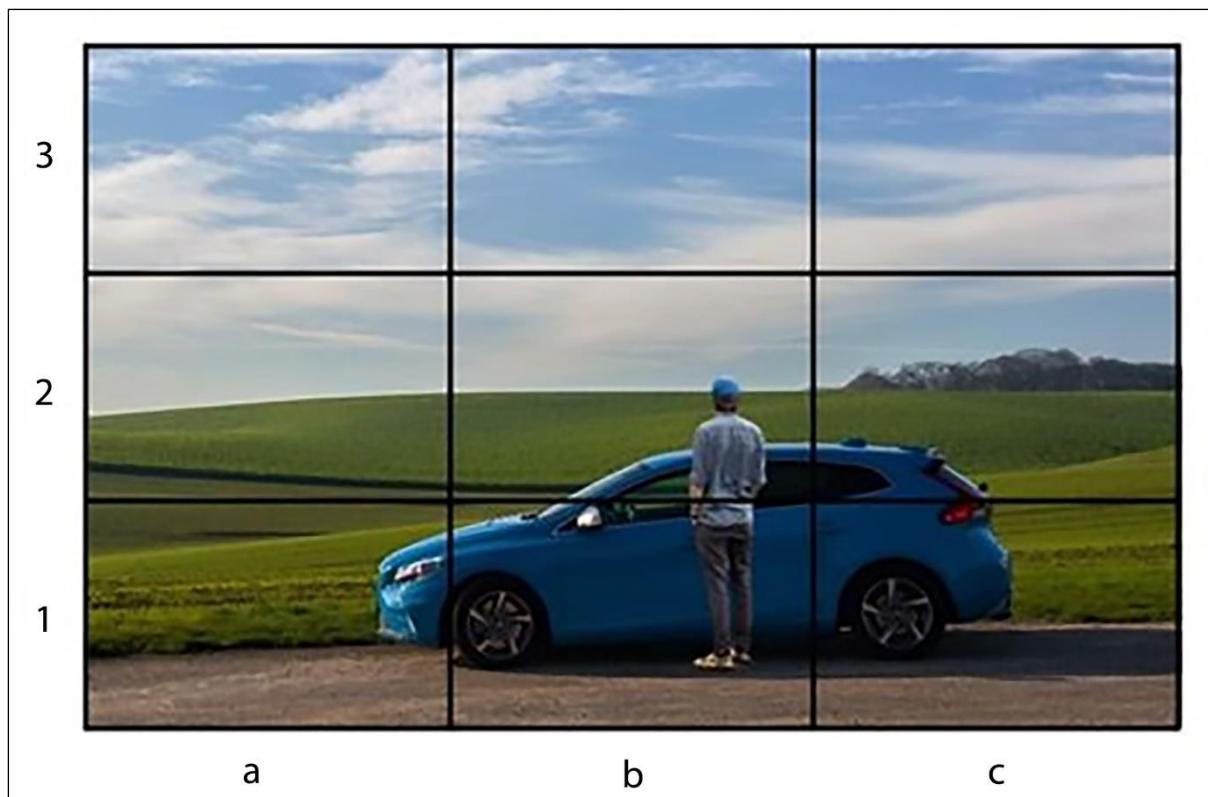
?

?

?

y=





y=

pc

bx

by

bh

bw

c1

c2

c3

pc

bx

by

bh

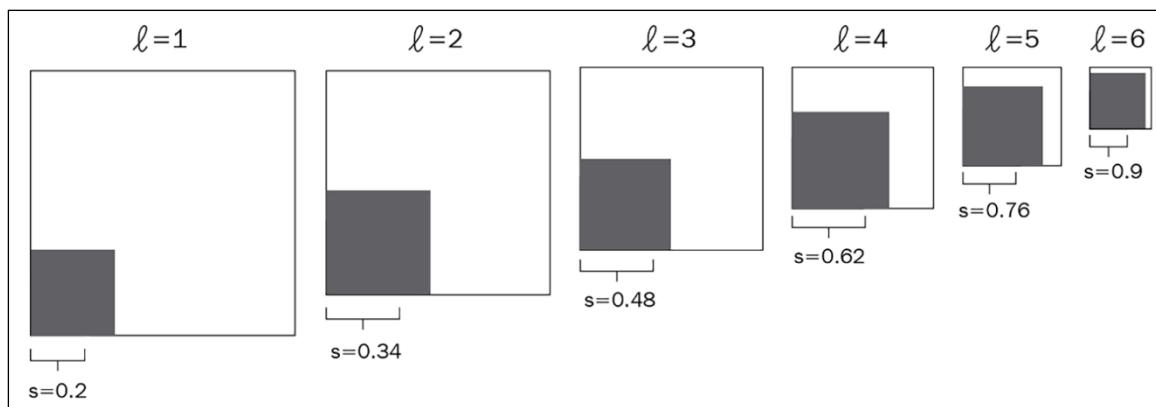
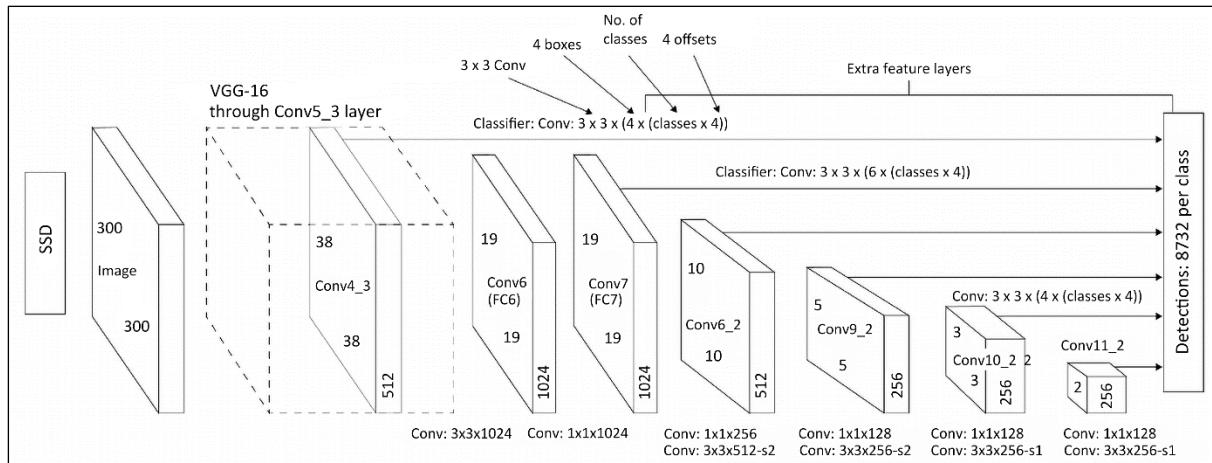
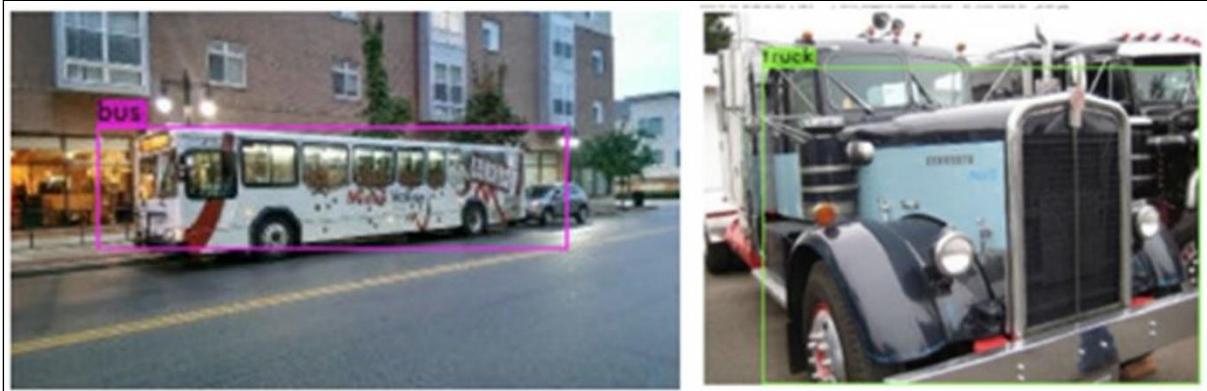
bw

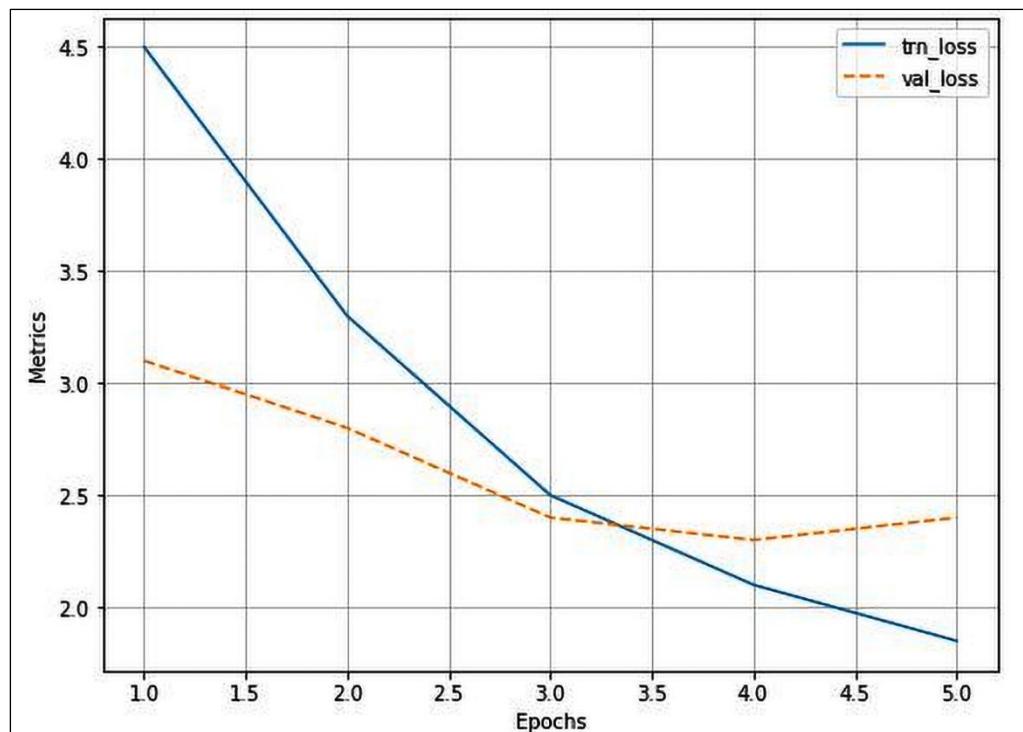
c1

c2

c3

data/person.jpg: Predicted in 54.532000 milli-seconds.
 dog: 99%
 person: 100%
 horse: 98%

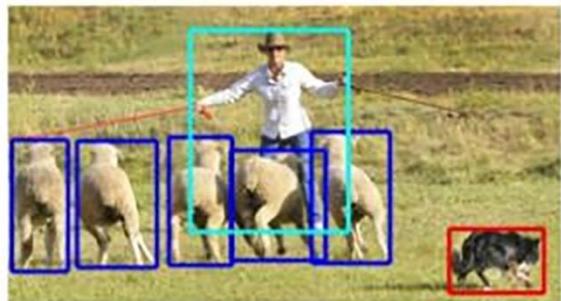




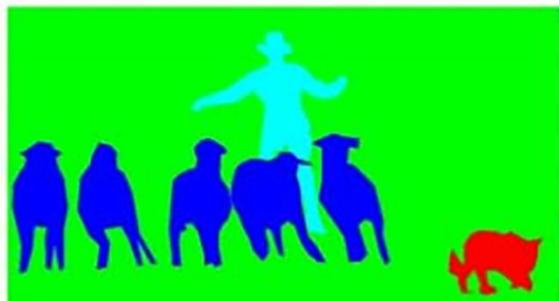
Chapter 9:



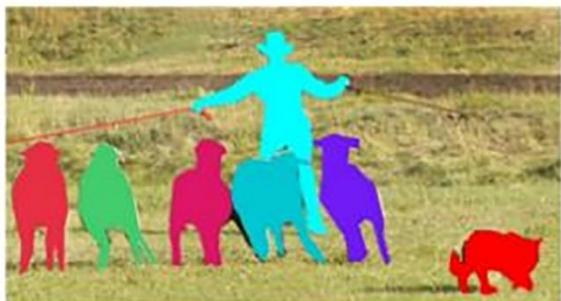
(a) Image classification



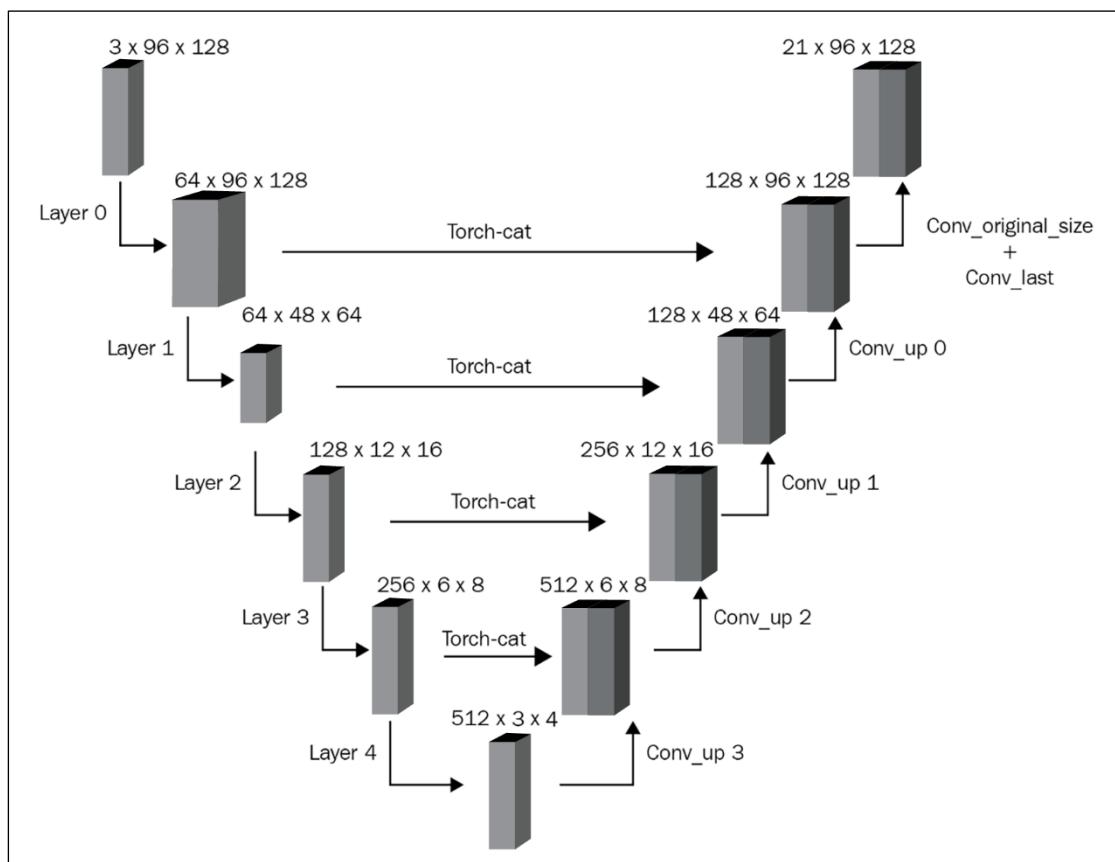
(b) Object localization



(c) Semantic segmentation



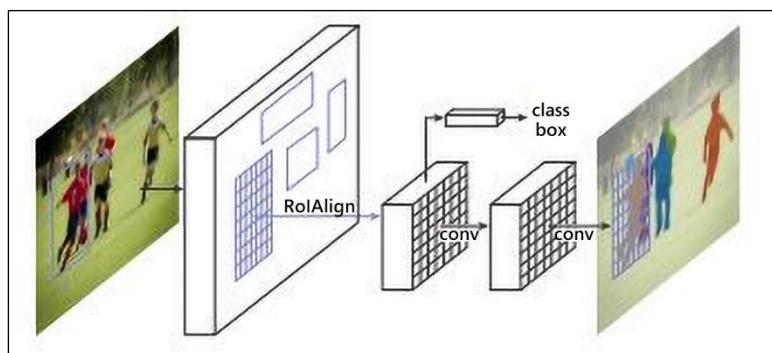
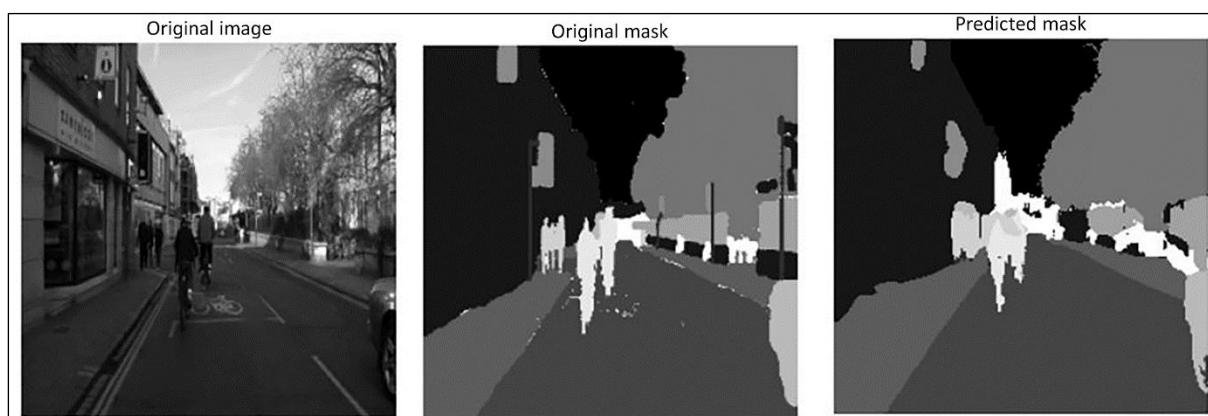
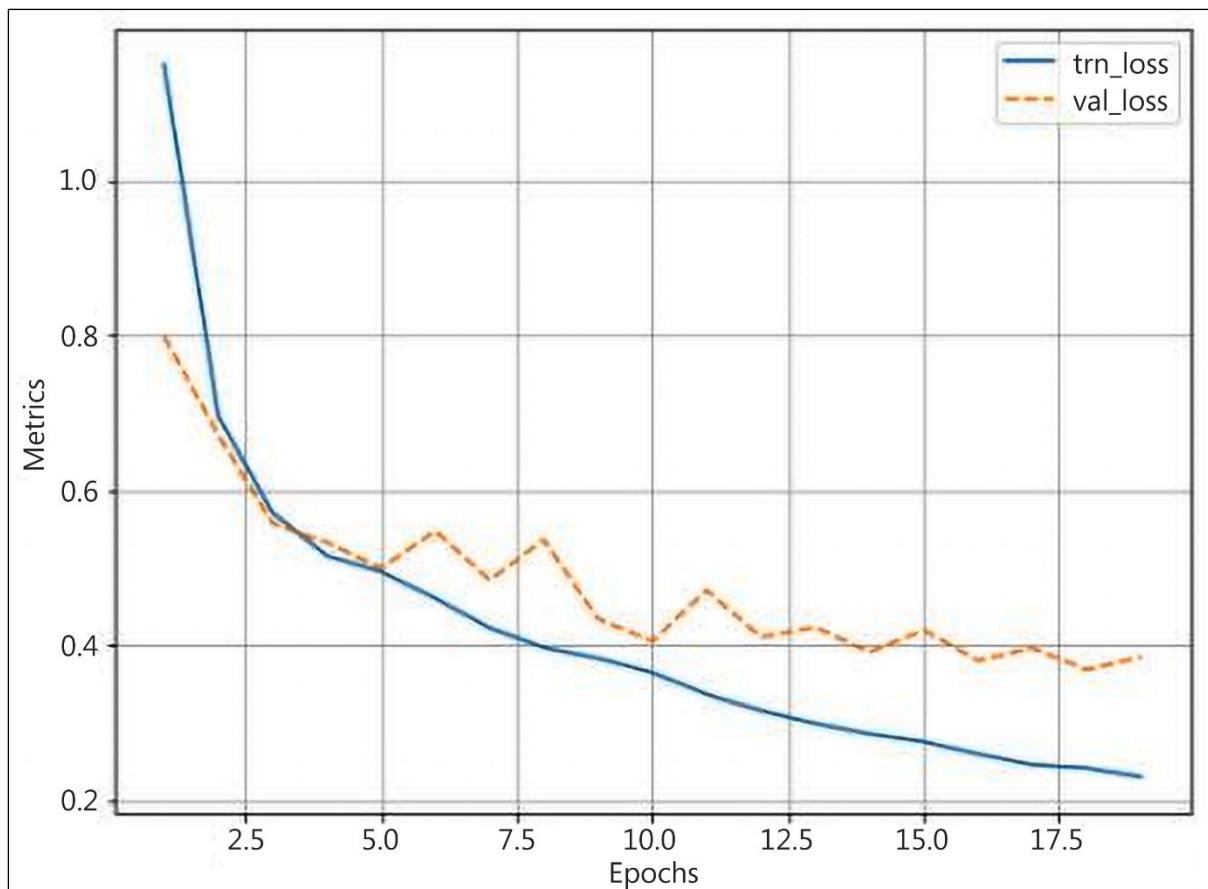
(d) Instance segmentation

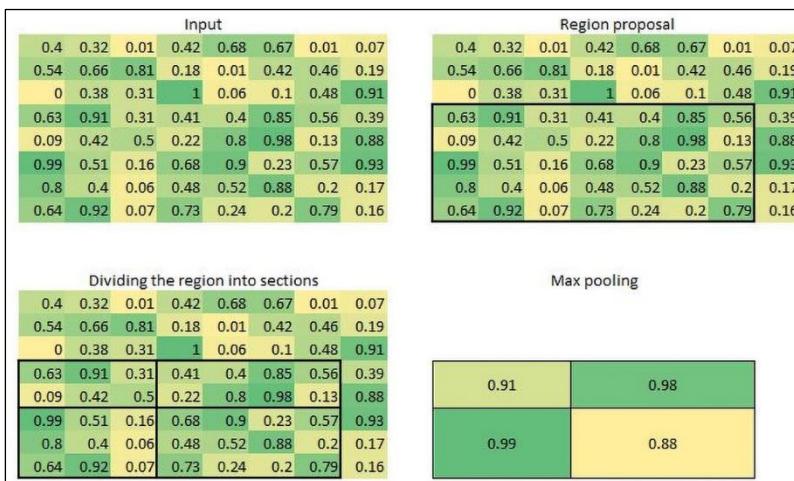
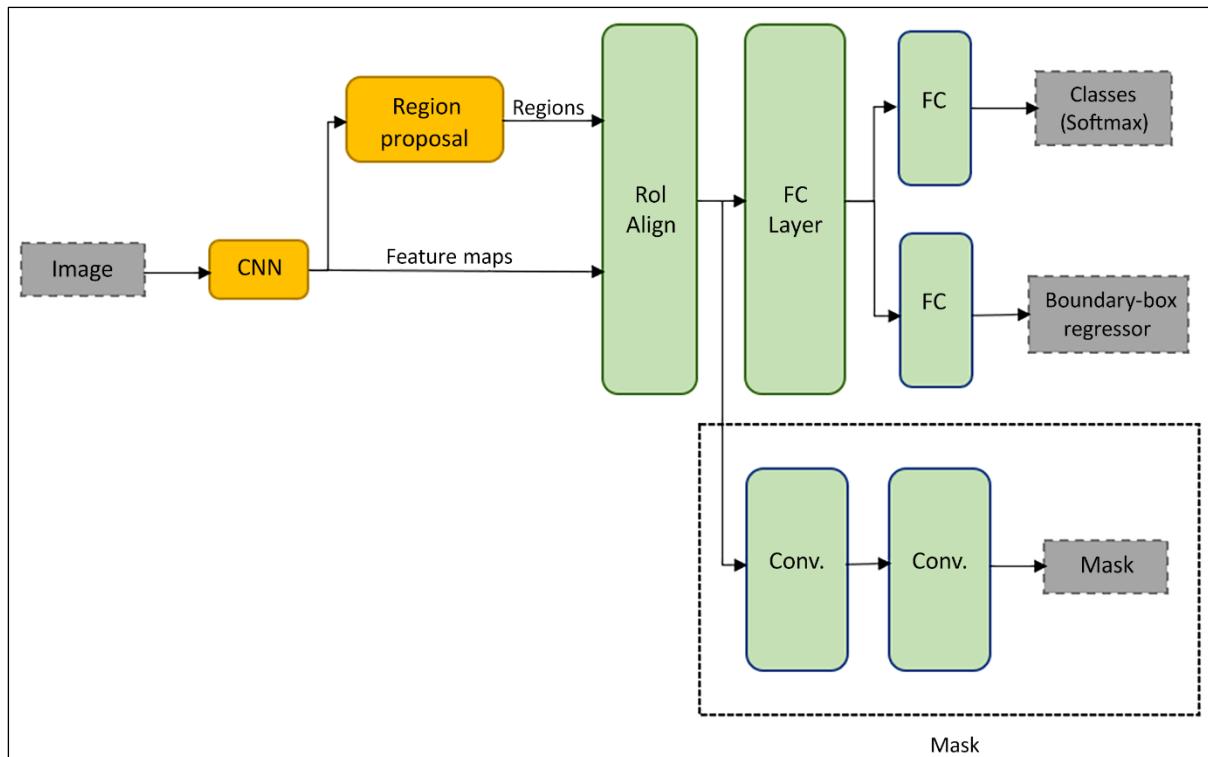


<u>Input array</u>	<u>Input array adjusted for stride</u>																																																																																																		
<table border="1"> <tr><td>1</td><td>1</td><td>1</td></tr> <tr><td>1</td><td>1</td><td>1</td></tr> <tr><td>1</td><td>1</td><td>1</td></tr> </table>	1	1	1	1	1	1	1	1	1	<table border="1"> <tr><td>1</td><td>0</td><td>1</td><td>0</td><td>1</td></tr> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> <tr><td>1</td><td>0</td><td>1</td><td>0</td><td>1</td></tr> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> <tr><td>1</td><td>0</td><td>1</td><td>0</td><td>1</td></tr> </table>	1	0	1	0	1	0	0	0	0	0	1	0	1	0	1	0	0	0	0	0	1	0	1	0	1																																																																
1	1	1																																																																																																	
1	1	1																																																																																																	
1	1	1																																																																																																	
1	0	1	0	1																																																																																															
0	0	0	0	0																																																																																															
1	0	1	0	1																																																																																															
0	0	0	0	0																																																																																															
1	0	1	0	1																																																																																															
<u>Input array adjusted for stride and padding</u>																																																																																																			
<table border="1"> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> </table>	0	0	0	0	0	0	0	0	1	0	1	0	1	0	0	0	0	0	0	0	0	0	1	0	1	0	1	0	0	0	0	0	0	0	0	0	1	0	1	0	1	0	0	0	0	0	0	0	0	<table border="1"> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr> </table>	0	0	0	0	0	0	0	0	1	0	1	0	1	0	0	0	0	0	0	0	0	0	1	0	1	0	1	0	0	0	0	0	0	0	0	0	1	0	1	0	1	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0																																																																																													
0	1	0	1	0	1	0																																																																																													
0	0	0	0	0	0	0																																																																																													
0	1	0	1	0	1	0																																																																																													
0	0	0	0	0	0	0																																																																																													
0	1	0	1	0	1	0																																																																																													
0	0	0	0	0	0	0																																																																																													
0	0	0	0	0	0	0																																																																																													
0	1	0	1	0	1	0																																																																																													
0	0	0	0	0	0	0																																																																																													
0	1	0	1	0	1	0																																																																																													
0	0	0	0	0	0	0																																																																																													
0	1	0	1	0	1	0																																																																																													
0	0	0	0	0	0	0																																																																																													
<u>Filter/kernel</u>	<u>Output array</u>																																																																																																		
<table border="1"> <tr><td>1</td><td>1</td></tr> <tr><td>1</td><td>1</td></tr> </table>	1	1	1	1	<table border="1"> <tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr> <tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr> <tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr> <tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr> <tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr> <tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr> </table>	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1																																																										
1	1																																																																																																		
1	1																																																																																																		
1	1	1	1	1	1																																																																																														
1	1	1	1	1	1																																																																																														
1	1	1	1	1	1																																																																																														
1	1	1	1	1	1																																																																																														
1	1	1	1	1	1																																																																																														
1	1	1	1	1	1																																																																																														

```
help(nn.ConvTranspose2d)
| Args:
|     in_channels (int): Number of channels in the input image
|     out_channels (int): Number of channels produced by the convolution
|     kernel_size (int or tuple): Size of the convolving kernel
|     stride (int or tuple, optional): Stride of the convolution. Default: 1
|     padding (int or tuple, optional): ``dilation * (kernel_size - 1) - padding`` zero-padding
|         will be added to both sides of each dimension in the input. Default: 0
|     output_padding (int or tuple, optional): Additional size added to one side
|         of each dimension in the output shape. Default: 0
|     groups (int, optional): Number of blocked connections from input channels to output channels. Default: 1
|     bias (bool, optional): If ``True``, adds a learnable bias to the output. Default: ``True``
|     dilation (int or tuple, optional): Spacing between kernel elements. Default: 1
```



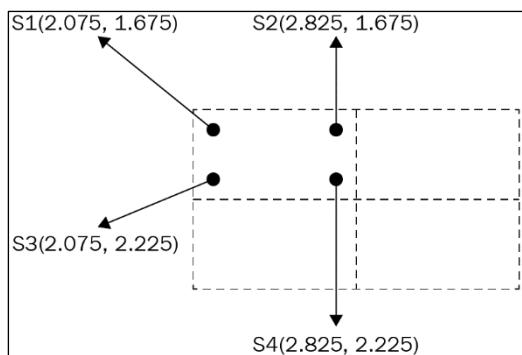




0.4	0.08	0.73	0.57	0.13
0.88	0.13	0.32	0.64	0.15
0.98	0.66	0.16	0.16	0.25
0.97	0.45	0.08	0.08	0.18
0.69	0.88	0.9	0.9	0.87

A (1.7, 1.4)		B (4.7, 1.4)		
0.4	0.08	0.73	0.57	0.13
0.88	0.13	0.32	0.64	0.15
0.98	0.66	0.16	0.16	0.25
0.97	0.45	0.08	0.08	0.18
0.69	0.88	0.9	0.9	0.87

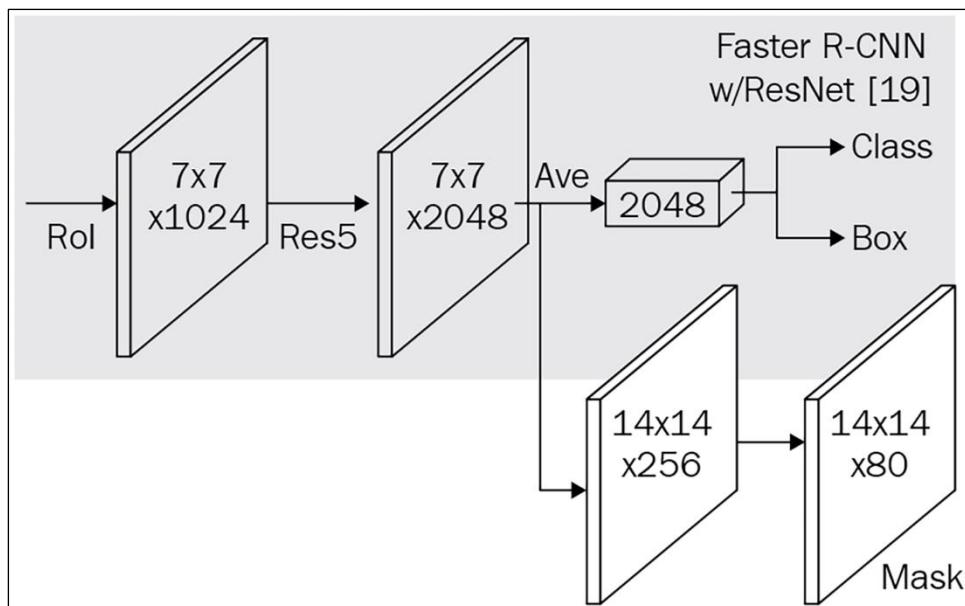
C (1.7, 3.6) D (4.7, 3.6)



0.4	0.08	0.73	0.57	0.13
0.88	0.13	P (0.32)	Q (0.64)	0.15
0.98	0.66	R (0.16)	S (0.12)	0.25
0.97	0.45	0.08	0.08	0.18
0.69	0.88	0.9	0.9	0.87

0.21778	0.27553	
0.14896	0.21852	

0.2152	0.2335
0.3763	0.3562



```
=====
Tensor Shape: torch.Size([3, 512, 683])      Min: 0.000      Max: 1.000      Mean: 0.555      dtype: torch.float32
=====
Dict Of 6 items
=====
BOXES:
Tensor Shape: torch.Size([2, 4])      Min: 233.000      Max: 480.000      Mean: 362.750      dtype: torch.float32
=====
LABELS:
Tensor Shape: torch.Size([2])      Min: 1.000      Max: 1.000      Mean: 1.000      dtype: torch.int64
=====
MASKS:
Tensor Shape: torch.Size([2, 512, 683])      Min: 0.000      Max: 1.000      Mean: 0.002      dtype: torch.uint8
=====
IMAGE_ID:
Tensor Shape: torch.Size([1])      Min: 0.000      Max: 0.000      Mean: 0.000      dtype: torch.int64
=====
AREA:
Tensor Shape: torch.Size([2])      Min: 864.000      Max: 935.000      Mean: 899.500      dtype: torch.float32
=====
... ... 1 more items


```

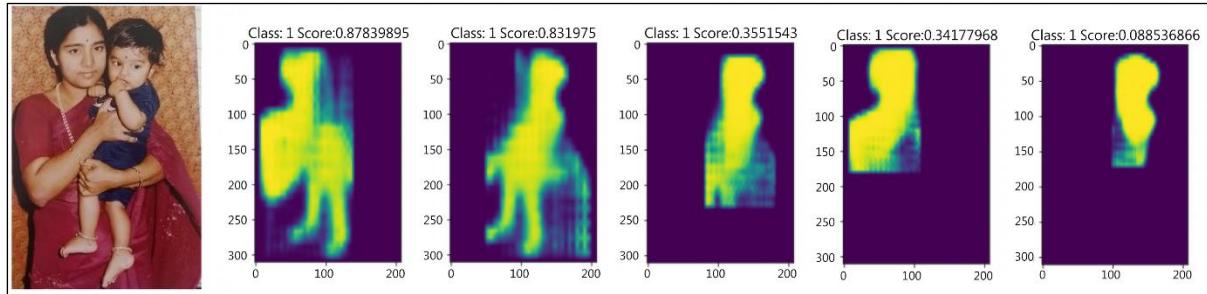
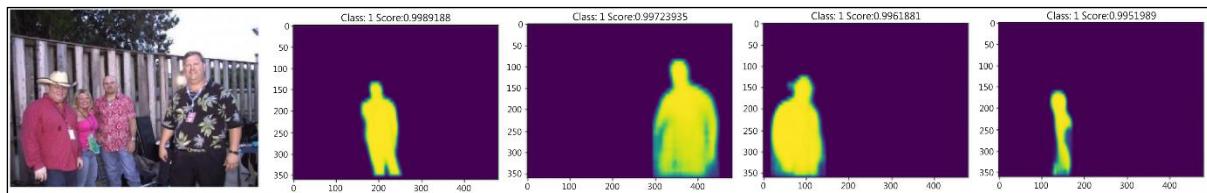
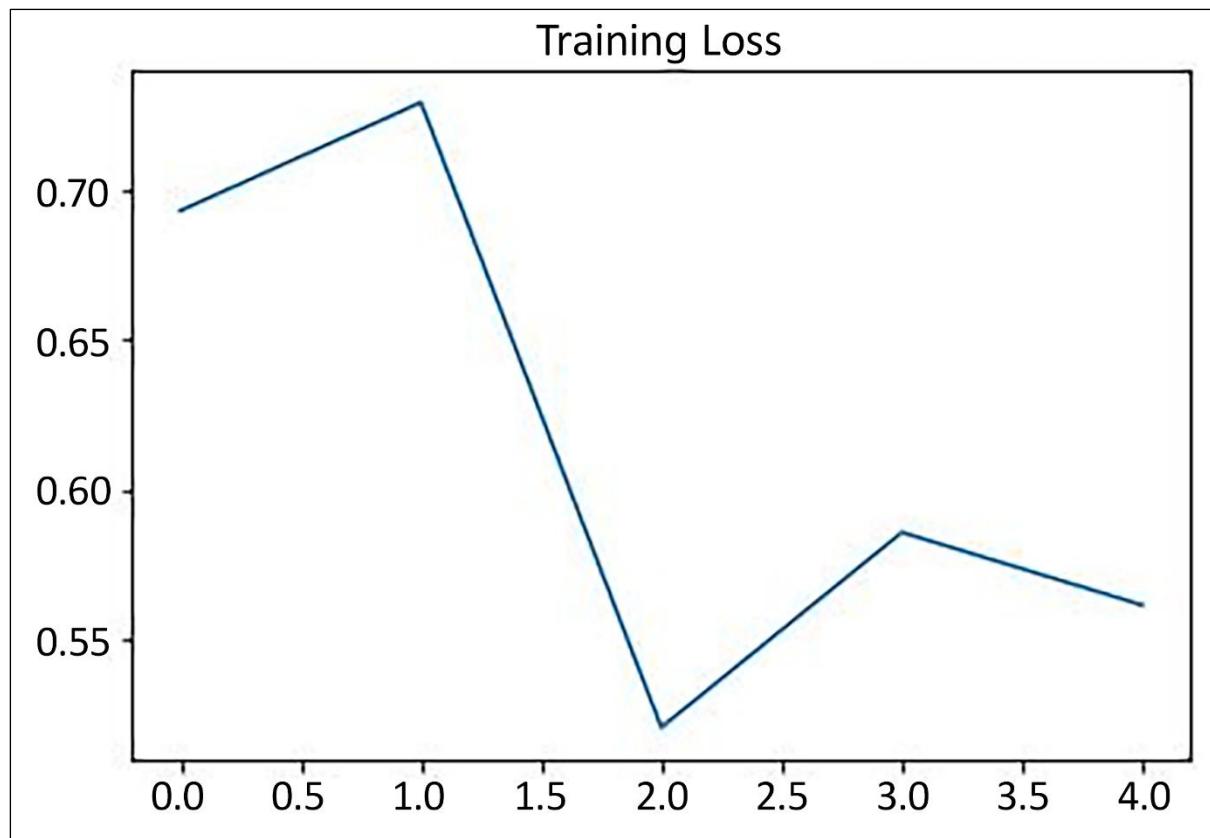
```

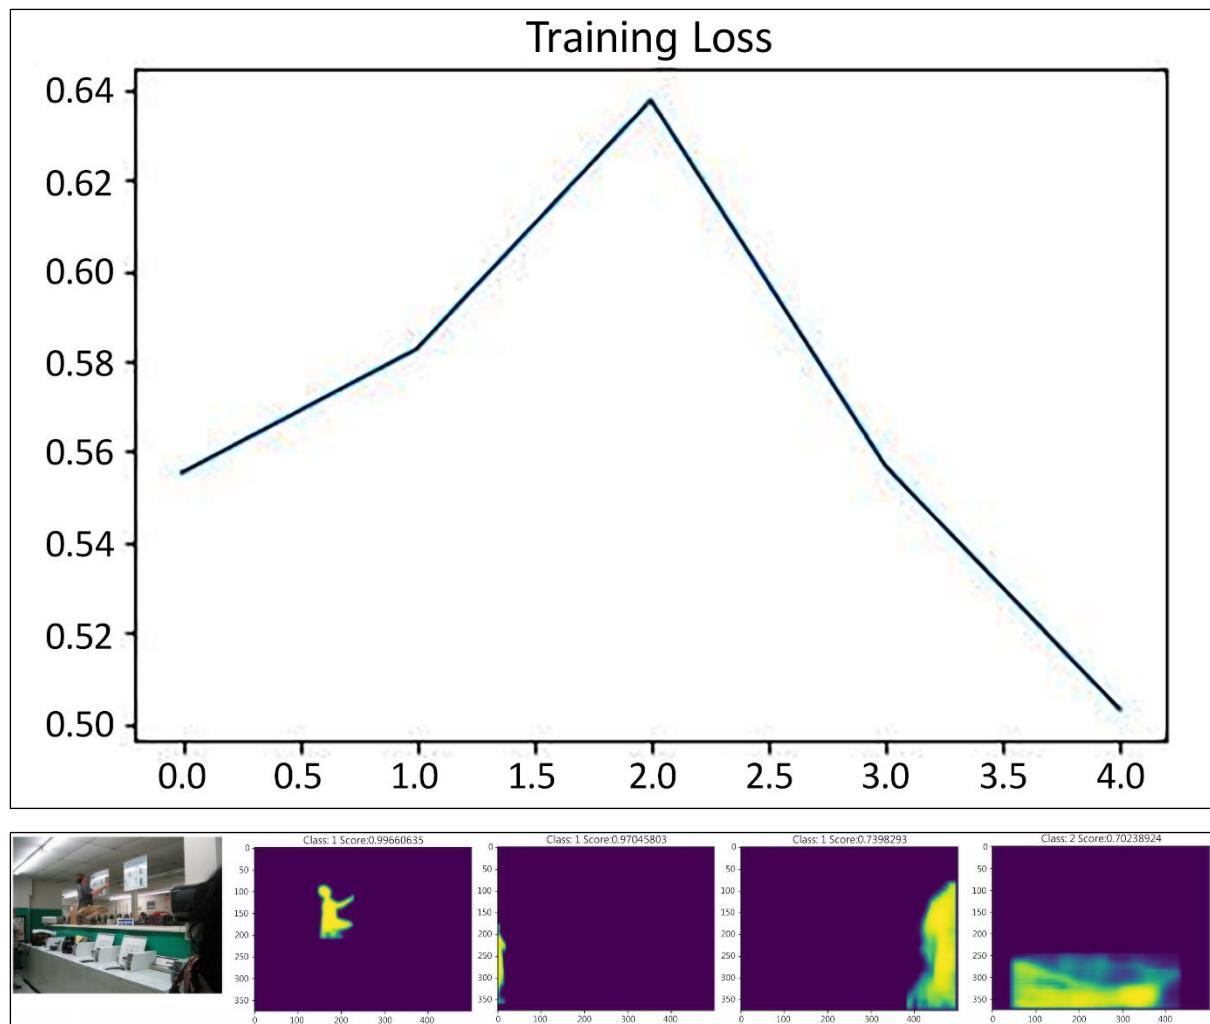
(roi_heads): RoIHeads(
  (box_roi_pool): MultiScaleRoIAlign()
  (box_head): TwoMLPHead(
    (fc6): Linear(in_features=12544, out_features=1024, bias=True)
    (fc7): Linear(in_features=1024, out_features=1024, bias=True)
  )
  (box_predictor): FastRCNNPredictor(
    (cls_score): Linear(in_features=1024, out_features=2, bias=True)
    (bbox_pred): Linear(in_features=1024, out_features=8, bias=True)
  )
  (mask_roi_pool): MultiScaleRoIAlign()
  (mask_head): MaskRCNNHeads(
    (mask_fcn1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (relu1): ReLU(inplace=True)
    (mask_fcn2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (relu2): ReLU(inplace=True)
    (mask_fcn3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (relu3): ReLU(inplace=True)
    (mask_fcn4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (relu4): ReLU(inplace=True)
  )
  (mask_predictor): MaskRCNNPredictor(
    (conv5_mask): ConvTranspose2d(256, 256, kernel_size=(2, 2), stride=(2, 2))
    (relu): ReLU(inplace=True)
    (mask_fcn_logits): Conv2d(256, 2, kernel_size=(1, 1), stride=(1, 1))
  )
)

```

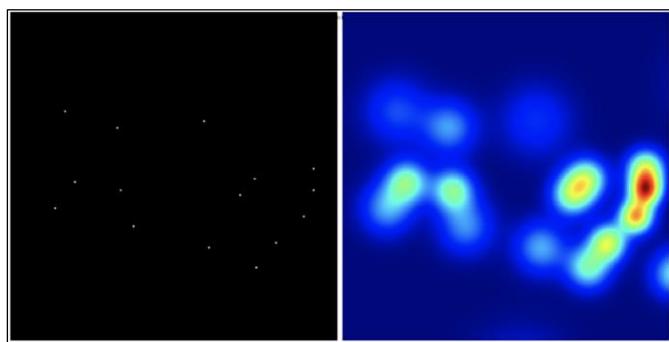
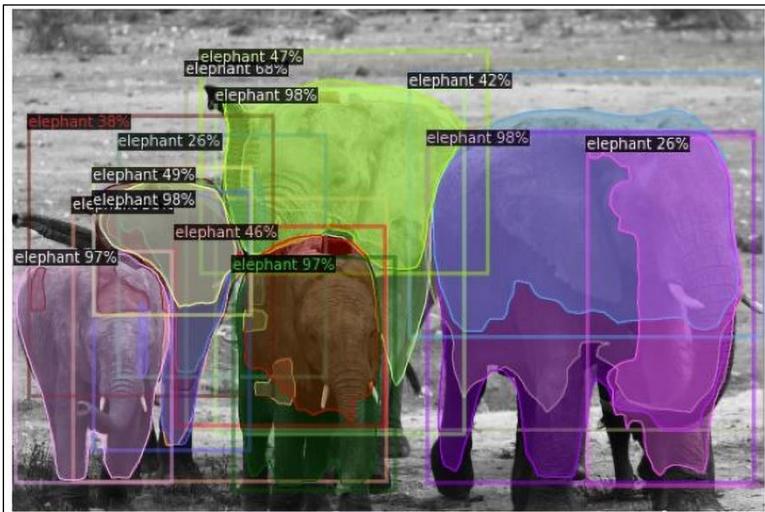
```

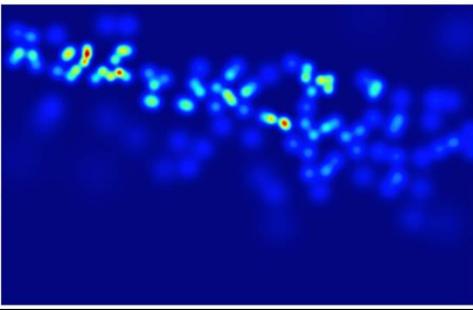
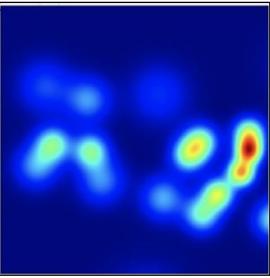
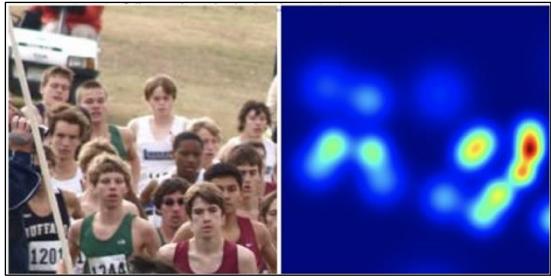
Dict Of 4 items
=====
BOXES:
Tensor Shape: torch.Size([100, 4])      Min: 0.000      Max: 1024.000      Mean: 385.767      dtype: torch.float32
=====
LABELS:
Tensor Shape: torch.Size([100])      Min: 1.000      Max: 1.000      Mean: 1.000      dtype: torch.int64
=====
SCORES:
Tensor Shape: torch.Size([100])      Min: 0.491      Max: 0.648      Mean: 0.531      dtype: torch.float32
=====
MASKS:
Tensor Shape: torch.Size([100, 1, 692, 1024])  Min: 0.000      Max: 1.000      Mean: 0.012      dtype: torch.float32
=====
```



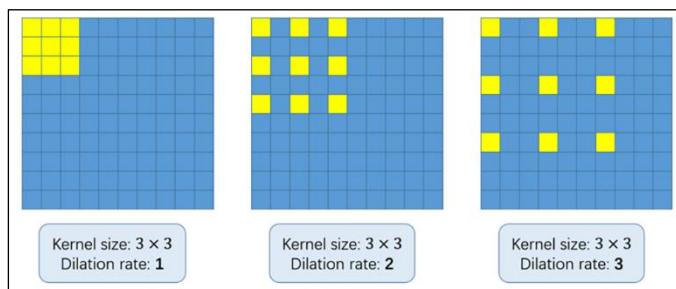


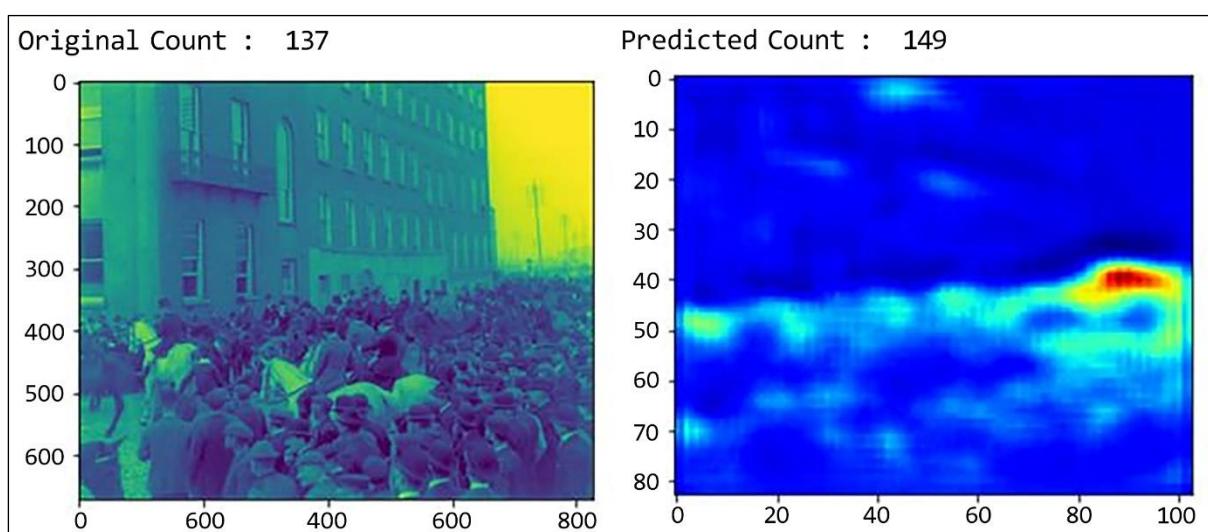
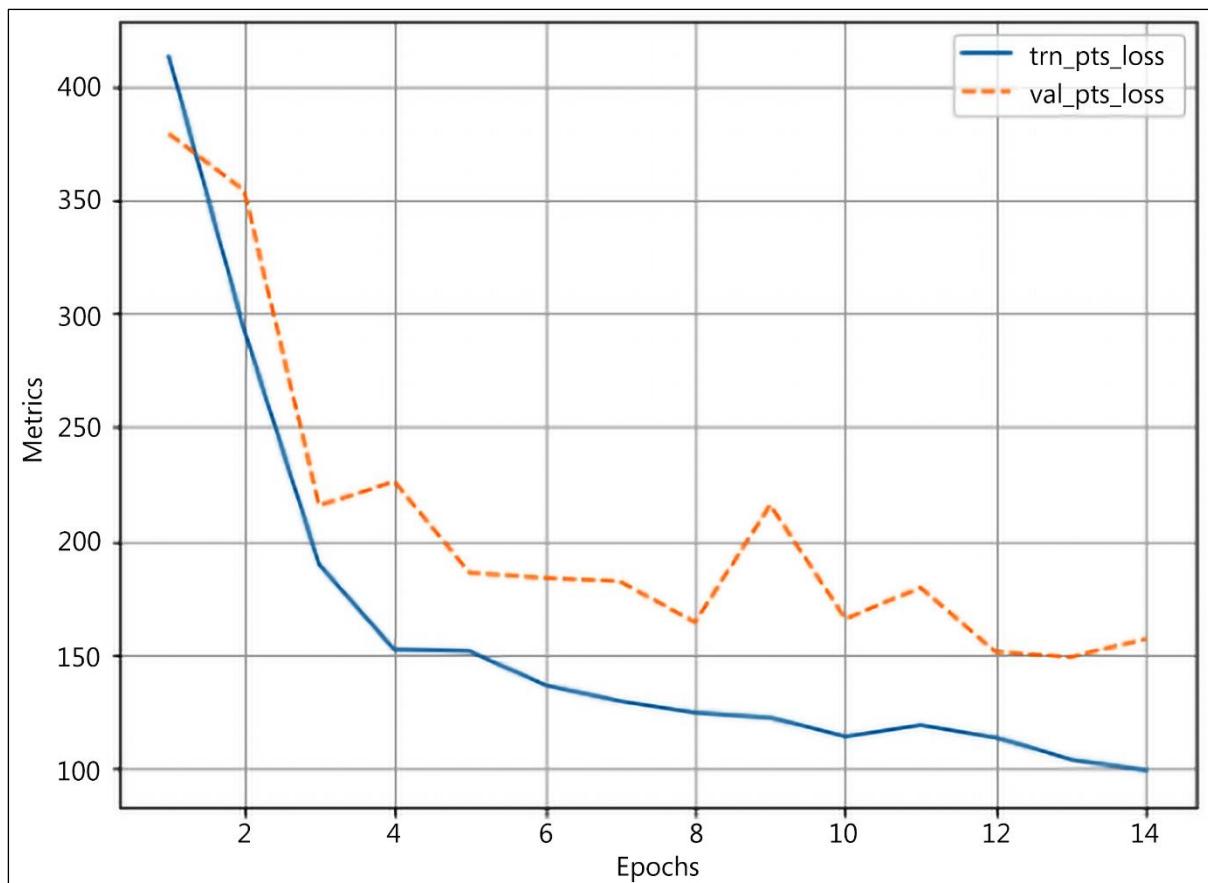
Chapter 10:



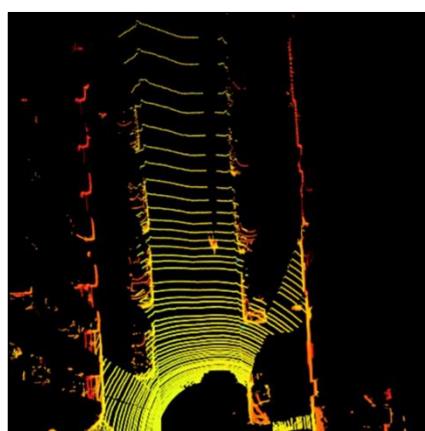
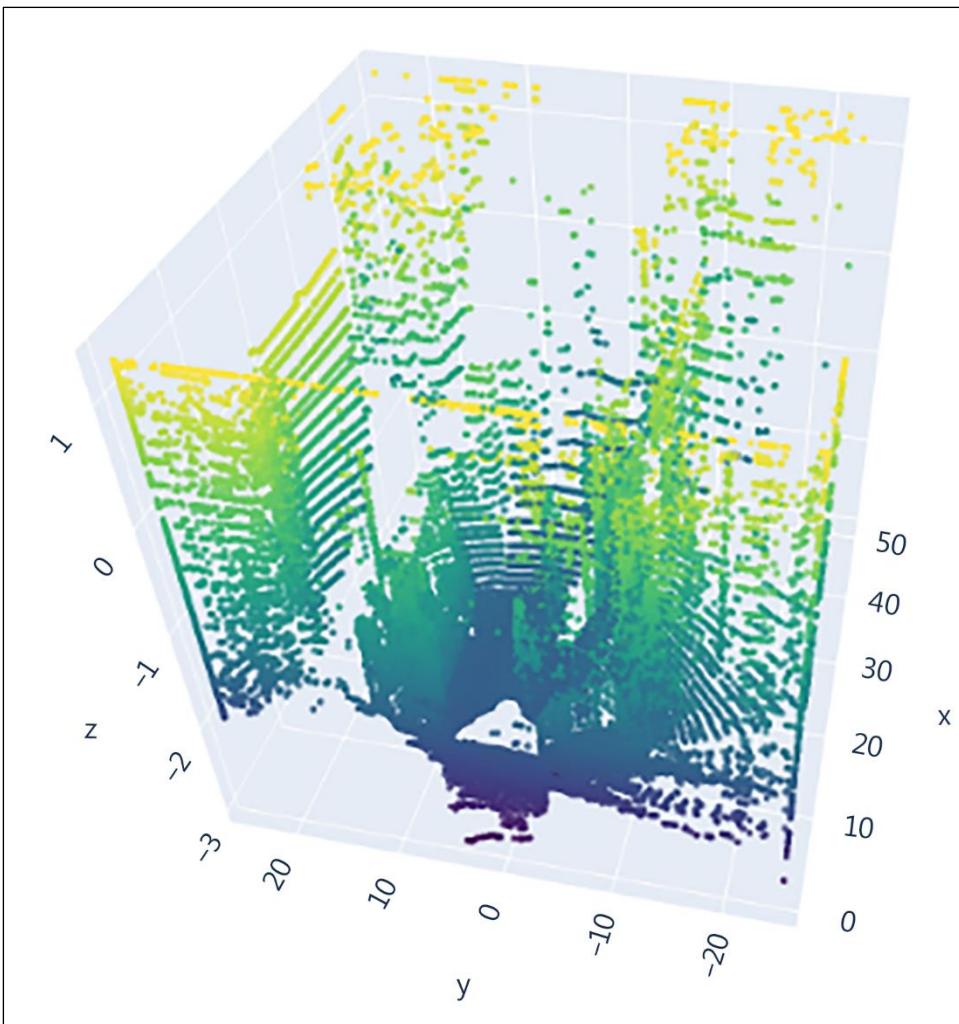


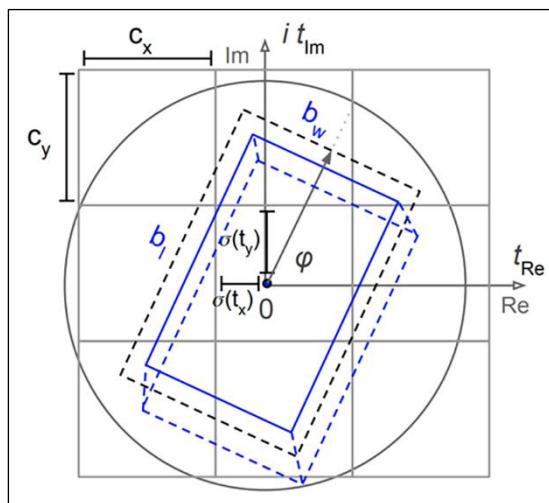
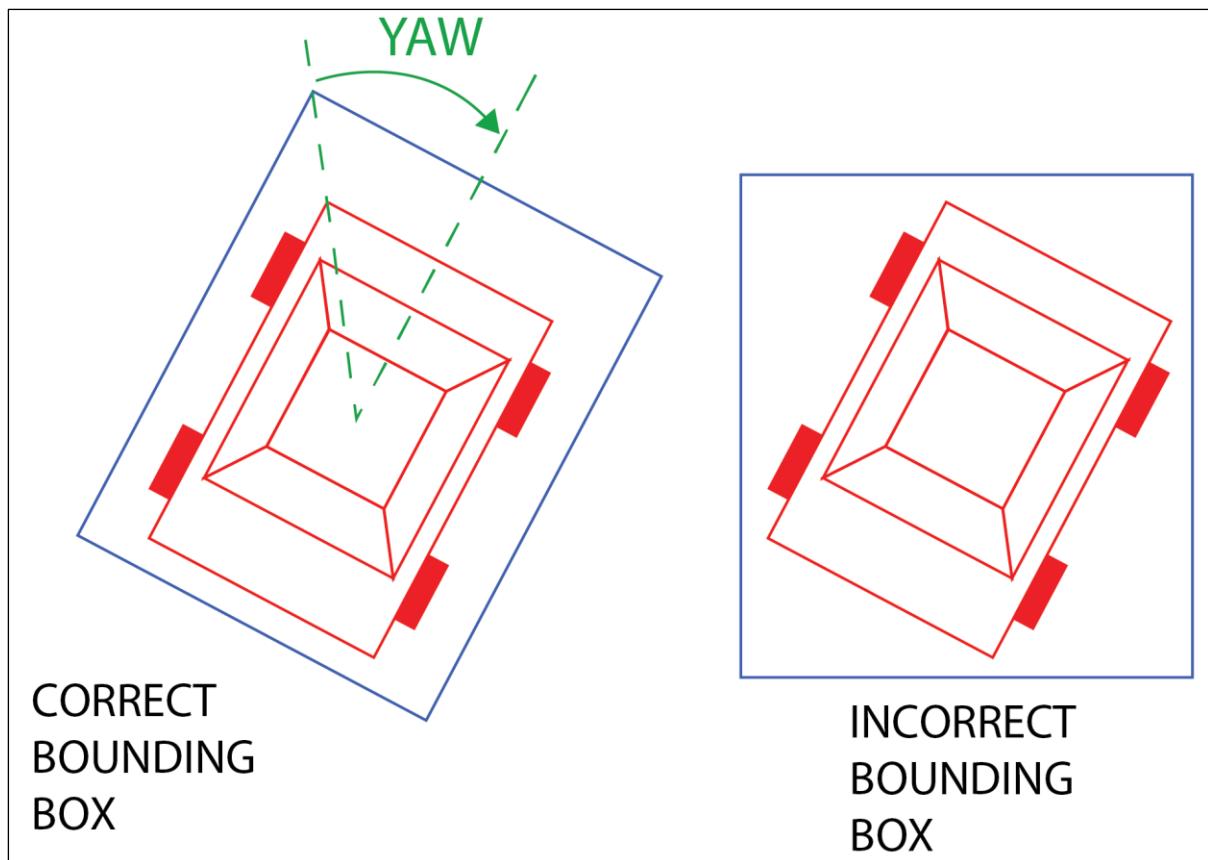
Configurations of CSRNet			
A	B	C	D
input(unfixed-resolution color image)			
front-end (fine-tuned from VGG-16)			
conv3-64-1			
conv3-64-1			
max-pooling			
conv3-128-1			
conv3-128-1			
max-pooling			
conv3-256-1			
conv3-256-1			
conv3-256-1			
max-pooling			
conv3-512-1			
conv3-512-1			
conv3-512-1			
back-end (four different configurations)			
conv3-512-1	conv3-512-2	conv3-512-2	conv3-512-4
conv3-512-1	conv3-512-2	conv3-512-2	conv3-512-4
conv3-512-1	conv3-512-2	conv3-512-2	conv3-512-4
conv3-256-1	conv3-256-2	conv3-256-4	conv3-256-4
conv3-128-1	conv3-128-2	conv3-128-4	conv3-128-4
conv3-64-1	conv3-64-2	conv3-64-4	conv3-64-4
conv1-1-1			





```
array([[62.502,  8.628,  2.343,  0.     ],
       [62.468,  8.824,  2.342,  0.     ],
       [66.793, 10.832,  2.497,  0.     ],
       ...,
       [ 3.75 , -1.418, -1.753,  0.25 ],
       [ 3.759, -1.409, -1.756,  0.32 ],
       [ 3.767, -1.398, -1.758,  0.     ]], dtype=float32)
```



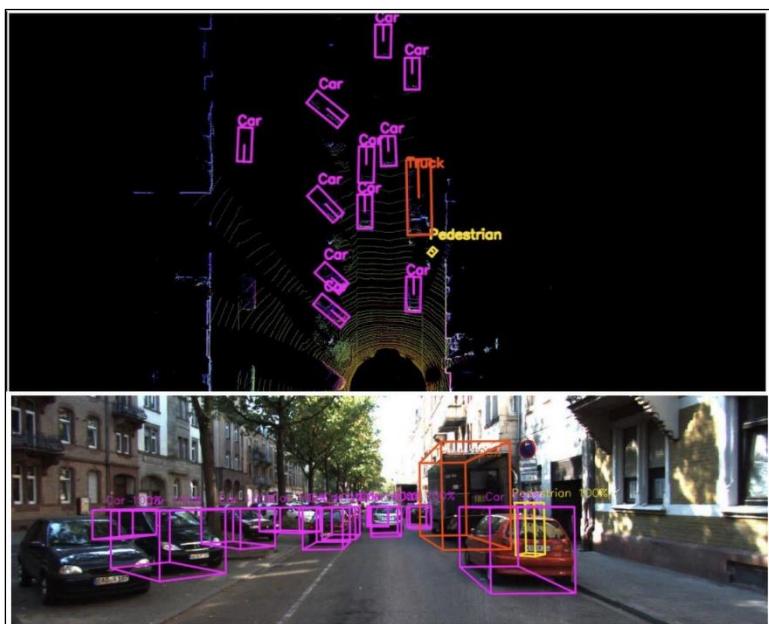


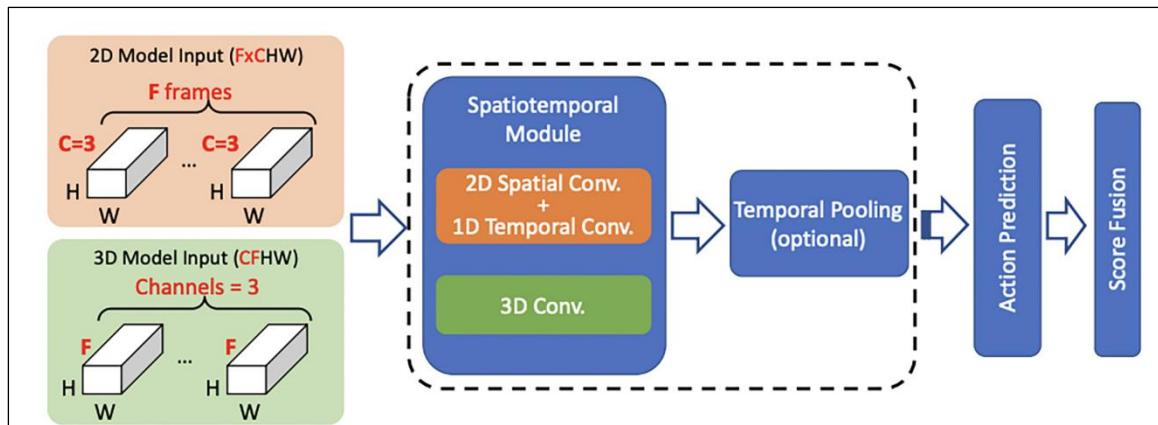
```

tree dataset
dataset
└── classes_names.txt
    └── ImageSets
        ├── test.txt
        ├── train.txt
        └── val.txt
    └── testing
        ├── calib
        ├── image_2
        └── velodyne
    └── training
        ├── calib
        ├── image_2
        ├── label_2
        └── velodyne

```

COLUMN	Example	Description	Range
type	Pedestrian	Class	Car/Pedestrian
truncation	0	Is object leaving image boundaries	0/1
occlusion	0	Is object occluded (0=fully visible, 1=partial, 2=mostly, 3=unknown)	0,1,2,3
alpha	-0.2	Observation angle	-pi to pi
x1	712.4	bbox	Image-coordinates
y1	143	bbox	Image-coordinates
x2	810.73	bbox	Image-coordinates
y2	307.92	bbox	Image-coordinates
h	1.89	height	meters
w	0.48	width	meters
l	1.2	length	meters
x	1.84	object location from camera	meters
y	1.47	object location from camera	meters
z	8.41	object location from camera	meters
ry	0.01	rotation of object around its own y-axis	-pi to pi

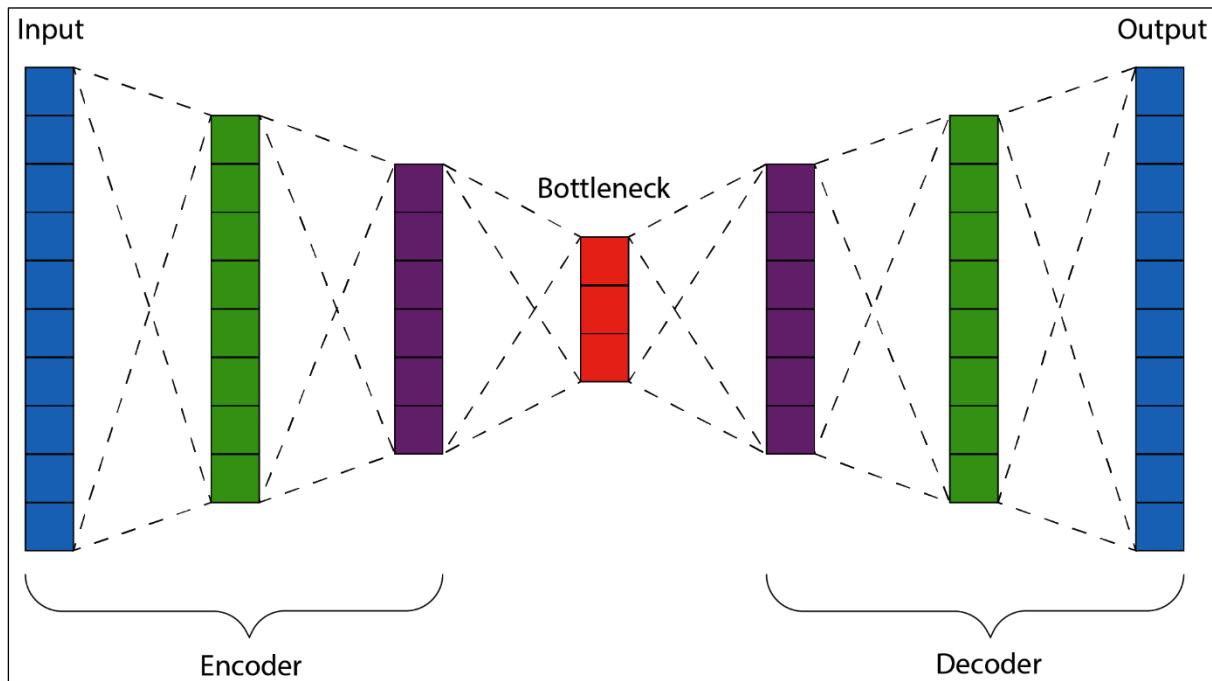




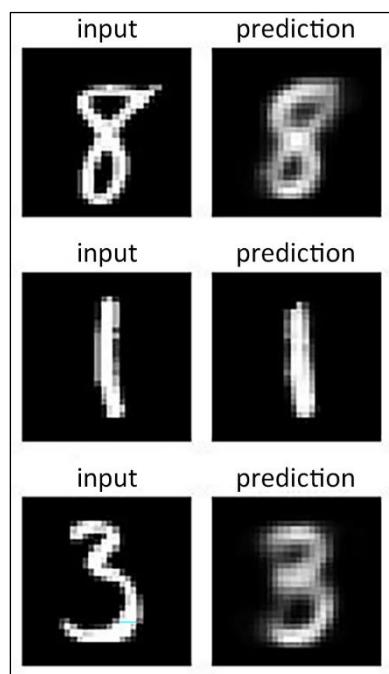
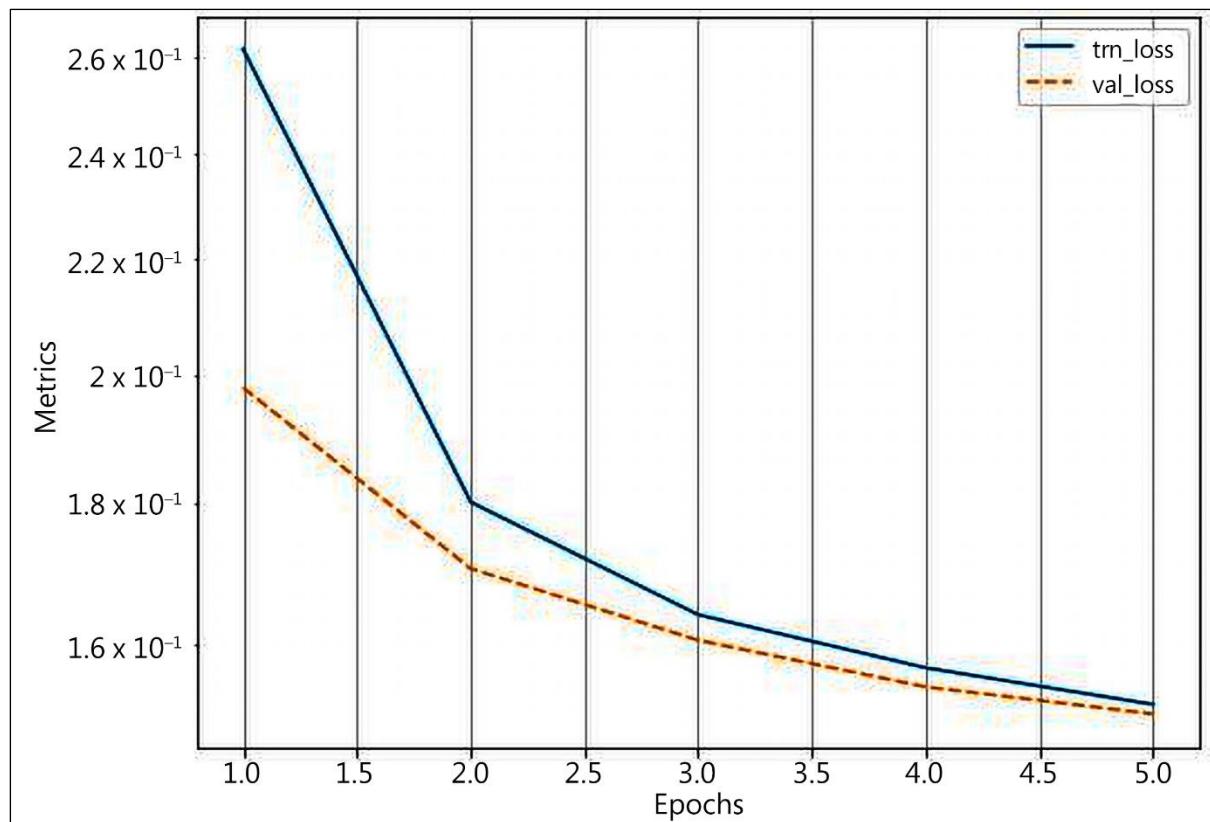
The top-5 labels with corresponding scores are:
 arm wrestling: 1.0
 rock scissors paper: 6.434453414527752e-09
 shaking hands: 2.7599860175087088e-09
 clapping: 1.3454612979302283e-09
 massaging feet: 5.555100823784187e-10

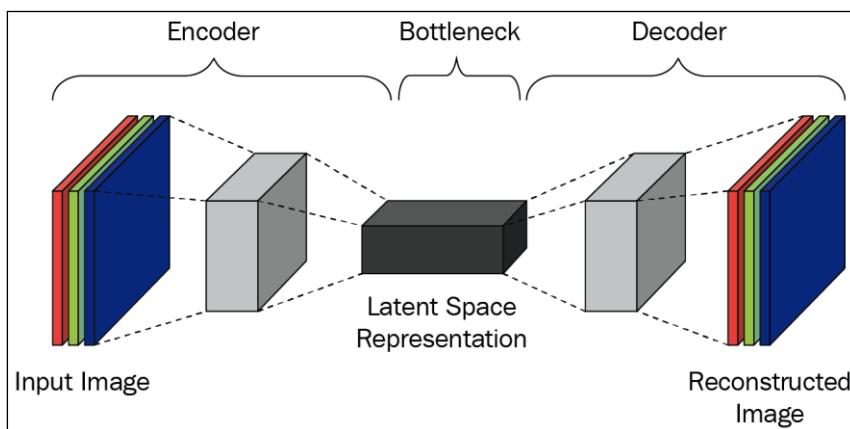
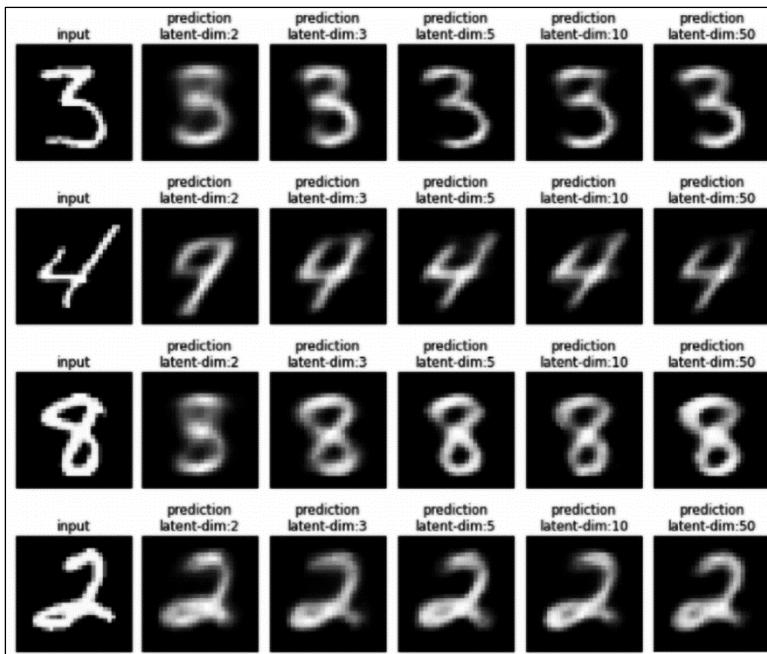
D32_1gwq35E.mp4	0
iRuyZSKhHRg.mp4	1
oXy-e_P_cAI.mp4	0
34XczvTaRiI.mp4	1
h2YqqUhnR34.mp4	0
O46YA8tI530.mp4	0
kFC3KY2bOP8.mp4	1

Chapter 11:

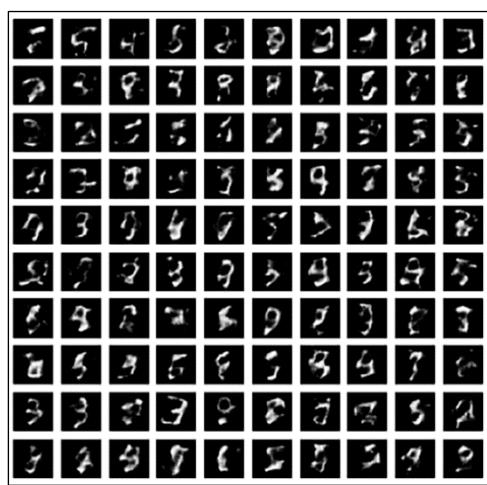
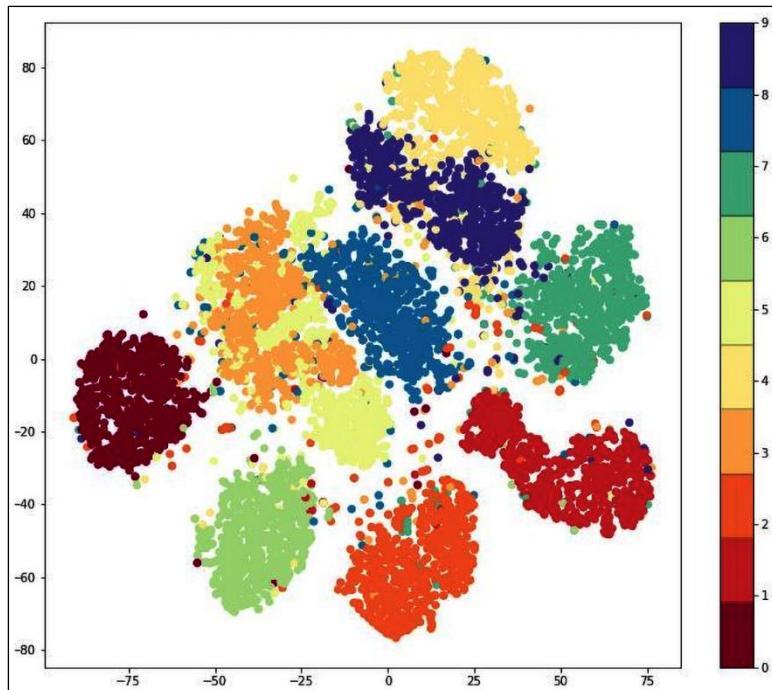
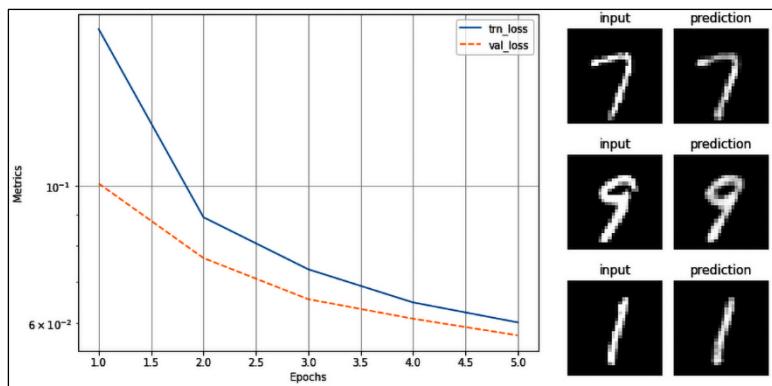


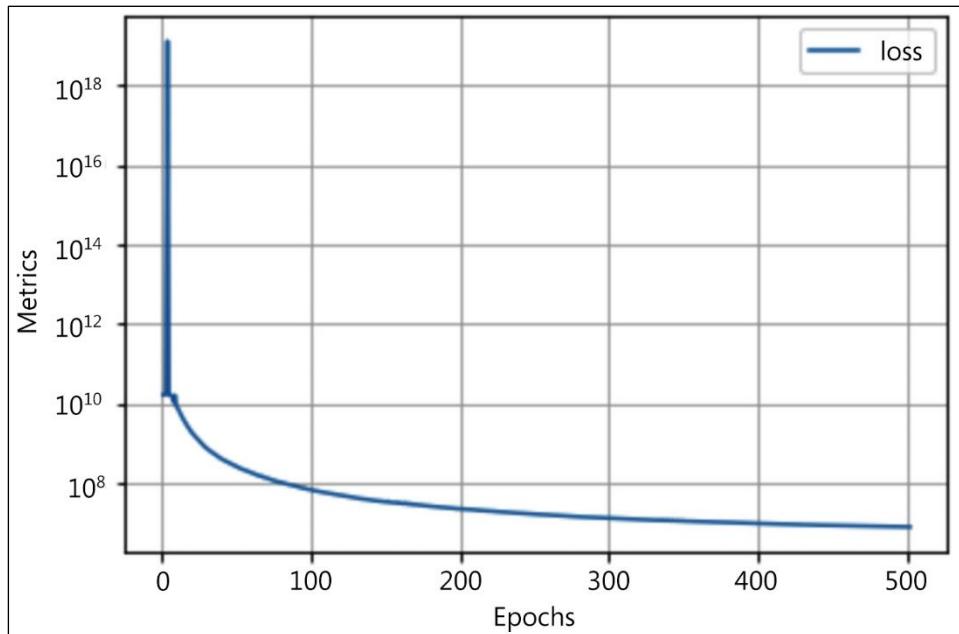
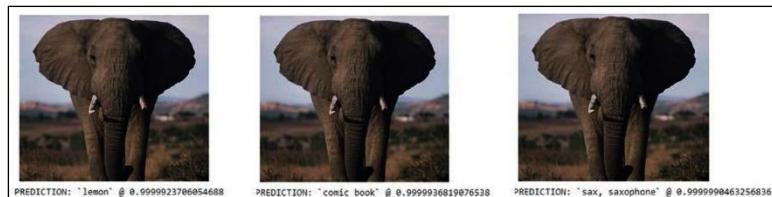
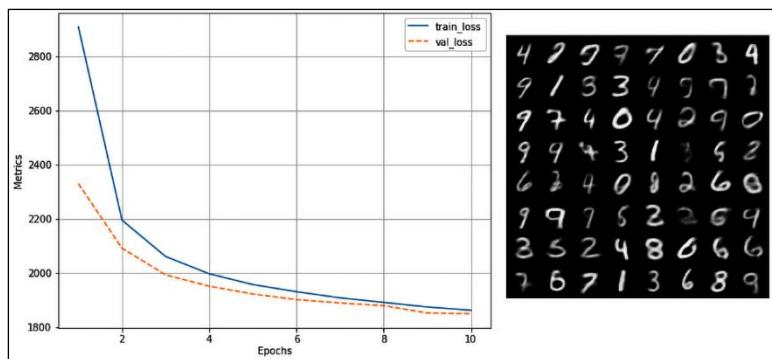
Layer (type:depth-idx)	Output Shape	Param #
<hr/>		
Sequential: 1-1		--
└ Linear: 2-1	[-1, 3]	100,480
└ ReLU: 2-2	[-1, 128]	--
└ Linear: 2-3	[-1, 64]	8,256
└ ReLU: 2-4	[-1, 64]	--
└ Linear: 2-5	[-1, 3]	195
Sequential: 1-2		--
└ Linear: 2-6	[-1, 784]	256
└ ReLU: 2-7	[-1, 64]	--
└ Linear: 2-8	[-1, 128]	8,320
└ ReLU: 2-9	[-1, 128]	--
└ Linear: 2-10	[-1, 784]	101,136
└ Tanh: 2-11	[-1, 784]	--
<hr/>		
Total params: 218,643		
Trainable params: 218,643		
Non-trainable params: 0		
Total mult-adds (M): 0.43		

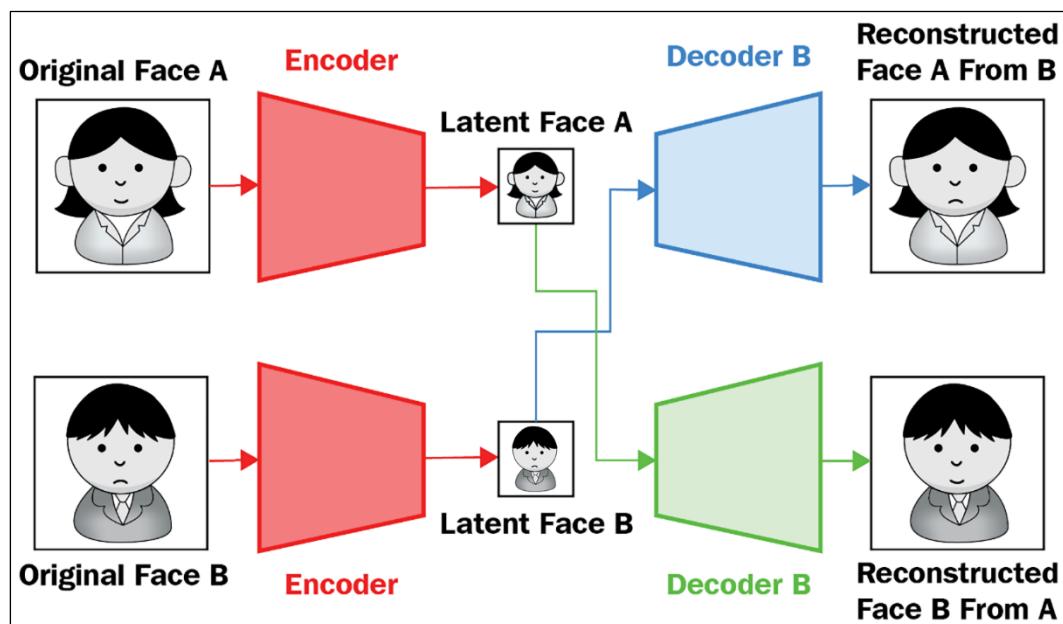
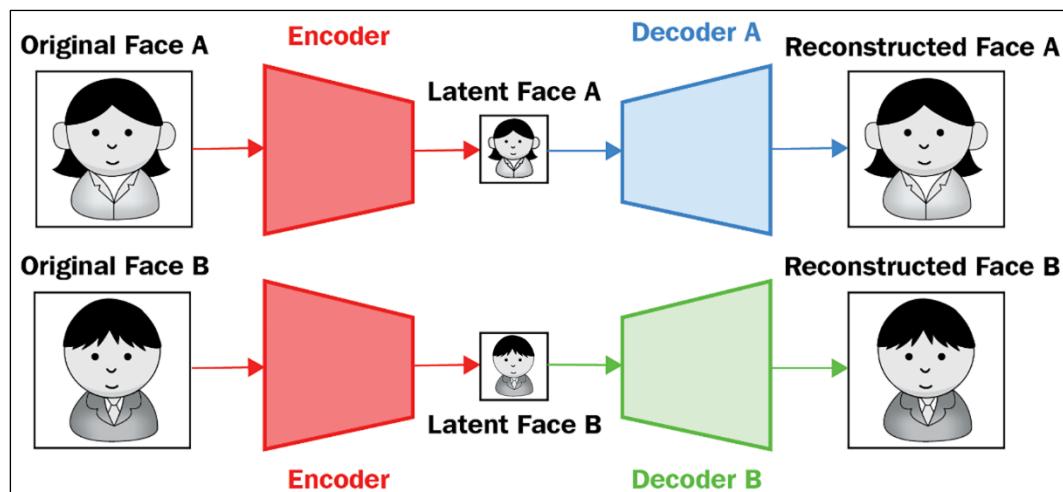


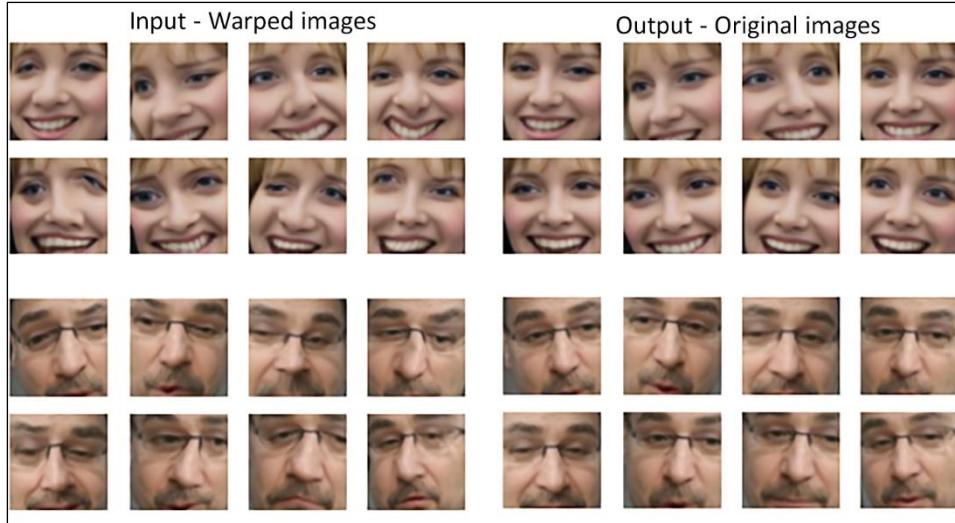


Layer (type:depth-idx)	Output Shape	Param #
<hr/>		
└ Sequential: 1-1	[-1, 64, 2, 2]	--
└ Conv2d: 2-1	[-1, 32, 10, 10]	320
└ ReLU: 2-2	[-1, 32, 10, 10]	--
└ MaxPool2d: 2-3	[-1, 32, 5, 5]	--
└ Conv2d: 2-4	[-1, 64, 3, 3]	18,496
└ ReLU: 2-5	[-1, 64, 3, 3]	--
└ MaxPool2d: 2-6	[-1, 64, 2, 2]	--
└ Sequential: 1-2	[-1, 1, 28, 28]	--
└ ConvTranspose2d: 2-7	[-1, 32, 5, 5]	18,464
└ ReLU: 2-8	[-1, 32, 5, 5]	--
└ ConvTranspose2d: 2-9	[-1, 16, 15, 15]	12,816
└ ReLU: 2-10	[-1, 16, 15, 15]	--
└ ConvTranspose2d: 2-11	[-1, 1, 28, 28]	65
└ Tanh: 2-12	[-1, 1, 28, 28]	--
<hr/>		
Total params: 50,161		
Trainable params: 50,161		
Non-trainable params: 0		
Total mult-adds (M): 3.64		

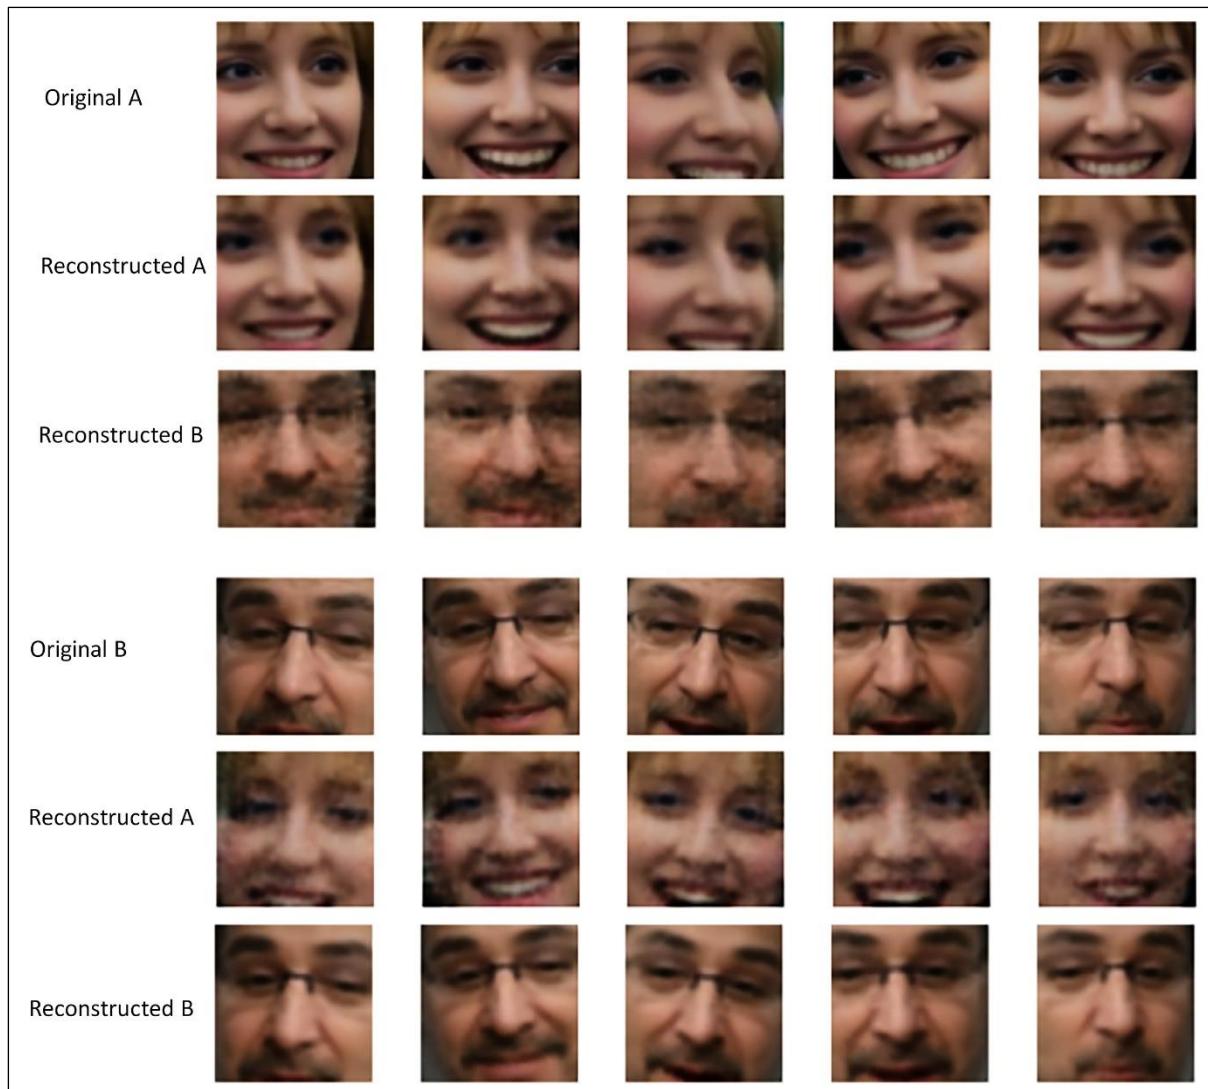


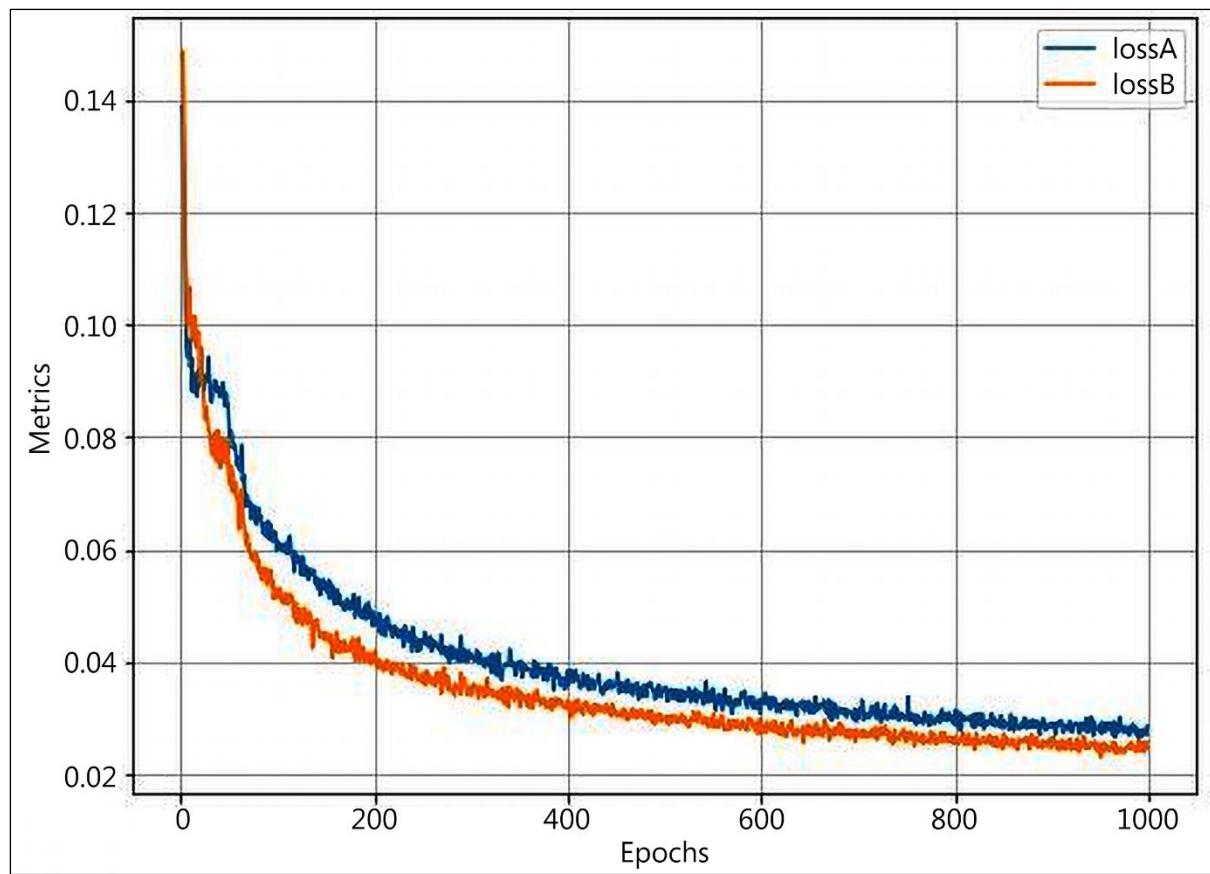




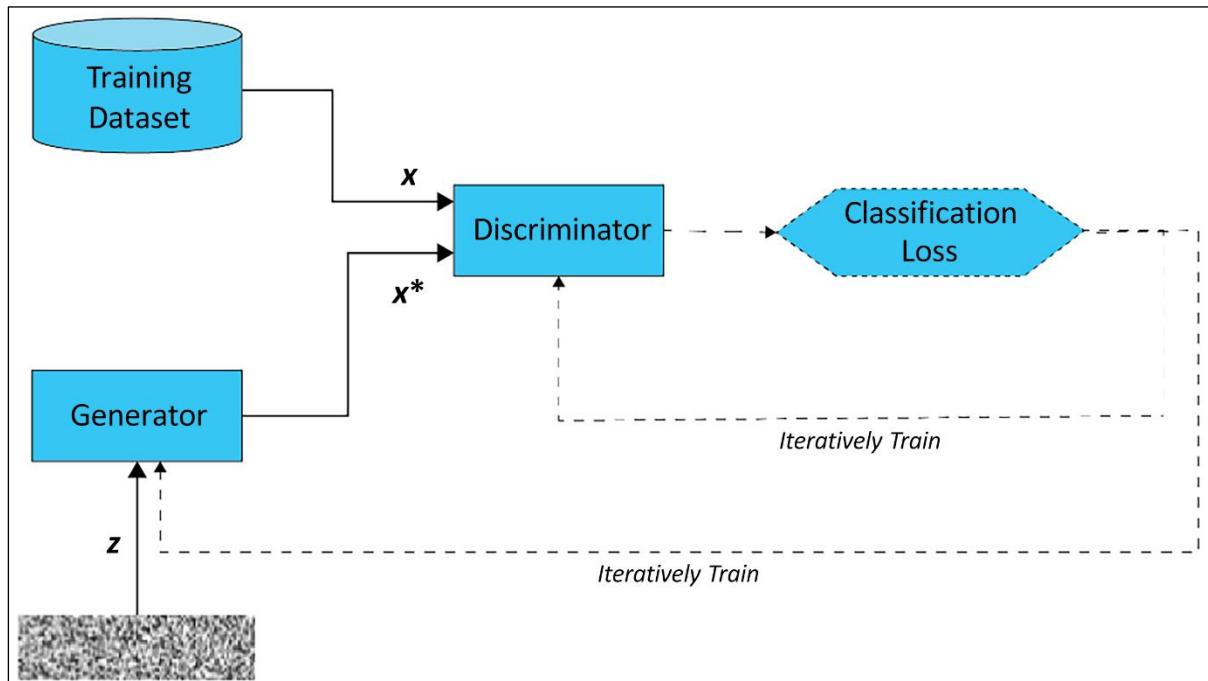
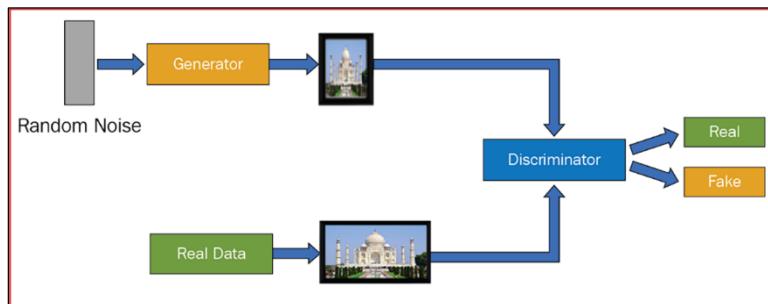


Layer (type:depth-idx)	Output Shape	Param #
Sequential: 1-1	[-1, 512, 8, 8]	--
└ Sequential: 2-1	[-1, 128, 32, 32]	--
└ Conv2d: 3-1	[-1, 128, 32, 32]	9,728
└ LeakyReLU: 3-2	[-1, 128, 32, 32]	--
└ Sequential: 2-2	[-1, 256, 16, 16]	--
└ Conv2d: 3-3	[-1, 256, 16, 16]	819,456
└ LeakyReLU: 3-4	[-1, 256, 16, 16]	--
└ Sequential: 2-3	[-1, 512, 8, 8]	--
└ Conv2d: 3-5	[-1, 512, 8, 8]	3,277,312
└ LeakyReLU: 3-6	[-1, 512, 8, 8]	--
└ Sequential: 2-4	[-1, 1024, 4, 4]	--
└ Conv2d: 3-7	[-1, 1024, 4, 4]	13,108,224
└ LeakyReLU: 3-8	[-1, 1024, 4, 4]	--
└ Flatten: 2-5	[-1, 16384]	--
└ Linear: 2-6	[-1, 1024]	16,778,240
└ Linear: 2-7	[-1, 16384]	16,793,600
└ Reshape: 2-8	[-1, 1024, 4, 4]	--
└ Sequential: 2-9	[-1, 512, 8, 8]	--
└ ConvTranspose2d: 3-9	[-1, 512, 8, 8]	2,097,664
└ LeakyReLU: 3-10	[-1, 512, 8, 8]	--
└ Sequential: 1-2	[-1, 3, 64, 64]	--
└ Sequential: 2-10	[-1, 256, 16, 16]	--
└ ConvTranspose2d: 3-11	[-1, 256, 16, 16]	524,544
└ LeakyReLU: 3-12	[-1, 256, 16, 16]	--
└ Sequential: 2-11	[-1, 128, 32, 32]	--
└ ConvTranspose2d: 3-13	[-1, 128, 32, 32]	131,200
└ LeakyReLU: 3-14	[-1, 128, 32, 32]	--
└ Sequential: 2-12	[-1, 64, 64, 64]	--
└ ConvTranspose2d: 3-15	[-1, 64, 64, 64]	32,832
└ LeakyReLU: 3-16	[-1, 64, 64, 64]	--
└ Conv2d: 2-13	[-1, 3, 64, 64]	1,731
└ Sigmoid: 2-14	[-1, 3, 64, 64]	--
Total params: 53,574,531		
Trainable params: 53,574,531		
Non-trainable params: 0		
Total mult-adds (G): 1.29		





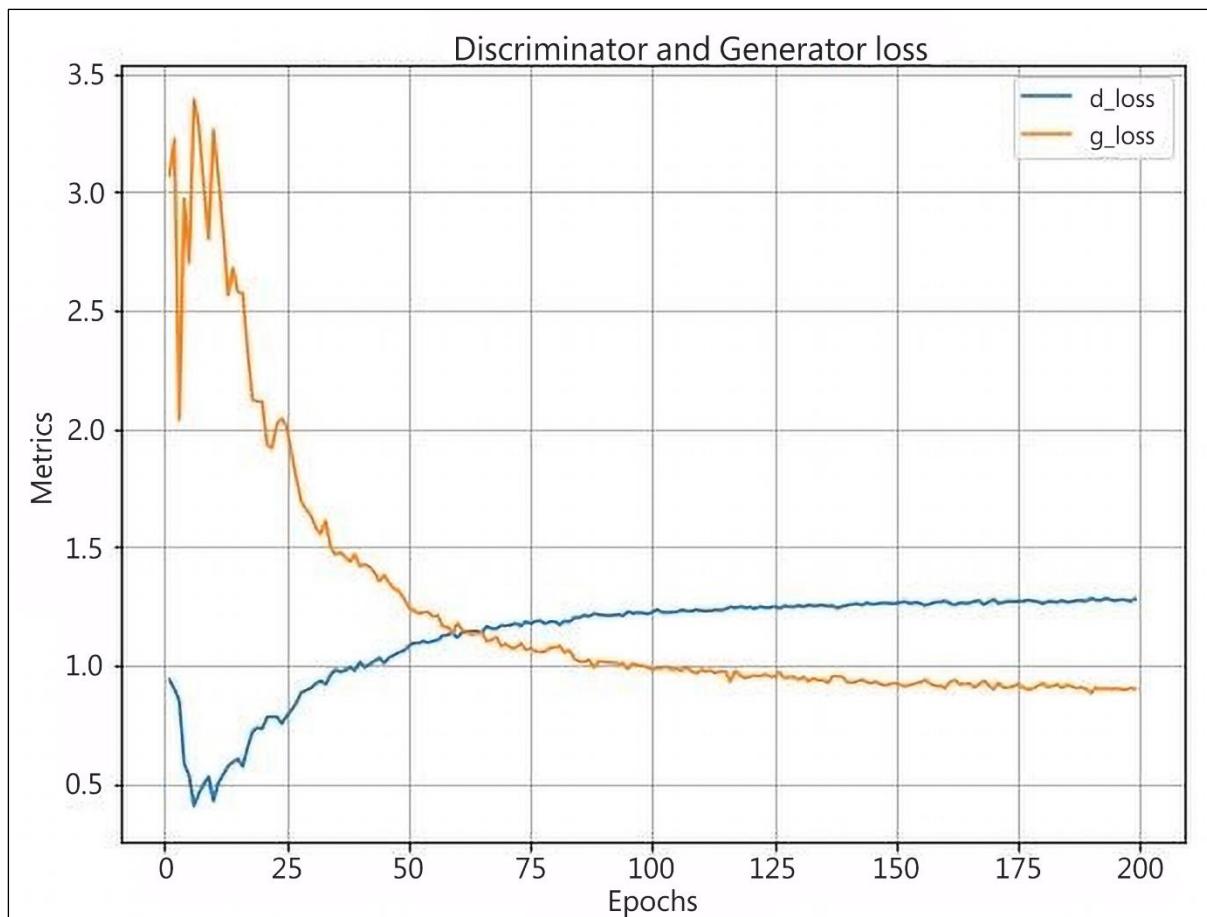
Chapter 12:



Layer (type:depth-idx)	Output Shape	Param #
Sequential: 1-1	[-1, 1]	--
└ Linear: 2-1	[-1, 1024]	803,840
└ LeakyReLU: 2-2	[-1, 1024]	--
└ Dropout: 2-3	[-1, 1024]	--
└ Linear: 2-4	[-1, 512]	524,800
└ LeakyReLU: 2-5	[-1, 512]	--
└ Dropout: 2-6	[-1, 512]	--
└ Linear: 2-7	[-1, 256]	131,328
└ LeakyReLU: 2-8	[-1, 256]	--
└ Dropout: 2-9	[-1, 256]	--
└ Linear: 2-10	[-1, 1]	257
└ Sigmoid: 2-11	[-1, 1]	--

Total params: 1,460,225
 Trainable params: 1,460,225
 Non-trainable params: 0
 Total mult-adds (M): 2.92

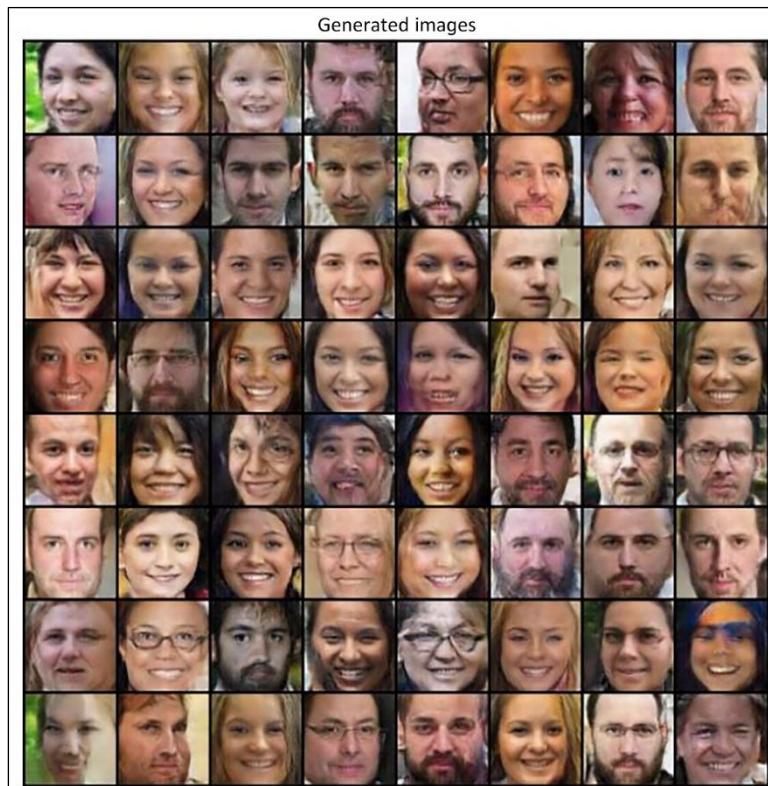
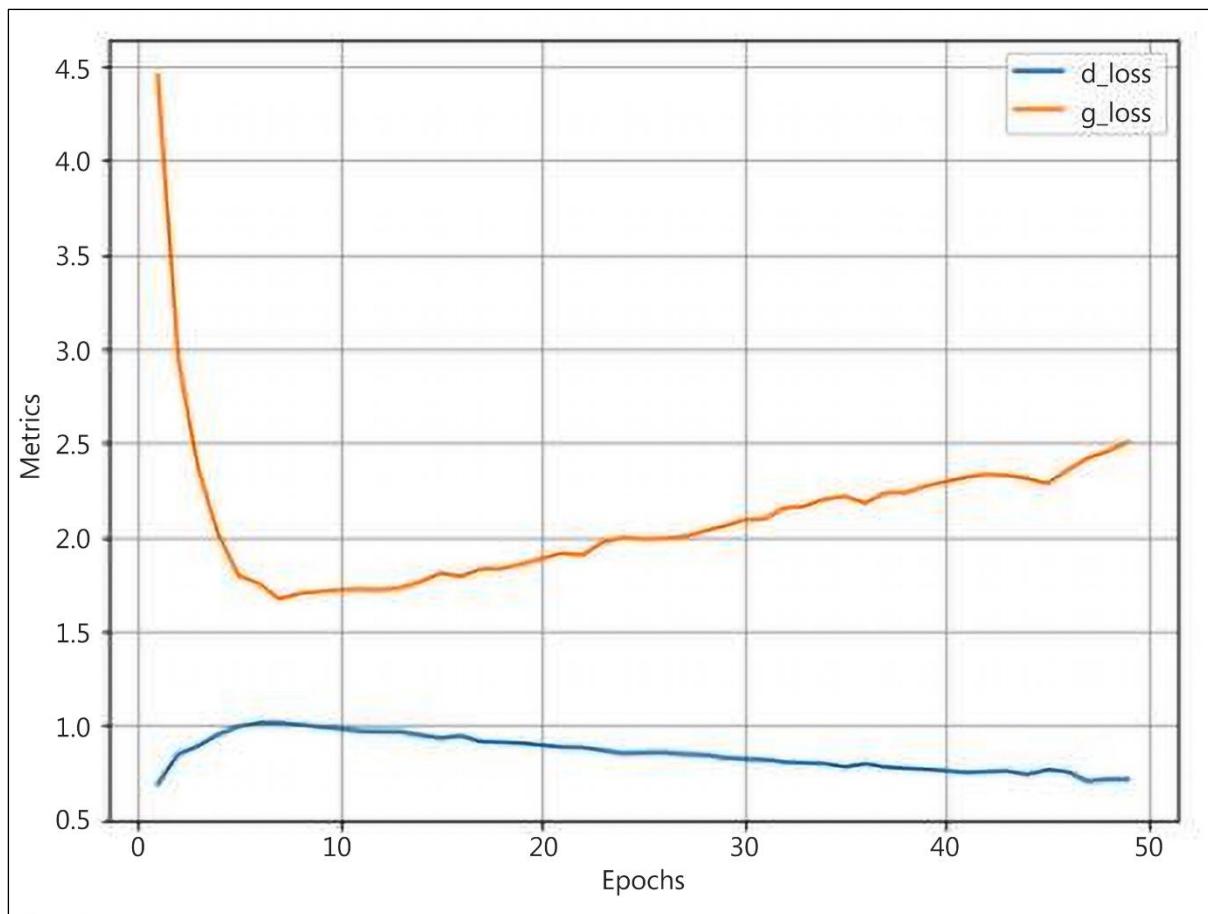
Layer (type:depth-idx)	Output Shape	Param #
Sequential: 1-1	[-1, 784]	--
└ Linear: 2-1	[-1, 256]	25,856
└ LeakyReLU: 2-2	[-1, 256]	--
└ Linear: 2-3	[-1, 512]	131,584
└ LeakyReLU: 2-4	[-1, 512]	--
└ Linear: 2-5	[-1, 1024]	525,312
└ LeakyReLU: 2-6	[-1, 1024]	--
└ Linear: 2-7	[-1, 784]	803,600
└ Tanh: 2-8	[-1, 784]	--
Total params: 1,486,352		
Trainable params: 1,486,352		
Non-trainable params: 0		
Total mult-adds (M): 2.97		





Layer (type:depth-idx)	Output Shape	Param #
Sequential: 1-1	[-1, 1, 1, 1]	--
└ Conv2d: 2-1	[-1, 64, 32, 32]	3,072
└ LeakyReLU: 2-2	[-1, 64, 32, 32]	--
└ Conv2d: 2-3	[-1, 128, 16, 16]	131,072
└ BatchNorm2d: 2-4	[-1, 128, 16, 16]	256
└ LeakyReLU: 2-5	[-1, 128, 16, 16]	--
└ Conv2d: 2-6	[-1, 256, 8, 8]	524,288
└ BatchNorm2d: 2-7	[-1, 256, 8, 8]	512
└ LeakyReLU: 2-8	[-1, 256, 8, 8]	--
└ Conv2d: 2-9	[-1, 512, 4, 4]	2,097,152
└ BatchNorm2d: 2-10	[-1, 512, 4, 4]	1,024
└ LeakyReLU: 2-11	[-1, 512, 4, 4]	--
└ Conv2d: 2-12	[-1, 1, 1, 1]	8,192
└ Sigmoid: 2-13	[-1, 1, 1, 1]	--
Total params: 2,765,568		
Trainable params: 2,765,568		
Non-trainable params: 0		
Total mult-adds (M): 106.58		

Layer (type:depth-idx)	Output Shape	Param #
Sequential: 1-1	[-1, 3, 64, 64]	--
└ ConvTranspose2d: 2-1	[-1, 512, 4, 4]	819,200
└ BatchNorm2d: 2-2	[-1, 512, 4, 4]	1,024
└ ReLU: 2-3	[-1, 512, 4, 4]	--
└ ConvTranspose2d: 2-4	[-1, 256, 8, 8]	2,097,152
└ BatchNorm2d: 2-5	[-1, 256, 8, 8]	512
└ ReLU: 2-6	[-1, 256, 8, 8]	--
└ ConvTranspose2d: 2-7	[-1, 128, 16, 16]	524,288
└ BatchNorm2d: 2-8	[-1, 128, 16, 16]	256
└ ReLU: 2-9	[-1, 128, 16, 16]	--
└ ConvTranspose2d: 2-10	[-1, 64, 32, 32]	131,072
└ BatchNorm2d: 2-11	[-1, 64, 32, 32]	128
└ ReLU: 2-12	[-1, 64, 32, 32]	--
└ ConvTranspose2d: 2-13	[-1, 3, 64, 64]	3,072
└ Tanh: 2-14	[-1, 3, 64, 64]	--
Total params: 3,576,704		
Trainable params: 3,576,704		
Non-trainable params: 0		
Total mult-adds (M): 431.92		



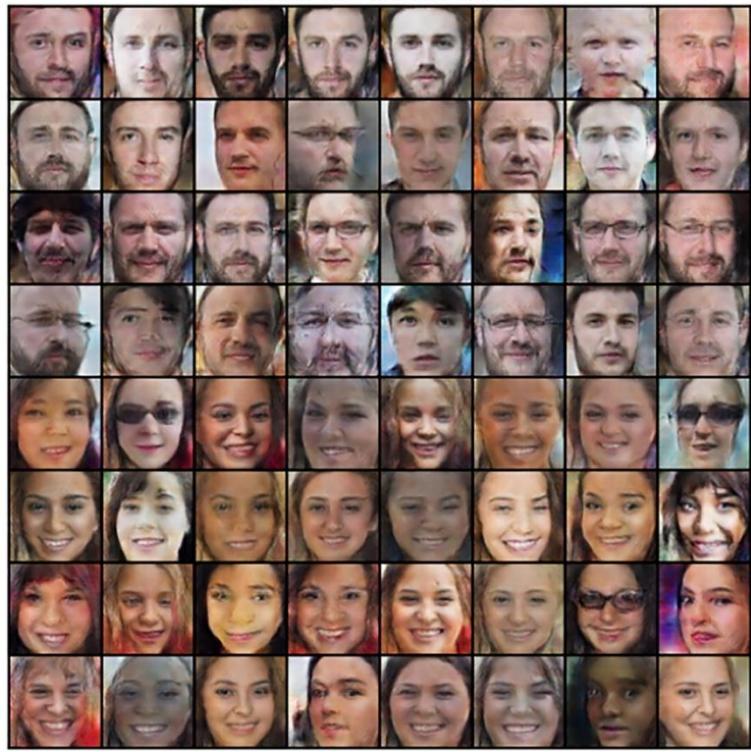
Layer (type:depth-idx)	Output Shape	Param #
-Sequential: 1-1	[-1, 256]	--
└Conv2d: 2-1	[-1, 64, 32, 32]	3,072
└LeakyReLU: 2-2	[-1, 64, 32, 32]	--
└Conv2d: 2-3	[-1, 128, 16, 16]	131,072
└BatchNorm2d: 2-4	[-1, 128, 16, 16]	256
└LeakyReLU: 2-5	[-1, 128, 16, 16]	--
└Conv2d: 2-6	[-1, 256, 8, 8]	524,288
└BatchNorm2d: 2-7	[-1, 256, 8, 8]	512
└LeakyReLU: 2-8	[-1, 256, 8, 8]	--
└Conv2d: 2-9	[-1, 512, 4, 4]	2,097,152
└BatchNorm2d: 2-10	[-1, 512, 4, 4]	1,024
└LeakyReLU: 2-11	[-1, 512, 4, 4]	--
└Conv2d: 2-12	[-1, 64, 2, 2]	524,288
└BatchNorm2d: 2-13	[-1, 64, 2, 2]	128
└LeakyReLU: 2-14	[-1, 64, 2, 2]	--
└Flatten: 2-15	[-1, 256]	--
-Embedding: 1-2	[-1, 32]	64
-Sequential: 1-3	[-1, 1]	--
└Linear: 2-16	[-1, 100]	28,900
└LeakyReLU: 2-17	[-1, 100]	--
└Linear: 2-18	[-1, 1]	101
└Sigmoid: 2-19	[-1, 1]	--

Total params: 3,310,857
Trainable params: 3,310,857
Non-trainable params: 0
Total mult-adds (M): 109.25

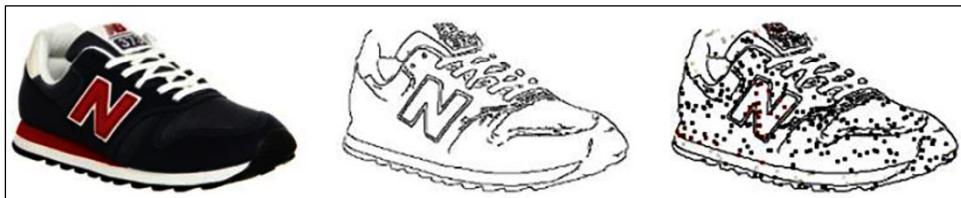
Layer (type:depth-idx)	Output Shape	Param #
-Embedding: 1-1	[-1, 32]	64
-Sequential: 1-2	[-1, 3, 64, 64]	--
└ConvTranspose2d: 2-1	[-1, 512, 4, 4]	1,081,344
└BatchNorm2d: 2-2	[-1, 512, 4, 4]	1,024
└ReLU: 2-3	[-1, 512, 4, 4]	--
└ConvTranspose2d: 2-4	[-1, 256, 8, 8]	2,097,152
└BatchNorm2d: 2-5	[-1, 256, 8, 8]	512
└ReLU: 2-6	[-1, 256, 8, 8]	--
└ConvTranspose2d: 2-7	[-1, 128, 16, 16]	524,288
└BatchNorm2d: 2-8	[-1, 128, 16, 16]	256
└ReLU: 2-9	[-1, 128, 16, 16]	--
└ConvTranspose2d: 2-10	[-1, 64, 32, 32]	131,072
└BatchNorm2d: 2-11	[-1, 64, 32, 32]	128
└ReLU: 2-12	[-1, 64, 32, 32]	--
└ConvTranspose2d: 2-13	[-1, 3, 64, 64]	3,072
└Tanh: 2-14	[-1, 3, 64, 64]	--

Total params: 3,838,912
Trainable params: 3,838,912
Non-trainable params: 0
Total mult-adds (M): 436.38

Generated male and female faces



Chapter 13:



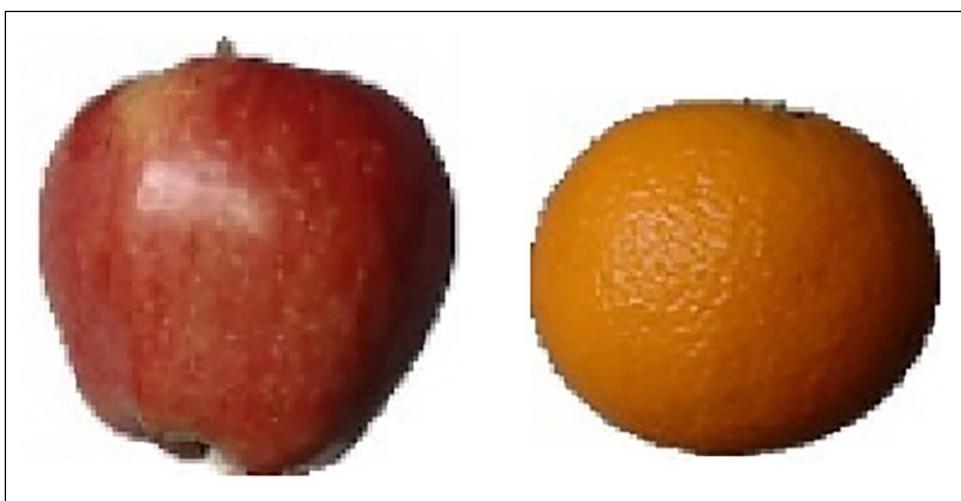
Layer (type)	Output Shape	Param #	Tr. Param #
UNetDown-1	[3, 64, 128, 128]	3,072	3,072
UNetDown-2	[3, 128, 64, 64]	131,072	131,072
UNetDown-3	[3, 256, 32, 32]	524,288	524,288
UNetDown-4	[3, 512, 16, 16]	2,097,152	2,097,152
UNetDown-5	[3, 512, 8, 8]	4,194,304	4,194,304
UNetDown-6	[3, 512, 4, 4]	4,194,304	4,194,304
UNetDown-7	[3, 512, 2, 2]	4,194,304	4,194,304
UNetDown-8	[3, 512, 1, 1]	4,194,304	4,194,304
UNetUp-9	[3, 1024, 2, 2]	4,194,304	4,194,304
UNetUp-10	[3, 1024, 4, 4]	8,388,608	8,388,608
UNetUp-11	[3, 1024, 8, 8]	8,388,608	8,388,608
UNetUp-12	[3, 1024, 16, 16]	8,388,608	8,388,608
UNetUp-13	[3, 512, 32, 32]	4,194,304	4,194,304
UNetUp-14	[3, 256, 64, 64]	1,048,576	1,048,576
UNetUp-15	[3, 128, 128, 128]	262,144	262,144
Upsample-16	[3, 128, 256, 256]	0	0
ZeroPad2d-17	[3, 128, 257, 257]	0	0
Conv2d-18	[3, 3, 256, 256]	6,147	6,147
Tanh-19	[3, 3, 256, 256]	0	0

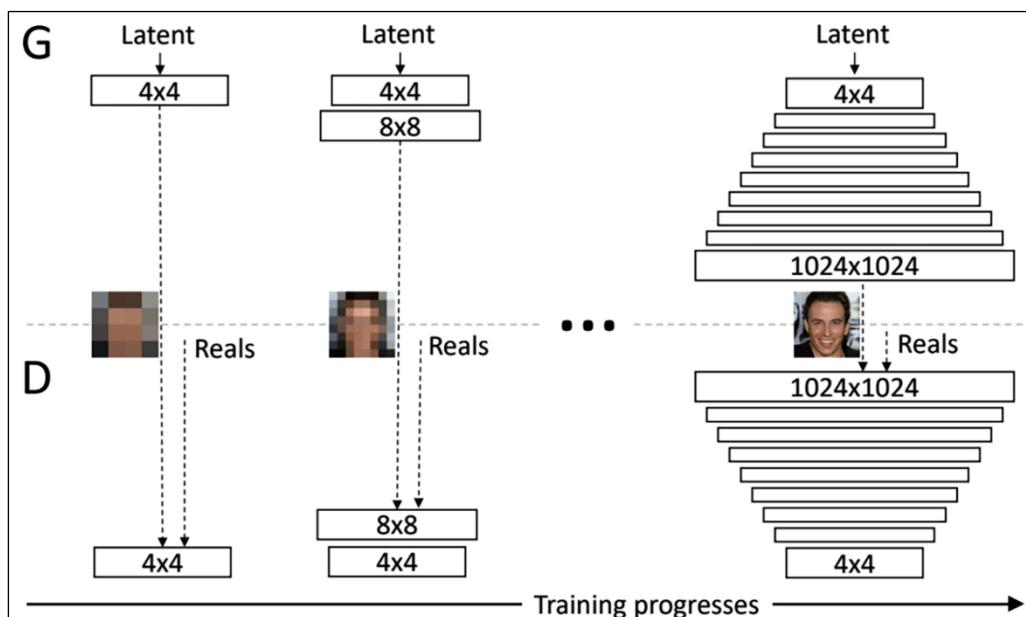
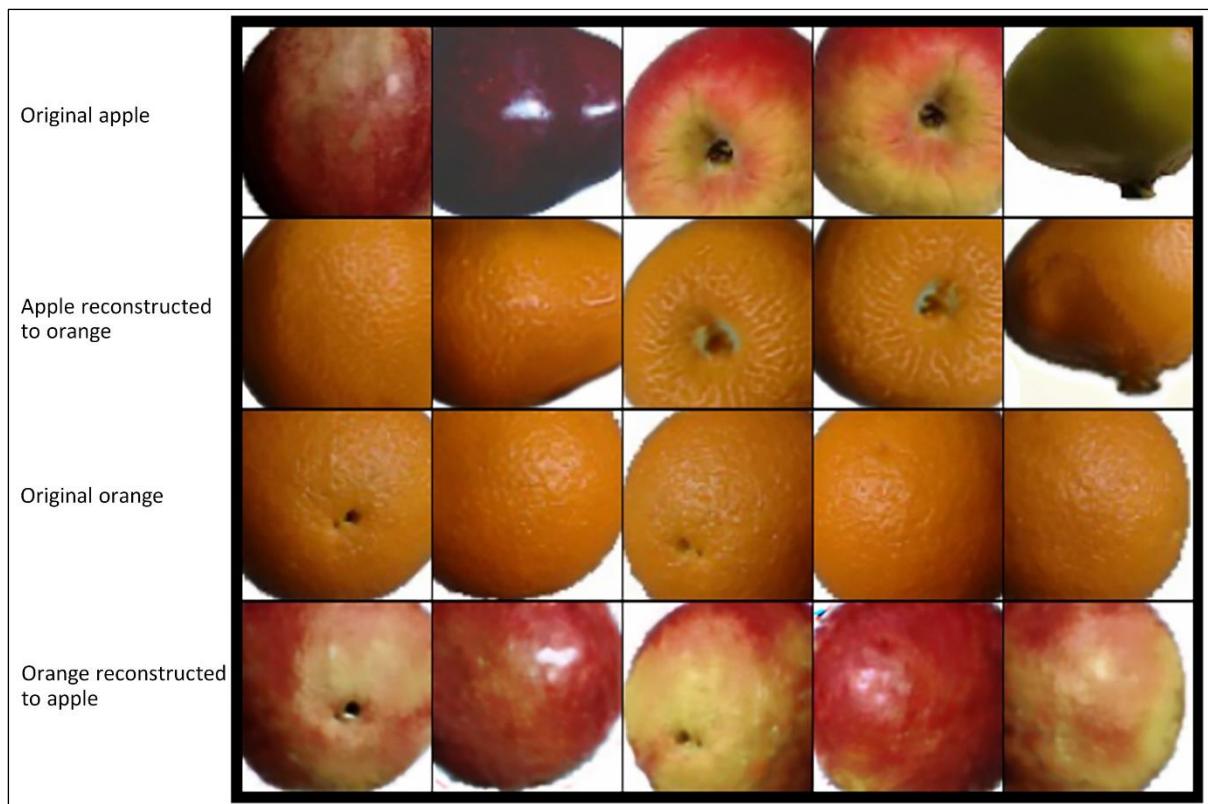
Total params: 54,404,099
 Trainable params: 54,404,099
 Non-trainable params: 0

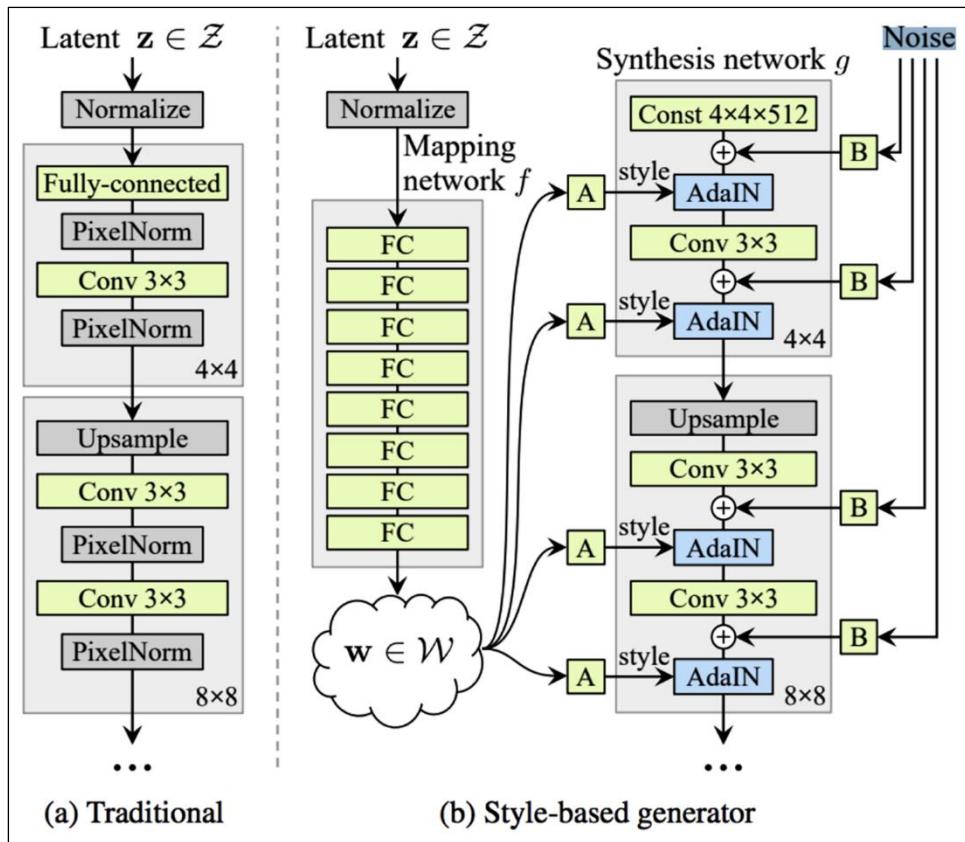
Layer (type)	Output Shape	Param #	Tr. Param #
Conv2d-1	[3, 64, 128, 128]	6,208	6,208
LeakyReLU-2	[3, 64, 128, 128]	0	0
Conv2d-3	[3, 128, 64, 64]	131,200	131,200
InstanceNorm2d-4	[3, 128, 64, 64]	0	0
LeakyReLU-5	[3, 128, 64, 64]	0	0
Conv2d-6	[3, 256, 32, 32]	524,544	524,544
InstanceNorm2d-7	[3, 256, 32, 32]	0	0
LeakyReLU-8	[3, 256, 32, 32]	0	0
Conv2d-9	[3, 512, 16, 16]	2,097,664	2,097,664
InstanceNorm2d-10	[3, 512, 16, 16]	0	0
LeakyReLU-11	[3, 512, 16, 16]	0	0
ZeroPad2d-12	[3, 512, 17, 17]	0	0
Conv2d-13	[3, 1, 16, 16]	8,192	8,192

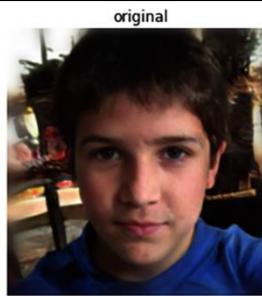
Total params: 2,767,808
 Trainable params: 2,767,808
 Non-trainable params: 0

Source::Generated::GroundTruth









Optimizing Latents.

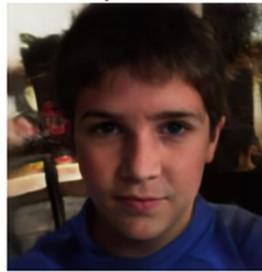
[2020-10-13 04:54:26,699][INFO] Loading pytorch model from `InterFaceGAN/models/pretrain/stylegan_ffhq.pth`.

[2020-10-13 04:54:26,982][INFO] Successfully loaded!

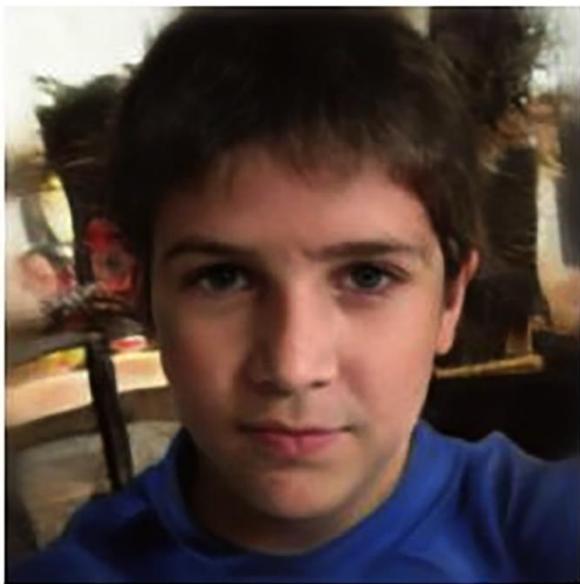
[2020-10-13 04:54:26,983][INFO] `lod` of the loaded model is 0.0.

Step: 999, Loss: 2.4222376346588135: 100% 1000/1000 [02:06<00:00, 7.89it/s]

synthesized



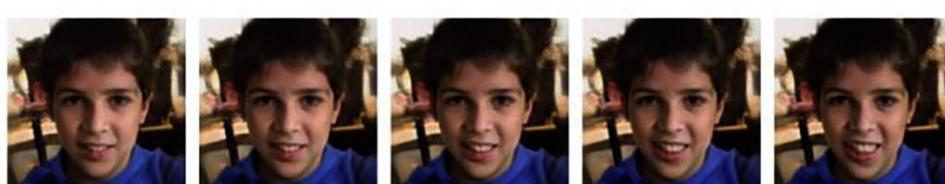
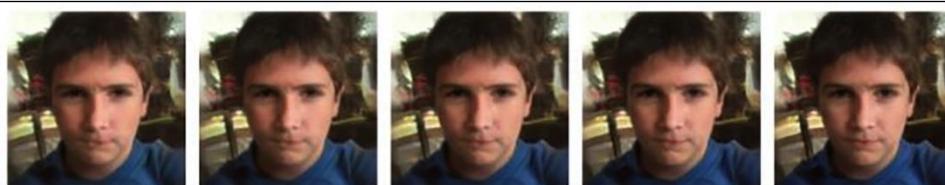
Transfer high level features

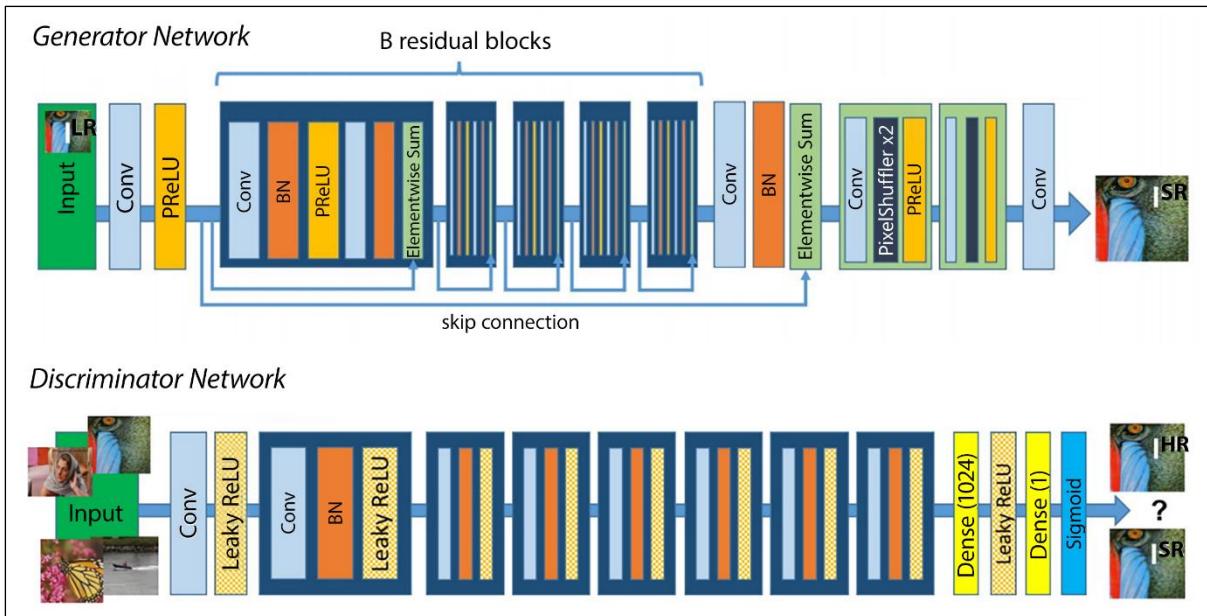


Transfer granular features



Transfer general image features

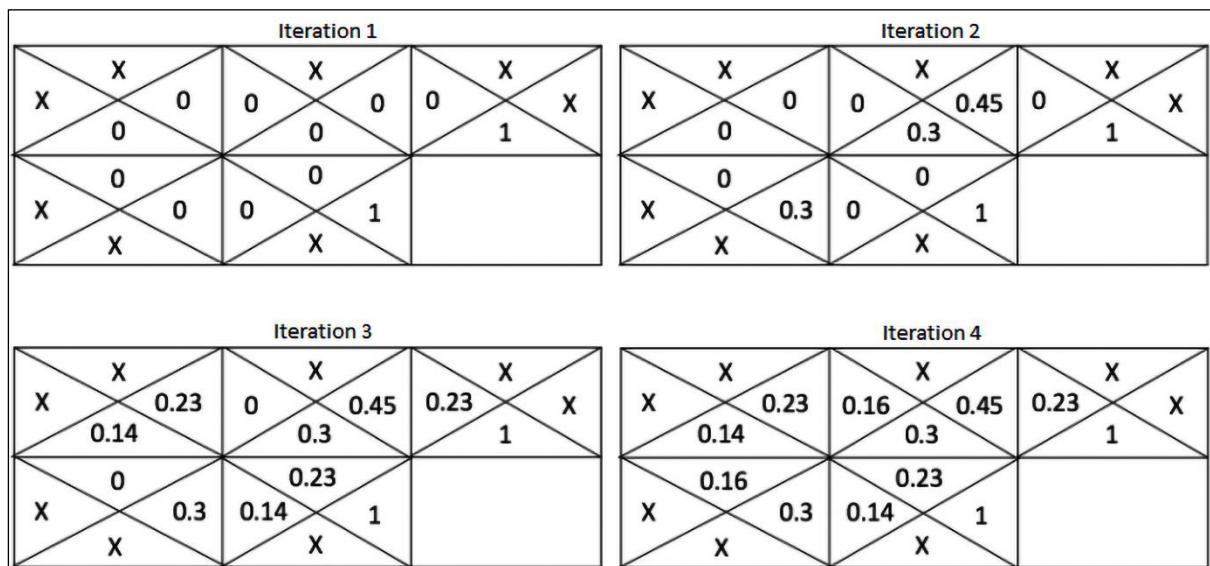
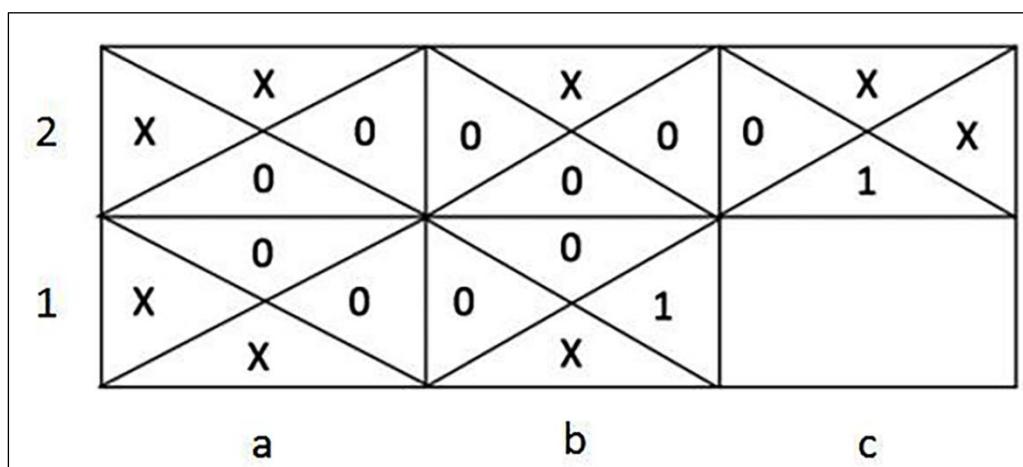




Chapter 14:

Start		
		+1

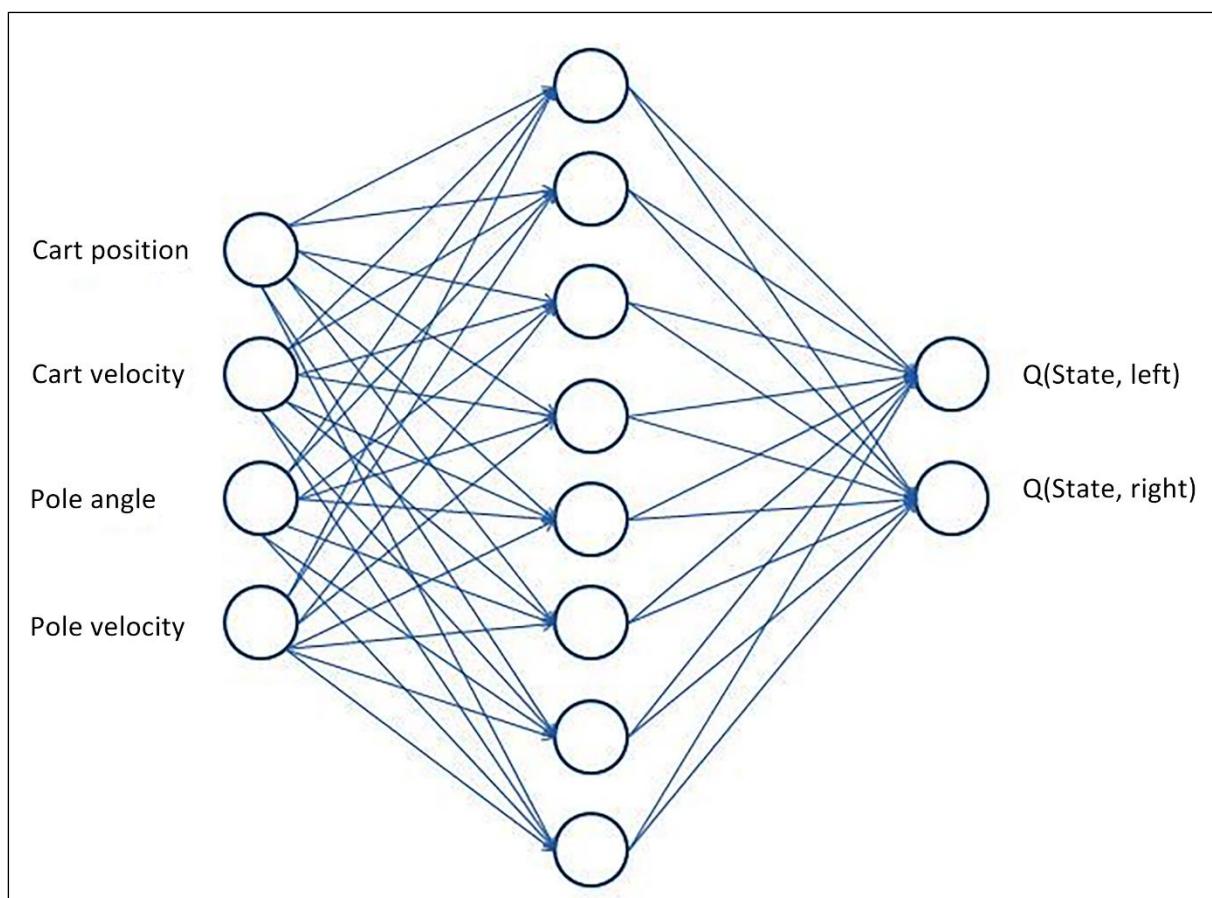
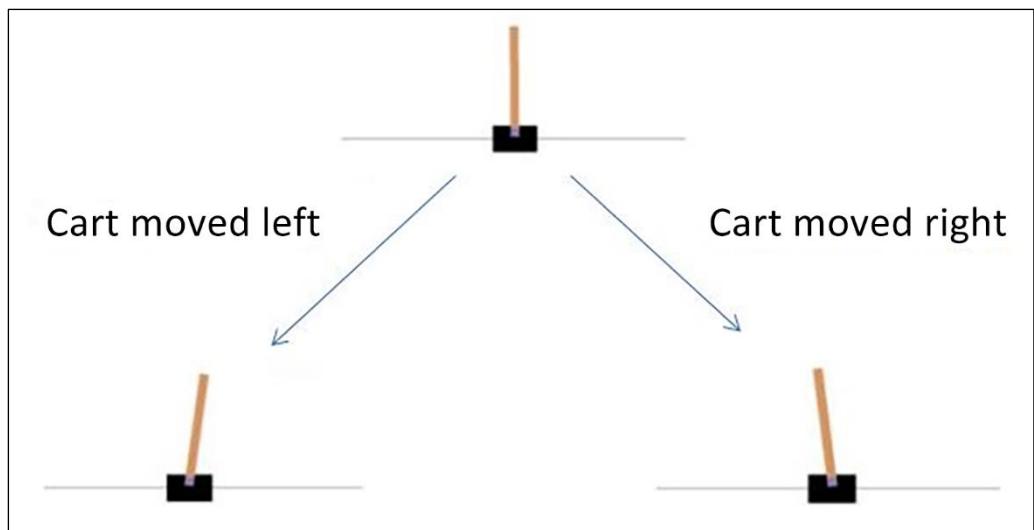
0.81	0.9	1
0.9	1	

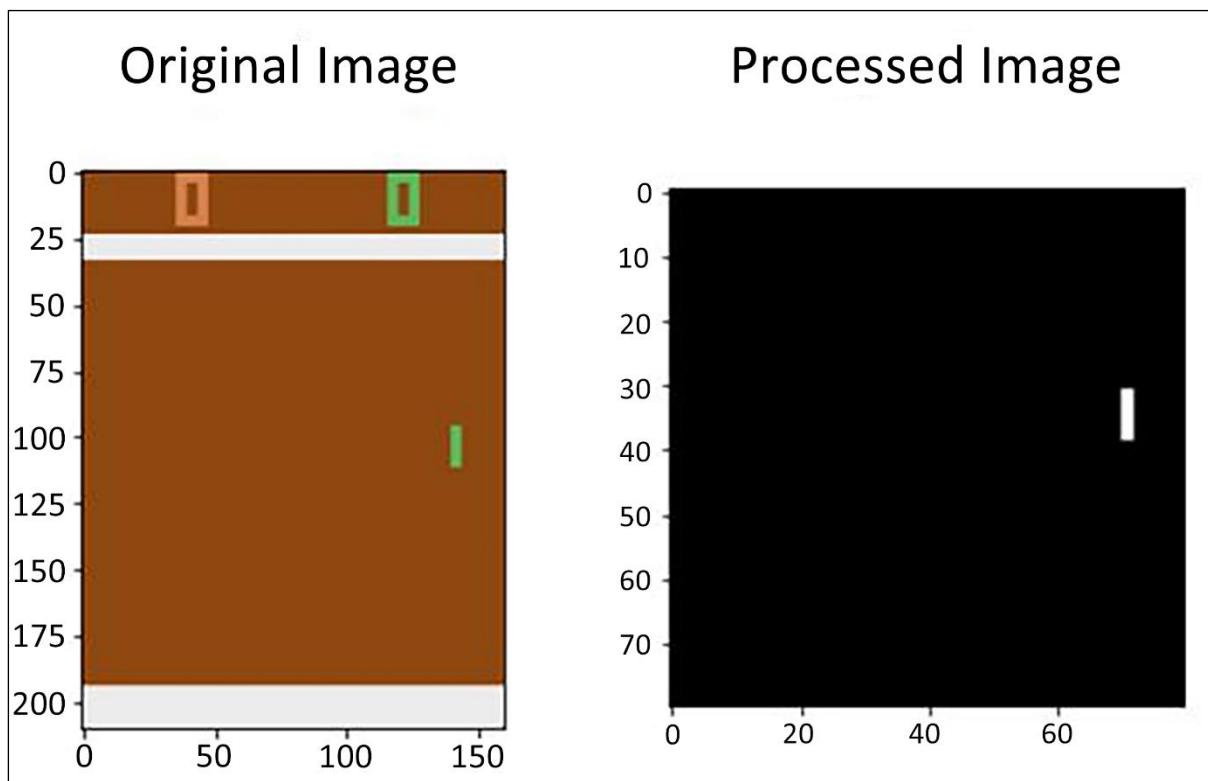
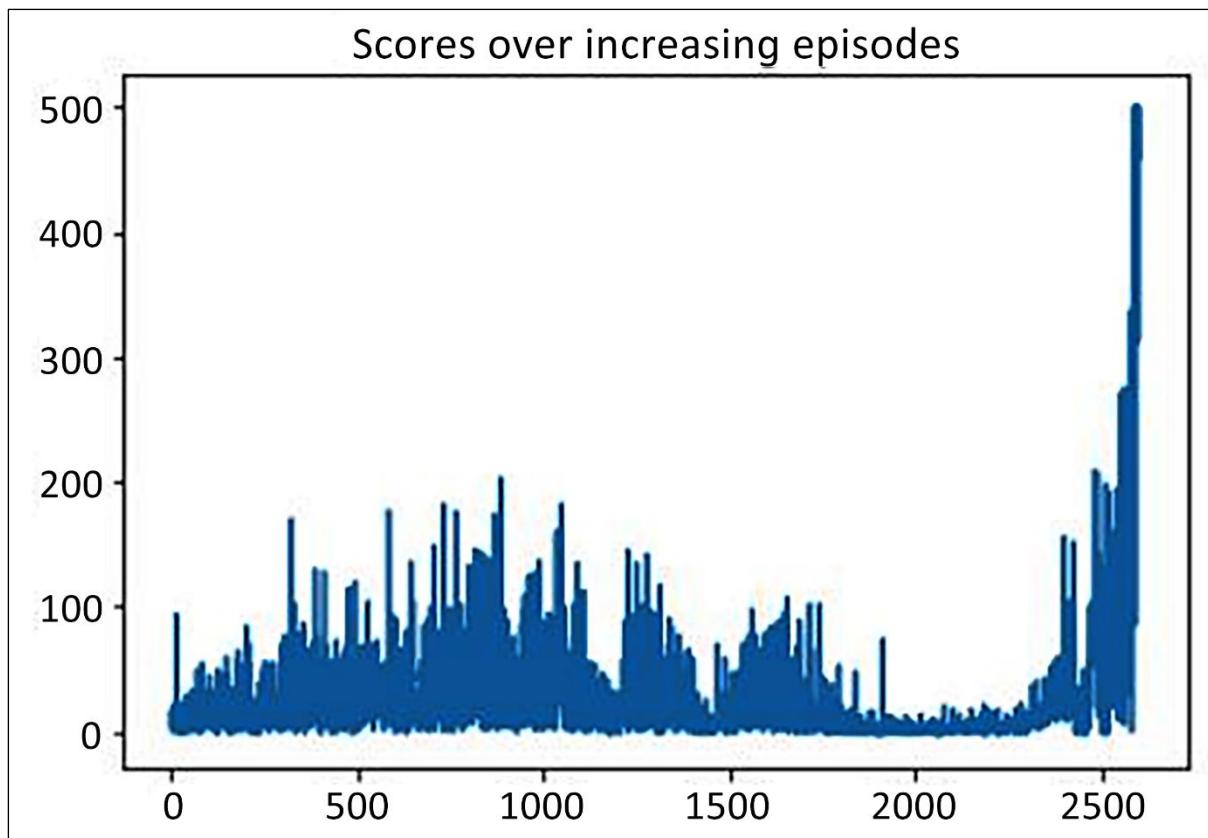


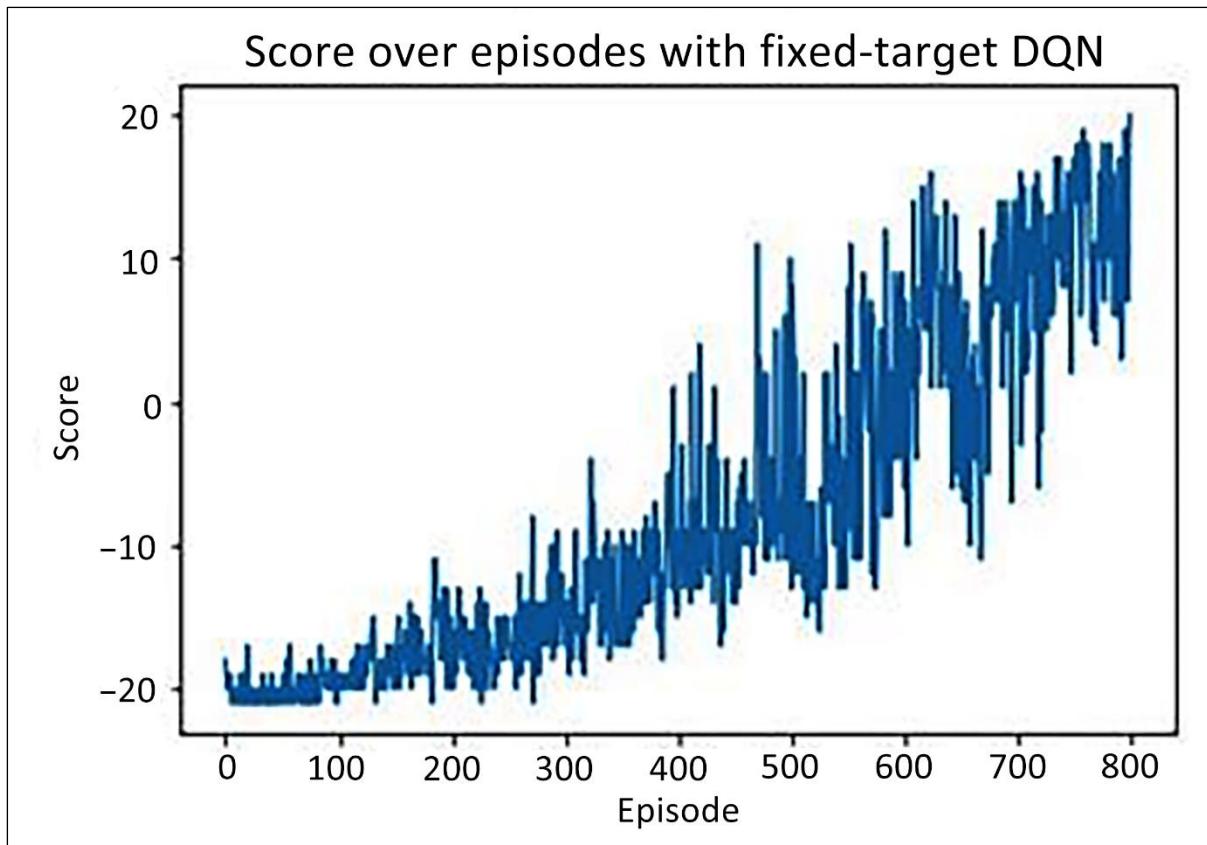


```
[[0.531441  0.59049  0.59049  0.531441 ],
 [0.531441  0.       0.6561   0.59049  ],
 [0.59049   0.729   0.59049  0.6561   ],
 [0.6561    0.       0.59049  0.59049  ],
 [0.59049   0.6561  0.       0.531441  ],
 [0.       0.       0.       0.       ],
 [0.       0.81    0.       0.6561  ],
 [0.       0.       0.       0.       ],
 [0.6561    0.       0.729   0.59049  ],
 [0.6561    0.81    0.81    0.       ],
 [0.729     0.9     0.       0.729   ],
 [0.       0.       0.       0.       ],
 [0.       0.       0.       0.       ],
 [0.       0.81    0.9     0.72899998],
 [0.80999997 0.9     1.       0.81    ],
 [0.       0.       0.       0.       ]]]
```

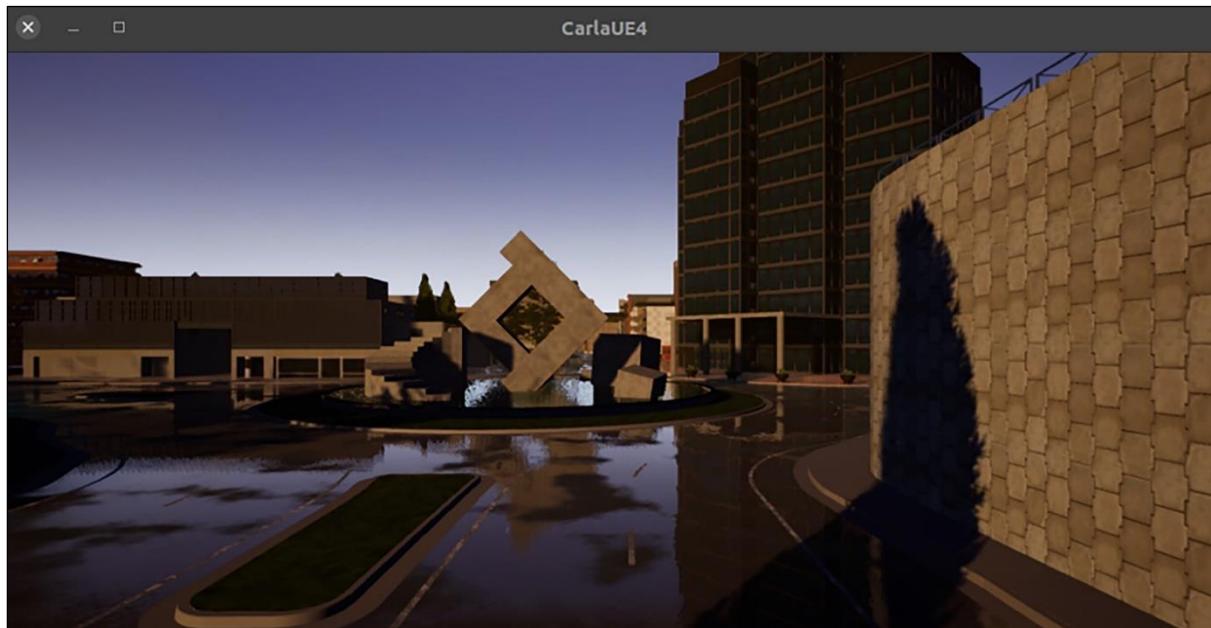


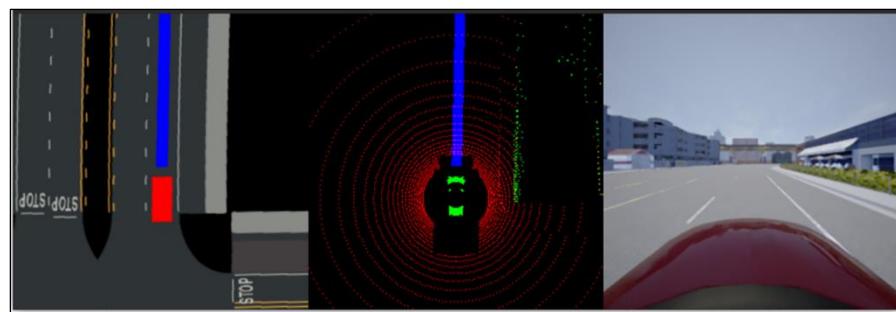




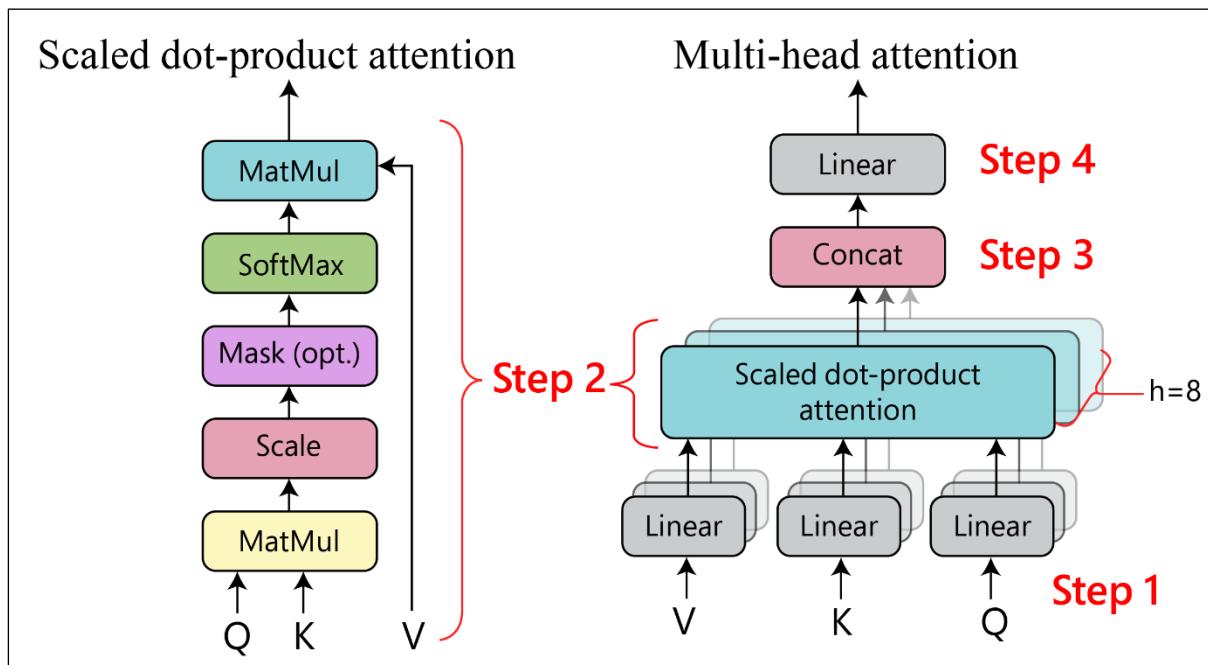
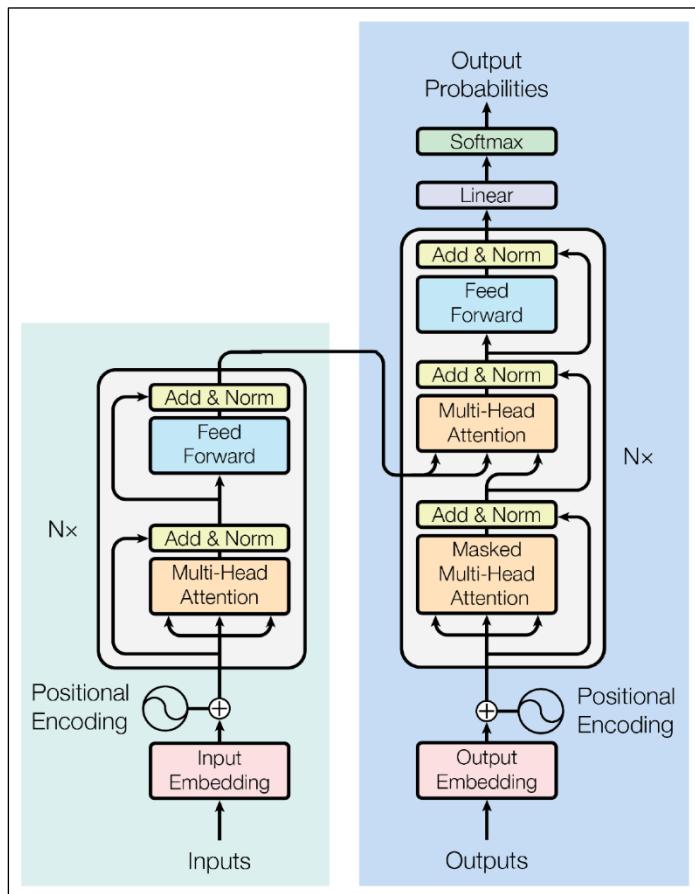


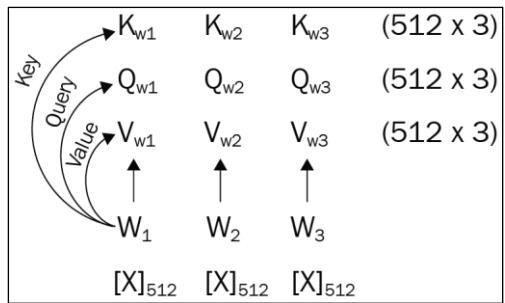
```
└ ipython -c "import carla; carla.__spec__"
out[1]: ModuleSpec(name='carla', loader=<_frozen_importlib_external.SourceFileLo
ader object at 0x7fb31646590>, origin='/home/yyr/anaconda3/lib/python3.7/site-p
ackages/carla-0.9.6-py3.5-linux-x86_64.egg/carla/__init__.py', submodule_search_
locations=['/home/yyr/anaconda3/lib/python3.7/site-packages/carla-0.9.6-py3.5-li
nux-x86_64.egg/carla'])
```





Chapter 15:

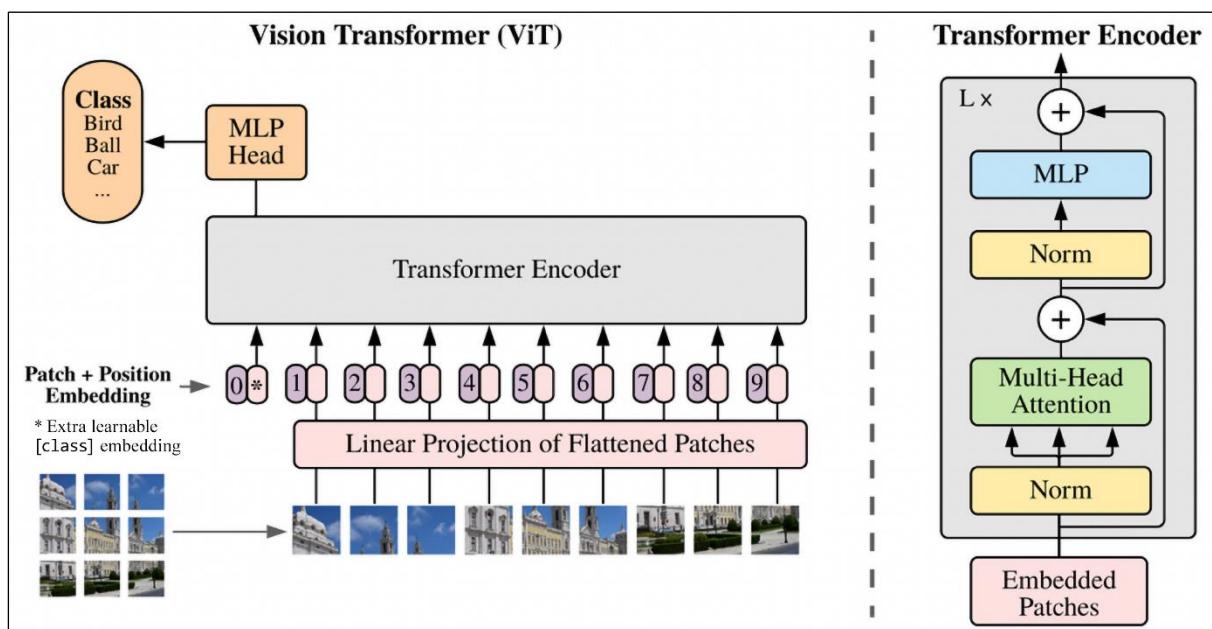


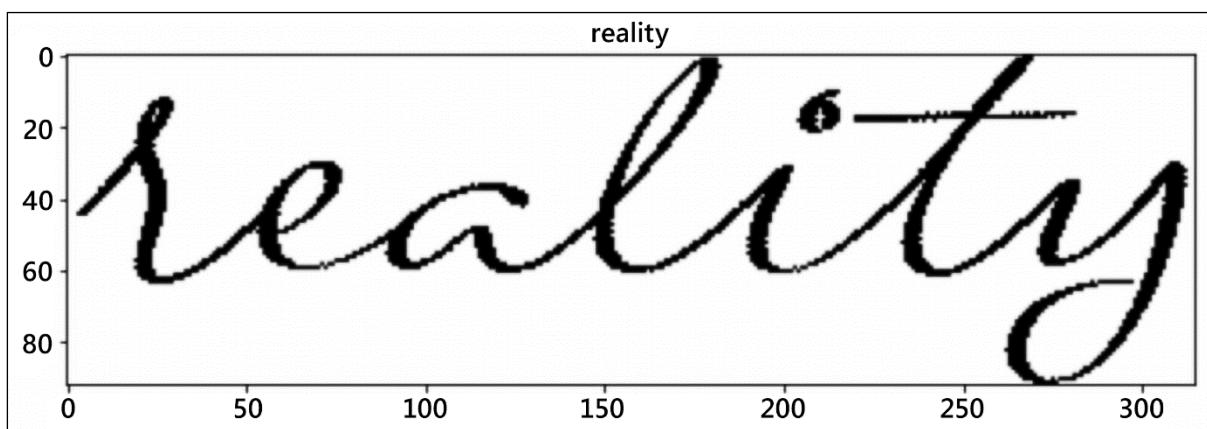
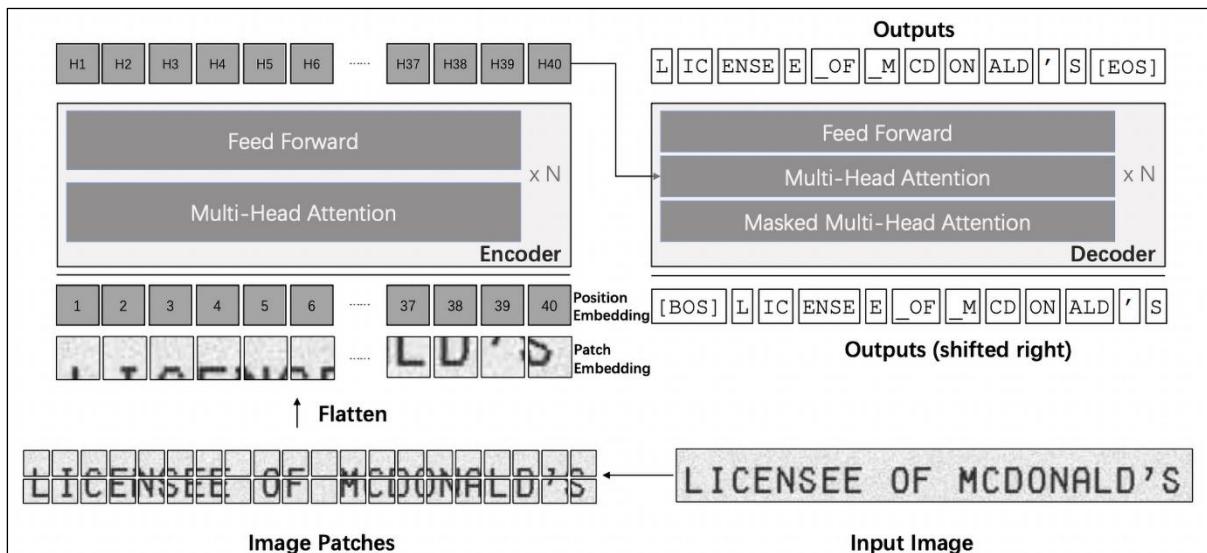
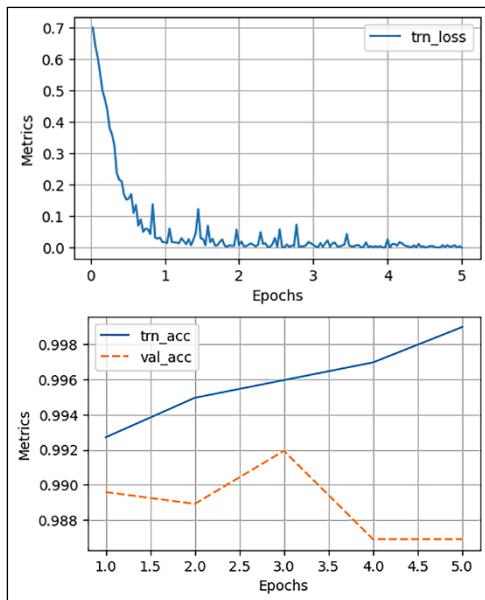


$$\begin{array}{llll}
 K_{w11} & K_{w21} & K_{w31} & (64 \times 3) = K_w \\
 Q_{w11} & Q_{w21} & Q_{w31} & (64 \times 3) = Q_w \\
 V_{w11} & V_{w21} & V_{w31} & (64 \times 3) = V_w
 \end{array}$$

$$\left(\begin{array}{cc} K_w^T & Q_w \\ (3 \times 64) & (64 \times 3) \\ (3 \times 3) & \end{array} \right) \rightarrow \left(\frac{K_w^T Q_w}{\sqrt{d_K}} \right) \rightarrow \text{Softmax} \left(\frac{K_w^T Q_w}{\sqrt{d_K}} \right)$$

$$\bigvee_w \text{Softmax} \left(\frac{K_w^T Q_w}{\sqrt{d_K}} \right) \rightarrow Z \quad (64 \times 3)$$



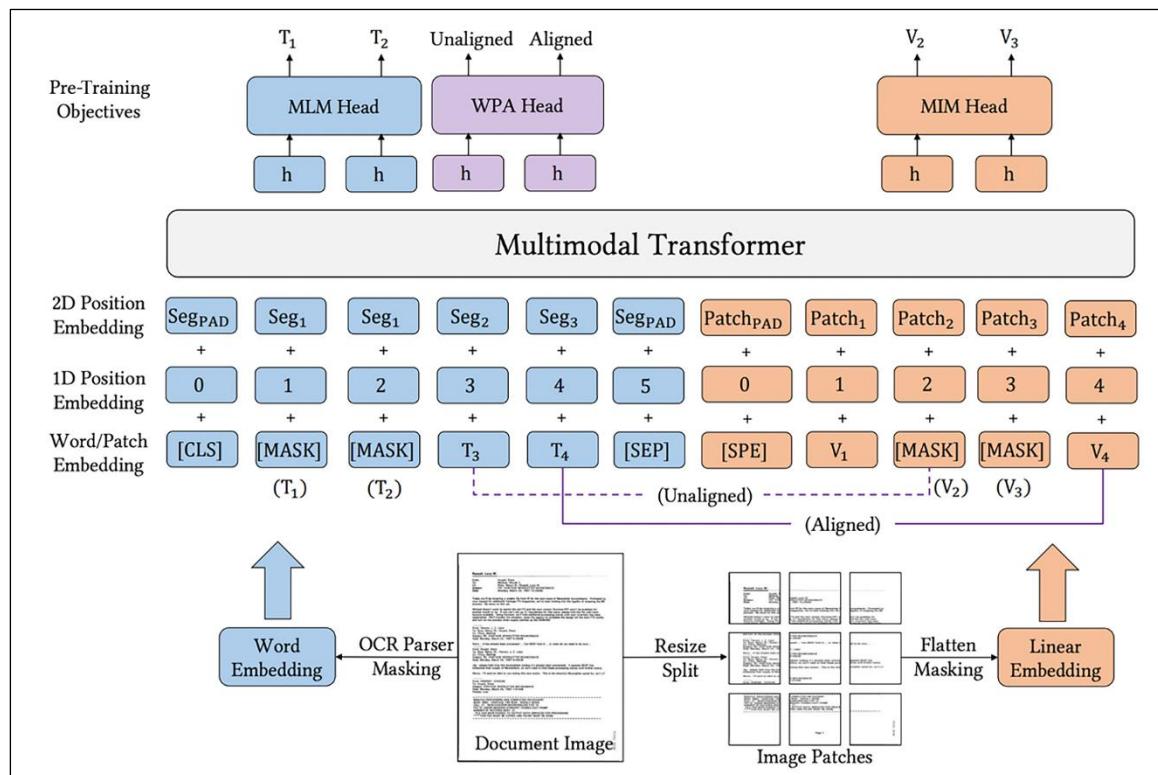
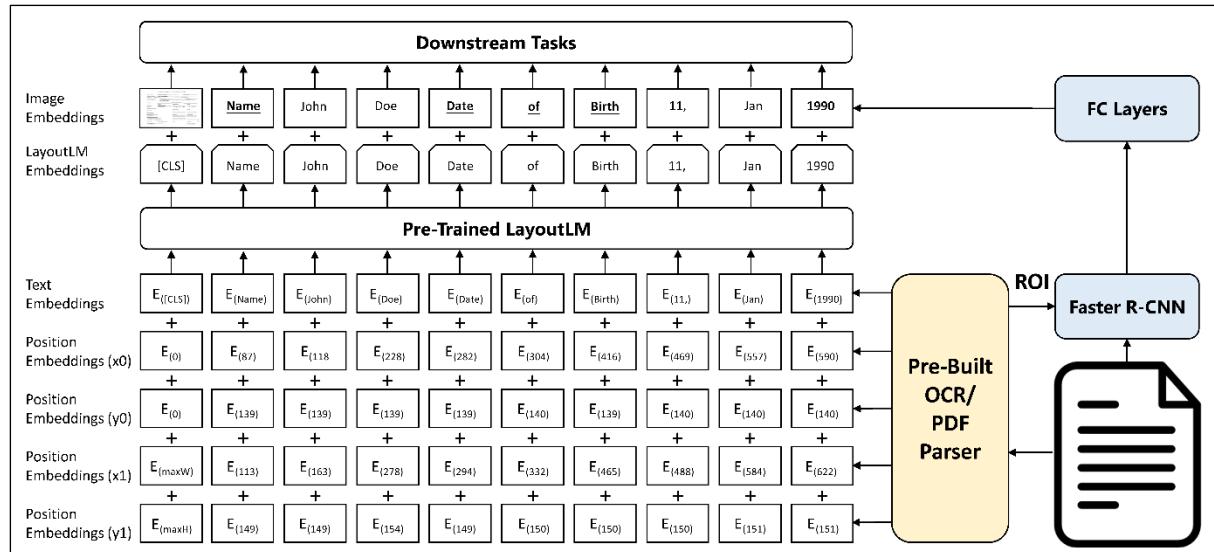


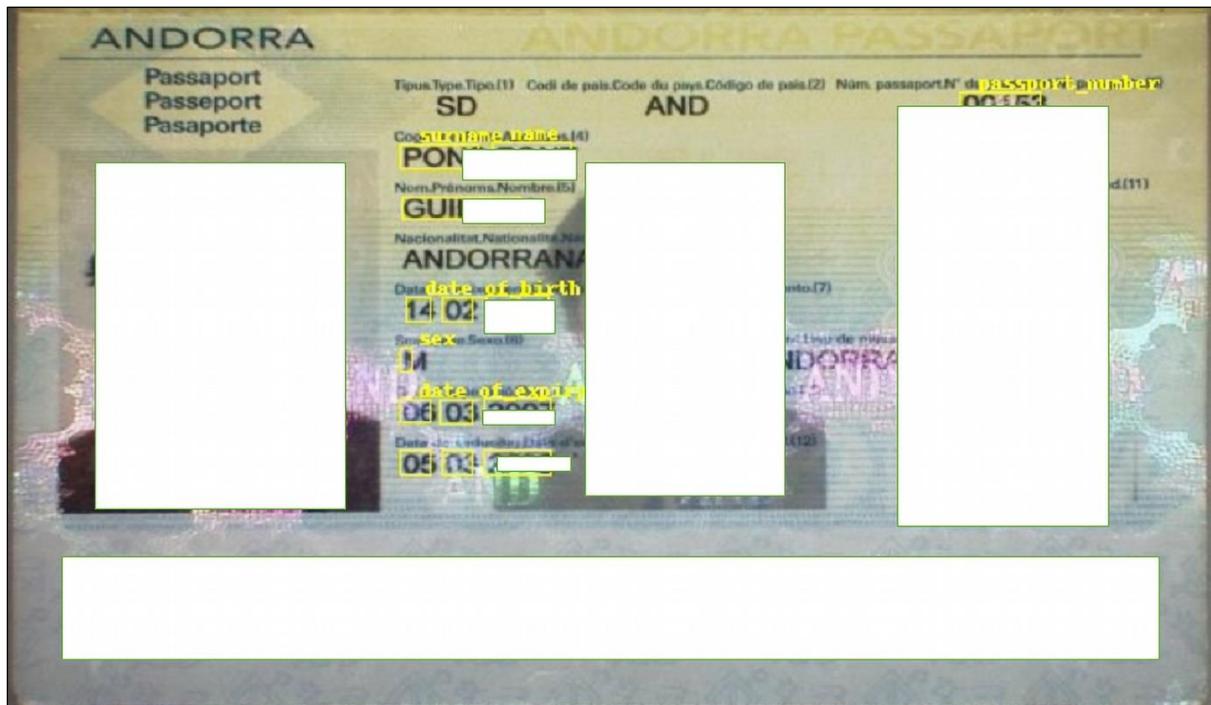
```
[ 'synthetic-data/already@dPkHBT.png',
  'synthetic-data/bring@OzNTFr.png',
  'synthetic-data/few@EhVOvU.png',
  'synthetic-data/research@cgo7rI.png',
  'synthetic-data/fast@bQ8Wkm.png' ]
```

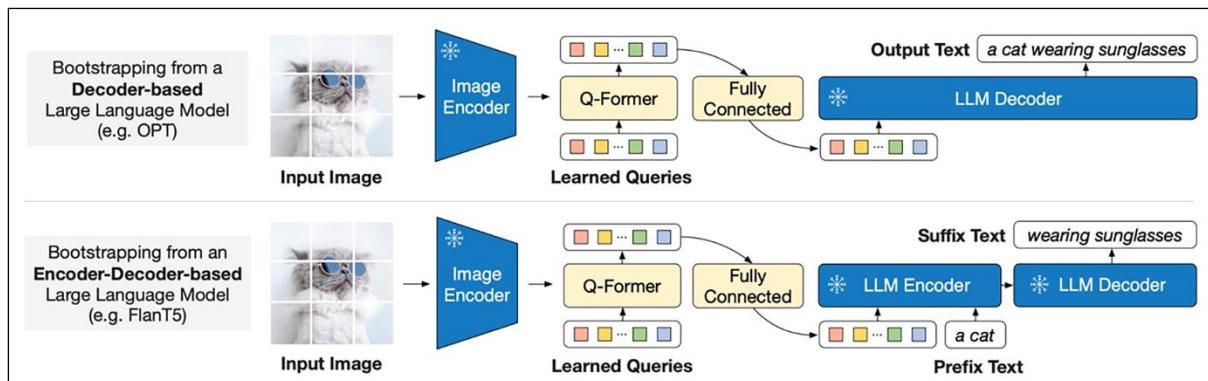
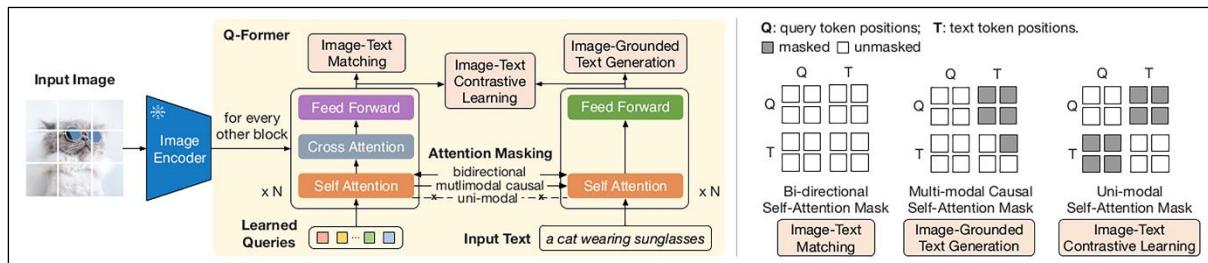
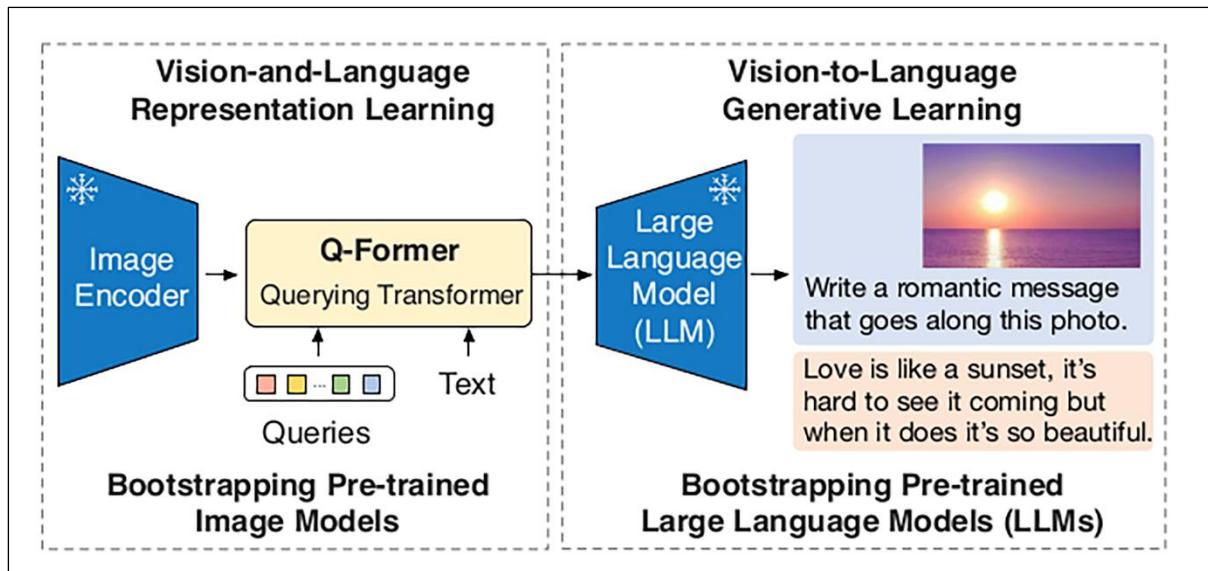
Step	Training Loss	Validation Loss	Cer
100	1.468400	1.485709	0.248651
200	1.201000	1.186254	0.245412
300	1.375000	1.082842	0.106513
400	1.034500	1.100040	0.167686
500	1.183800	1.077437	0.183879
600	0.777200	0.947086	0.169845
700	0.913500	1.018770	0.154372
800	0.759700	0.909072	0.077726
900	0.965700	0.894016	0.082044
1000	0.940400	0.830890	0.065491
1100	0.811900	0.762985	0.032746
1200	0.965600	0.873897	0.087801
1300	0.729500	0.851937	0.077726
1400	0.931800	0.811440	0.066931
1500	0.759000	0.789583	0.052177
1600	0.665700	0.761649	0.064772
1700	0.625800	0.721781	0.038143
1800	0.660300	0.713733	0.041742
1900	0.695700	0.720314	0.033105
2000	0.595400	0.661299	0.017992



Predicted Text: American

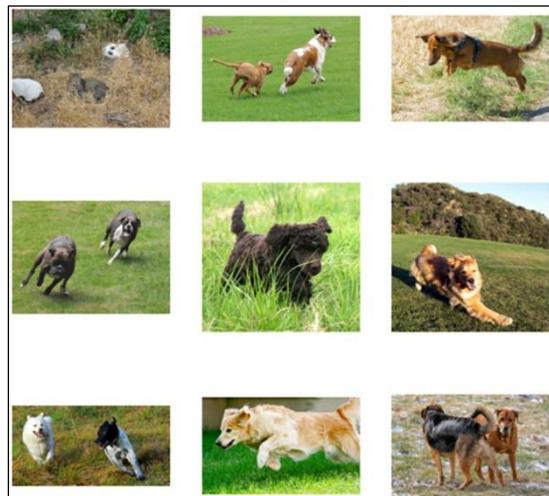
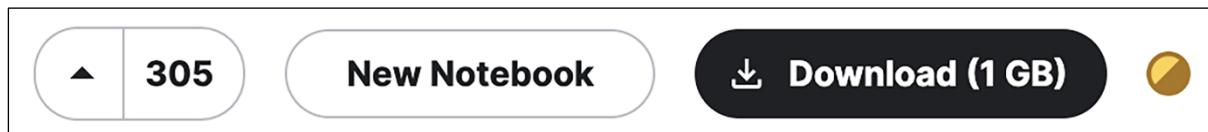
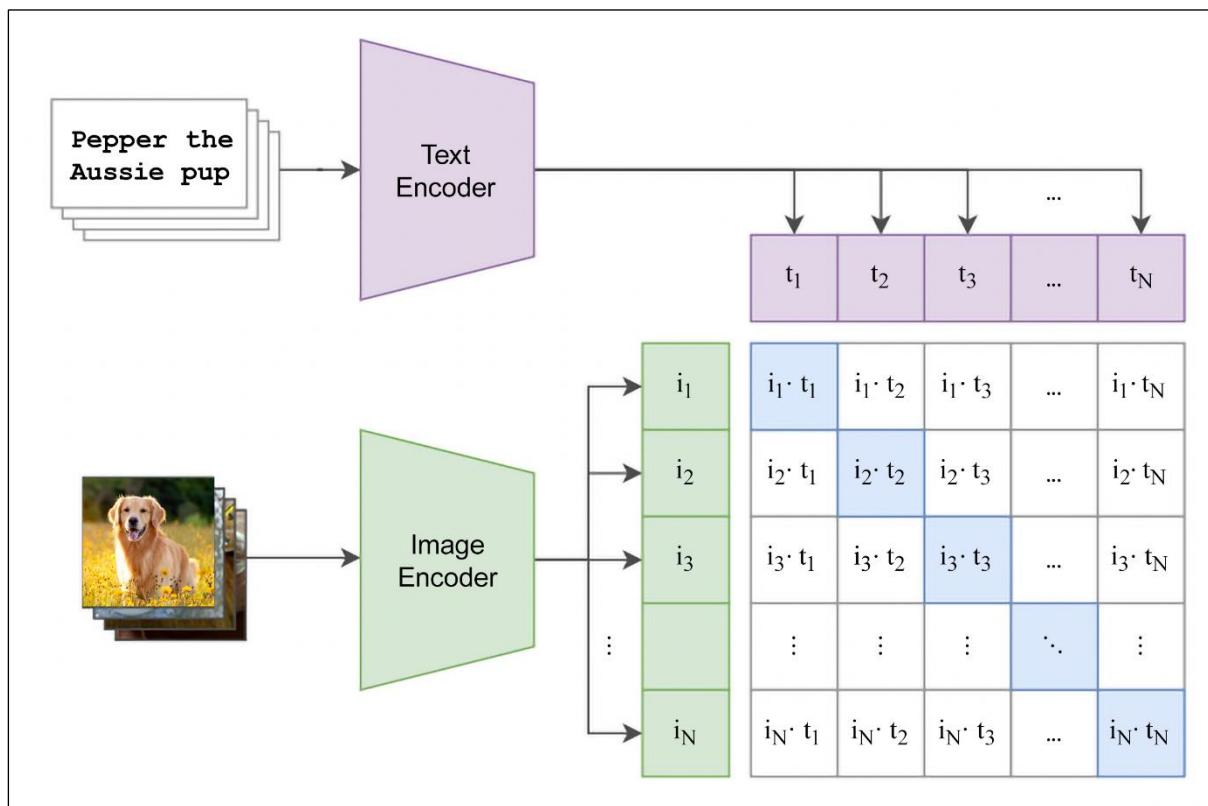


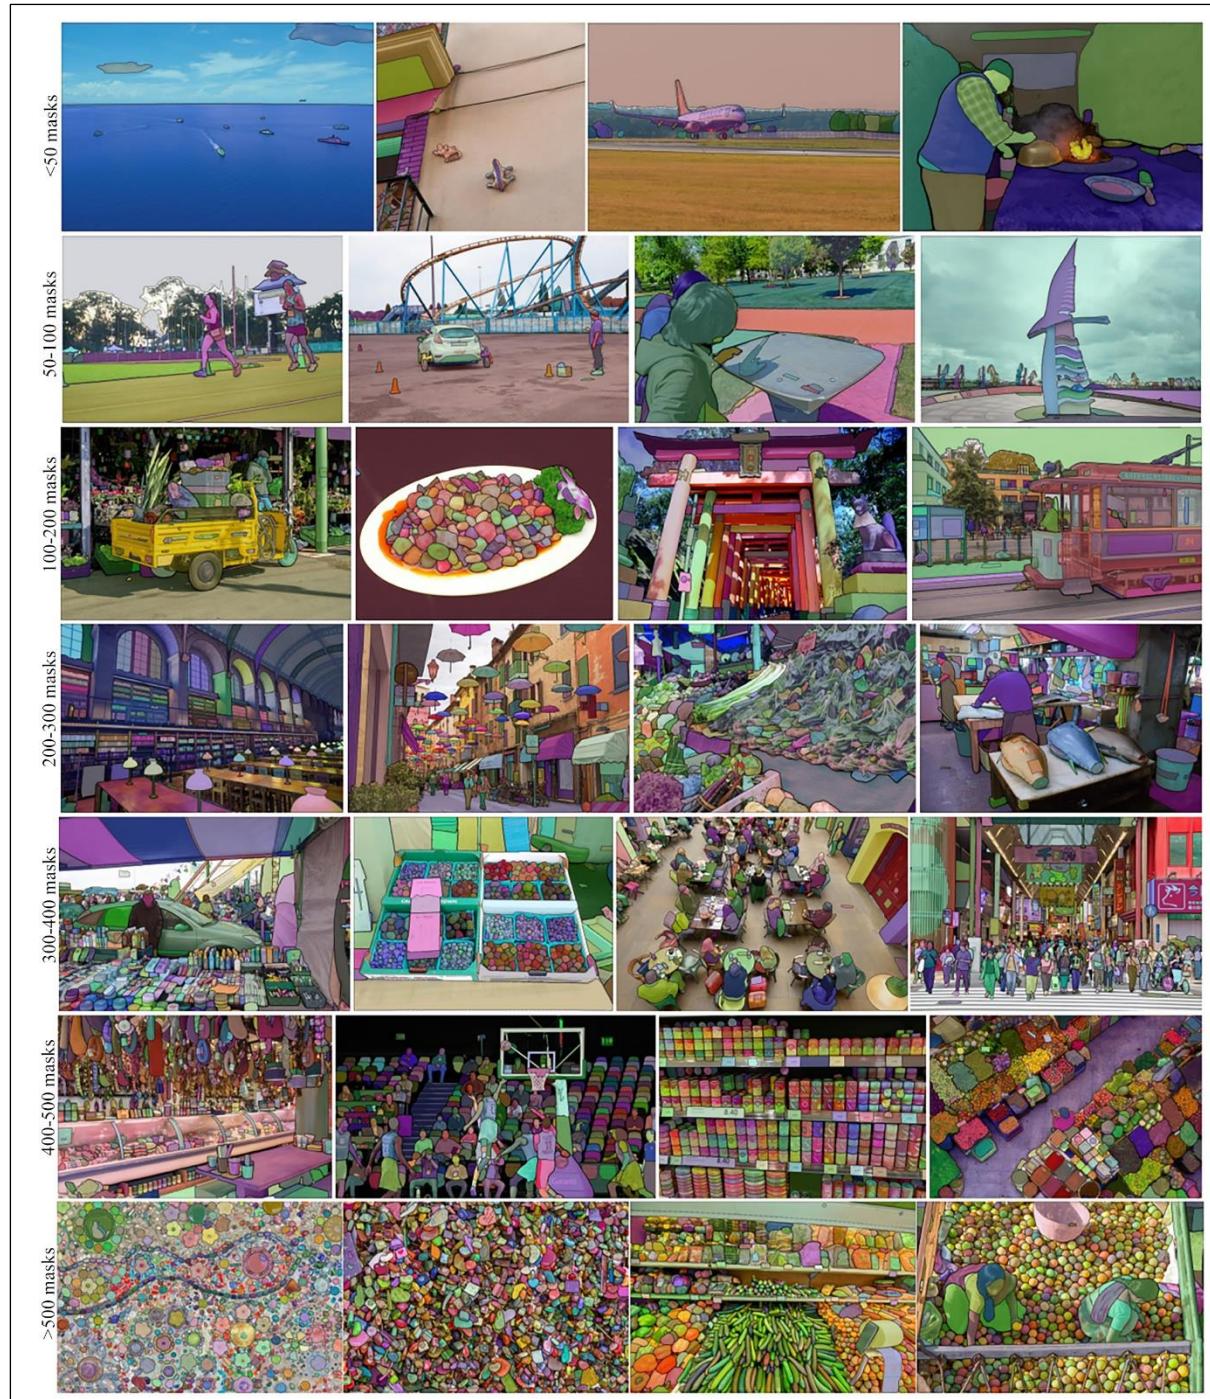
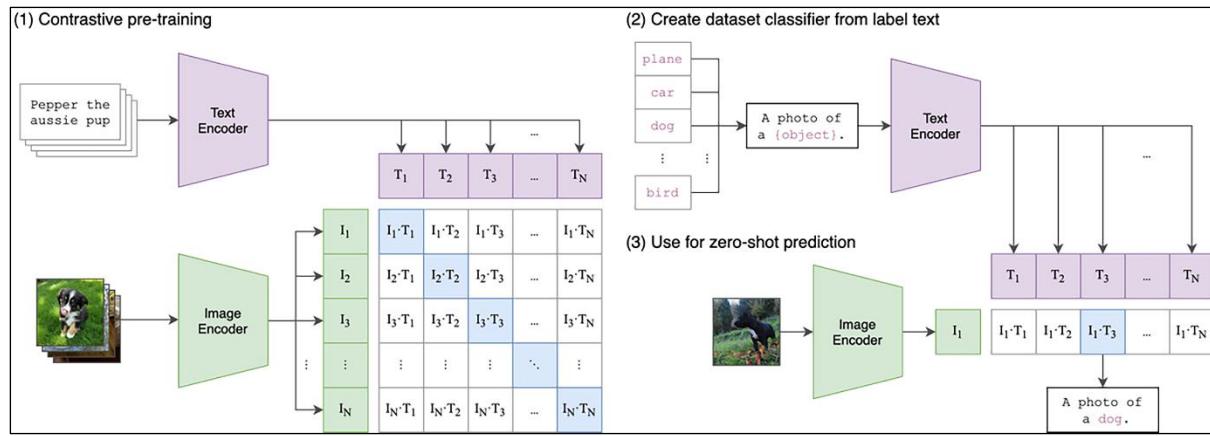


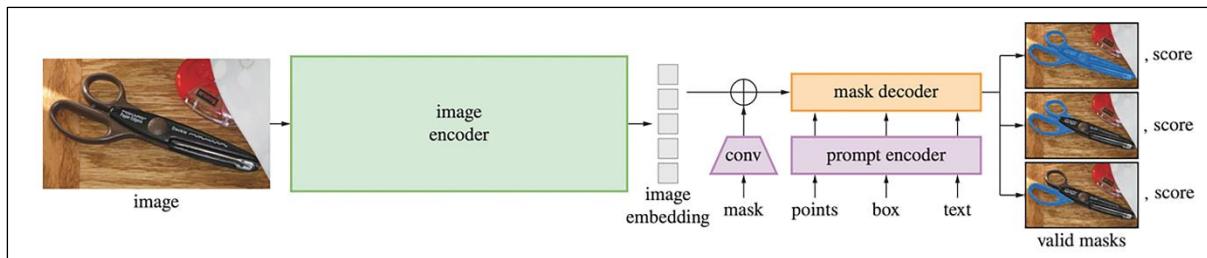




Chapter 16:







```

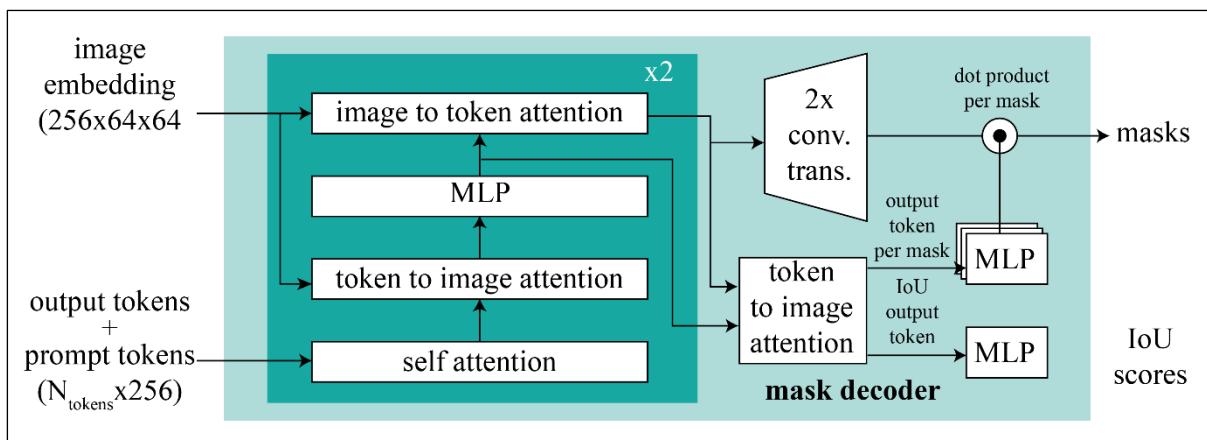
def forward(~
    self,
    points: Optional[Tuple[torch.Tensor, torch.Tensor]], ~
    boxes: Optional[torch.Tensor], ~
    masks: Optional[torch.Tensor], ~
):~
    bs = self._get_batch_size(points, boxes, masks)~
    sparse_embeddings = torch.empty((bs, 0, self.embed_dim), device=self._get_device())~
    if points is not None:~
        coords, labels = points~
        point_embeddings = self._embed_points(coords, labels, pad=(boxes is None))~
        sparse_embeddings = torch.cat([sparse_embeddings, point_embeddings], dim=1)~
    if boxes is not None:~
        box_embeddings = self._embed_boxes(boxes)~
        sparse_embeddings = torch.cat([sparse_embeddings, box_embeddings], dim=1)~
    if masks is not None:~
        dense_embeddings = self._embed_masks(masks)~
    else:~
        dense_embeddings = self.no_mask_embed.weight.reshape(1, -1, 1, 1).expand(~
            bs, -1, self.image_embedding_size[0], self.image_embedding_size[1]~
        )~
    ~
    return sparse_embeddings, dense_embeddings.

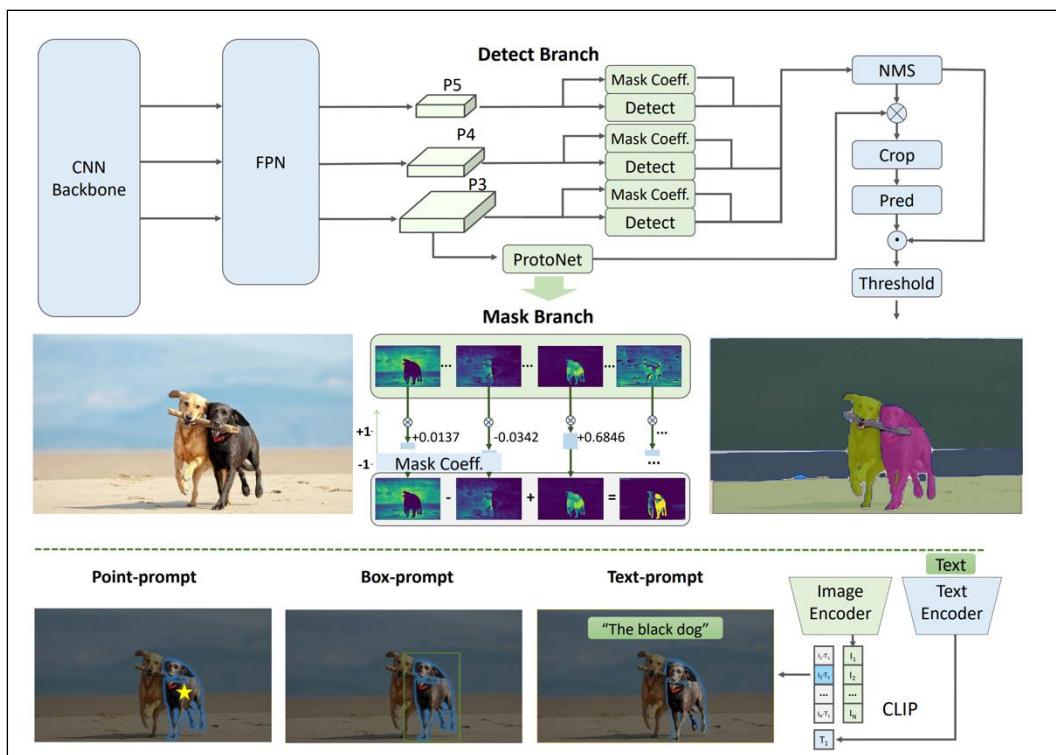
```

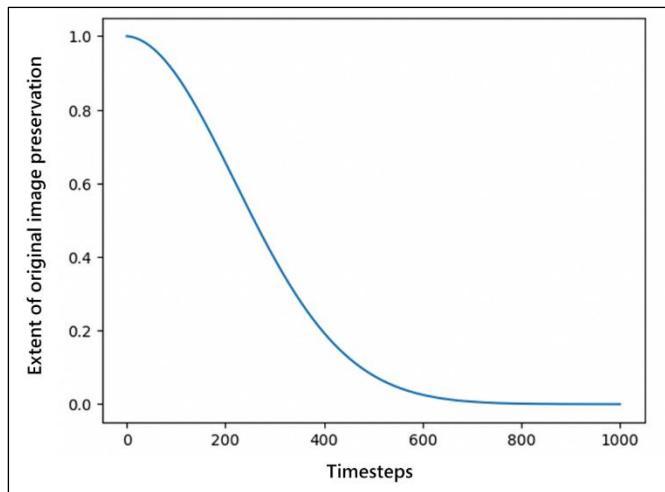
```

Module Name: PromptEncoder~
Input Kwargs: ~
    points:~
        0 - tensor[64, 1, 2] f64 n=128 (1Kb) x€[10.672, 672.328] μ=186.750 σ=209.371 cuda:0 - ID:#5b3b2d-~
        1 - tensor[64, 1] i32 x€[1, 1] μ=1.000 σ=0. cuda:0 - ID:#08f4ad-~
    boxes - NoneType~
    masks - NoneType~
~
Input Args:~
~
Outputs: ~
    sparse_embeddings - tensor[64, 2, 256] n=32768 (0.1Mb) x€[-1.352, 1.358] μ=-0.006 σ=0.530 cuda:0 - ID:#ca373f-~
    dense_embeddings - tensor[64, 256, 64, 64] n=67108864 (0.2Gb) x€[-0.218, 0.138] μ=-0.001 σ=0.031 grad cuda:0 - ID:#3c6848-

```

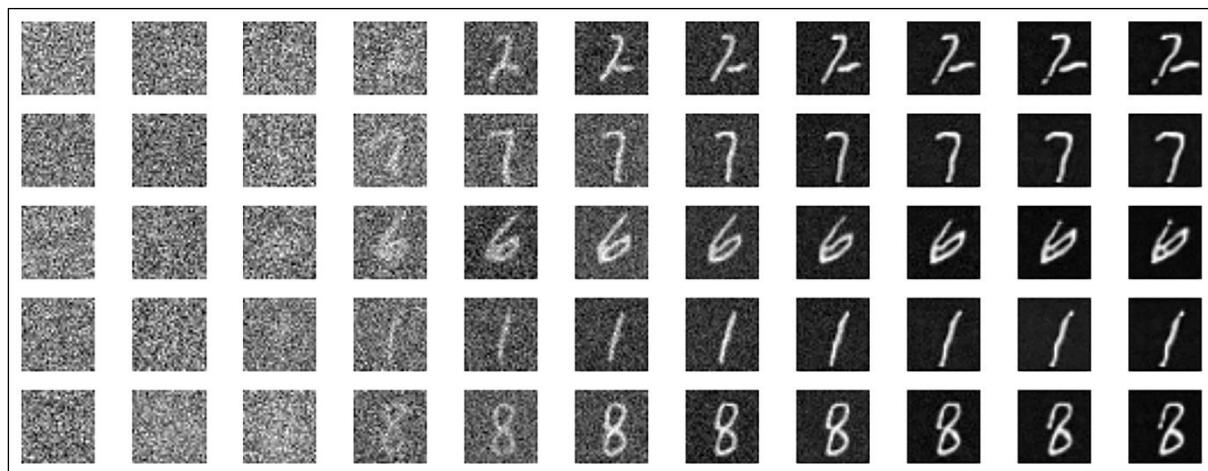
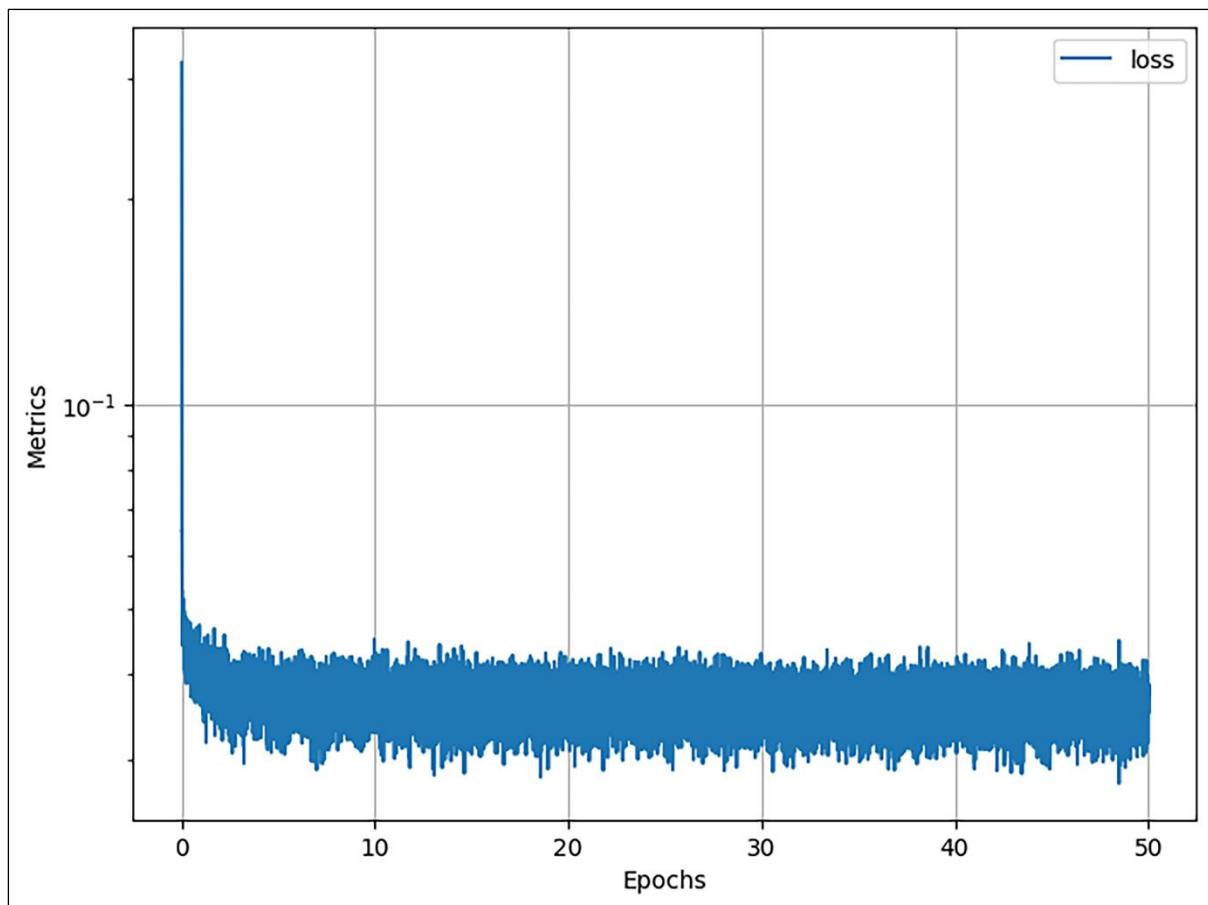


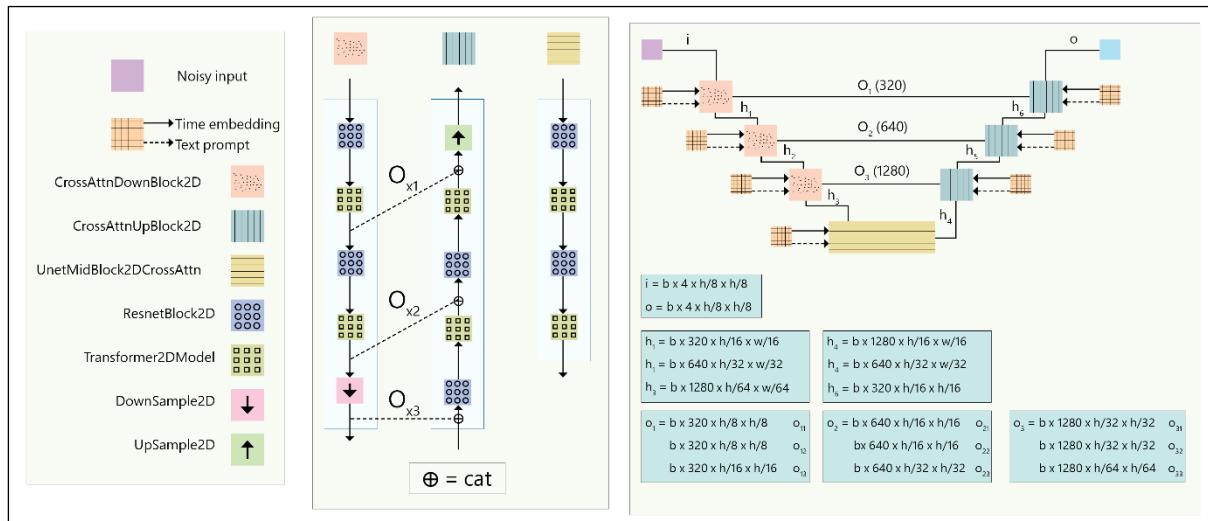
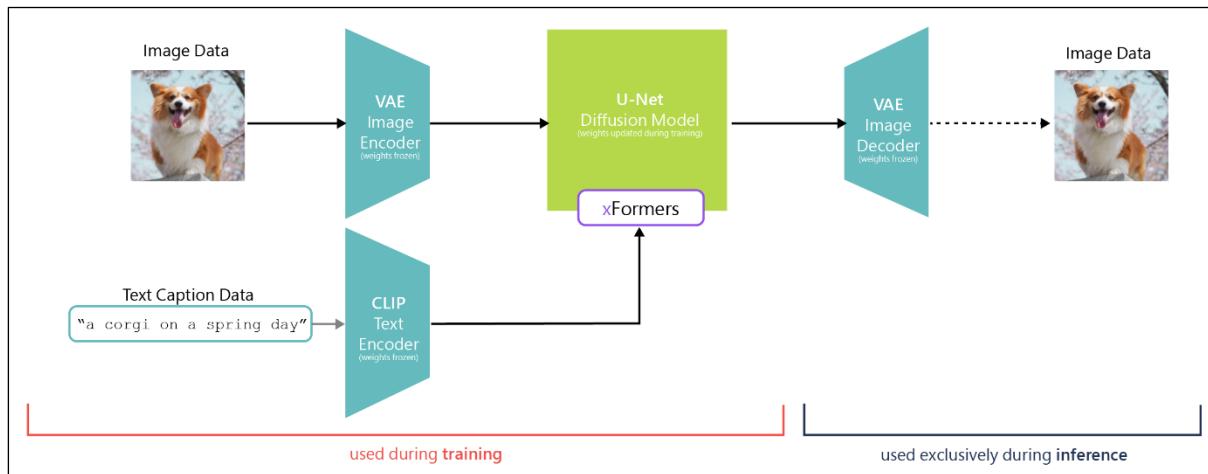




```
=====
Layer (type:depth-idx)           Param #
=====
UNet2DModel                      --
├Conv2d: 1-1                      320
├Timesteps: 1-2                   --
├TimestepEmbedding: 1-3
│└Linear: 2-1                     4,224
│└SiLU: 2-2                       --
│└Linear: 2-3                     16,512
├ModuleList: 1-4                  --
│└DownBlock2D: 2-4                32,000
│└AttnDownBlock2D: 2-5             119,680
│└AttnDownBlock2D: 2-6             460,544
│└AttnDownBlock2D: 2-7             1,215,744
├ModuleList: 1-5                  --
│└AttnUpBlock2D: 2-8              4,661,248
│└AttnUpBlock2D: 2-9              1,347,840
│└AttnUpBlock2D: 2-10             346,240
│└UpBlock2D: 2-11                 78,528
├UNetMidBlock2D: 1-6              --
│└ModuleList: 2-12                263,680
│└ModuleList: 2-13                2,428,416
├GroupNorm: 1-7                   64
├SiLU: 1-8                        --
└Conv2d: 1-9                      289
=====

Total params: 10,975,329
Trainable params: 10,975,329
Non-trainable params: 0
=====
```





```
1 ======+-----+  
2 Layer ·(type:depth-idx) ..... Param #  
3 ======+-----+  
4 UNet2DConditionModel .....  
5 |Conv2d: 1-1 ..... 11,840  
6 |Timesteps: 1-2 .....  
7 |TimestepEmbedding: 1-3 .....  
8 | |LoRACompatibleLinear: 2-1 ..... 410,880  
9 | |SiLU: 2-2 .....  
10 | |LoRACompatibleLinear: 2-3 ..... 1,639,680  
11 |ModuleList: 1-4 .....  
12 | |CrossAttnDownBlock2D: 2-4 ..... 10,852,160  
13 | |CrossAttnDownBlock2D: 2-5 ..... 37,473,920  
14 | |CrossAttnDownBlock2D: 2-6 ..... 141,303,040  
15 | |DownBlock2D: 2-7 ..... 62,277,120  
16 |ModuleList: 1-5 .....  
17 | |UpBlock2D: 2-8 ..... 162,241,280  
18 | |CrossAttnUpBlock2D: 2-9 ..... 260,296,960  
19 | |CrossAttnUpBlock2D: 2-10 ..... 72,396,160  
20 | |CrossAttnUpBlock2D: 2-11 ..... 19,302,080  
21 |UNetMidBlock2DCrossAttn: 1-6 .....  
22 | |ModuleList: 2-12 ..... 35,416,320  
23 | |ModuleList: 2-13 ..... 62,277,120  
24 |GroupNorm: 1-7 ..... 640  
25 |SiLU: 1-8 .....  
26 |Conv2d: 1-9 ..... 11,524  
27 ======+-----+  
28 Total ·params: 865,910,724  
29 Trainable ·params: 865,910,724  
30 Non-trainable ·params: 0  
31 ======+-----+
```

```
1 ======+-----+  
2 Layer ·(type:depth-idx) ..... Param #  
3 ======+-----+  
4 CrossAttnDownBlock2D .....  
5 |ModuleList: 1-1 .....  
6 | |Transformer2DModel: 2-1 ..... 2,710,080  
7 | |Transformer2DModel: 2-2 ..... 2,710,080  
8 |ModuleList: 1-2 .....  
9 | |ResnetBlock2D: 2-3 ..... 2,255,040  
10 | |ResnetBlock2D: 2-4 ..... 2,255,040  
11 |ModuleList: 1-3 .....  
12 | |Downsample2D: 2-5 ..... 921,920  
13 ======+-----+  
14 Total ·params: 10,852,160  
15 Trainable ·params: 10,852,160  
16 Non-trainable ·params: 0  
17 ======+-----+
```

```
1 hidden_states-->tensor[2,-320,-96,-96]·n=5898240 (22Mb) ·x€[-3.921,-3.216]·μ=0.004·σ=0.351·cuda:0·
2 temb-->tensor[2,-1280]·n=2560 (10Kb) ·x€[-2.014,-5.339]·μ=-0.003·σ=0.256·cuda:0·
3 encoder_hidden_states-->tensor[2,-77,-1024]·n=157696 (0.6Mb) ·x€[-6.571,-13.023]·μ=-0.169·σ=1.032·cuda:0
```

```
1 ....def forward(·
2 .....··self,·
3 .....··hidden_states:·torch.FloatTensor,·
4 .....··temb:·Optional[torch.FloatTensor] =·None,·
5 .....··encoder_hidden_states:·Optional[torch.FloatTensor] =·None,·
6 ....):·
7 .....··output_states =·()·
8 .....··for i, (resnet, attn)·in·enumerate(blocks):·
9 .....···hidden_states = resnet(hidden_states,·temb)·
10 .....···hidden_states = attn(·
11 .....····hidden_states,·
12 .....····encoder_hidden_states=encoder_hidden_states,·
13 .....···)@[0]·
14 .....···output_states = output_states +·(hidden_states,)·
15 .....··if self.downsamplers is not·None:·
16 .....···for downampler in self.downsamplers:·
17 .....····hidden_states =·downampler(hidden_states,·scale=lora_scale)·
18 .....···output_states = output_states +·(hidden_states,)·
19 .....··return hidden_states,·output_states
```

```
hidden_states-->tensor[2,-320,-48,-48]·n=1474560 (5.6Mb) ·x€[-13.873,-16.032]·μ=0.017·σ=2.290·cuda:0·
output_states-->·
.....1-->tensor[2,-320,-96,-96]·n=5898240 (22Mb) ·x€[-8.291,-6.062]·μ=-0.084·σ=0.641·cuda:0·
.....2-->tensor[2,-320,-96,-96]·n=5898240 (22Mb) ·x€[-9.585,-7.211]·μ=-0.082·σ=0.963·cuda:0·
.....3-->tensor[2,-320,-48,-48]·n=1474560 (5.6Mb) ·x€[-13.873,-16.032]·μ=0.017·σ=2.290·cuda:0·
```

```
1 hidden_states-->tensor[2,-1280,-12,-12]·n=368640 (1.4Mb) ·x€[-43.398,-32.103]·μ=-0.329·σ=4.600·cuda:0·
2 temb-->tensor[2,-1280]·n=2560 (10Kb) ·x€[-2.014,-5.339]·μ=-0.003·σ=0.256·cuda:0·
3 encoder_hidden_states-->tensor[2,-77,-1024]·n=157696 (0.6Mb) ·x€[-6.571,-13.023]·μ=-0.169·σ=1.032·cuda:0
```

```
1 .....def forward(‐
2 .....    self,‐
3 .....    hidden_states: torch.FloatTensor,‐
4 .....    temb: Optional[torch.FloatTensor] = None,‐
5 .....    encoder_hidden_states: Optional[torch.FloatTensor] = None,‐
6 .....):‐
7 .....    hidden_states = self.resnets[0](hidden_states, temb)‐
8 .....    for attn, resnet in zip(self.attentions, self.resnets[1:]):‐
9 .....        hidden_states = attn(‐
10 .....            hidden_states,‐
11 .....            encoder_hidden_states=encoder_hidden_states,‐
12 .....            )[0]‐
13 .....        hidden_states = resnet(hidden_states, temb)‐
14 .....    return hidden_states
```

```
1 tensor[2, 1280, 12, 12] n=368640 (1.4Mb) xE[-43.307, 34.261] μ=-0.404 σ=4.882 cuda:0
```

```
1. hidden_states -- tensor[2, 640, 96, 96] n=11796480 (45Mb) xE[-486.840, 68.294] μ=-0.102 σ=3.369 cuda:0
2. temb -- tensor[2, 1280] n=2560 (10Kb) xE[-2.014, -5.339] μ=-0.003 σ=0.256 cuda:0
3. res_hidden_states_tuple=
.....1 -- tensor[2, 320, 96, 96] n=5898240 (22Mb) xE[-3.706, 3.410] μ=0.003 σ=0.351 cuda:0
.....2 -- tensor[2, 320, 96, 96] n=5898240 (22Mb) xE[-8.291, 6.062] μ=-0.084 σ=0.641 cuda:0
.....3 -- tensor[2, 320, 96, 96] n=5898240 (22Mb) xE[-9.585, 7.211] μ=-0.082 σ=0.963 cuda:0
4. encoder_hidden_states -- tensor[2, 77, 1024] n=157696 (0.6Mb) xE[-6.571, 13.023] μ=-0.169 σ=1.032 cuda:0
```

```
...def forward(...
...    self,
...    hidden_states: torch.FloatTensor,
...    res_hidden_states_tuple: Tuple[torch.FloatTensor, ...],
...    temb: Optional[torch.FloatTensor] = None,
...    encoder_hidden_states: Optional[torch.FloatTensor] = None,
...):
...    for resnet, attn in zip(self.resnets, self.attentions):
...        # pop res hidden states
...        res_hidden_states = res_hidden_states_tuple[-1]
...        res_hidden_states_tuple = res_hidden_states_tuple[:-1]
...        hidden_states = torch.cat([hidden_states, res_hidden_states], dim=1)
...        hidden_states = resnet(hidden_states, temb, scale=lora_scale)
...        hidden_states = attn(
...            hidden_states,
...            encoder_hidden_states=encoder_hidden_states,
...            return_dict=False,
...        )[0]
...
...    if self.upsamplers is not None:
...        for upsample in self.upsamplers:
...            hidden_states = upsample(hidden_states, upsample_size, scale=lora_scale)
...
...    return hidden_states
```

```
tensor[2, 320, 96, 96] n=5898240 (22Mb) xE[-29.935, 52.374] μ=-0.834 σ=4.356 cuda:0
```

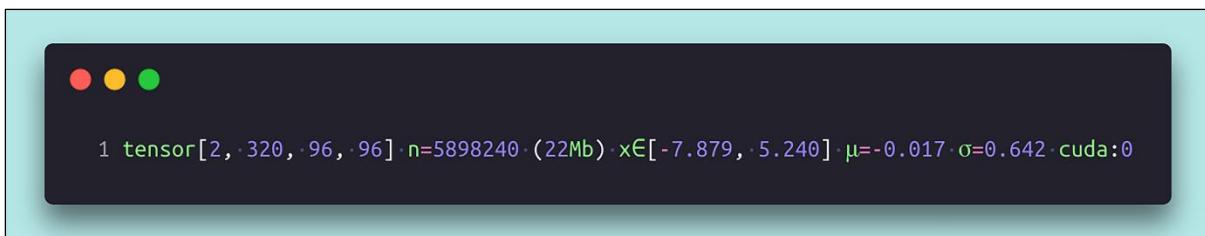
```
1 hidden_states - tensor[2, 320, 96, 96] n=5898240 (22Mb) xE[-10.410, 6.776] μ=-0.085 σ=0.782 cuda:0
2 encoder_hidden_states - tensor[2, 77, 1024] n=157696 (0.6Mb) xE[-6.571, 13.023] μ=-0.169 σ=1.032 cuda:0
```

```
1 .....def forward(‐
2 .....    self,‐
3 .....    hidden_states: torch.Tensor,‐
4 .....    encoder_hidden_states: Optional[torch.Tensor] = None,‐
5 .....):‐
6 .....    batch, _, height, width = hidden_states.shape‐
7 .....    residual = hidden_states‐
8 .....    hidden_states = self.norm(hidden_states)‐
9 .....    hidden_states = self.proj_in(hidden_states) #·nn.Conv2D‐
10 .....   inner_dim = hidden_states.shape[1]‐
11 .....  hidden_states = hidden_states.permute(0, -2, -3, -1).reshape(batch, height * width, inner_dim)‐
12 .....  for block in self.transformer_blocks:‐
13 .....      hidden_states = block(‐
14 .....          hidden_states,‐
15 .....          attention_mask=attention_mask,‐
16 .....          encoder_hidden_states=encoder_hidden_states,‐
17 .....      )‐
18 .....  hidden_states = self.proj_out(hidden_states) #·nn.Conv2D‐
19 .....  hidden_states = hidden_states.reshape(batch, height, width, inner_dim).permute(0, 3, 1, 2).contiguous()‐
20 .....  output = hidden_states + residual‐
21 .....  return Transformer2DModelOutput(sample=output)
```

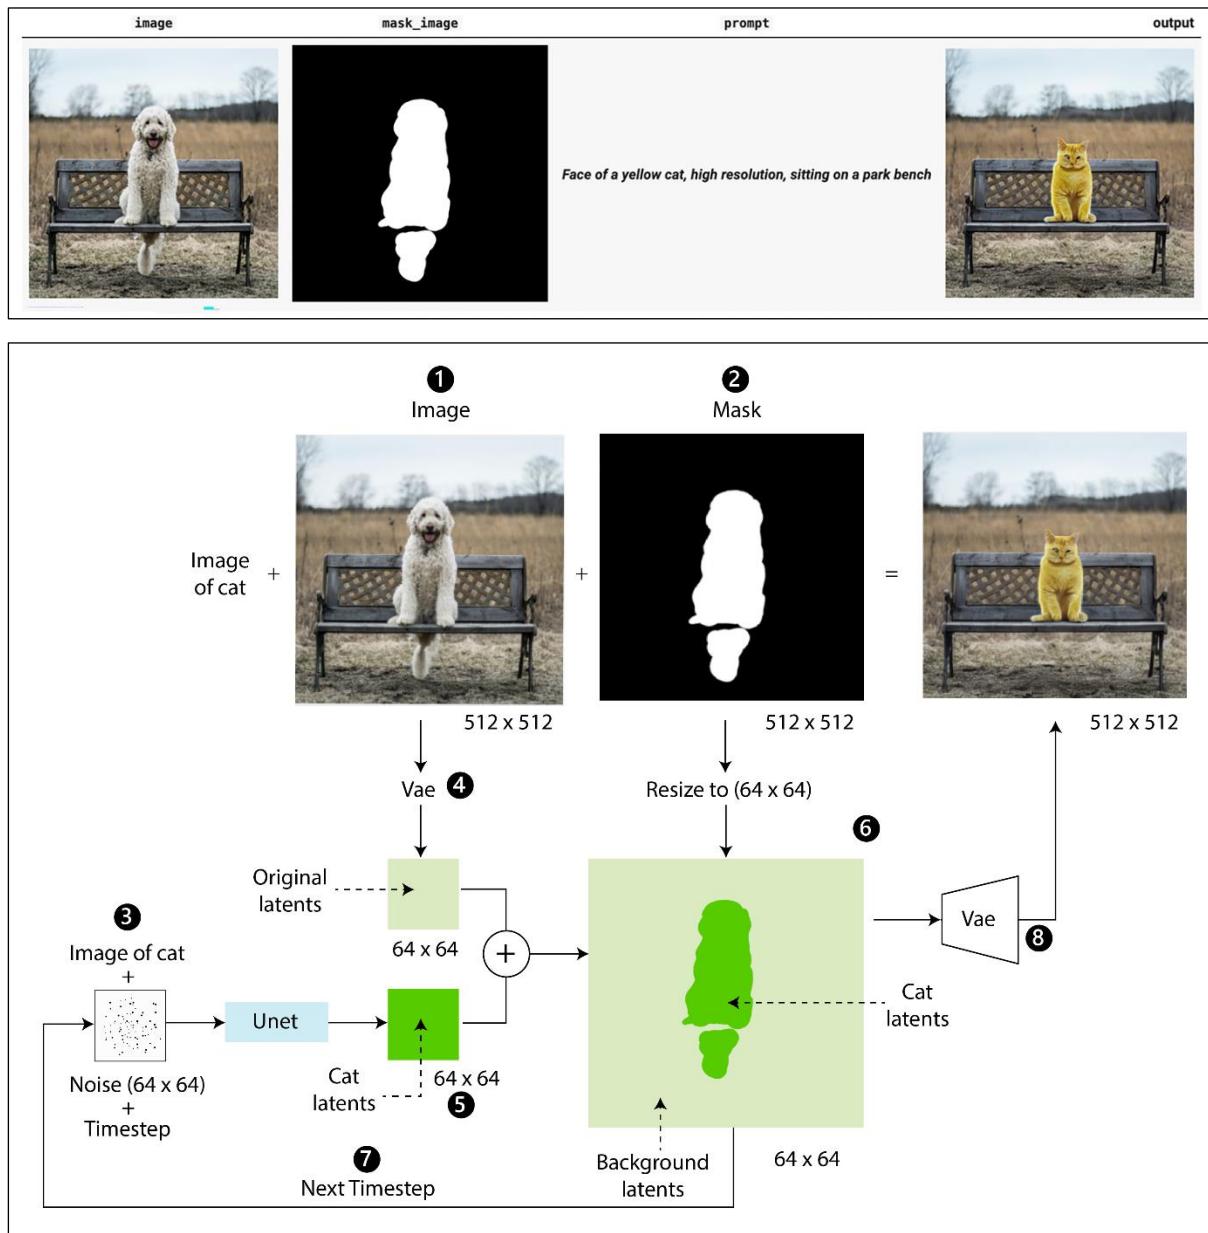
```
1 tensor[2, · 320, · 96, · 96] · n=5898240 · (22Mb) · x€[ · -30.745, · 52.179] · μ=-0.837 · σ=4.354 · cuda:0
```

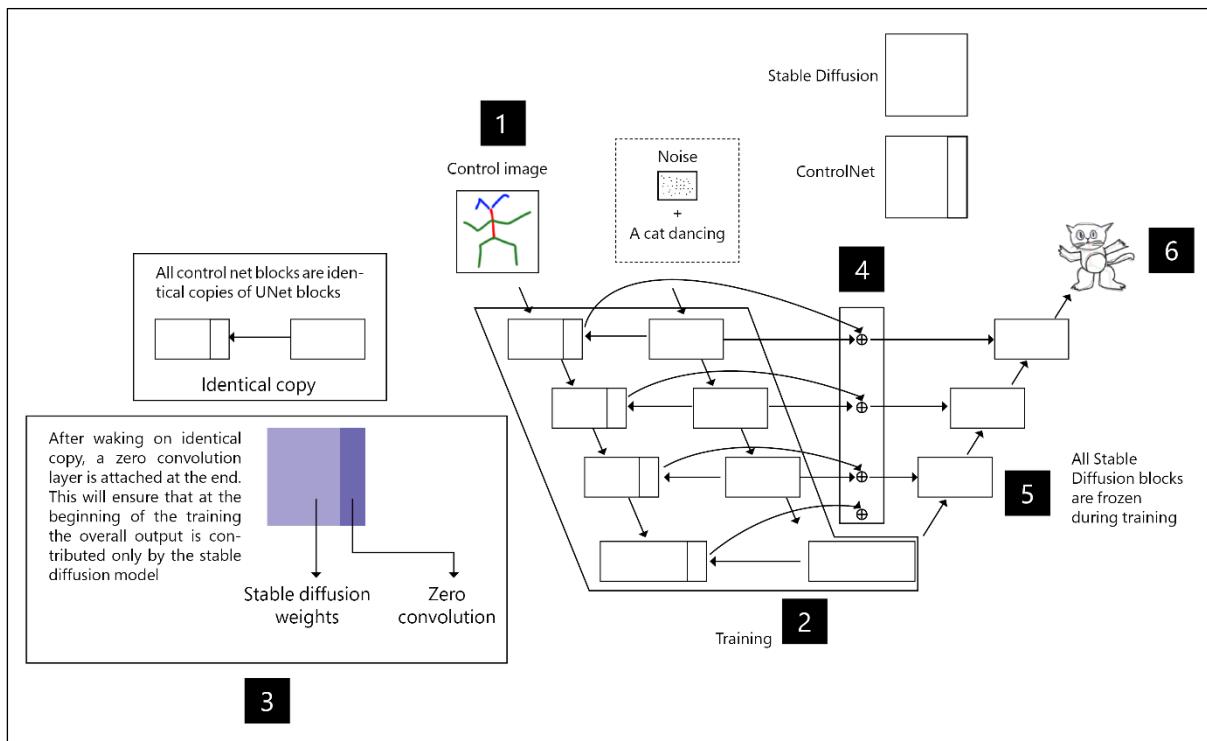
```
1 hidden_states · tensor[2, · 320, · 96, · 96] · n=5898240 · (22Mb) · x€[ · -3.619, · 3.284] · μ=0.004 · σ=0.348 · cuda:0‐
2 temb · tensor[2, · 1280] · n=2560 · (10Kb) · x€[ · -2.014, · 5.339] · μ=-0.003 · σ=0.256 · cuda:0
```

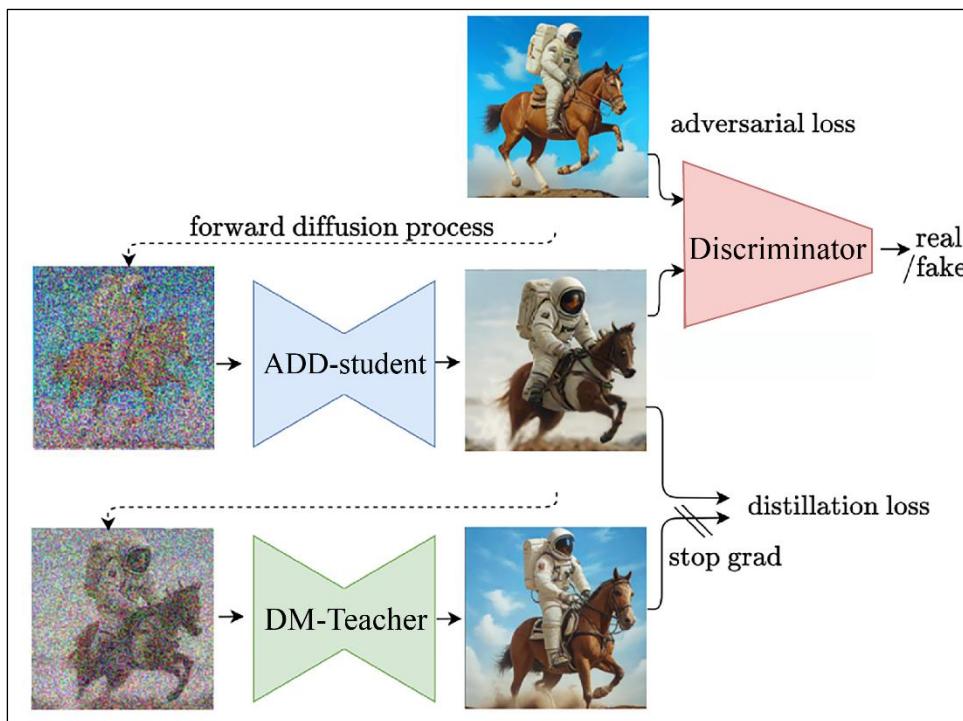
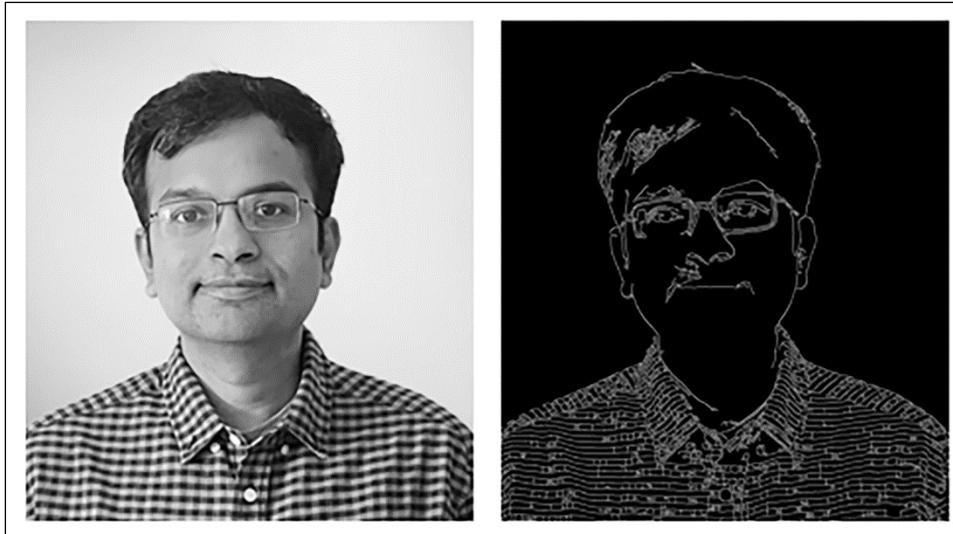
```
1 def forward(‐
2 .....    self,‐
3 .....    input_tensor: torch.FloatTensor,‐
4 .....    temb: torch.FloatTensor,‐
5 .....):‐
6 .....    hidden_states = input_tensor‐
7 .....    hidden_states = self.norm1(hidden_states)‐
8 .....    hidden_states = self.nonlinearity(hidden_states)‐
9 .....    hidden_states = self.conv1(hidden_states)‐
10 .....   if self.time_emb_proj is not None:‐
11 .....     if not self.skip_time_act:‐
12 .....       temb = self.nonlinearity(temb)‐
13 .....       temb = self.time_emb_proj(temb)[::, ::, None, None]‐
14 .....       hidden_states = hidden_states + temb‐
15 .....     hidden_states = self.nonlinearity(hidden_states)‐
16 .....     hidden_states = self.dropout(hidden_states)‐
17 .....     hidden_states = self.conv2(hidden_states, scale) if not USE_PEFT_BACKEND else self.conv2(hidden_states)‐
18 .....     output_tensor = (input_tensor + hidden_states)‐
19 .....     return output_tensor
```

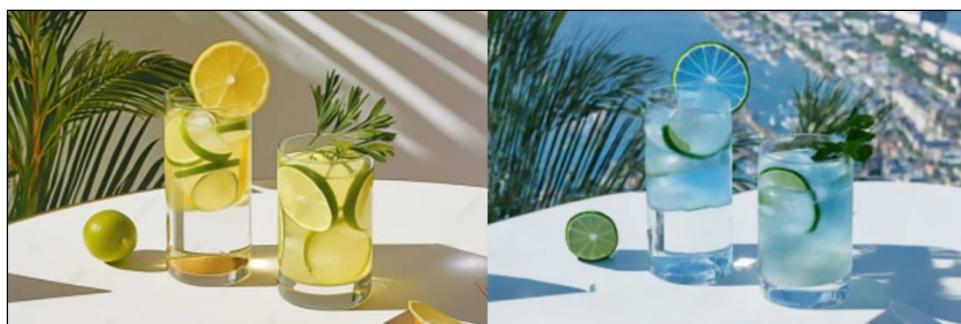
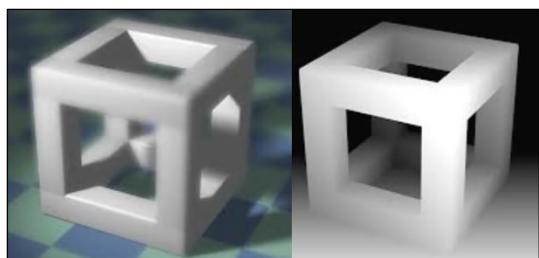


Chapter 17:

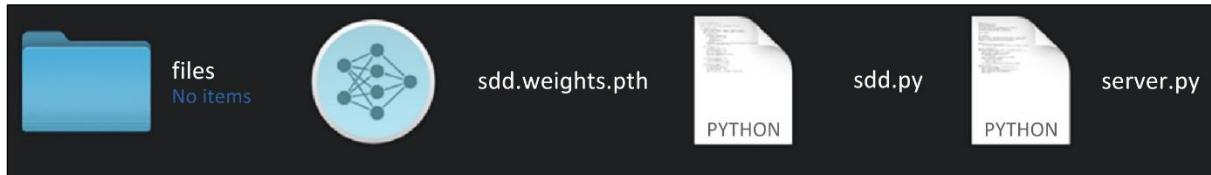




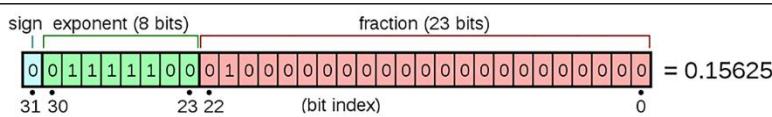




Chapter 18:



```
INFO:     Started server process [71575]
INFO:     Waiting for application startup.
INFO:     Application startup complete.
INFO:     Uvicorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
INFO:     127.0.0.1:62268 - "GET / HTTP/1.1" 404 Not Found
INFO:     127.0.0.1:62268 - "GET /favicon.ico HTTP/1.1" 404 Not Found
INFO:     127.0.0.1:62274 - "GET /docs HTTP/1.1" 200 OK
INFO:     127.0.0.1:62274 - "GET /openapi.json HTTP/1.1" 200 OK
```

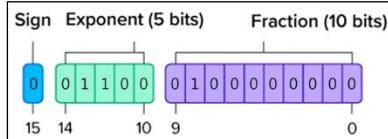


The real value assumed by a given 32-bit *binary32* data with a given *sign*, biased exponent *e* (the 8-bit unsigned integer), and a 23-bit *fraction* is

$$(-1)^{b_{31}} \times 2^{(b_{30}b_{29}\dots b_3)_2 - 127} \times (1.b_{22}b_{21}\dots b_0)_2,$$

which yields

$$\text{value} = (-1)^{\text{sign}} \times 2^{(E-127)} \times \left(1 + \sum_{i=1}^{23} b_{23-i} 2^{-i}\right).$$



```
⇒ [8/9] COPY . /app          0.8s
⇒ [9/9] RUN ls /app          0.6s
⇒ exporting to image         0.7s
⇒ ⇒ exporting layers         0.6s
⇒ ⇒ writing image sha256:1b921c0b5f03a6eab430e518e341987e54b3f1462ec55f5e8 0.0s
⇒ ⇒ naming to docker.io/library/sdd:latest 0.0s
```

```
❯ make docker-run
docker run -p 5000:5000 sdd:latest
INFO:     Started server process [1]
INFO:     Waiting for application startup.
INFO:     Application startup complete.
INFO:     Uvicorn running on http://0.0.0.0:5000 (Press CTRL+C to quit)
```

```
curl -X POST "http://127.0.0.1:5000/predict" -H "accept: application/json" -H "Content-Type: multipart/form-data" -F "file=@/tmp/non-defect.png;type=image/png"
```

